

Optimistic Stack Allocation and Dynamic Heapification for Managed Runtimes

Software Engineering Research in India (SERI 2024)



Aditya Anand

Indian Institute of Technology Bombay

20th July 2024

Content of the slides

Aditya Anand^{*}, Solai Adithya[†], Swapnil Rustagi[†], Priyam Seth[†], Vijay Sundaresan[#], Daryl Maier[#], V Krishna Nandivada⁺ and **Manas Thakur**^{*}. “**Optimistic Stack Allocation and Dynamic Heapification in Managed Runtimes**”, *PLDI 2024*.

^{*}IIT Bombay, [†]IIT Mandi, [#]IBM Canada, ⁺IIT Madras



Objects in Java

Objects in Java

- Managed runtime for Java allocate objects on the heap.

Objects in Java

- Managed runtime for Java allocate objects on the heap.

- `A a = new A(); // On heap`

Objects in Java

- Managed runtime for Java allocate objects on the heap.

```
• A a = new A(); // On heap
```

- Benefits:

Objects in Java

- Managed runtime for Java allocate objects on the heap.

```
• A a = new A(); // On heap
```

- Benefits:
 - Unburden programmer from making complex allocation-deallocation decisions and reduce the possibility of harmful memory bugs.

Objects in Java

- Managed runtime for Java allocate objects on the heap.

```
• A a = new A(); // On heap
```

- Benefits:
 - Unburden programmer from making complex allocation-deallocation decisions and reduce the possibility of harmful memory bugs.
 - Automatic garbage collection.

Objects in Java

- Managed runtime for Java allocate objects on the heap.

```
• A a = new A(); // On heap
```

- Benefits:
 - Unburden programmer from making complex allocation-deallocation decisions and reduce the possibility of harmful memory bugs.
 - Automatic garbage collection.
- Challenges:

Objects in Java

- Managed runtime for Java allocate objects on the heap.

```
• A a = new A(); // On heap
```

- Benefits:
 - Unburden programmer from making complex allocation-deallocation decisions and reduce the possibility of harmful memory bugs.
 - Automatic garbage collection.
- Challenges:
 - Access time is high.

Objects in Java

- Managed runtime for Java allocate objects on the heap.

```
• A a = new A(); // On heap
```

- Benefits:
 - Unburden programmer from making complex allocation-deallocation decisions and reduce the possibility of harmful memory bugs.
 - Automatic garbage collection.
- Challenges:
 - Access time is high.
 - Garbage collection is an overhead.

Stack Allocation

Stack Allocation

- Memory allocated on stack:

Stack Allocation

- Memory allocated on stack:
 - Less access time.

Stack Allocation

- Memory allocated on stack:
 - Less access time.
 - Get freed up as soon as the allocating method returns.

Stack Allocation

- Memory allocated on stack:
 - Less access time.
 - Get freed up as soon as the allocating method returns.
- Escape Analysis

Stack Allocation

- Memory allocated on stack:
 - Less access time.
 - Get freed up as soon as the allocating method returns.
- **Escape Analysis**
- In case of Java:

Stack Allocation

- Memory allocated on stack:
 - Less access time.
 - Get freed up as soon as the allocating method returns.
- **Escape Analysis**
- In case of Java:
 - Escape analysis is performed: Just-in-time (JIT) compilation

Stack Allocation

- Memory allocated on stack:
 - Less access time.
 - Get freed up as soon as the allocating method returns.
- **Escape Analysis**
- In case of Java:
 - Escape analysis is performed: Just-in-time (JIT) compilation — **Imprecise**

Stack Allocation

- Memory allocated on stack:
 - Less access time.
 - Get freed up as soon as the allocating method returns.
- **Escape Analysis**
- In case of Java:
 - Escape analysis is performed: Just-in-time (JIT) compilation — **Imprecise**
 - Very few objects get allocated on stack.

Static Analysis for Stack Allocation

Static Analysis for Stack Allocation

- Perform precise (context-, flow-, field-sensitive) escape analysis statically.

Static Analysis for Stack Allocation

- Perform precise (context-, flow-, field-sensitive) escape analysis statically.
- Use statically generated escape analysis result to optimistically allocate objects on stack at runtime.

Static Analysis for Stack Allocation

- Perform precise (context-, flow-, field-sensitive) escape analysis statically.
- Use statically generated escape analysis result to **optimistically** allocate objects on stack at runtime.

Static Analysis for Stack Allocation

- Perform precise (context-, flow-, field-sensitive) escape analysis statically.
- Use statically generated escape analysis result to **optimistically** allocate objects on stack at runtime.
- **Challenges:**

Static Analysis for Stack Allocation

- Perform precise (context-, flow-, field-sensitive) escape analysis statically.
- Use statically generated escape analysis result to **optimistically** allocate objects on stack at runtime.
- **Challenges:**
 - Dynamic Features: **Dynamic Class Loading (DCL)**, **Hot-Code Replacement (HCR)** allows code changes.

Static Analysis for Stack Allocation

- Perform precise (context-, flow-, field-sensitive) escape analysis statically.
- Use statically generated escape analysis result to **optimistically** allocate objects on stack at runtime.
- **Challenges:**
 - Dynamic Features: **Dynamic Class Loading** (DCL), **Hot-Code Replacement** (HCR) allows code changes.
 - An object that was stack allocated based on static-analysis results, might start escaping at run-time.

Static Analysis for Stack Allocation

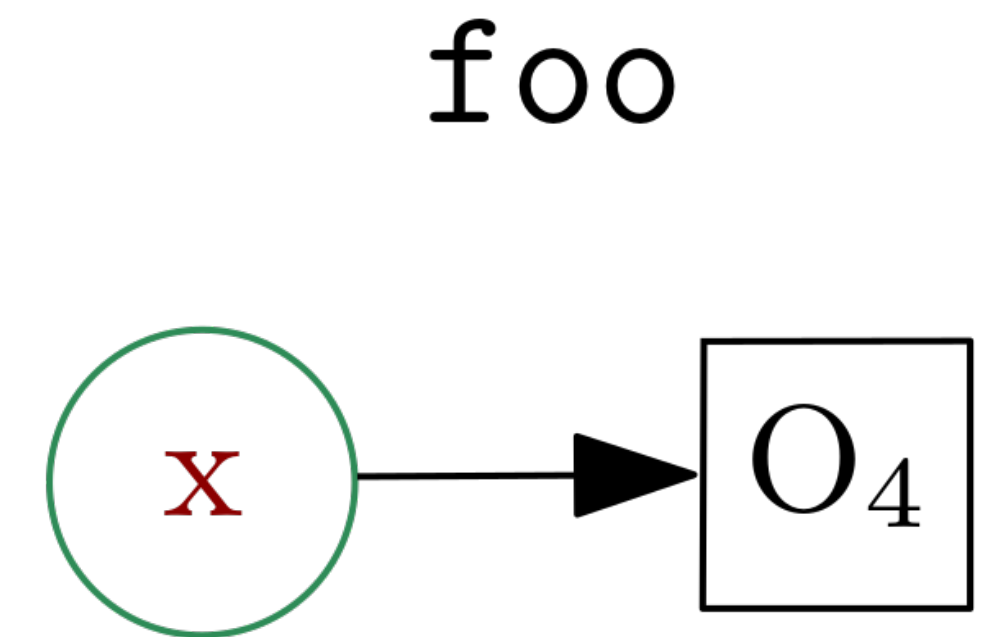
- Perform precise (context-, flow-, field-sensitive) escape analysis statically.
- Use statically generated escape analysis result to **optimistically** allocate objects on stack at runtime.
- **Challenges:**
 - Dynamic Features: **Dynamic Class Loading** (DCL), **Hot-Code Replacement** (HCR) allows code changes.
 - An object that was stack allocated based on static-analysis results, might start escaping at run-time.
- How to safely allocate objects on stack in a managed runtime?

Motivating Example

```
1. class A {
2.     A f;
3.     void foo(A q, A r) {
4.         A x = new A(); // O4
5.         A y = new A(); // O5
6.         x.f = new A(); // O6
7.         A p = x.f;
8.         bar(p, y);
9.         r.zar(p, q);
10.    } /* method foo */
11.    void zar(A p, A q) { . . . }
12.    void bar(A p1, A p2) {
13.        p1.f = p2;
14.    } /* method bar */
15. } /* class A */
```

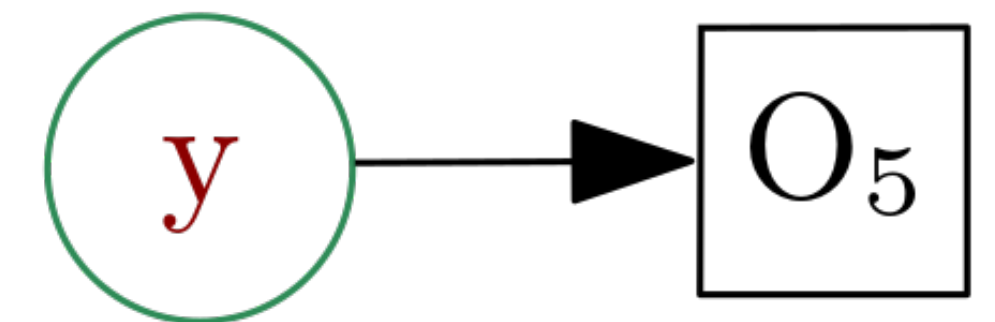
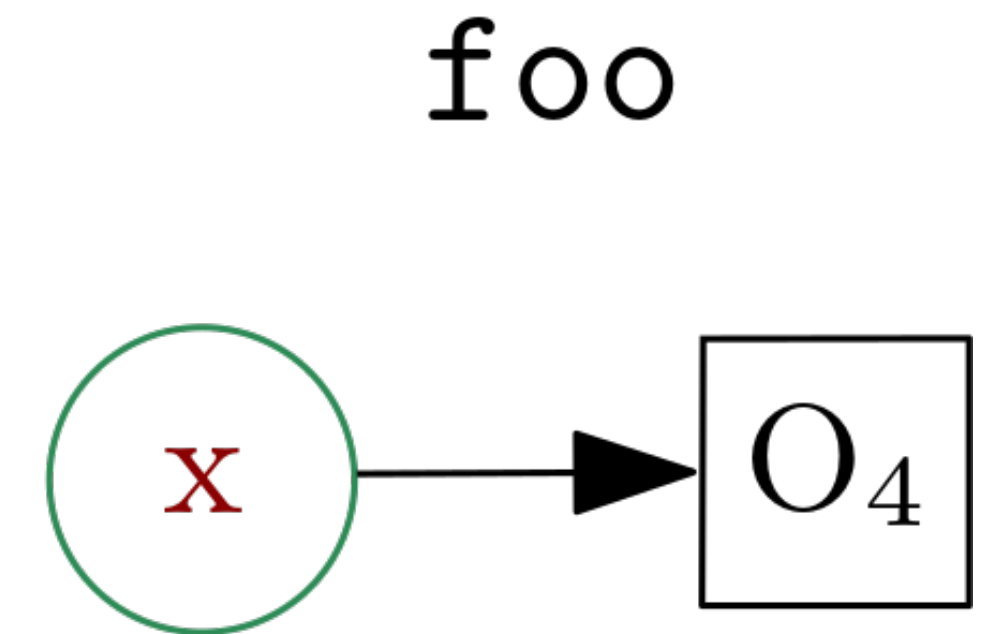
Motivating Example

```
1. class A {  
2.     A f;  
3.     void foo(A q, A r) {  
4.         A x = new A(); // O4  
5.         A y = new A(); // O5  
6.         x.f = new A(); // O6  
7.         A p = x.f;  
8.         bar(p, y);  
9.         r.zar(p, q);  
10.    } /* method foo */  
11.    void zar(A p, A q) { . . . }  
12.    void bar(A p1, A p2) {  
13.        p1.f = p2;  
14.    } /* method bar */  
15. } /* class A */
```



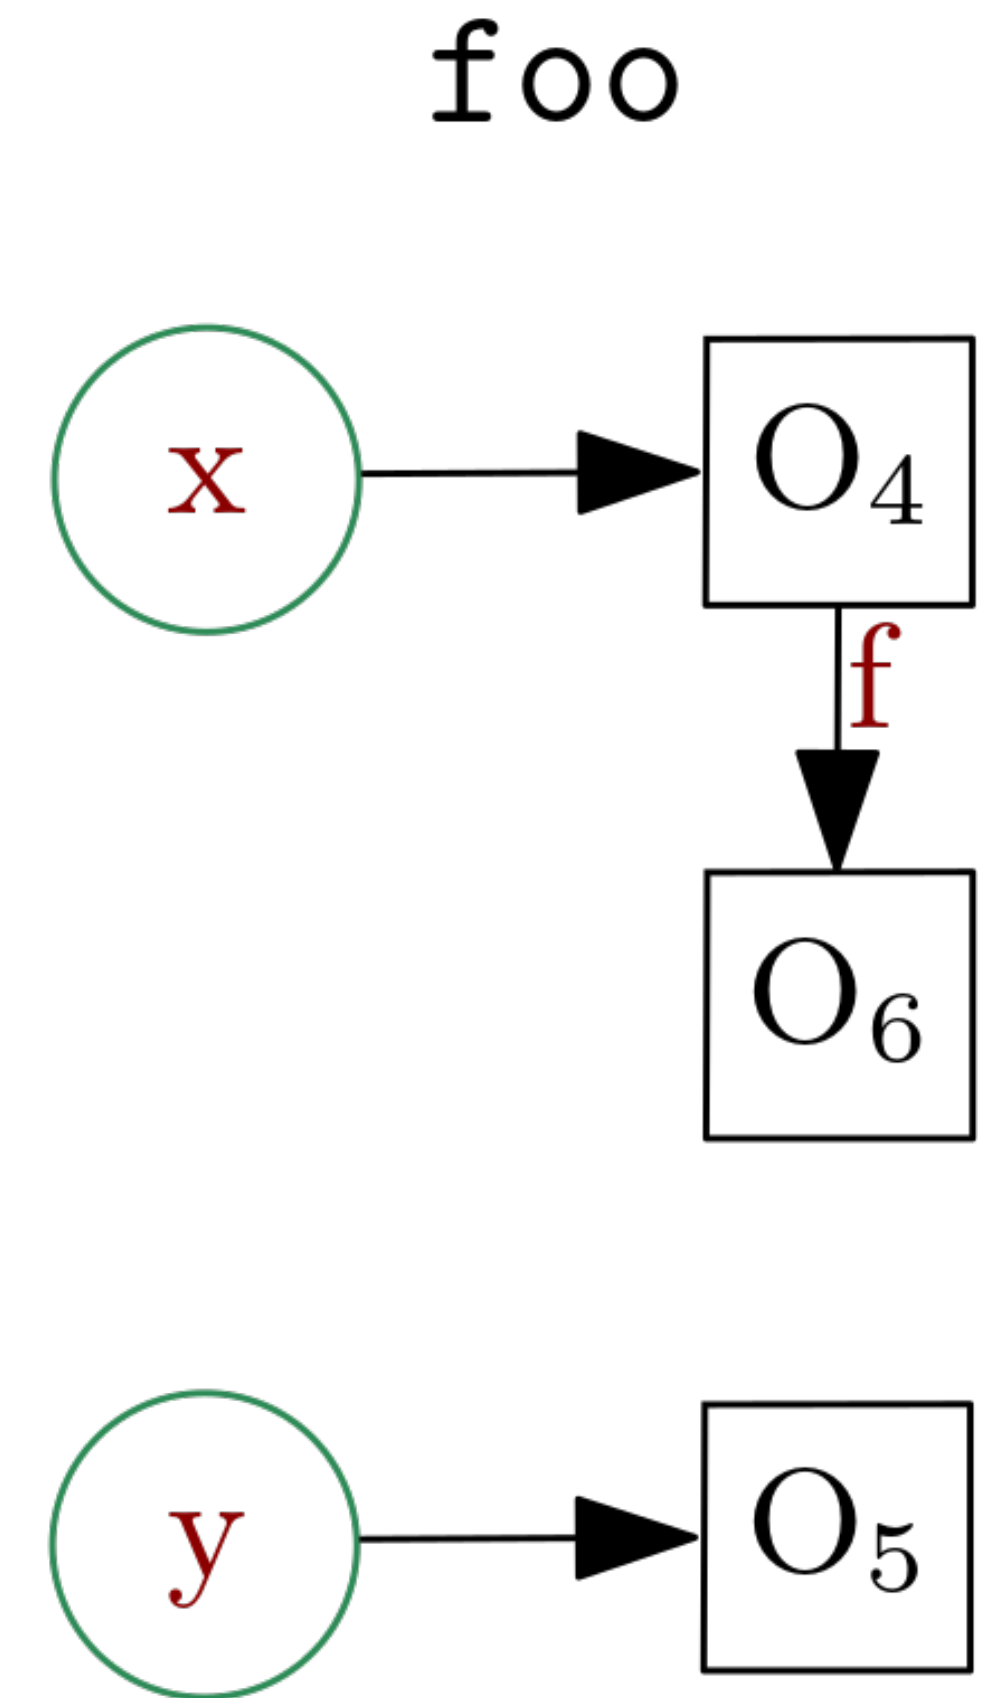
Motivating Example

```
1. class A {
2.     A f;
3.     void foo(A q, A r) {
4.         A x = new A(); // O4
5.         A y = new A(); // O5
6.         x.f = new A(); // O6
7.         A p = x.f;
8.         bar(p, y);
9.         r.zar(p, q);
10.    } /* method foo */
11.    void zar(A p, A q) { . . . }
12.    void bar(A p1, A p2) {
13.        p1.f = p2;
14.    } /* method bar */
15. } /* class A */
```



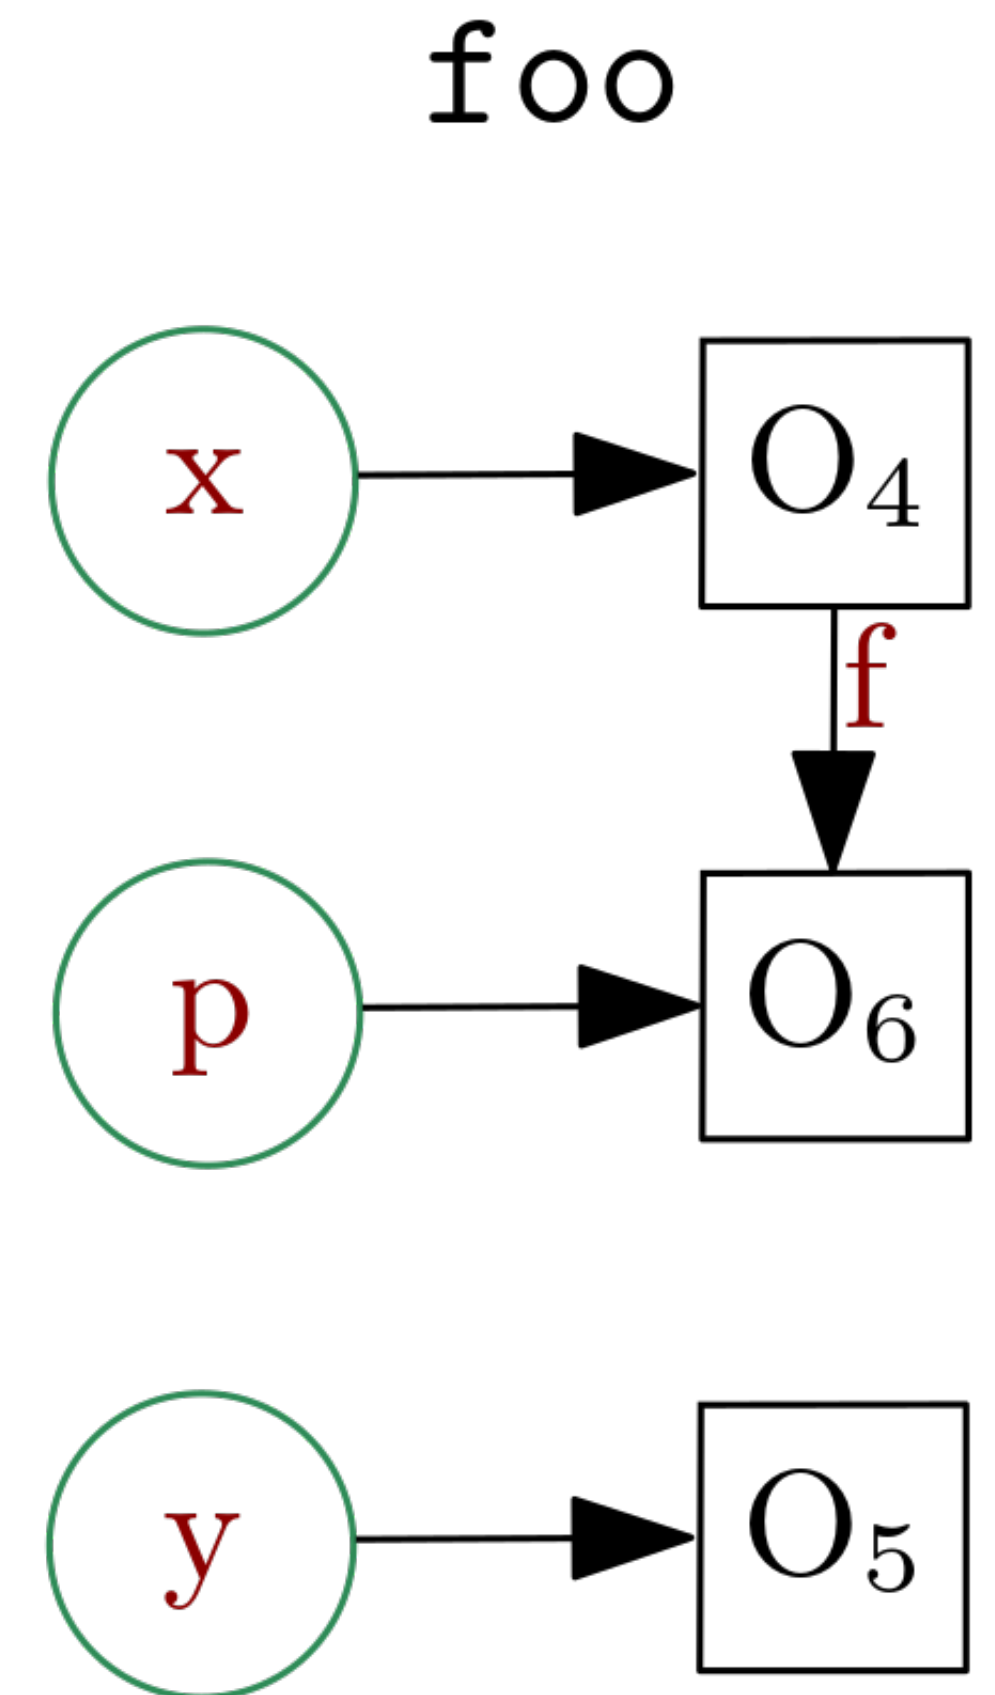
Motivating Example

```
1. class A {  
2.     A f;  
3.     void foo(A q, A r) {  
4.         A x = new A(); // O4  
5.         A y = new A(); // O5  
6.         x.f = new A(); // O6  
7.         A p = x.f;  
8.         bar(p, y);  
9.         r.zar(p, q);  
10.    } /* method foo */  
11. void zar(A p, A q) { . . . }  
12. void bar(A p1, A p2) {  
13.     p1.f = p2;  
14. } /* method bar */  
15. } /* class A */
```



Motivating Example

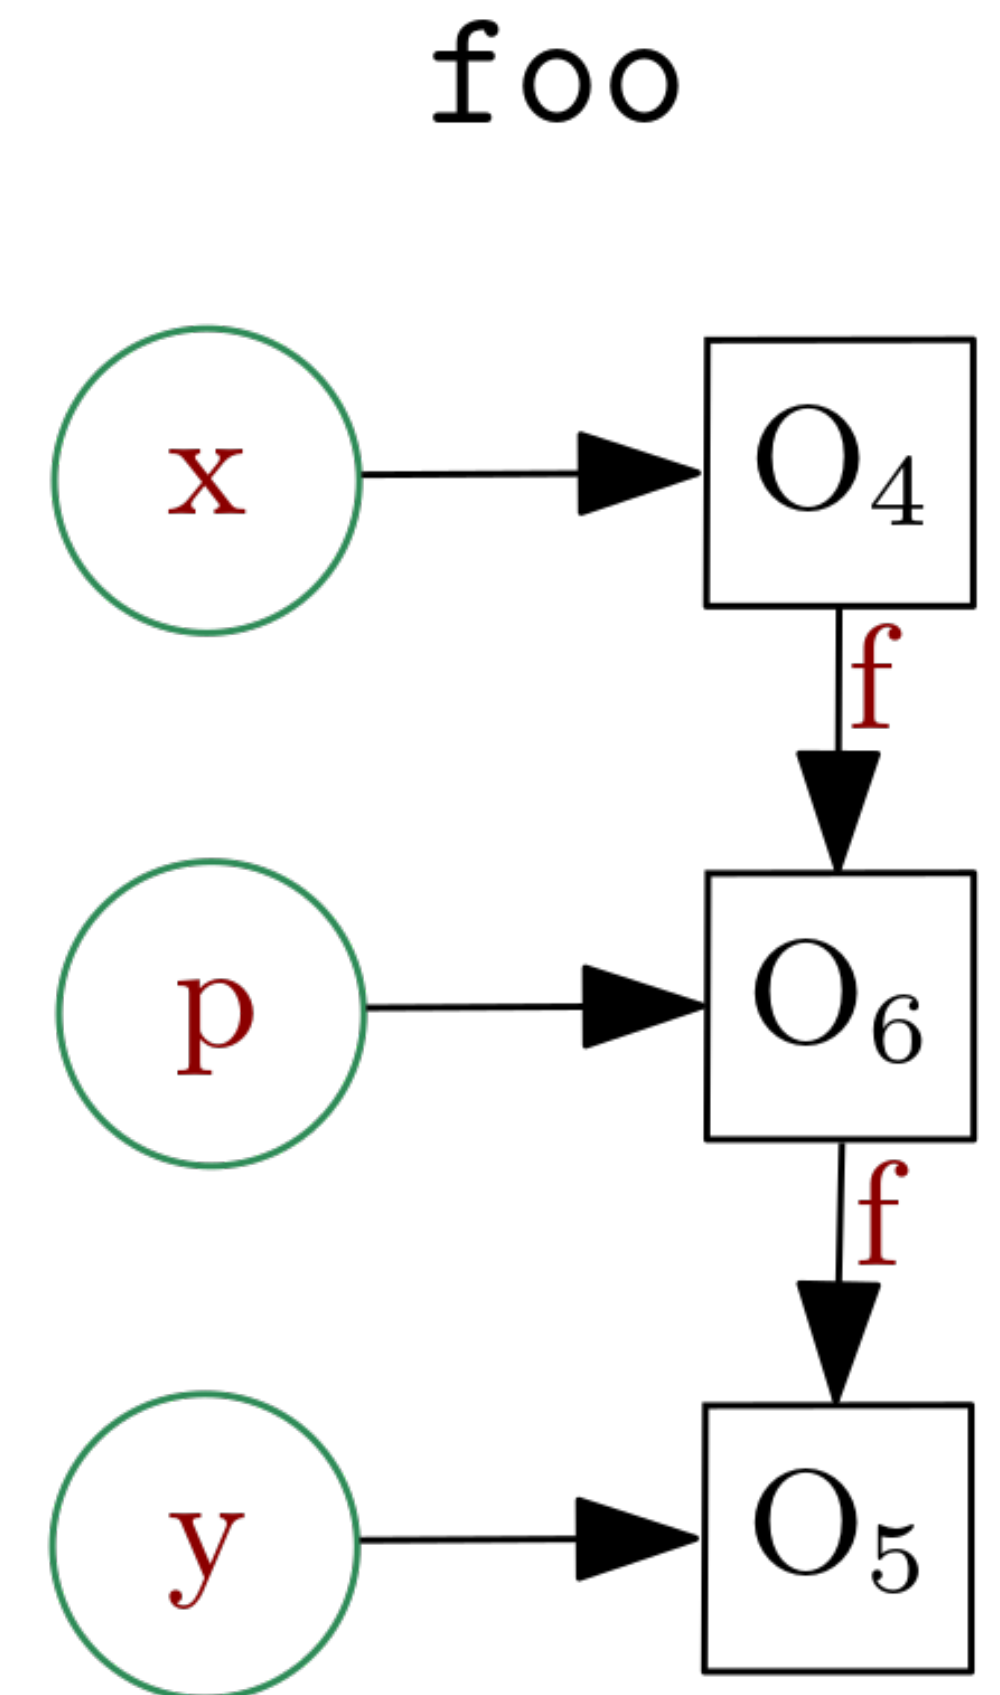
```
1. class A {  
2.     A f;  
3.     void foo(A q, A r) {  
4.         A x = new A(); // O4  
5.         A y = new A(); // O5  
6.         x.f = new A(); // O6  
7.         A p = x.f;  
8.         bar(p, y);  
9.         r.zar(p, q);  
10.    } /* method foo */  
11.    void zar(A p, A q) { . . . }  
12.    void bar(A p1, A p2) {  
13.        p1.f = p2;  
14.    } /* method bar */  
15. } /* class A */
```



Motivating Example

```
1. class A {  
2.     A f;  
3.     void foo(A q, A r) {  
4.         A x = new A(); // O4  
5.         A y = new A(); // O5  
6.         x.f = new A(); // O6  
7.         A p = x.f;  
8.         bar(p, y);  
9.         r.zar(p, q);  
10.    } /* method foo */
```

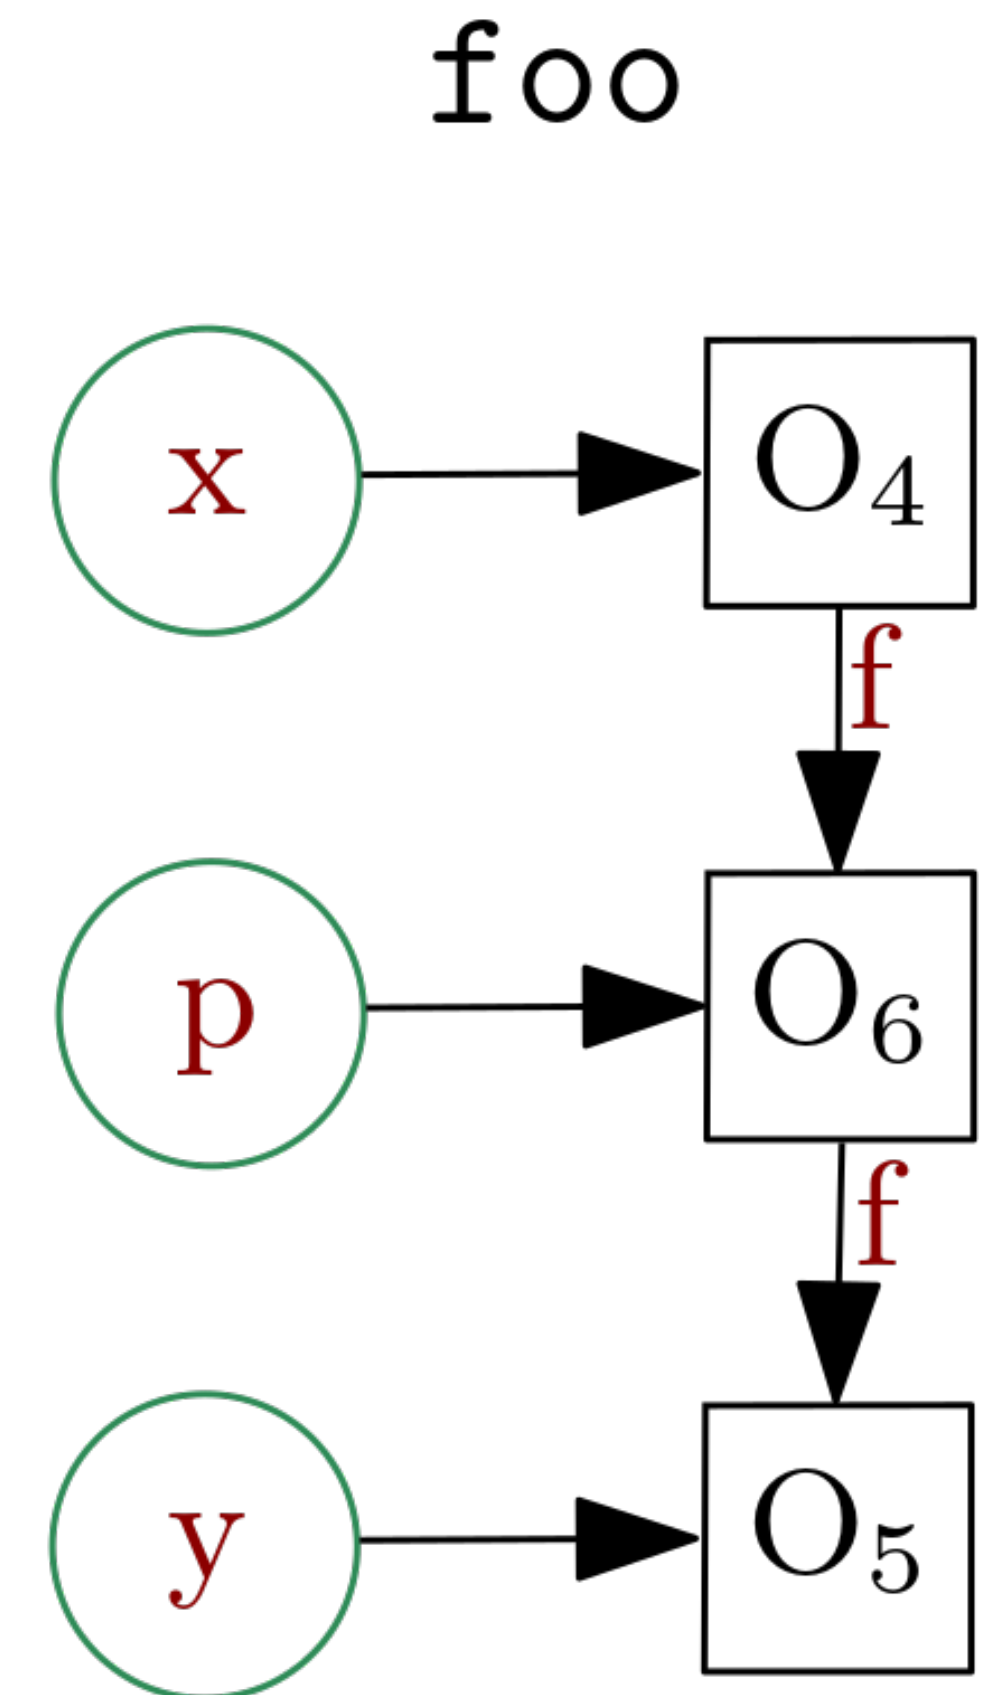
```
11. void zar(A p, A q) { . . . }  
12. void bar(A p1, A p2) {  
13.     p1.f = p2;  
14. } /* method bar */  
15. } /* class A */
```



Motivating Example

```
1. class A {  
2.     A f;  
3.     void foo(A q, A r) {  
4.         A x = new A(); // O4  
5.         A y = new A(); // O5  
6.         x.f = new A(); // O6  
7.         A p = x.f;  
8.         bar(p, y);  
9.         r.zar(p, q);  
10.    } /* method foo */
```

```
11. void zar(A p, A q) { . . . }  
12. void bar(A p1, A p2) {  
13.     p1.f = p2;  
14. } /* method bar */  
15. } /* class A */
```

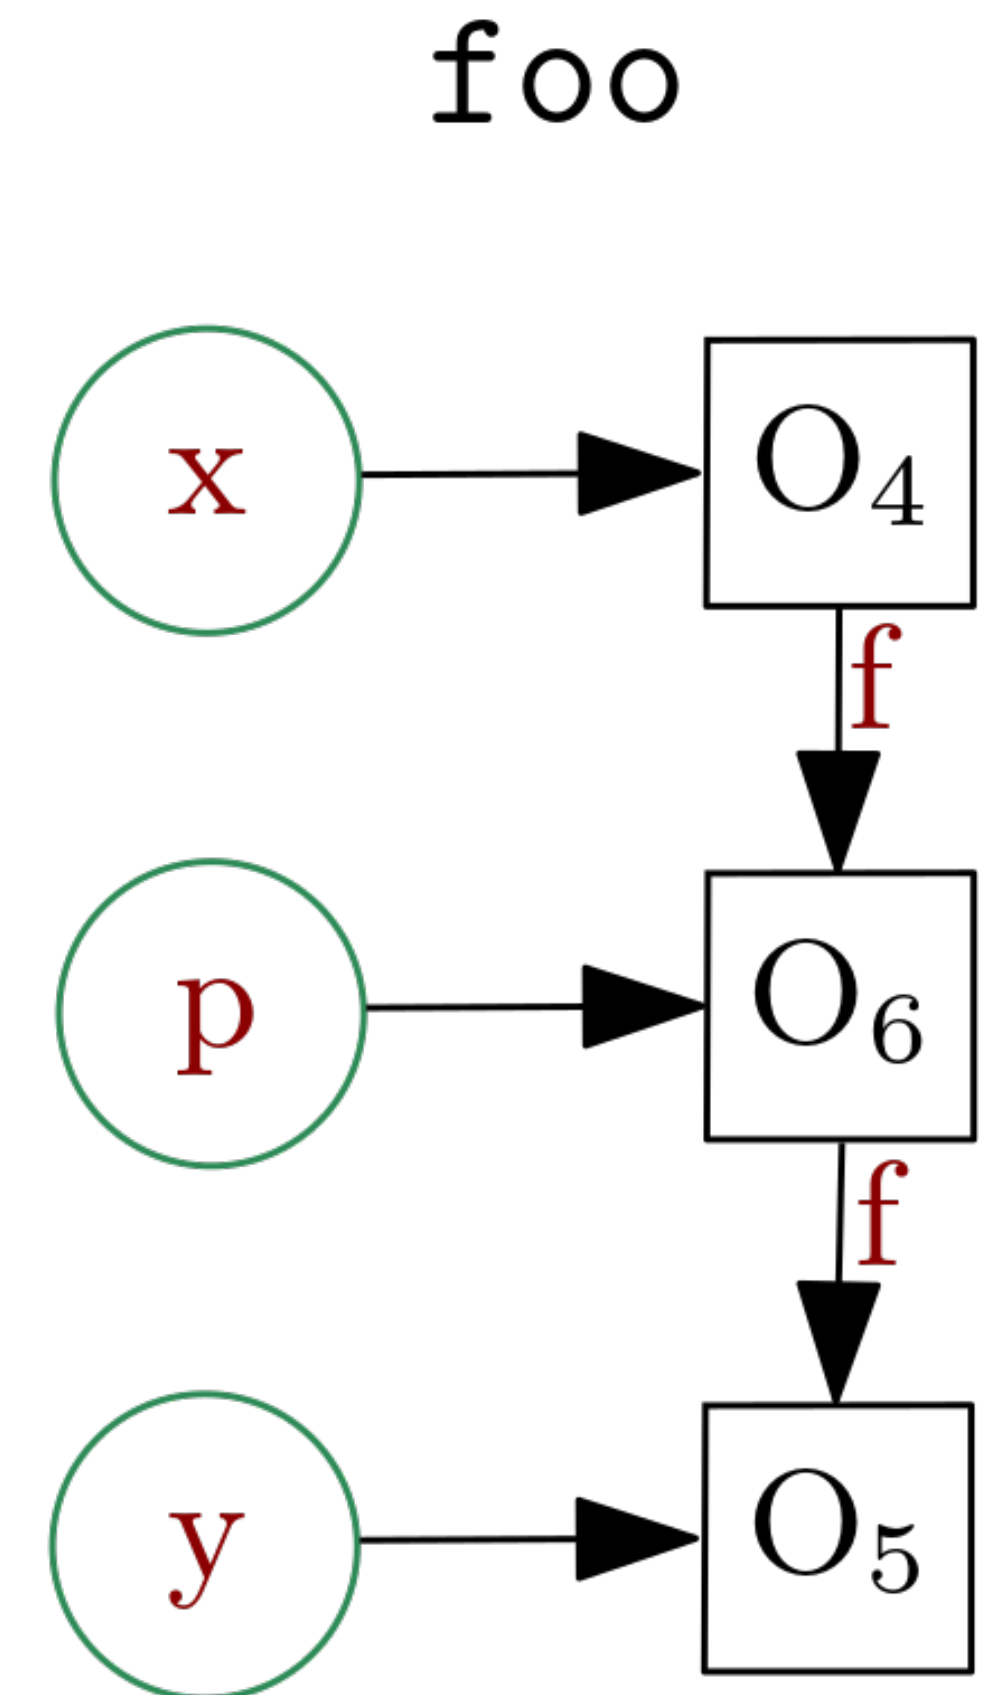


Motivating Example

```
1. class A {  
2.   A f;  
3.   void foo(A q, A r) {  
4.     A x = new A(); // O4  
5.     A y = new A(); // O5  
6.     x.f = new A(); // O6  
7.     A p = x.f;  
8.     bar(p, y);  
9.     r.zar(p, q);  
10.  } /* method foo */
```

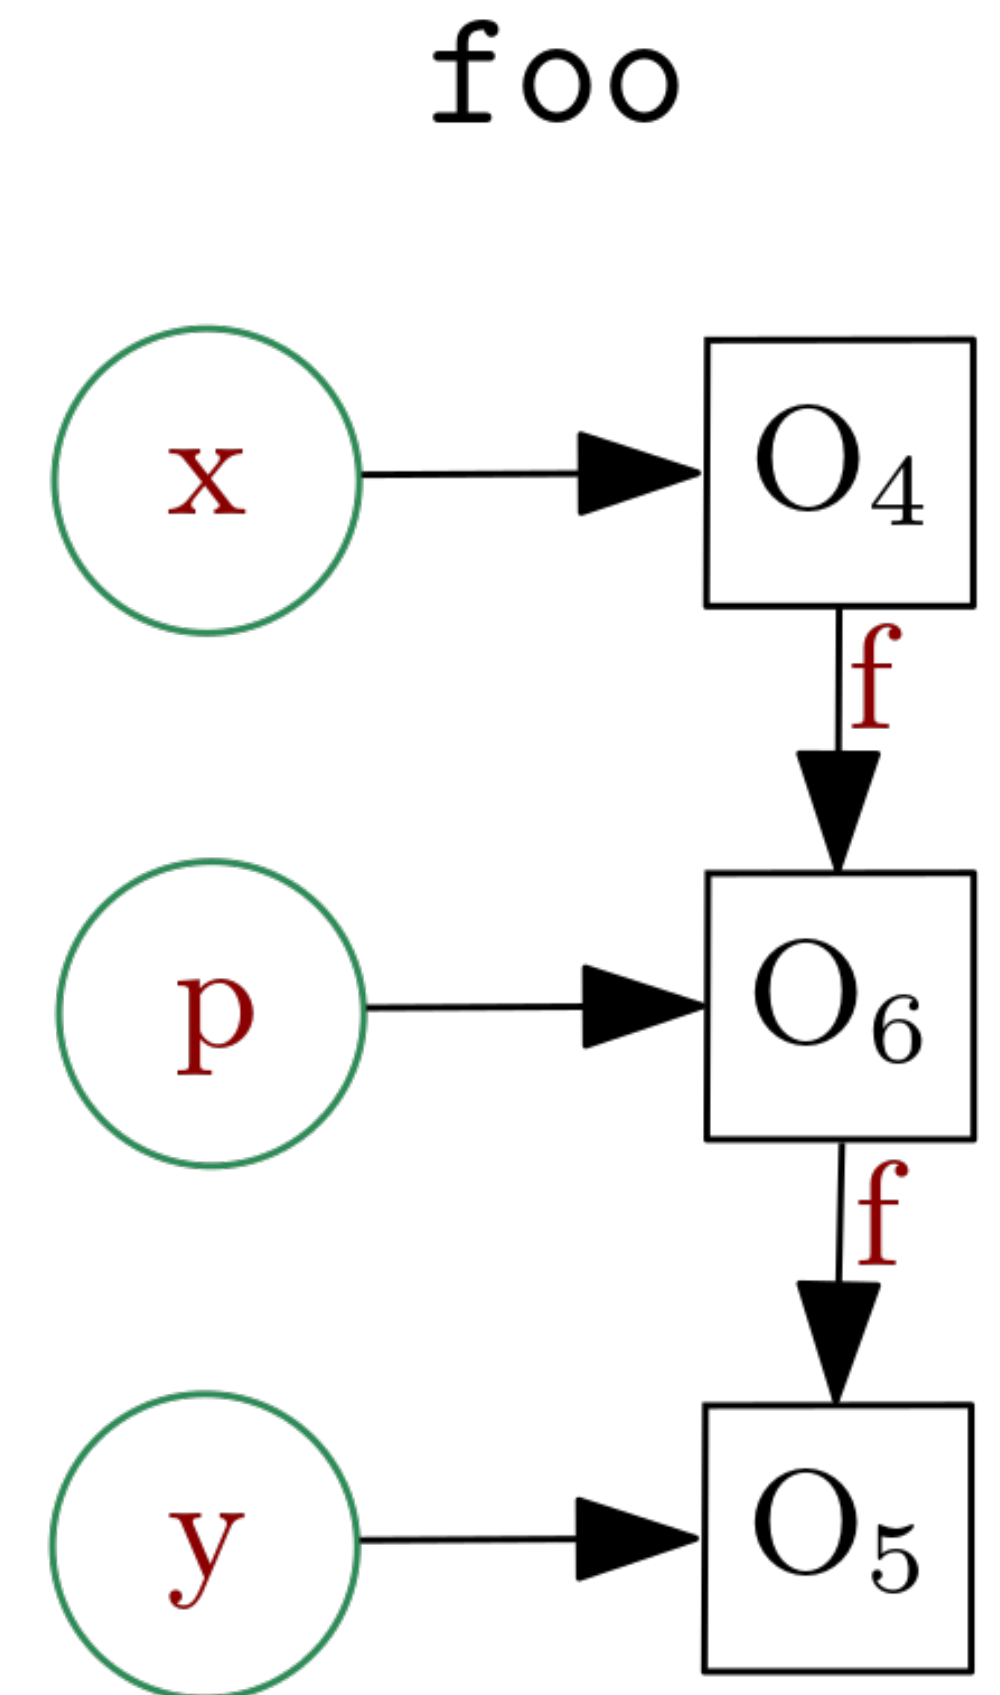
```
11. void zar(A p, A q) { . . . }  
12. void bar(A p1, A p2) {  
13.   p1.f = p2;  
14. } /* method bar */  
15. } /* class A */
```

Stack Allocate
O₄, O₅ and O₆



Motivating Example

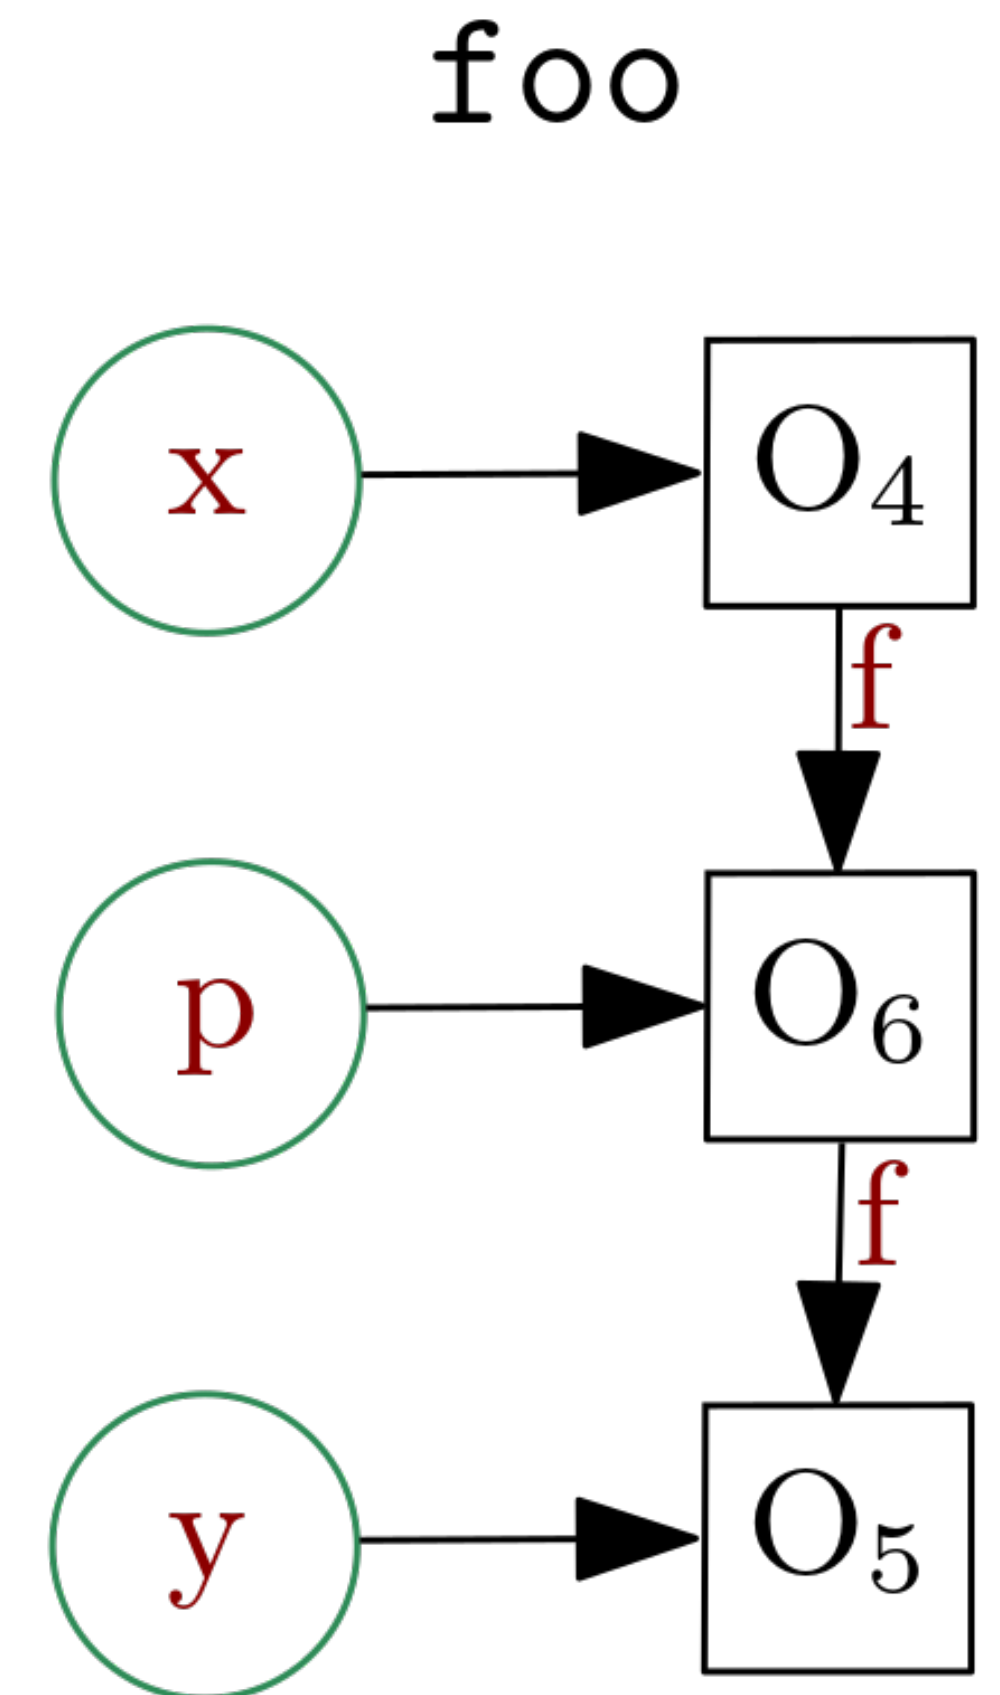
```
1. class A {
2.     A f;
3.     void foo(A q, A r) {
4.         A x = new A(); // O4
5.         A y = new A(); // O5
6.         x.f = new A(); // O6
7.         A p = x.f;
8.         bar(p, y);
9.         r.zar(p, q);
10.    } /* method foo */
11.    void zar(A p, A q) { . . . }
12.    void bar(A p1, A p2) {
13.        p1.f = p2;
14.    } /* method bar */
15. } /* class A */
16. class B extends A
17.     void zar(A p, A q) {
18.         q.f = p;
19.     } /* method zar */
20. } /* class B */
```



Motivating Example

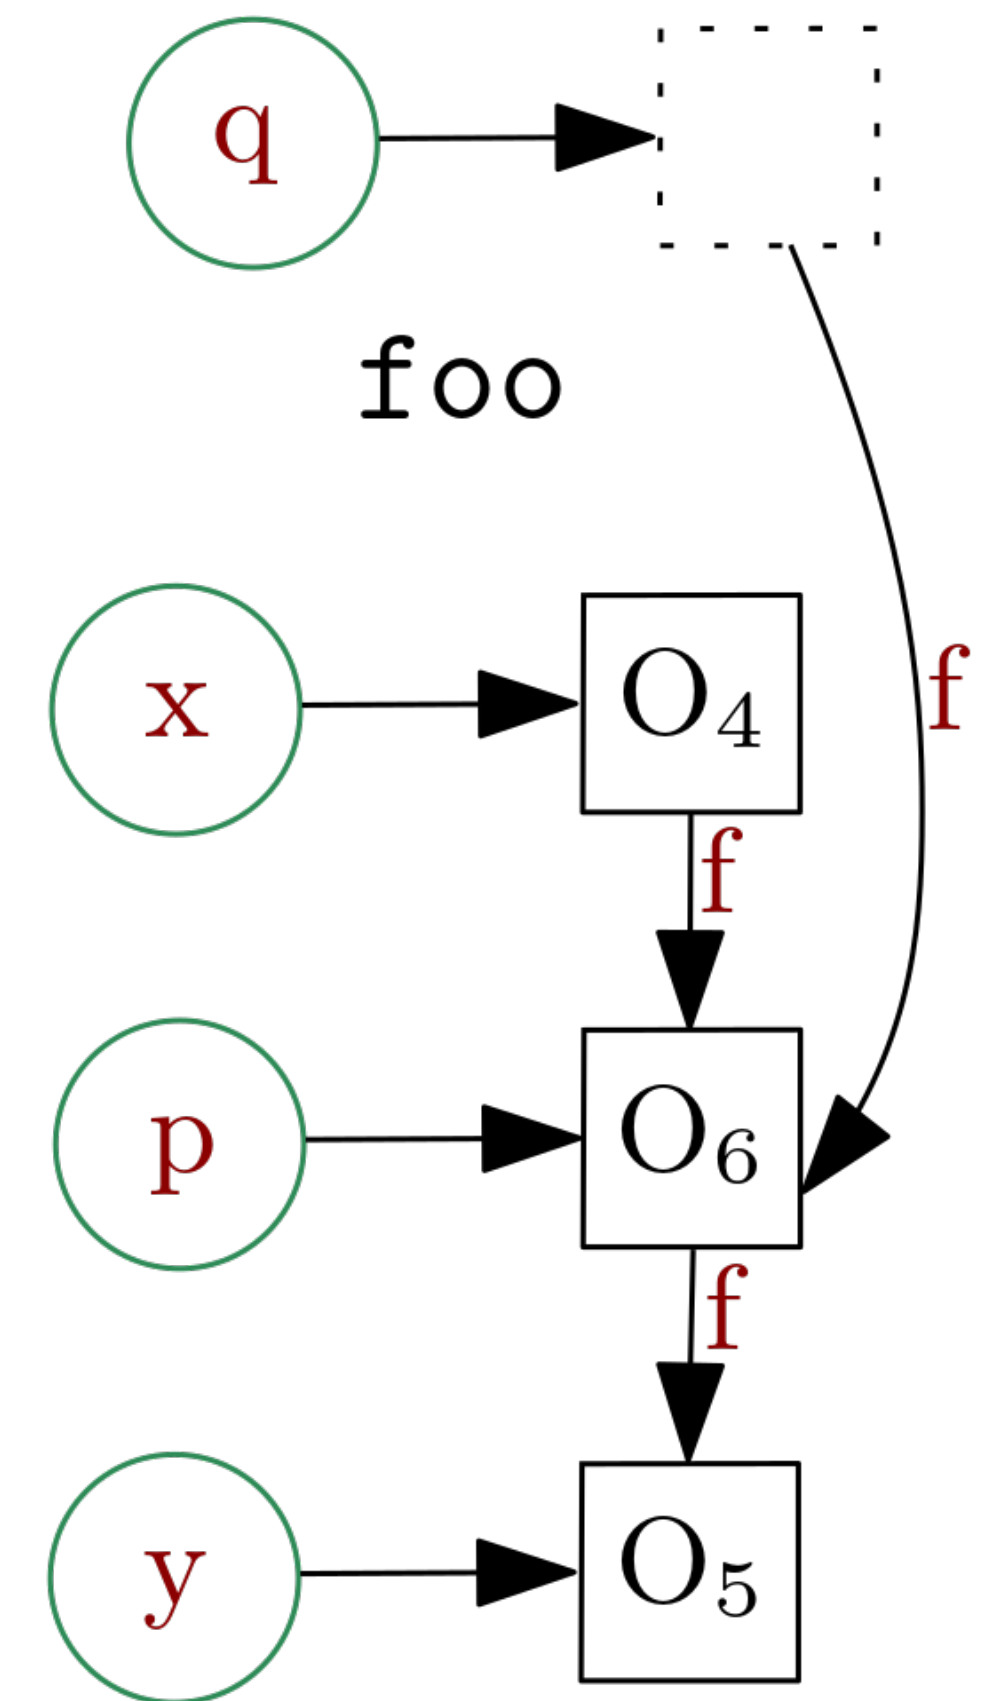
```
1. class A {
2.     A f;
3.     void foo(A q, A r) {
4.         A x = new A(); // O4
5.         A y = new A(); // O5
6.         x.f = new A(); // O6
7.         A p = x.f;
8.         bar(p, y);
9.         r.zar(p, q);
10.    } /* method foo */
11.    void zar(A p, A q) { . . . }
12.    void bar(A p1, A p2) {
13.        p1.f = p2;
14.    } /* method bar */
15. } /* class A */
16. class B extends A
17.     void zar(A p, A q) {
18.         q.f = p;
19.     } /* method zar */
20. } /* class B */
```

Dynamically
loaded



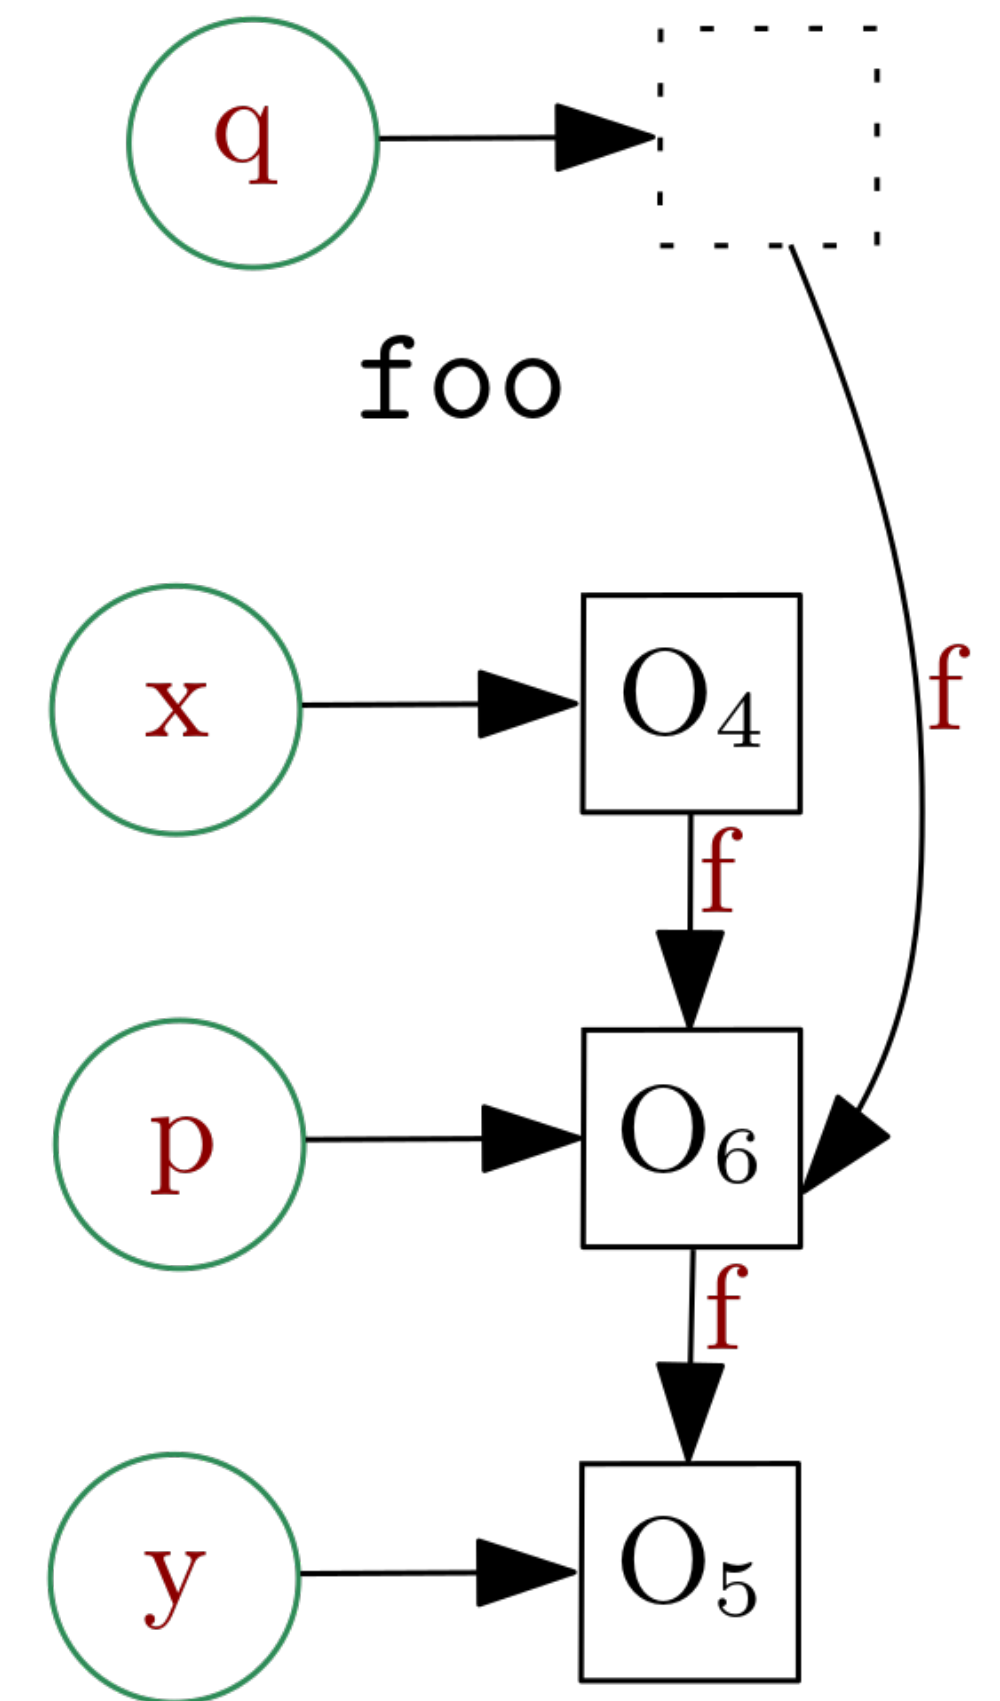
Motivating Example

```
1. class A {
2.     A f;
3.     void foo(A q, A r) {
4.         A x = new A(); // O4
5.         A y = new A(); // O5
6.         x.f = new A(); // O6
7.         A p = x.f;
8.         bar(p, y);
9.         r.zar(p, q);
10.    } /* method foo */
11.    void zar(A p, A q) { . . . }
12.    void bar(A p1, A p2) {
13.        p1.f = p2;
14.    } /* method bar */
15. } /* class A */
16. class B extends A
17.     void zar(A p, A q) {
18.         q.f = p;
19.     } /* method zar */
20. } /* class B */
```



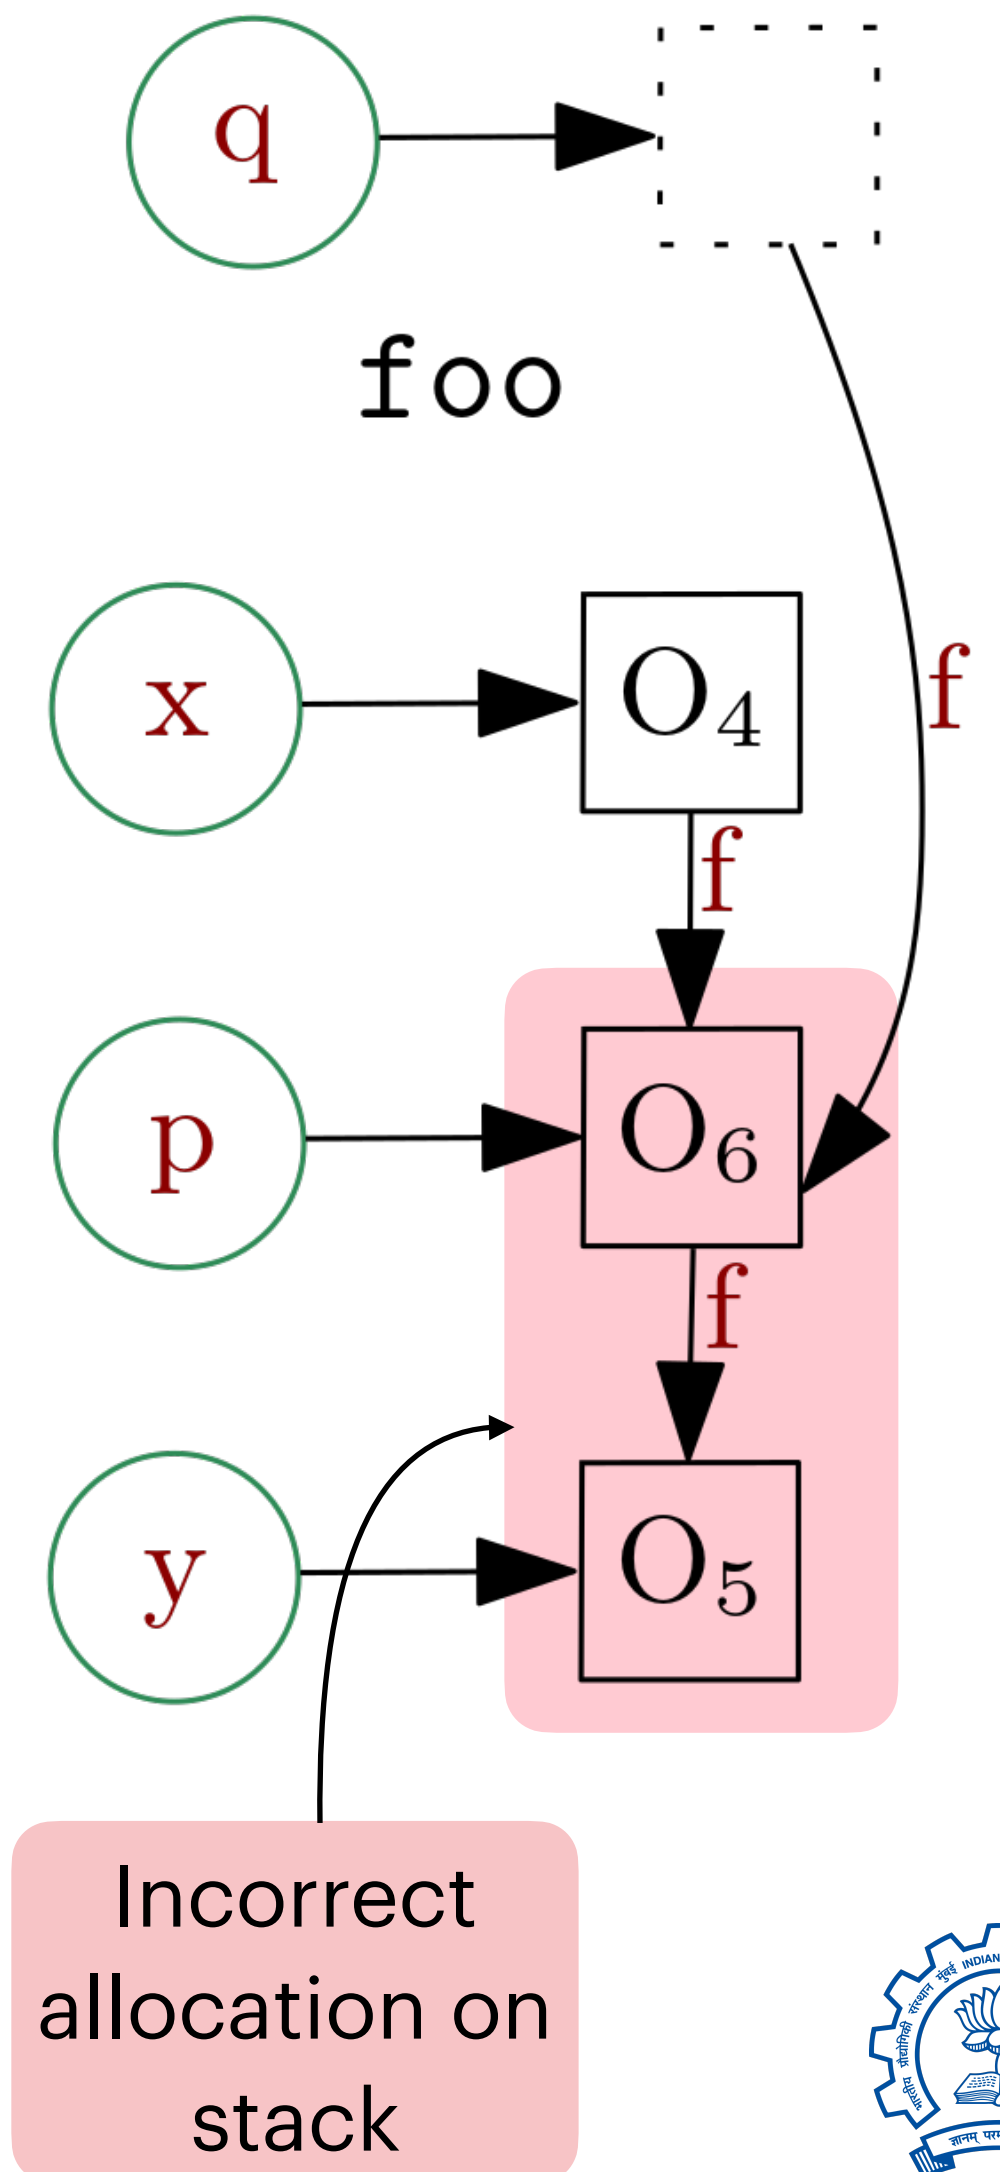
Motivating Example

```
1. class A {
2.     A f;
3.     void foo(A q, A r) {
4.         A x = new A(); // O4
5.         A y = new A(); // O5
6.         x.f = new A(); // O6
7.         A p = x.f;
8.         bar(p, y);
9.         r.zar(p, q);
10.    } /* method foo */
11.    void zar(A p, A q) { . . . }
12.    void bar(A p1, A p2) {
13.        p1.f = p2;
14.    } /* method bar */
15. } /* class A */
16. class B extends A
17.     void zar(A p, A q) {
18.         q.f = p;
19.     } /* method zar */
20. } /* class B */
```



Motivating Example

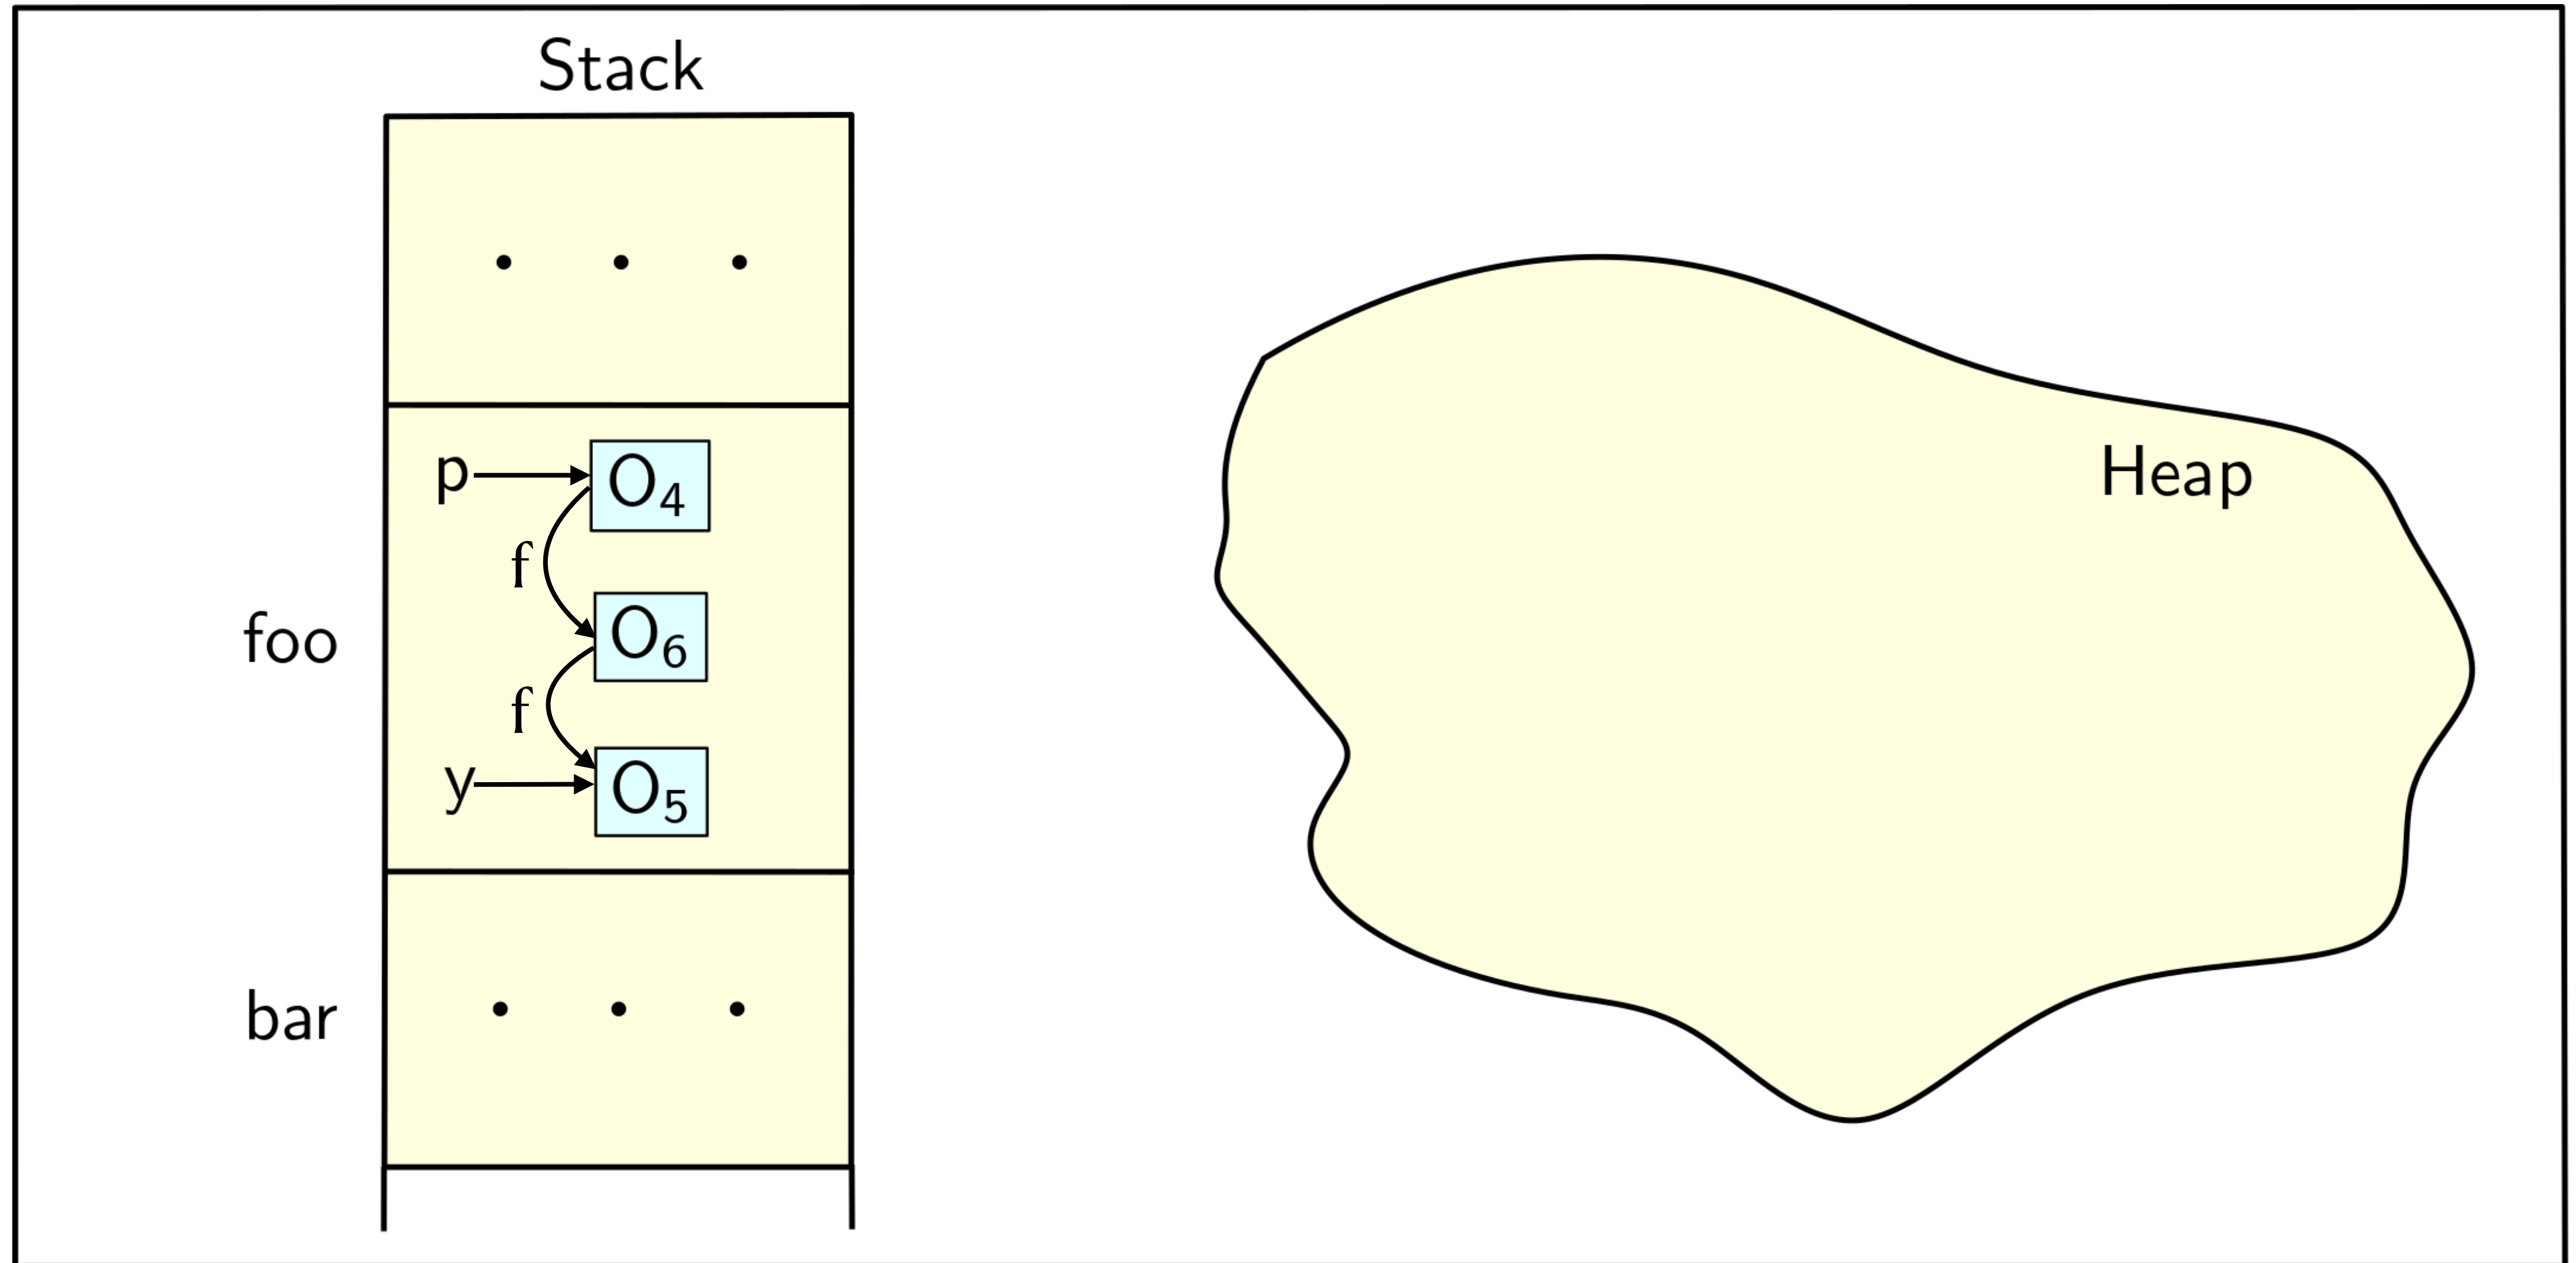
```
1. class A {
2.     A f;
3.     void foo(A q, A r) {
4.         A x = new A(); // O4
5.         A y = new A(); // O5
6.         x.f = new A(); // O6
7.         A p = x.f;
8.         bar(p, y);
9.         r.zar(p, q);
10.    } /* method foo */
11.    void zar(A p, A q) { . . . }
12.    void bar(A p1, A p2) {
13.        p1.f = p2;
14.    } /* method bar */
15. } /* class A */
16. class B extends A
17.     void zar(A p, A q) {
18.         q.f = p;
19.     } /* method zar */
20. } /* class B */
```



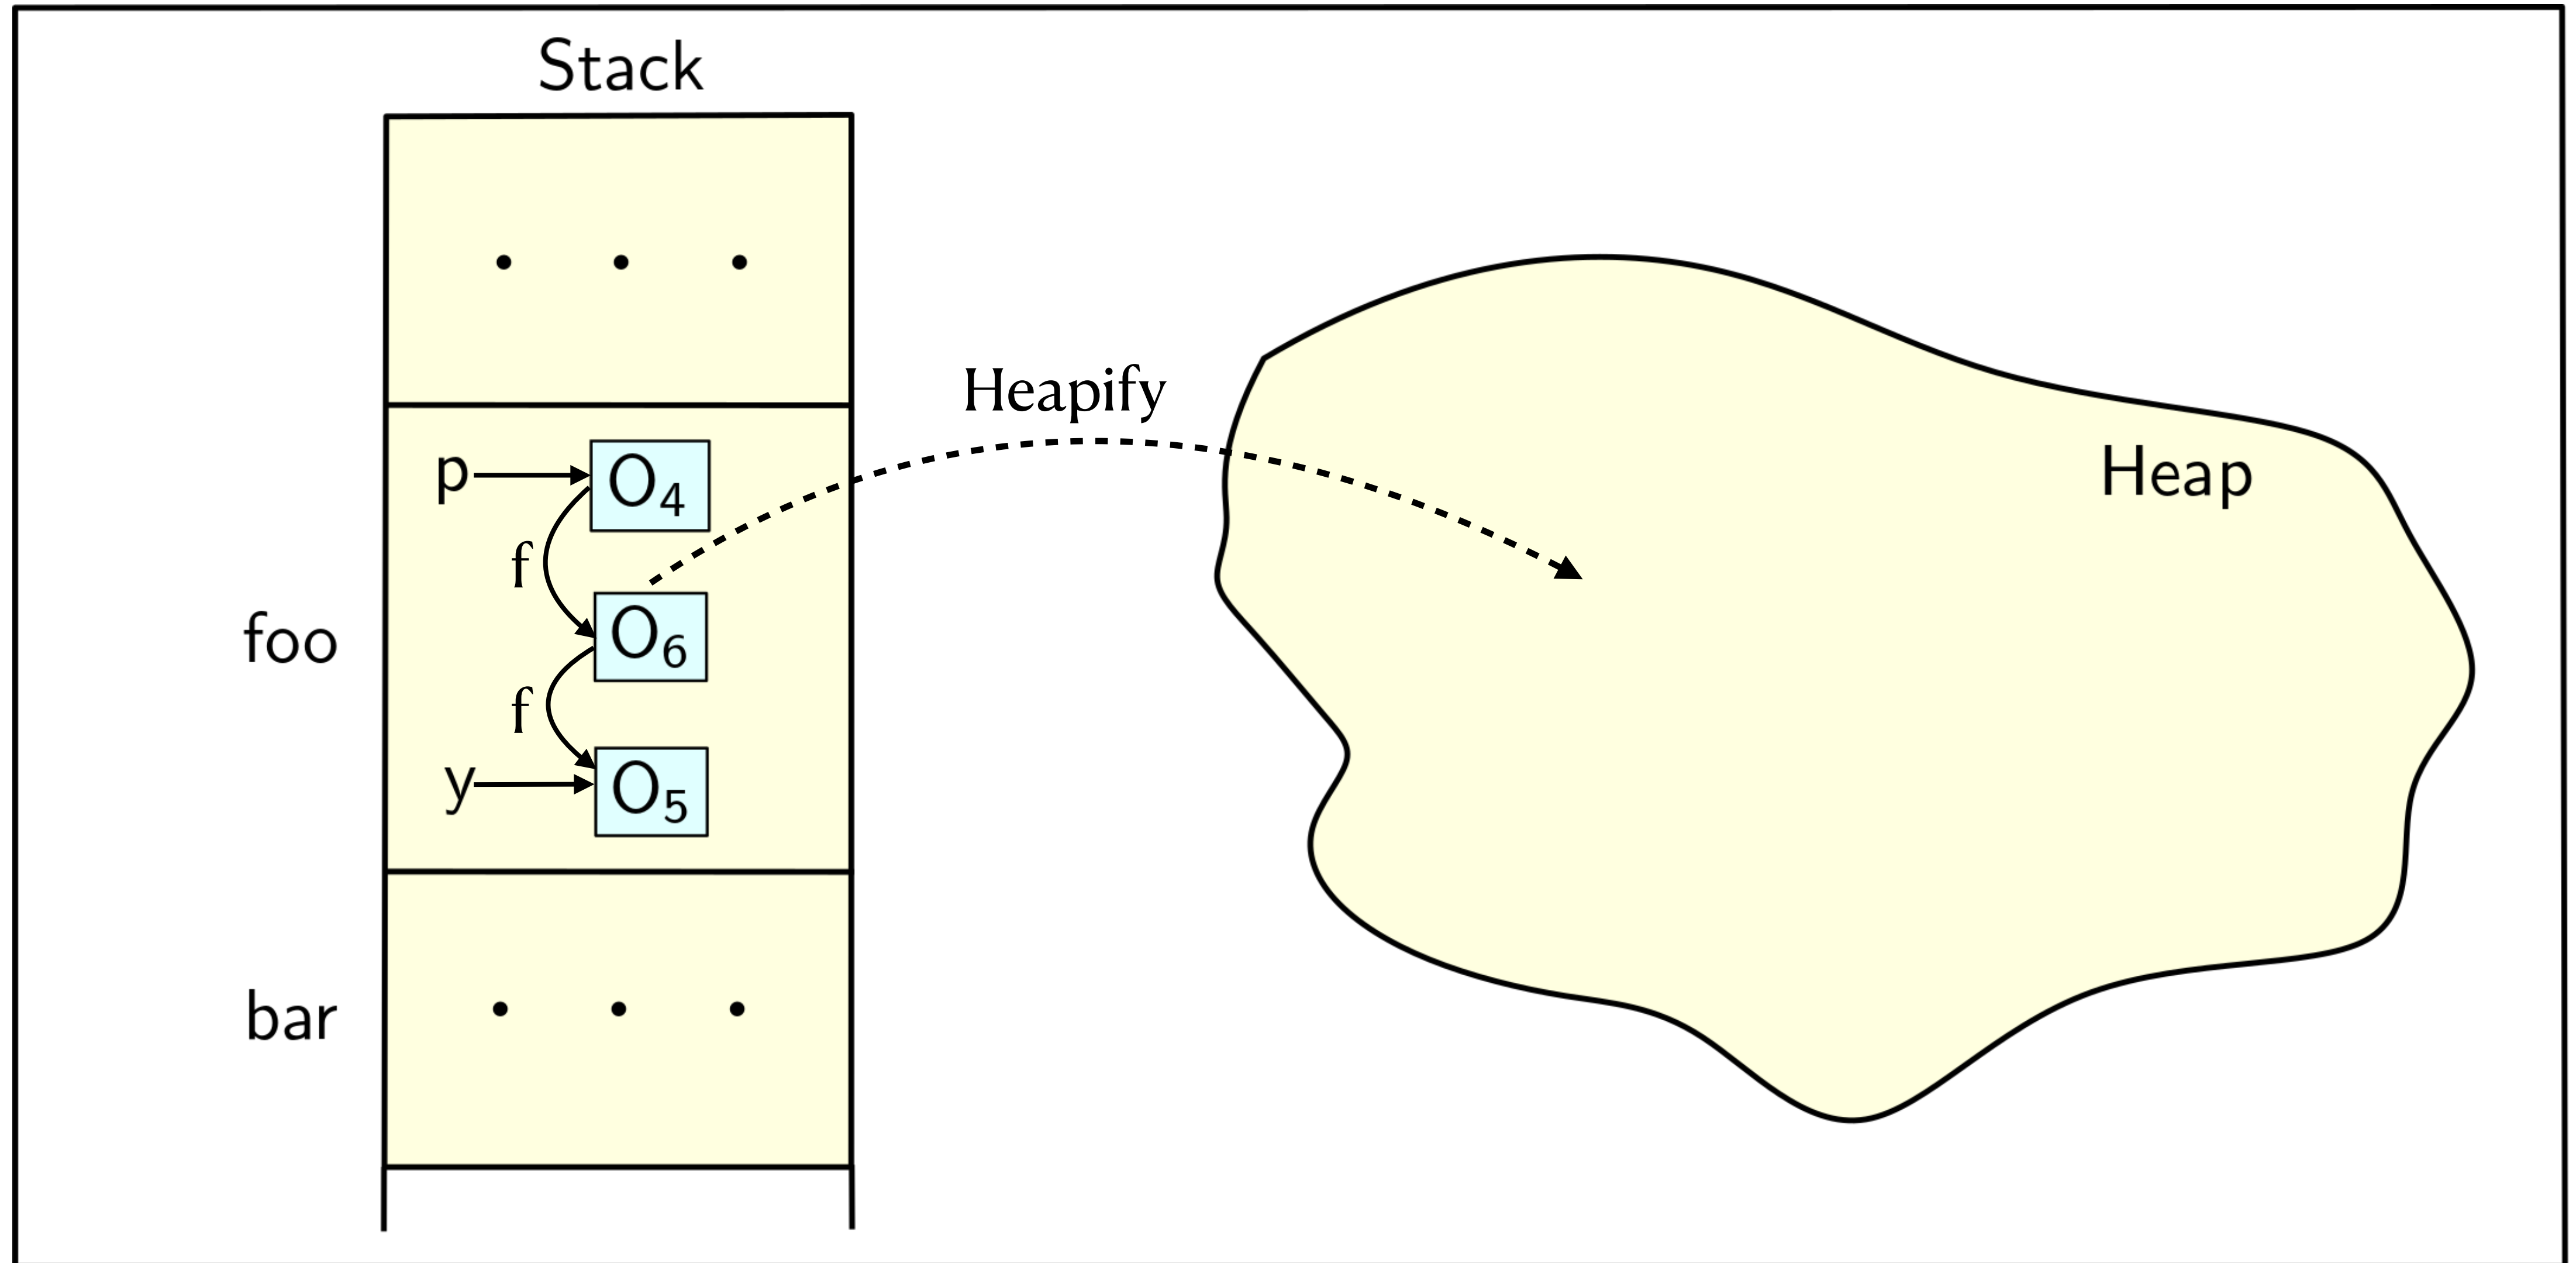


Dynamic Heapification

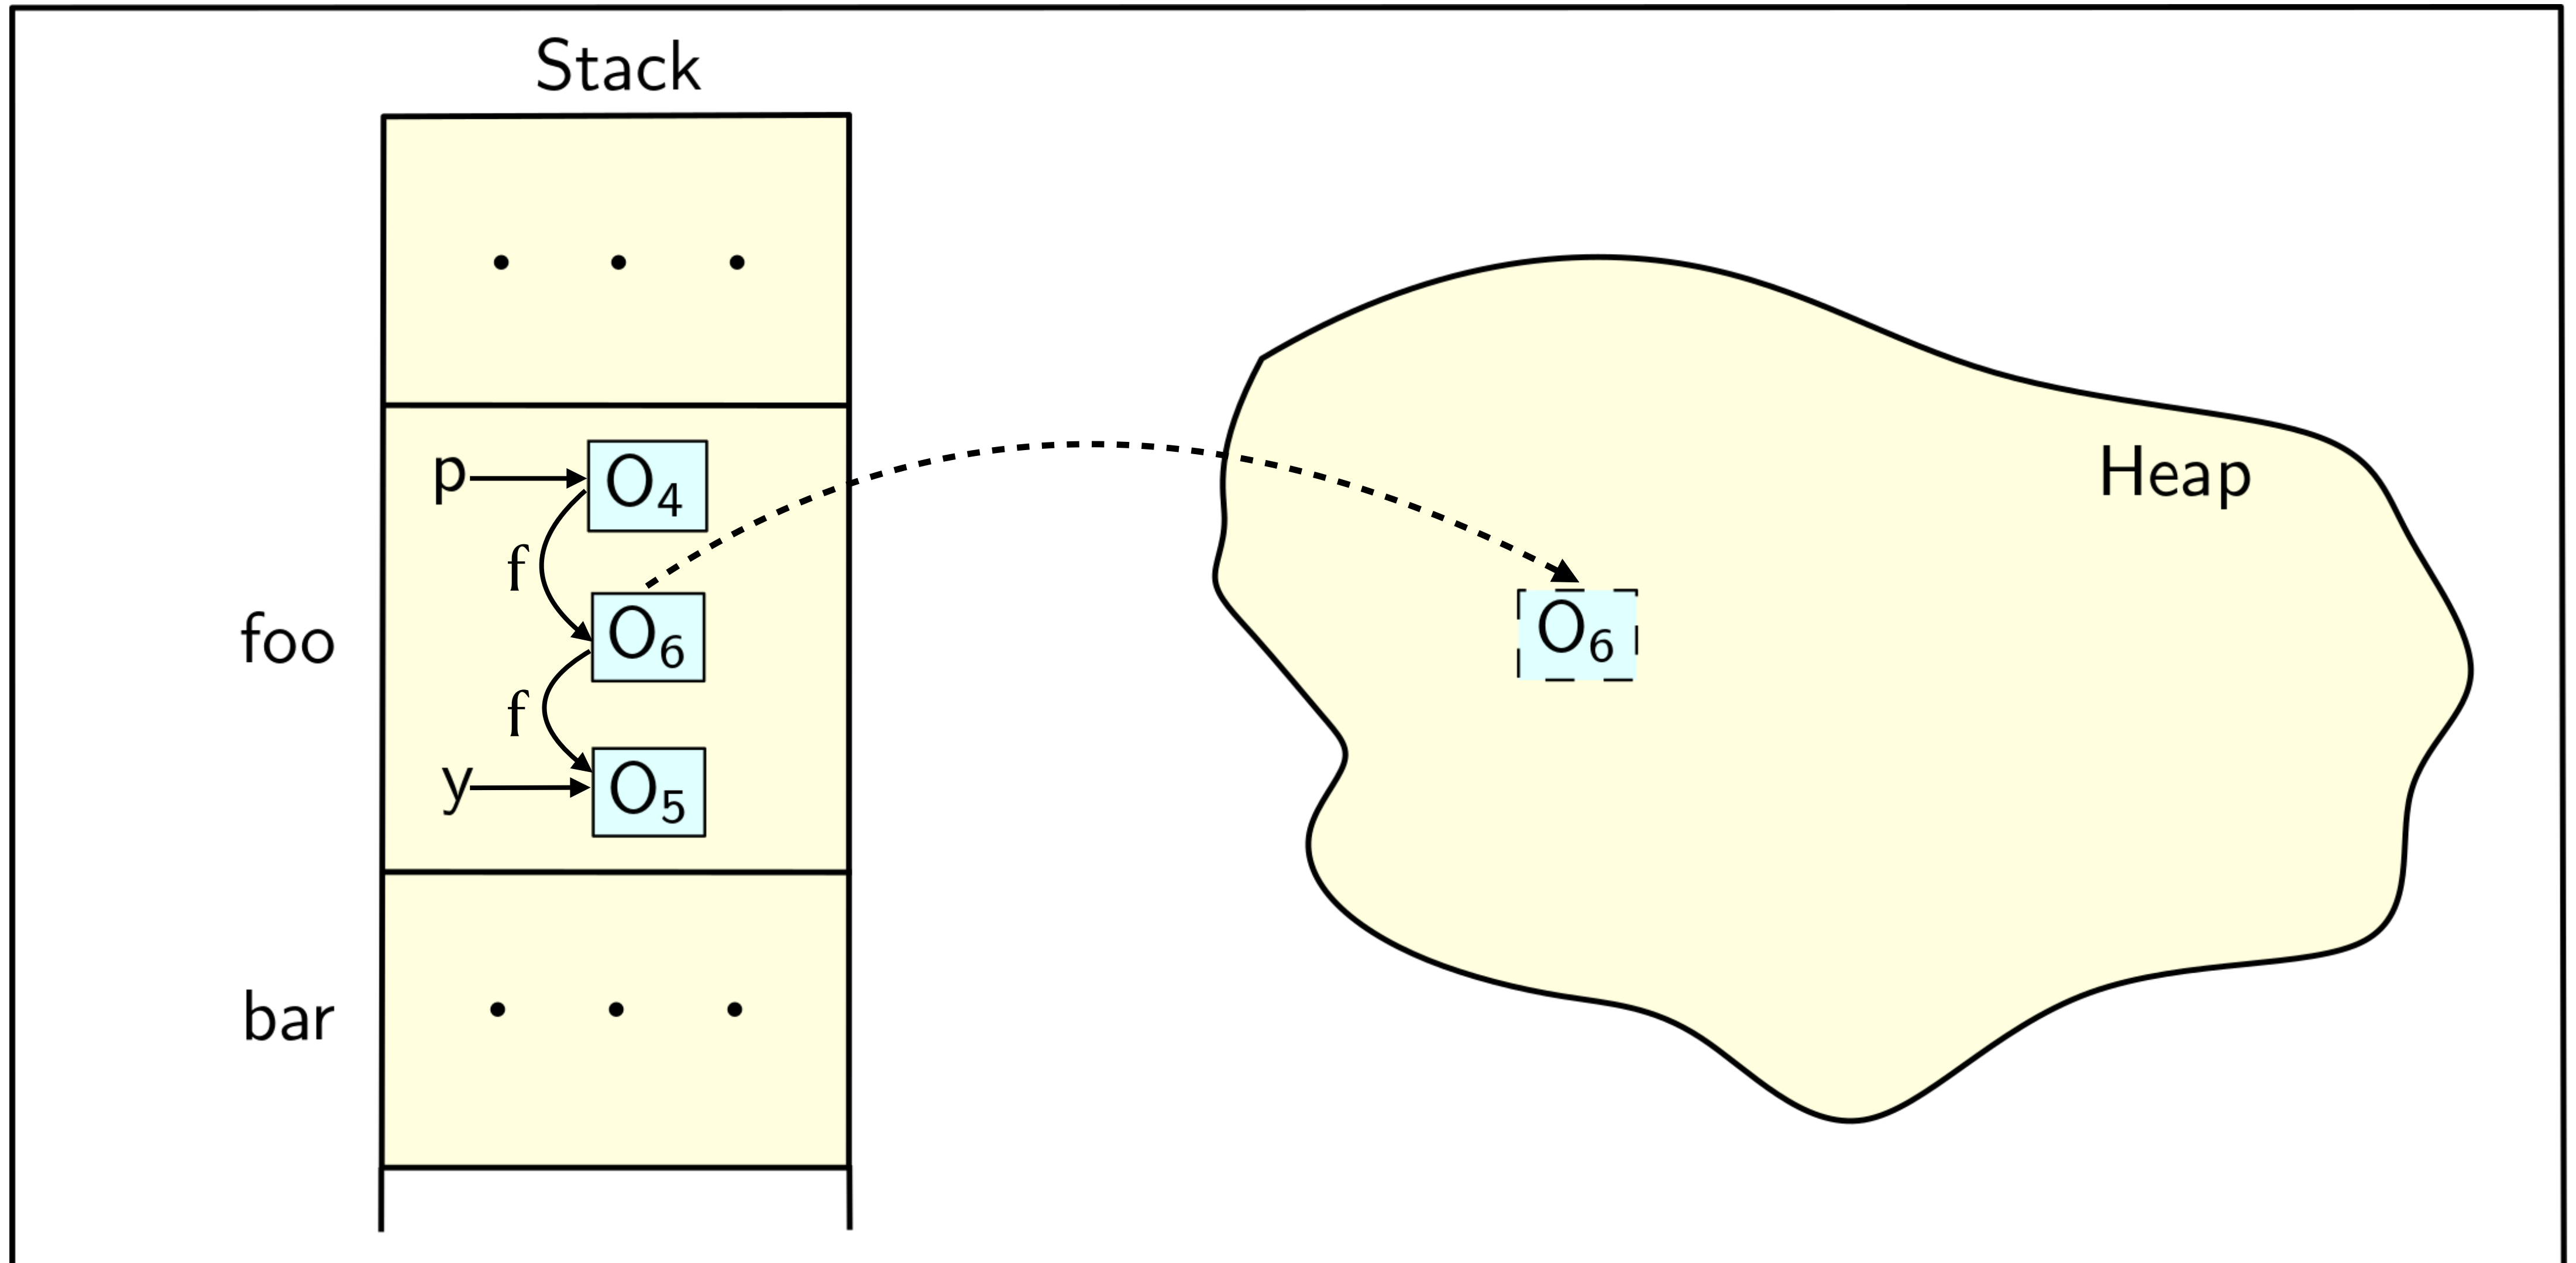
Heapification



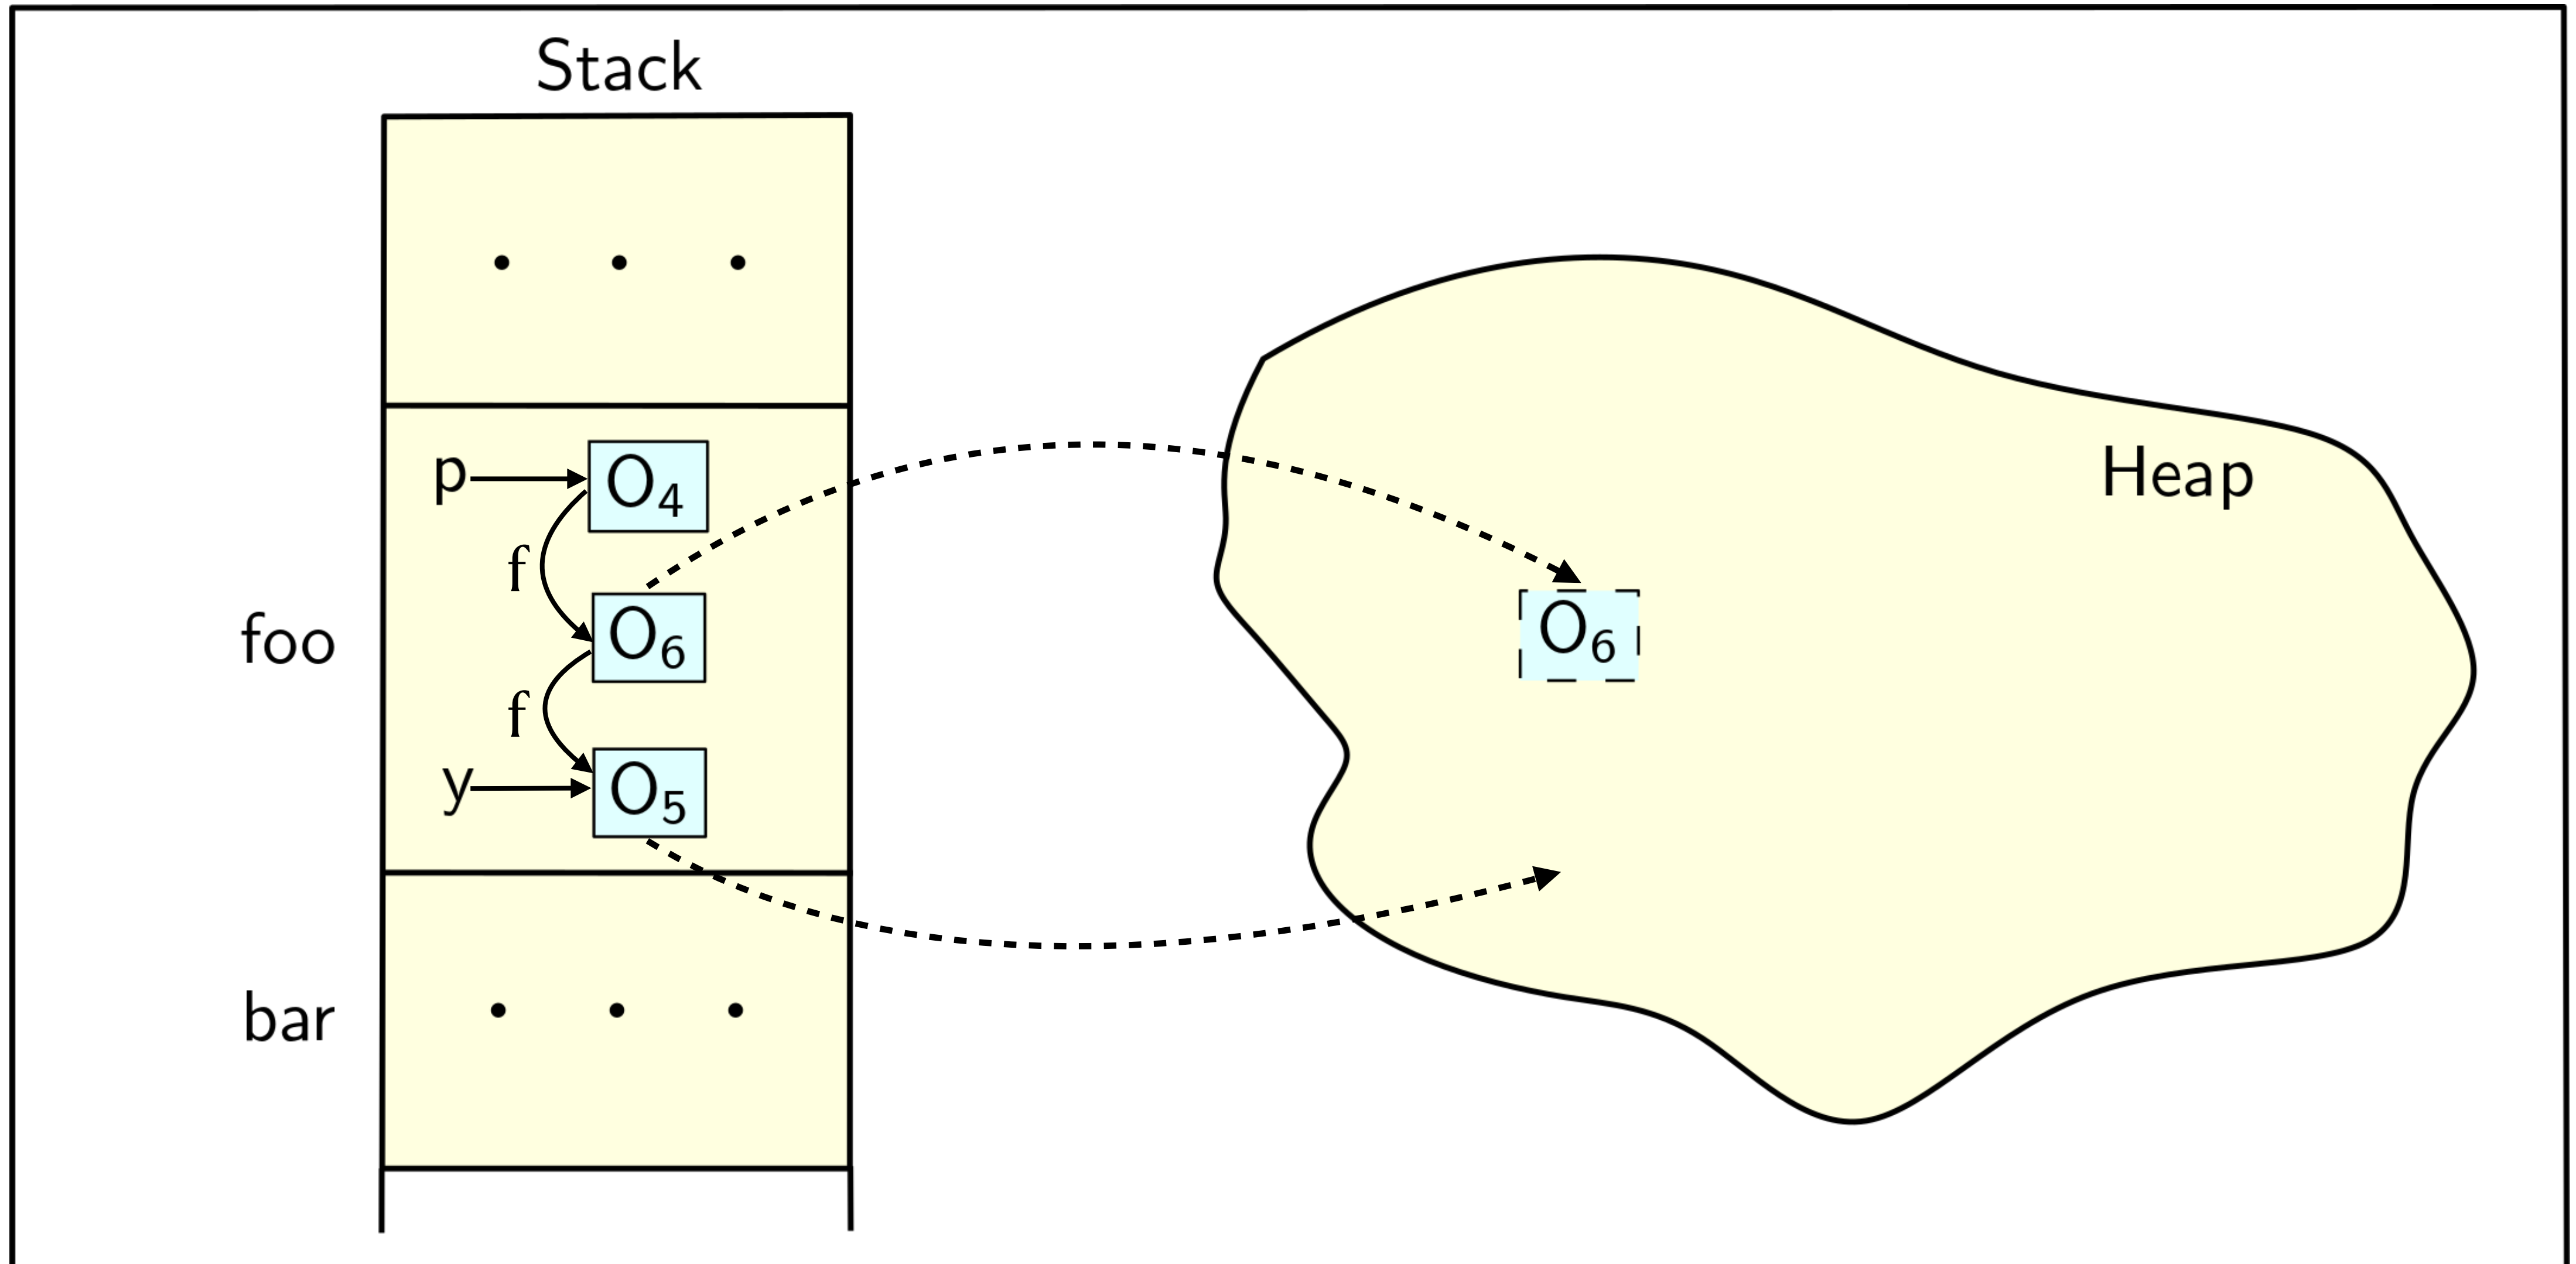
Heapification



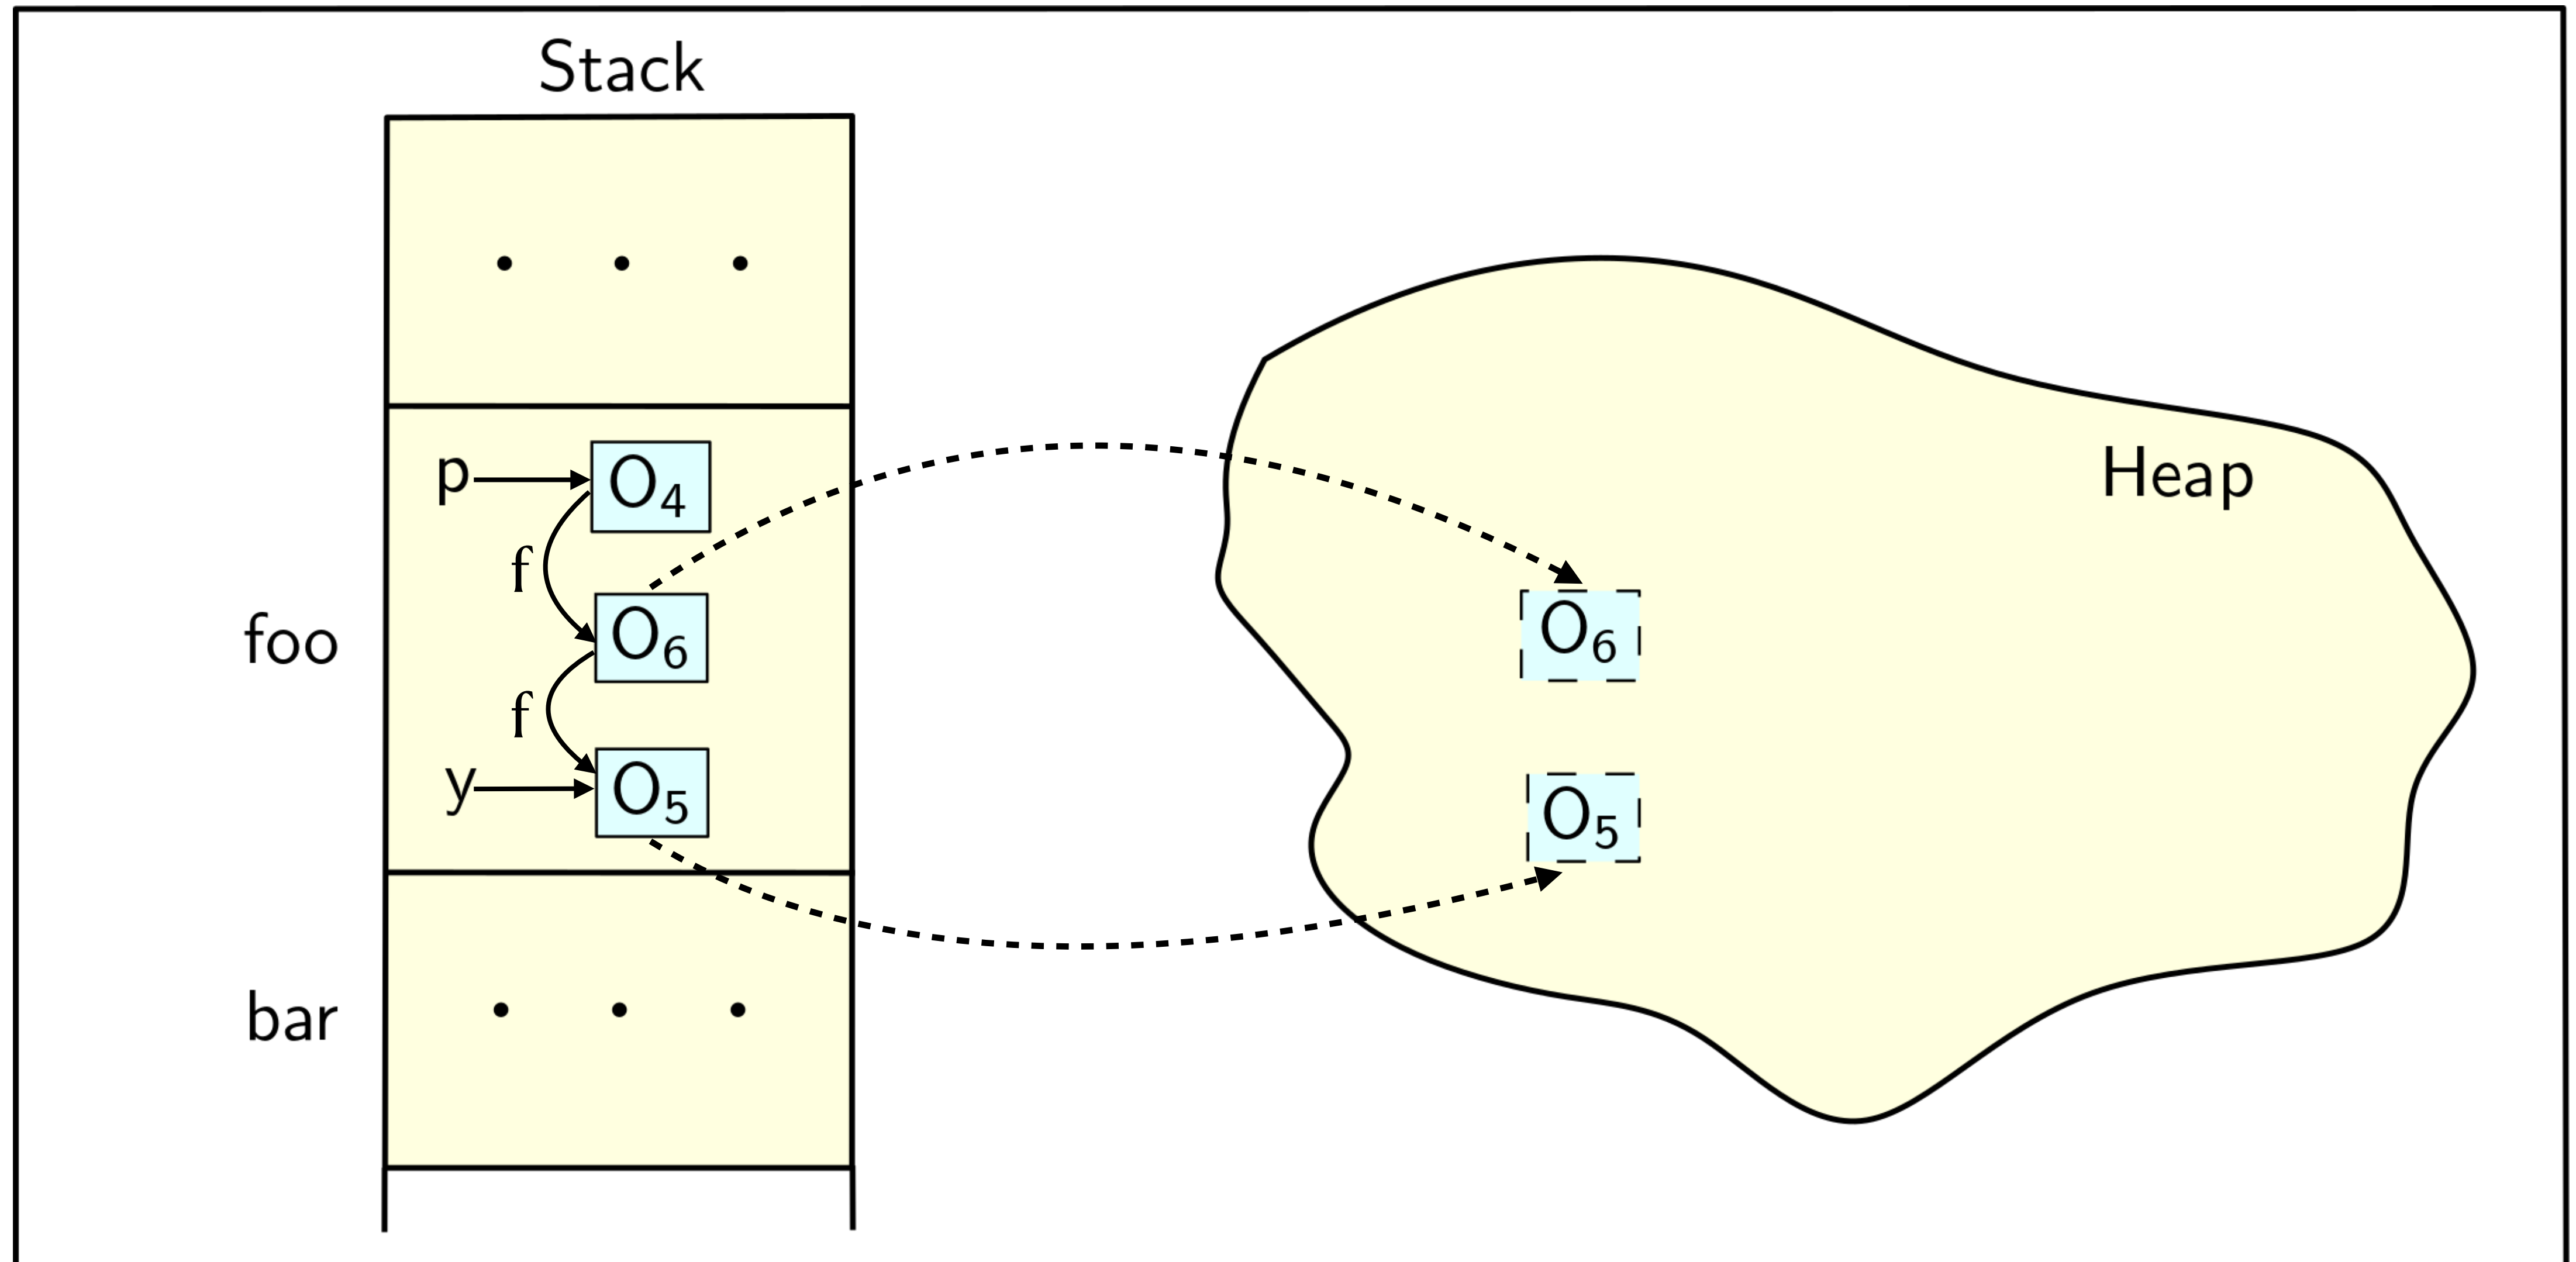
Heapification



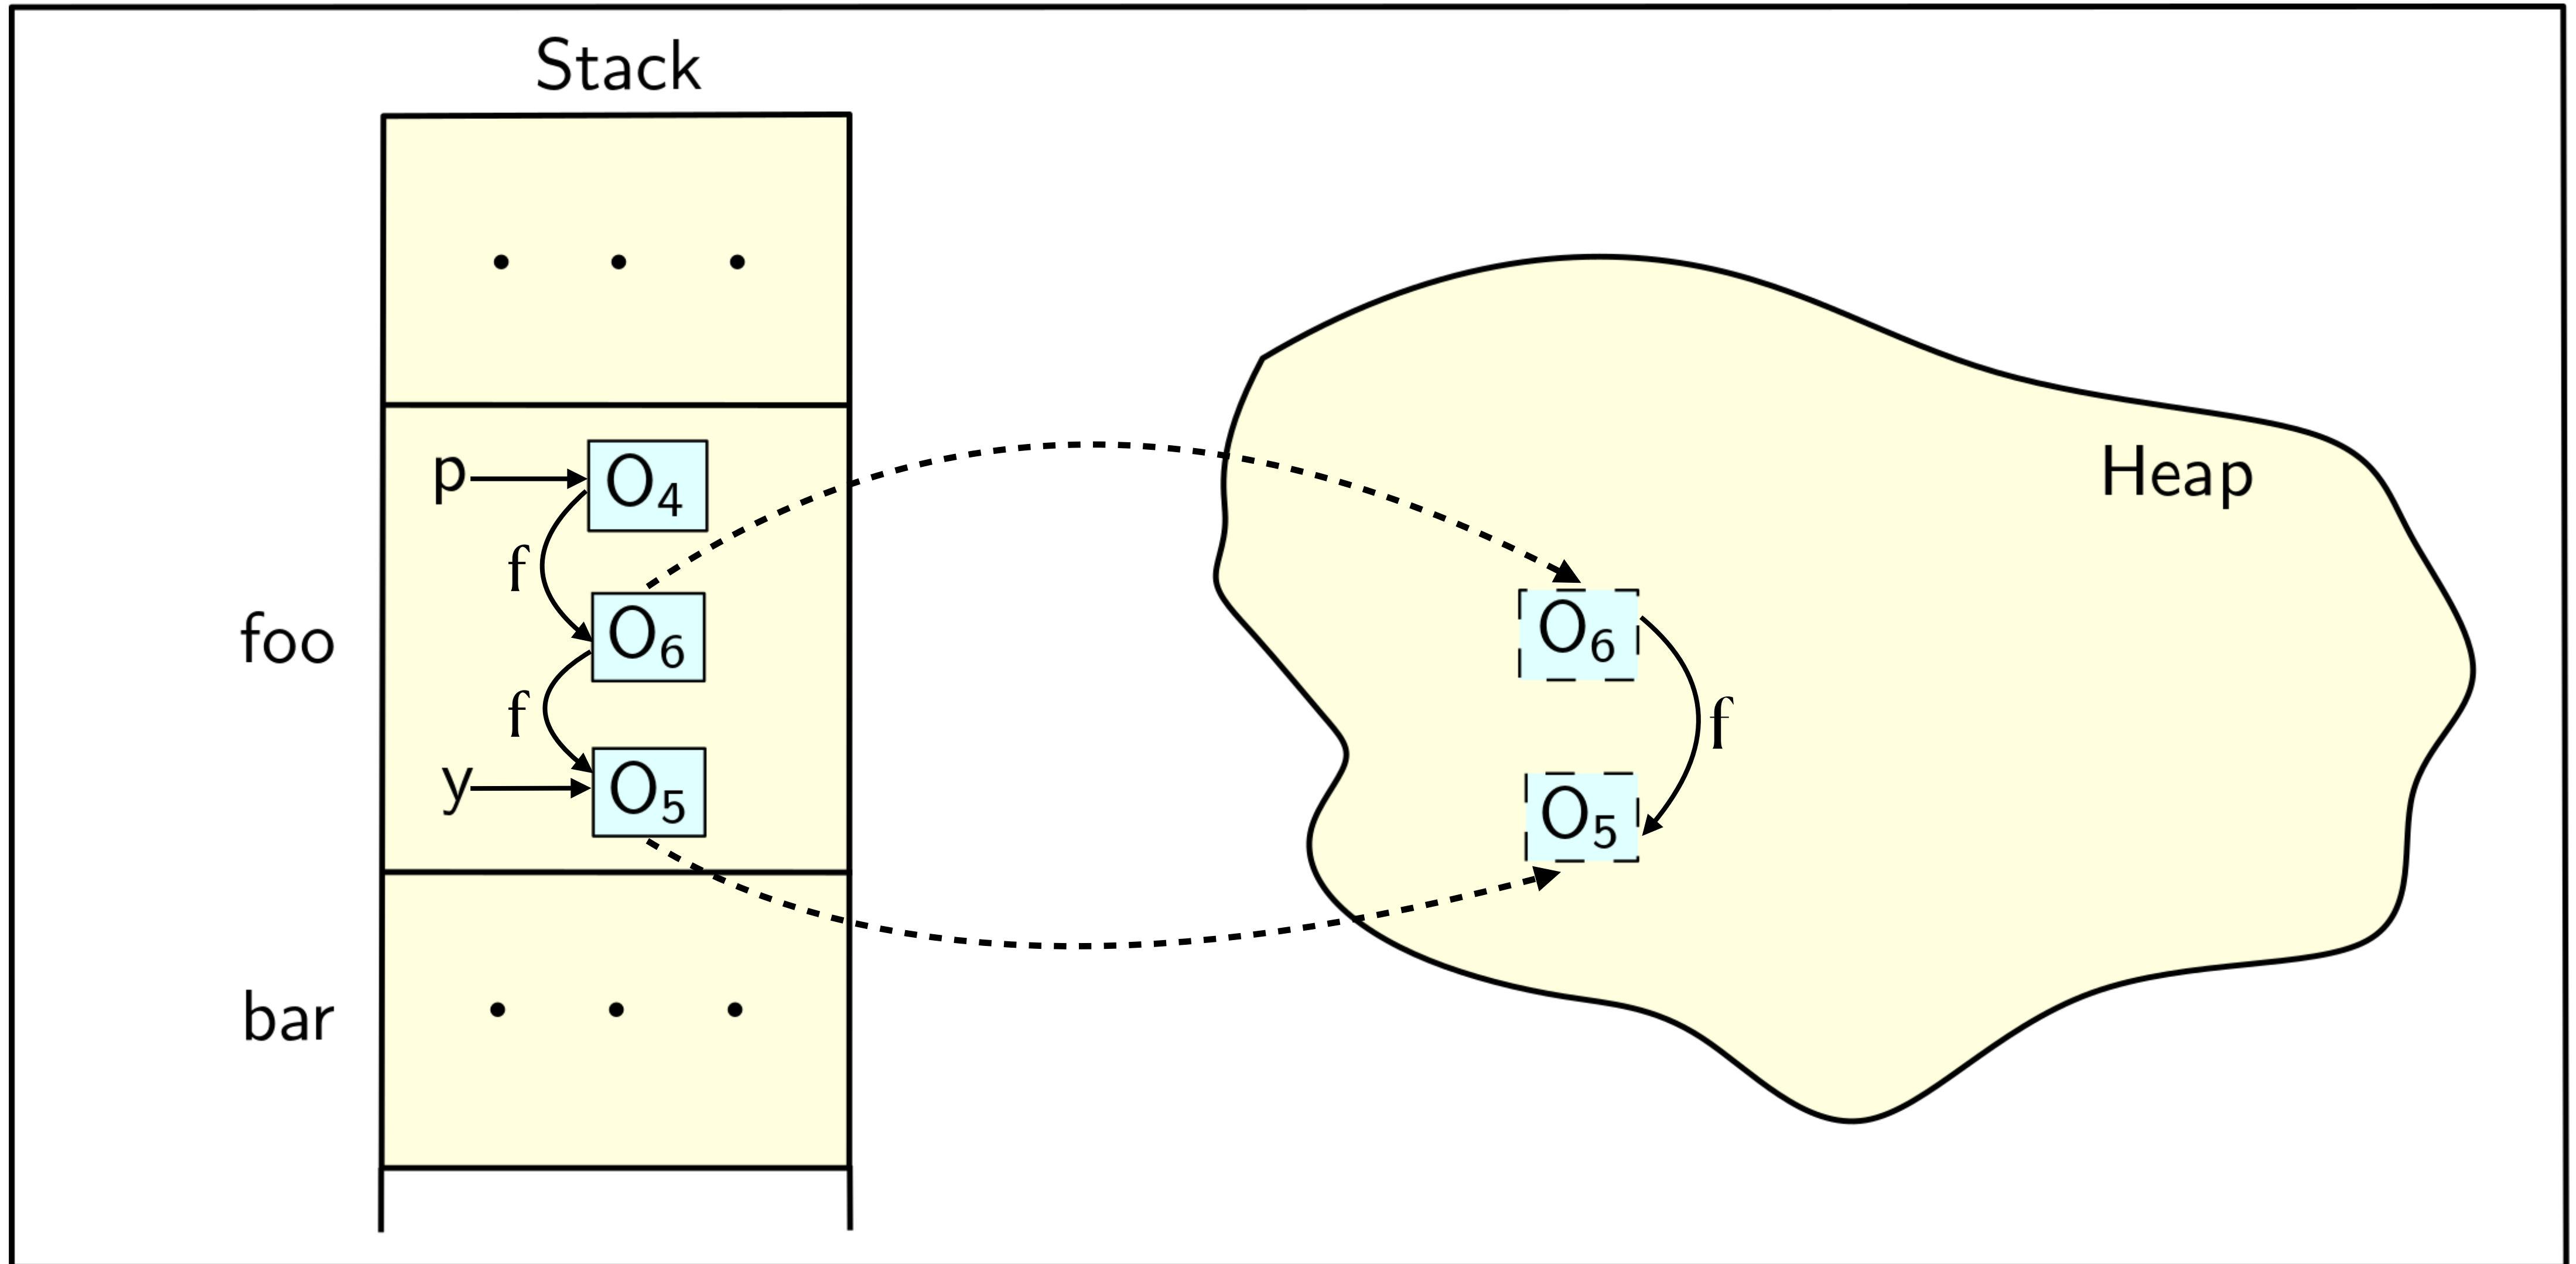
Heapification



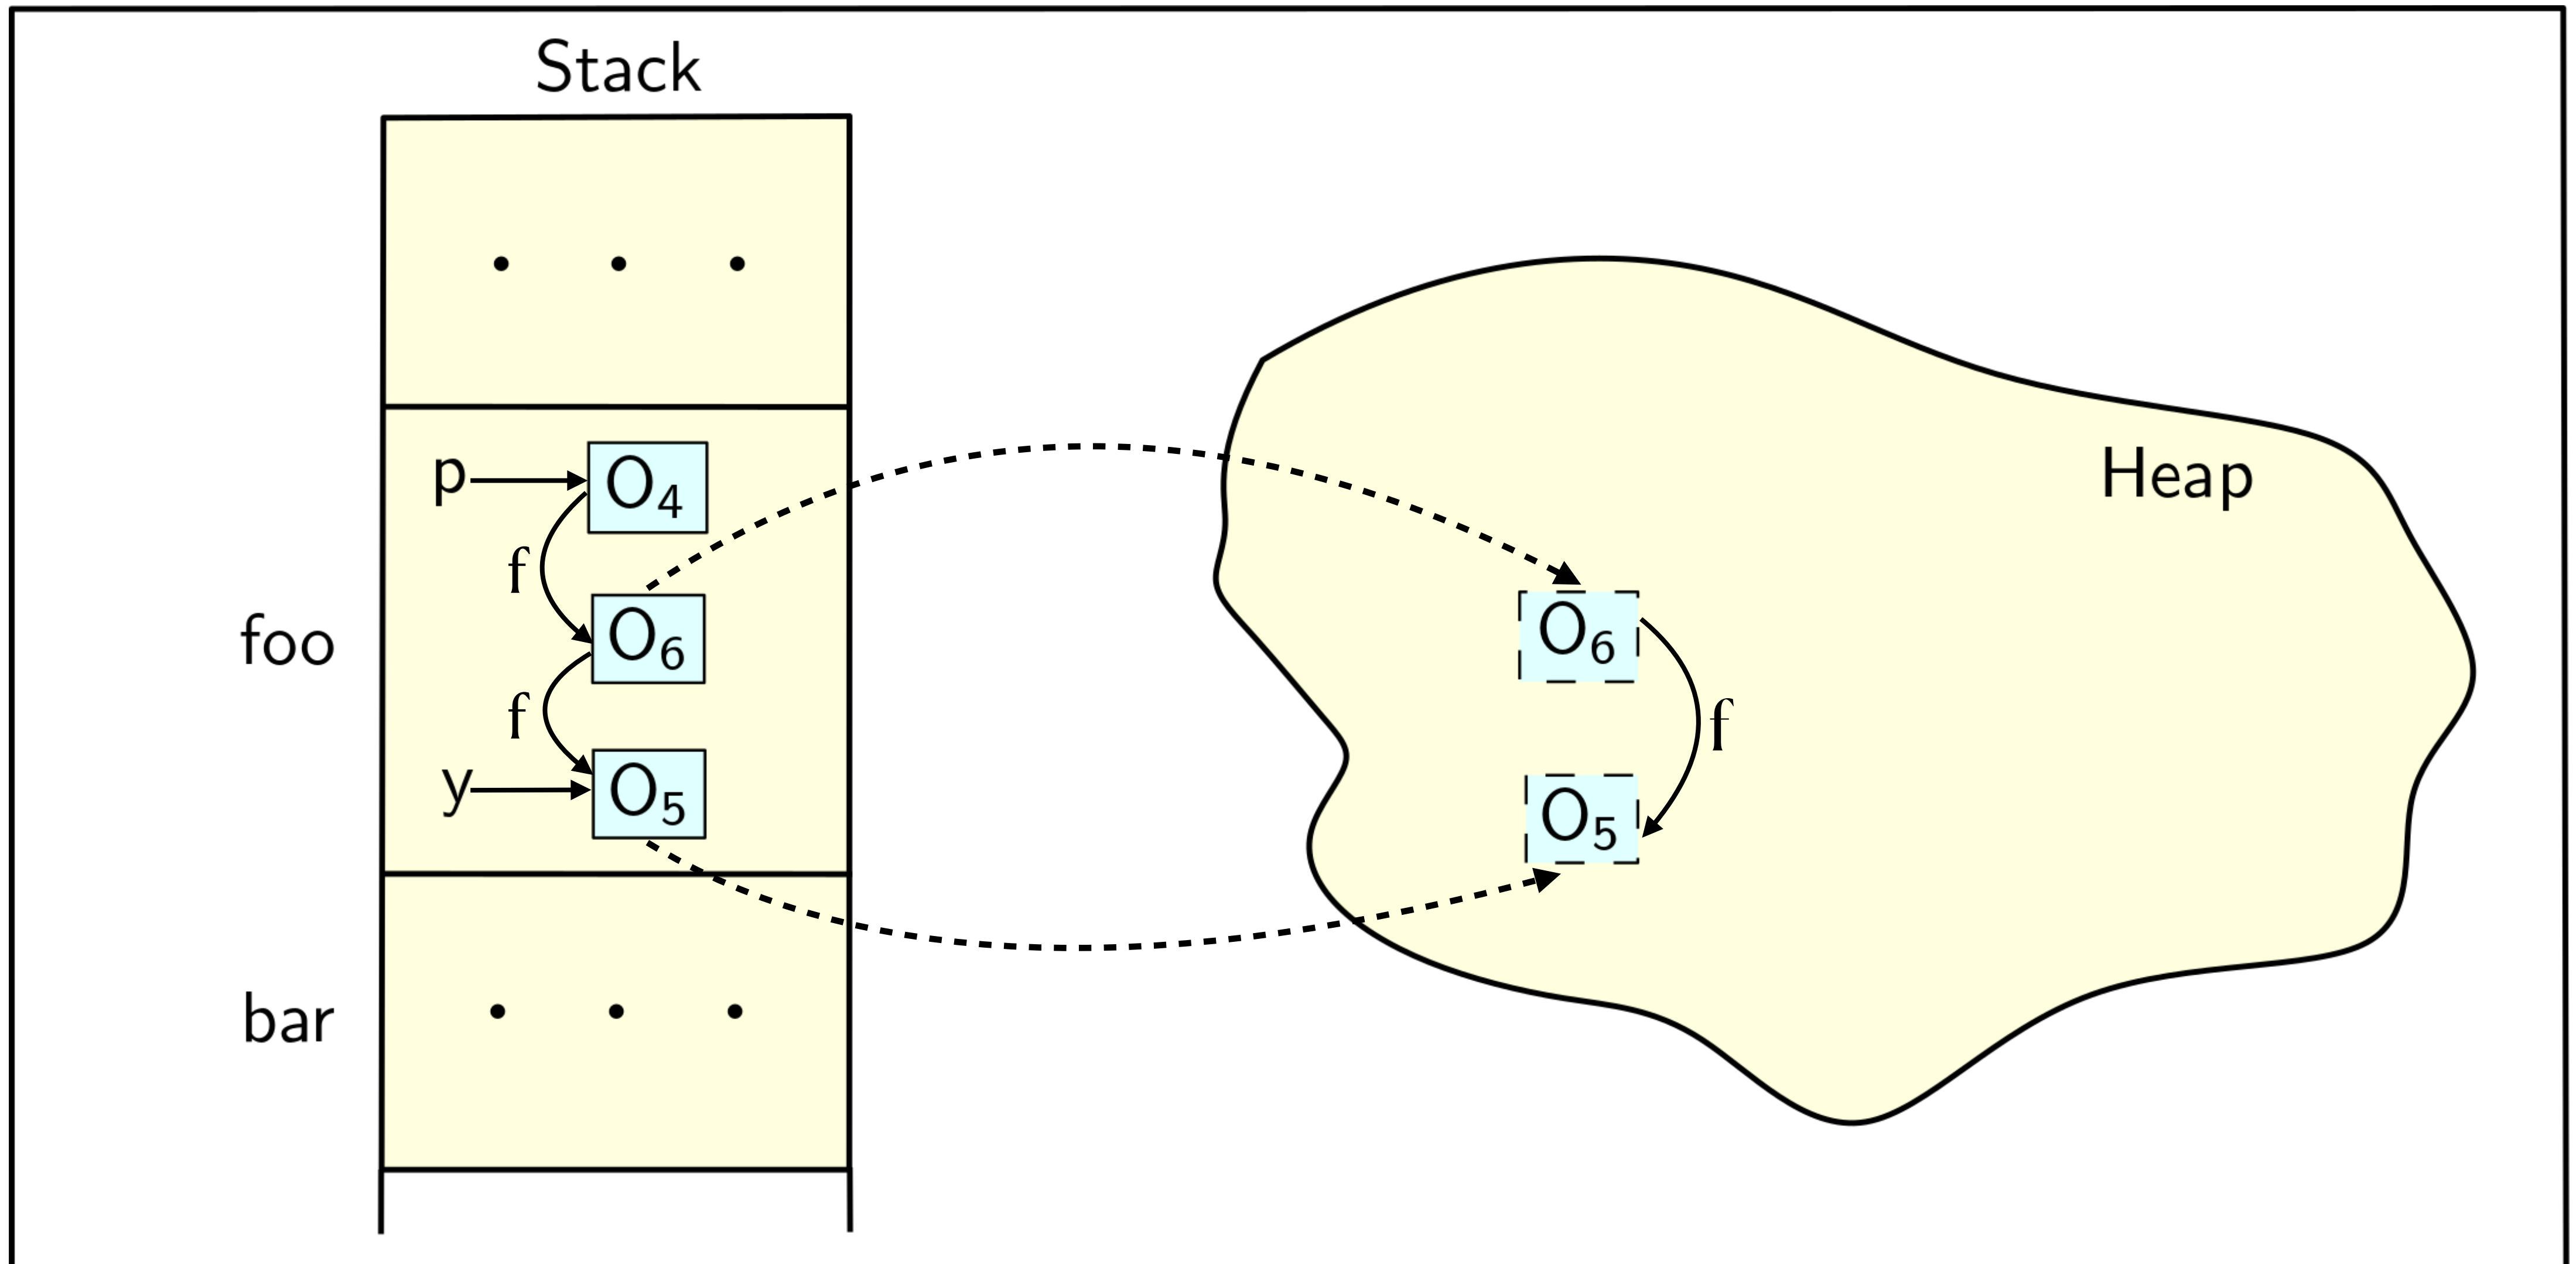
Heapification



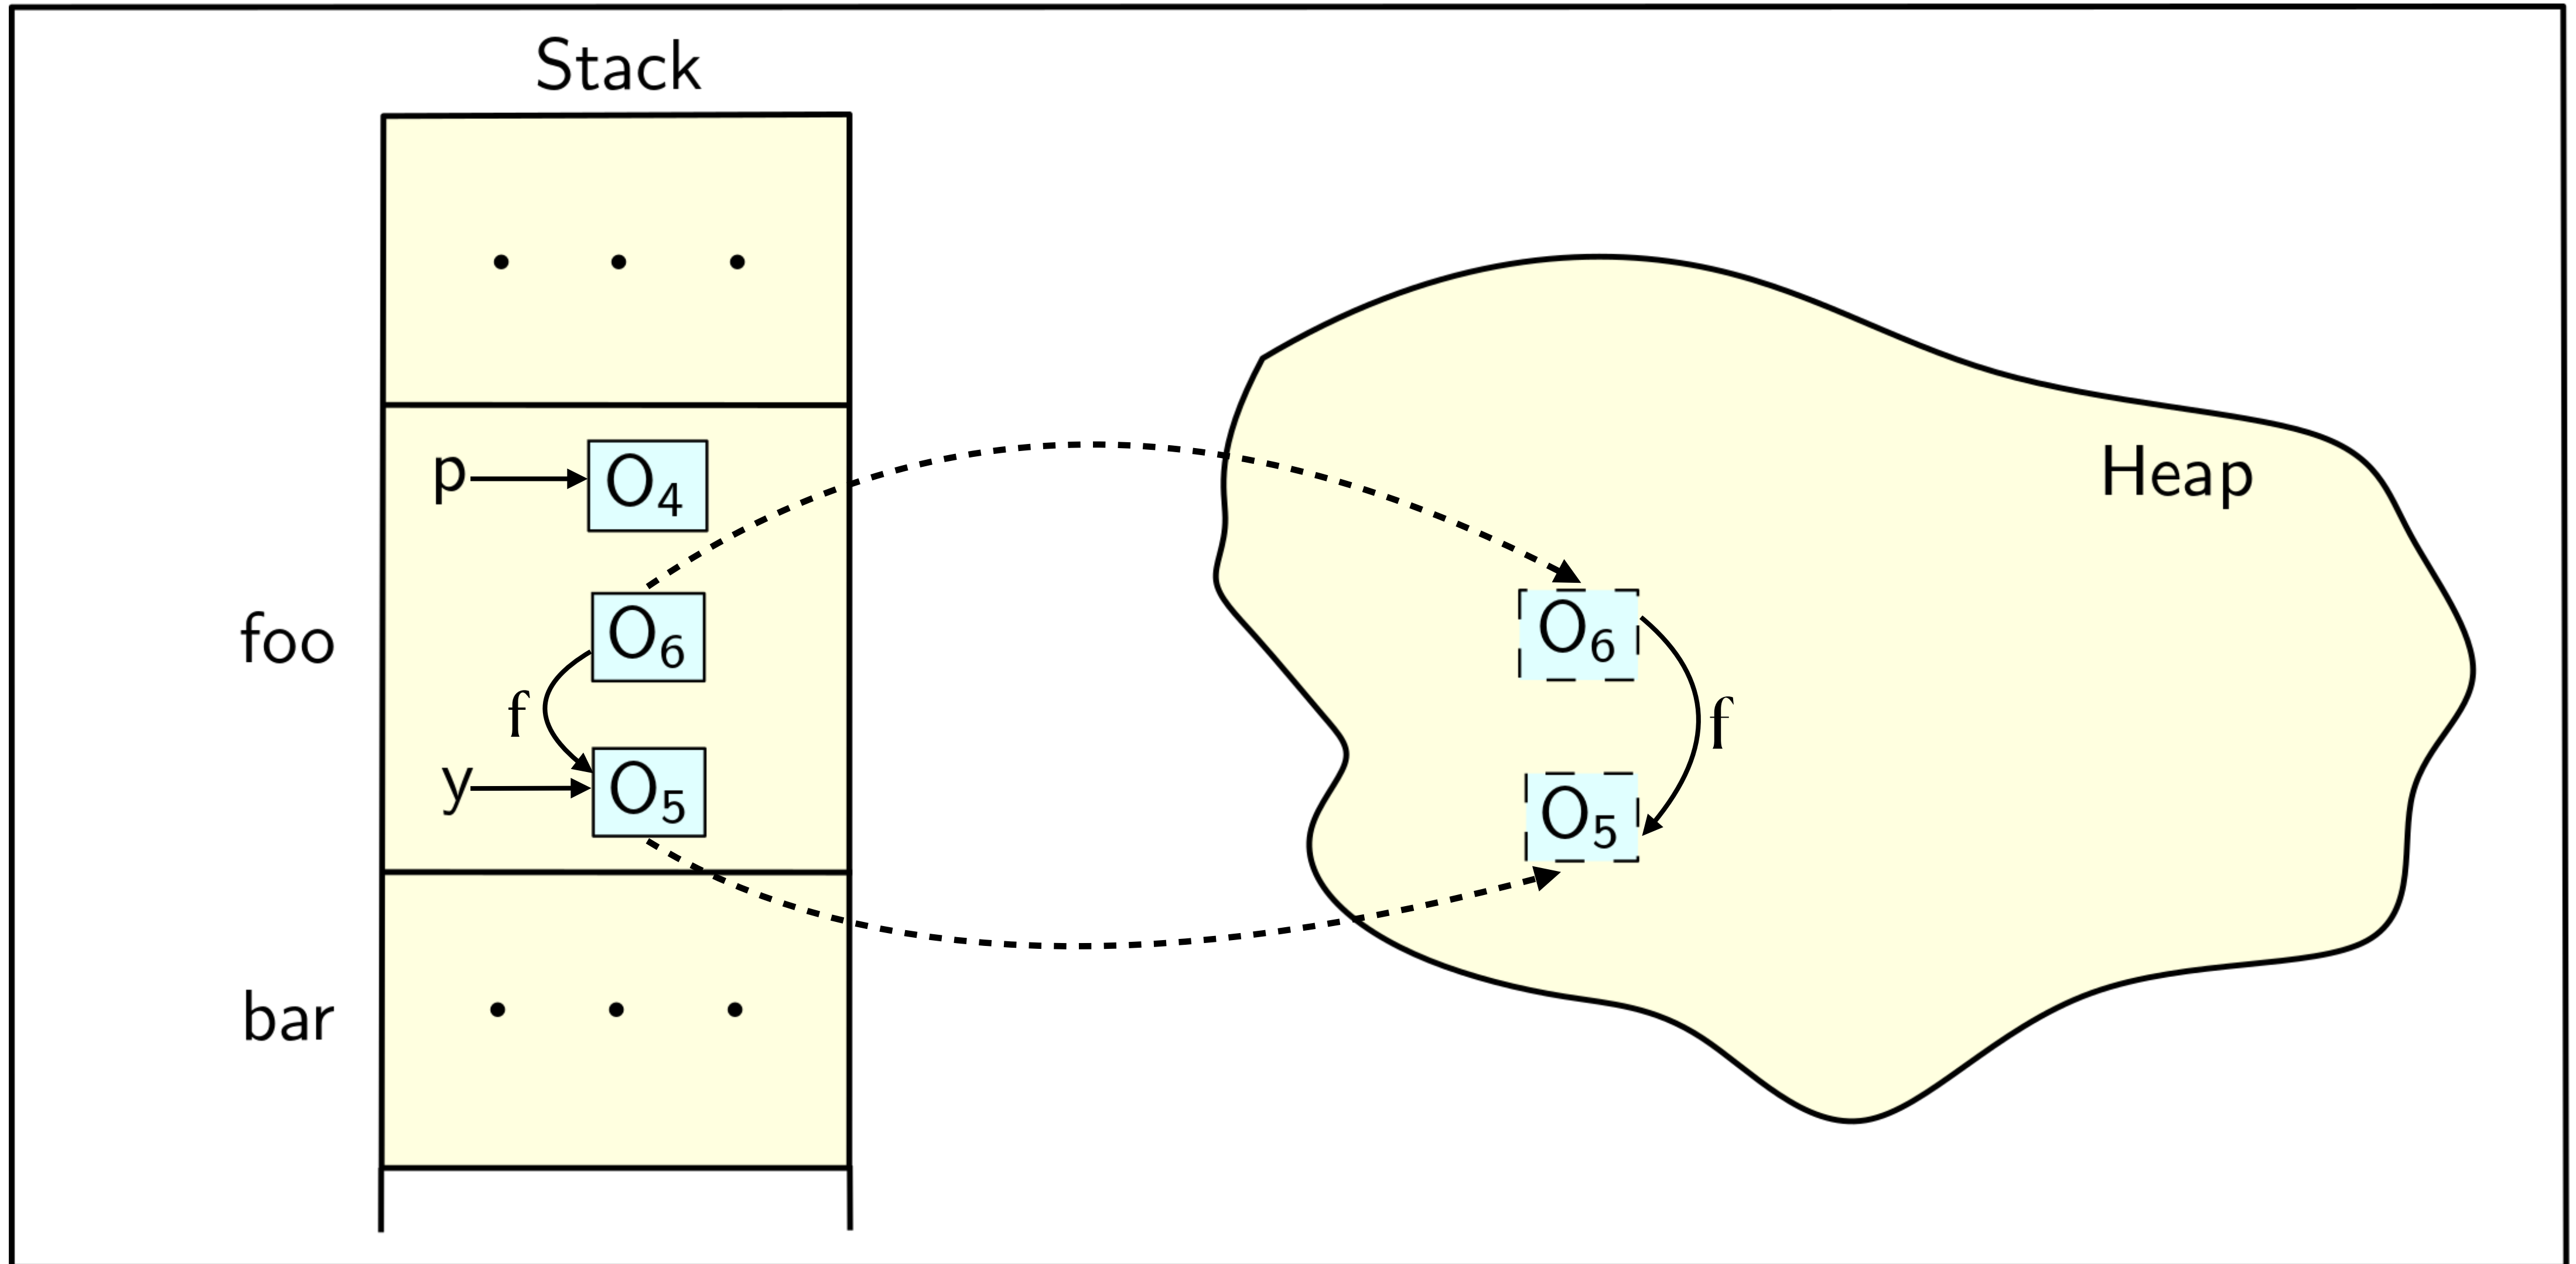
Heapification



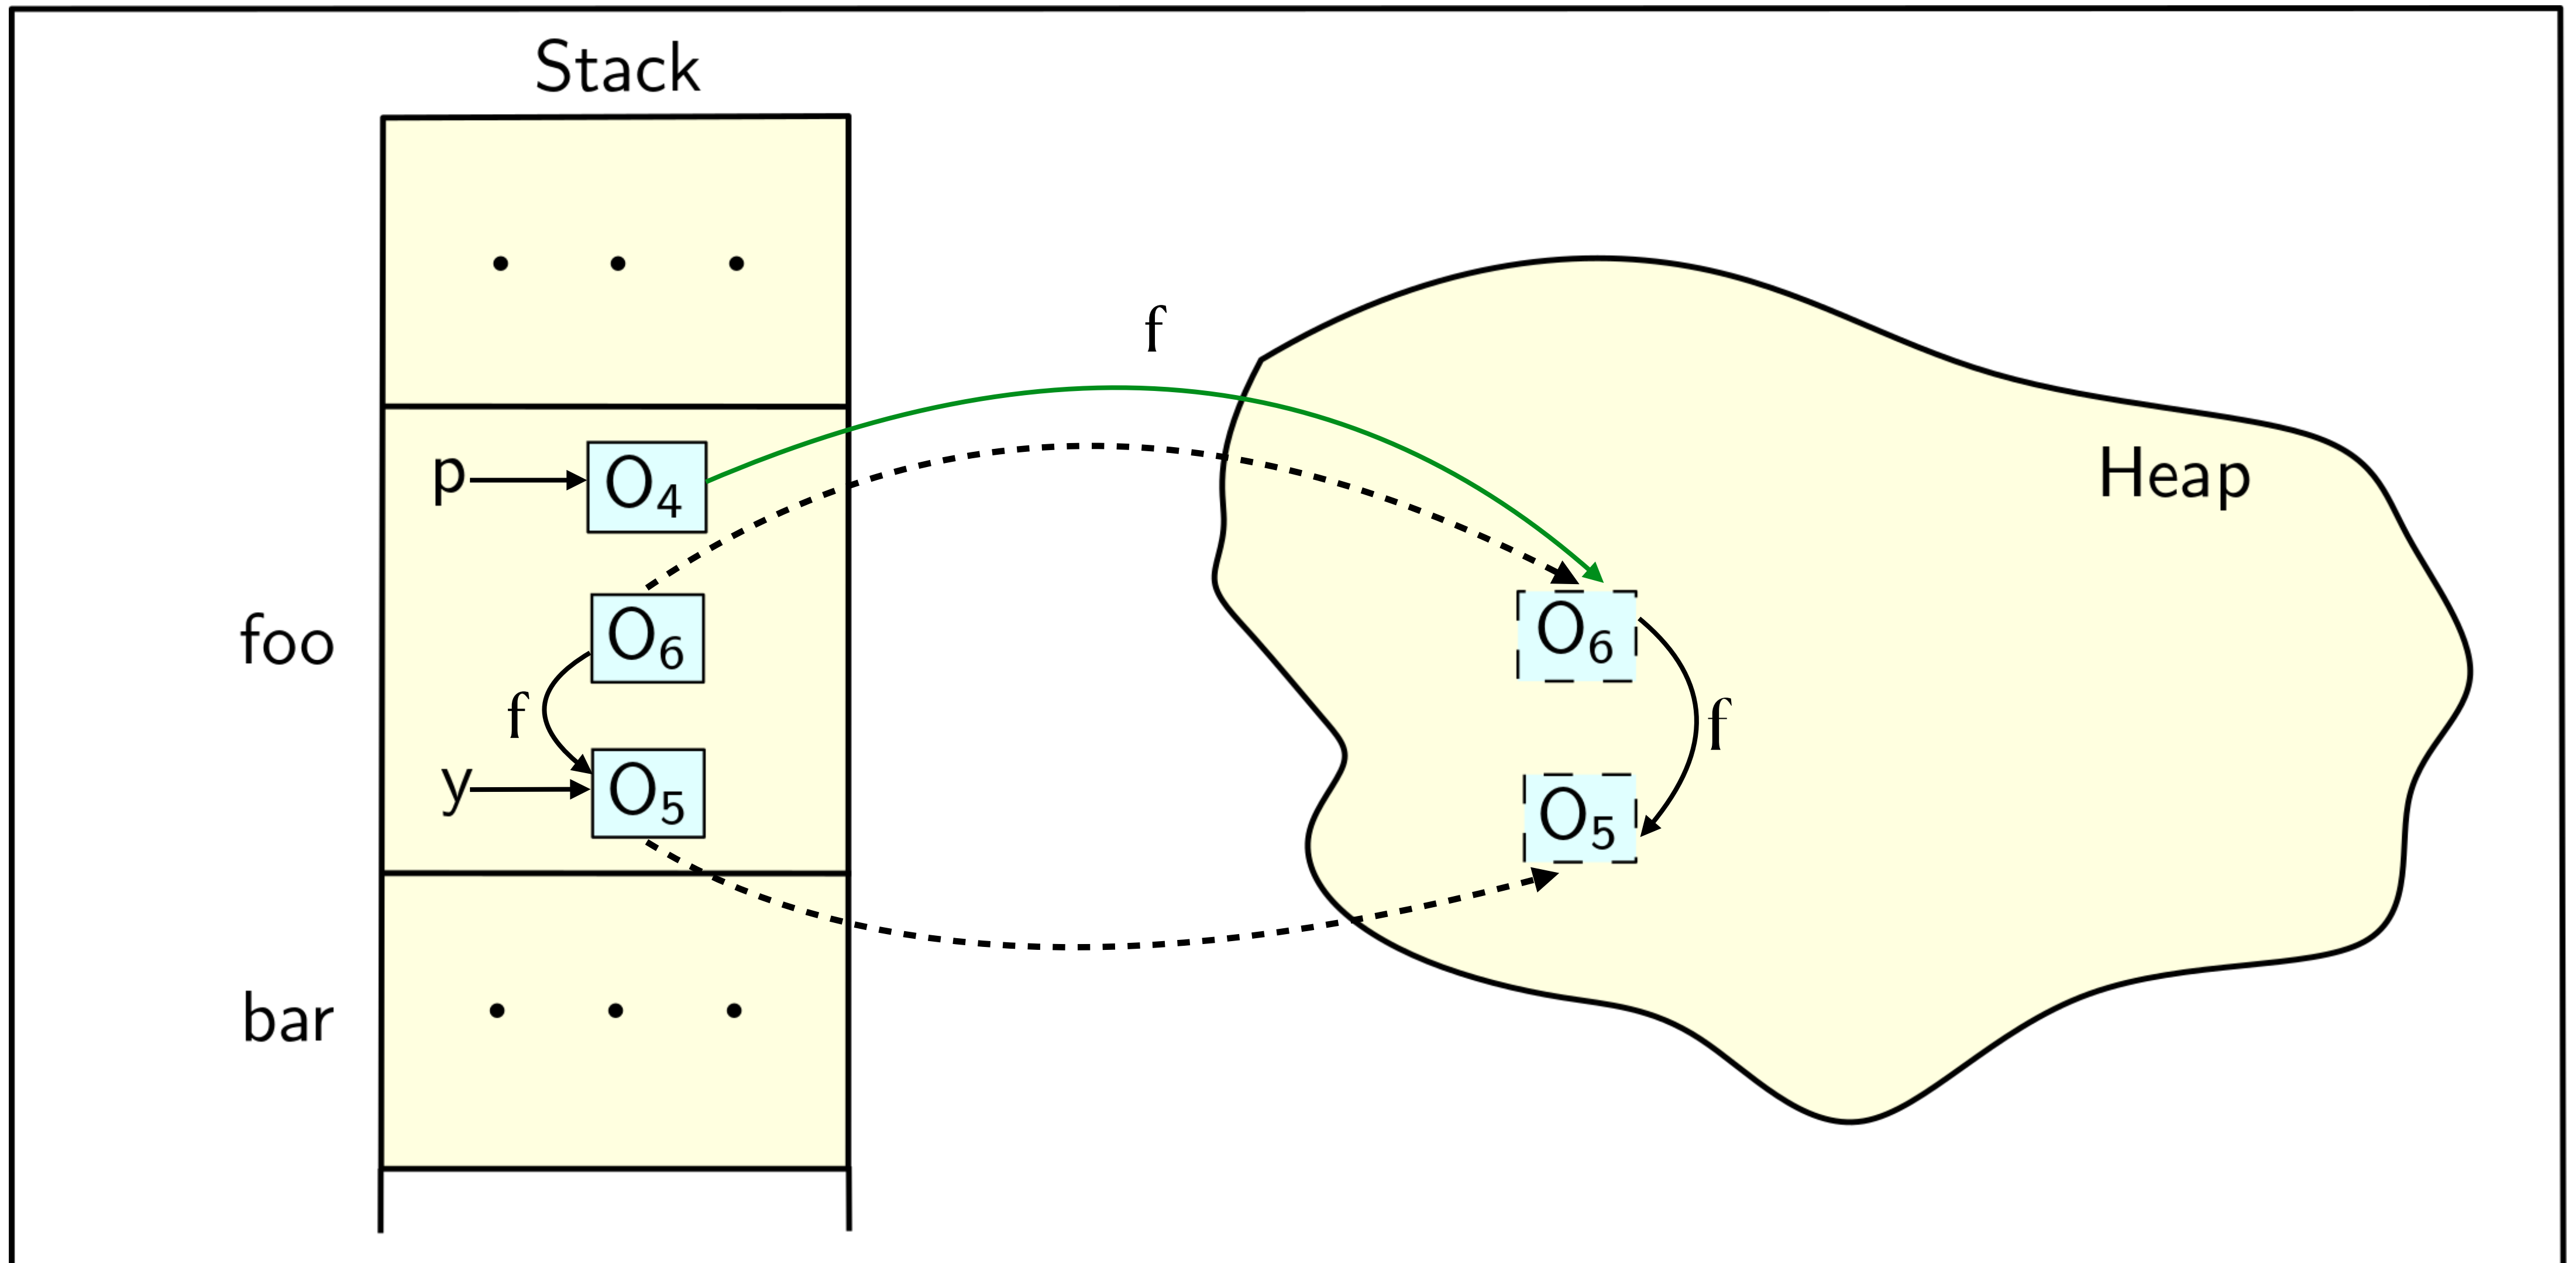
Heapification



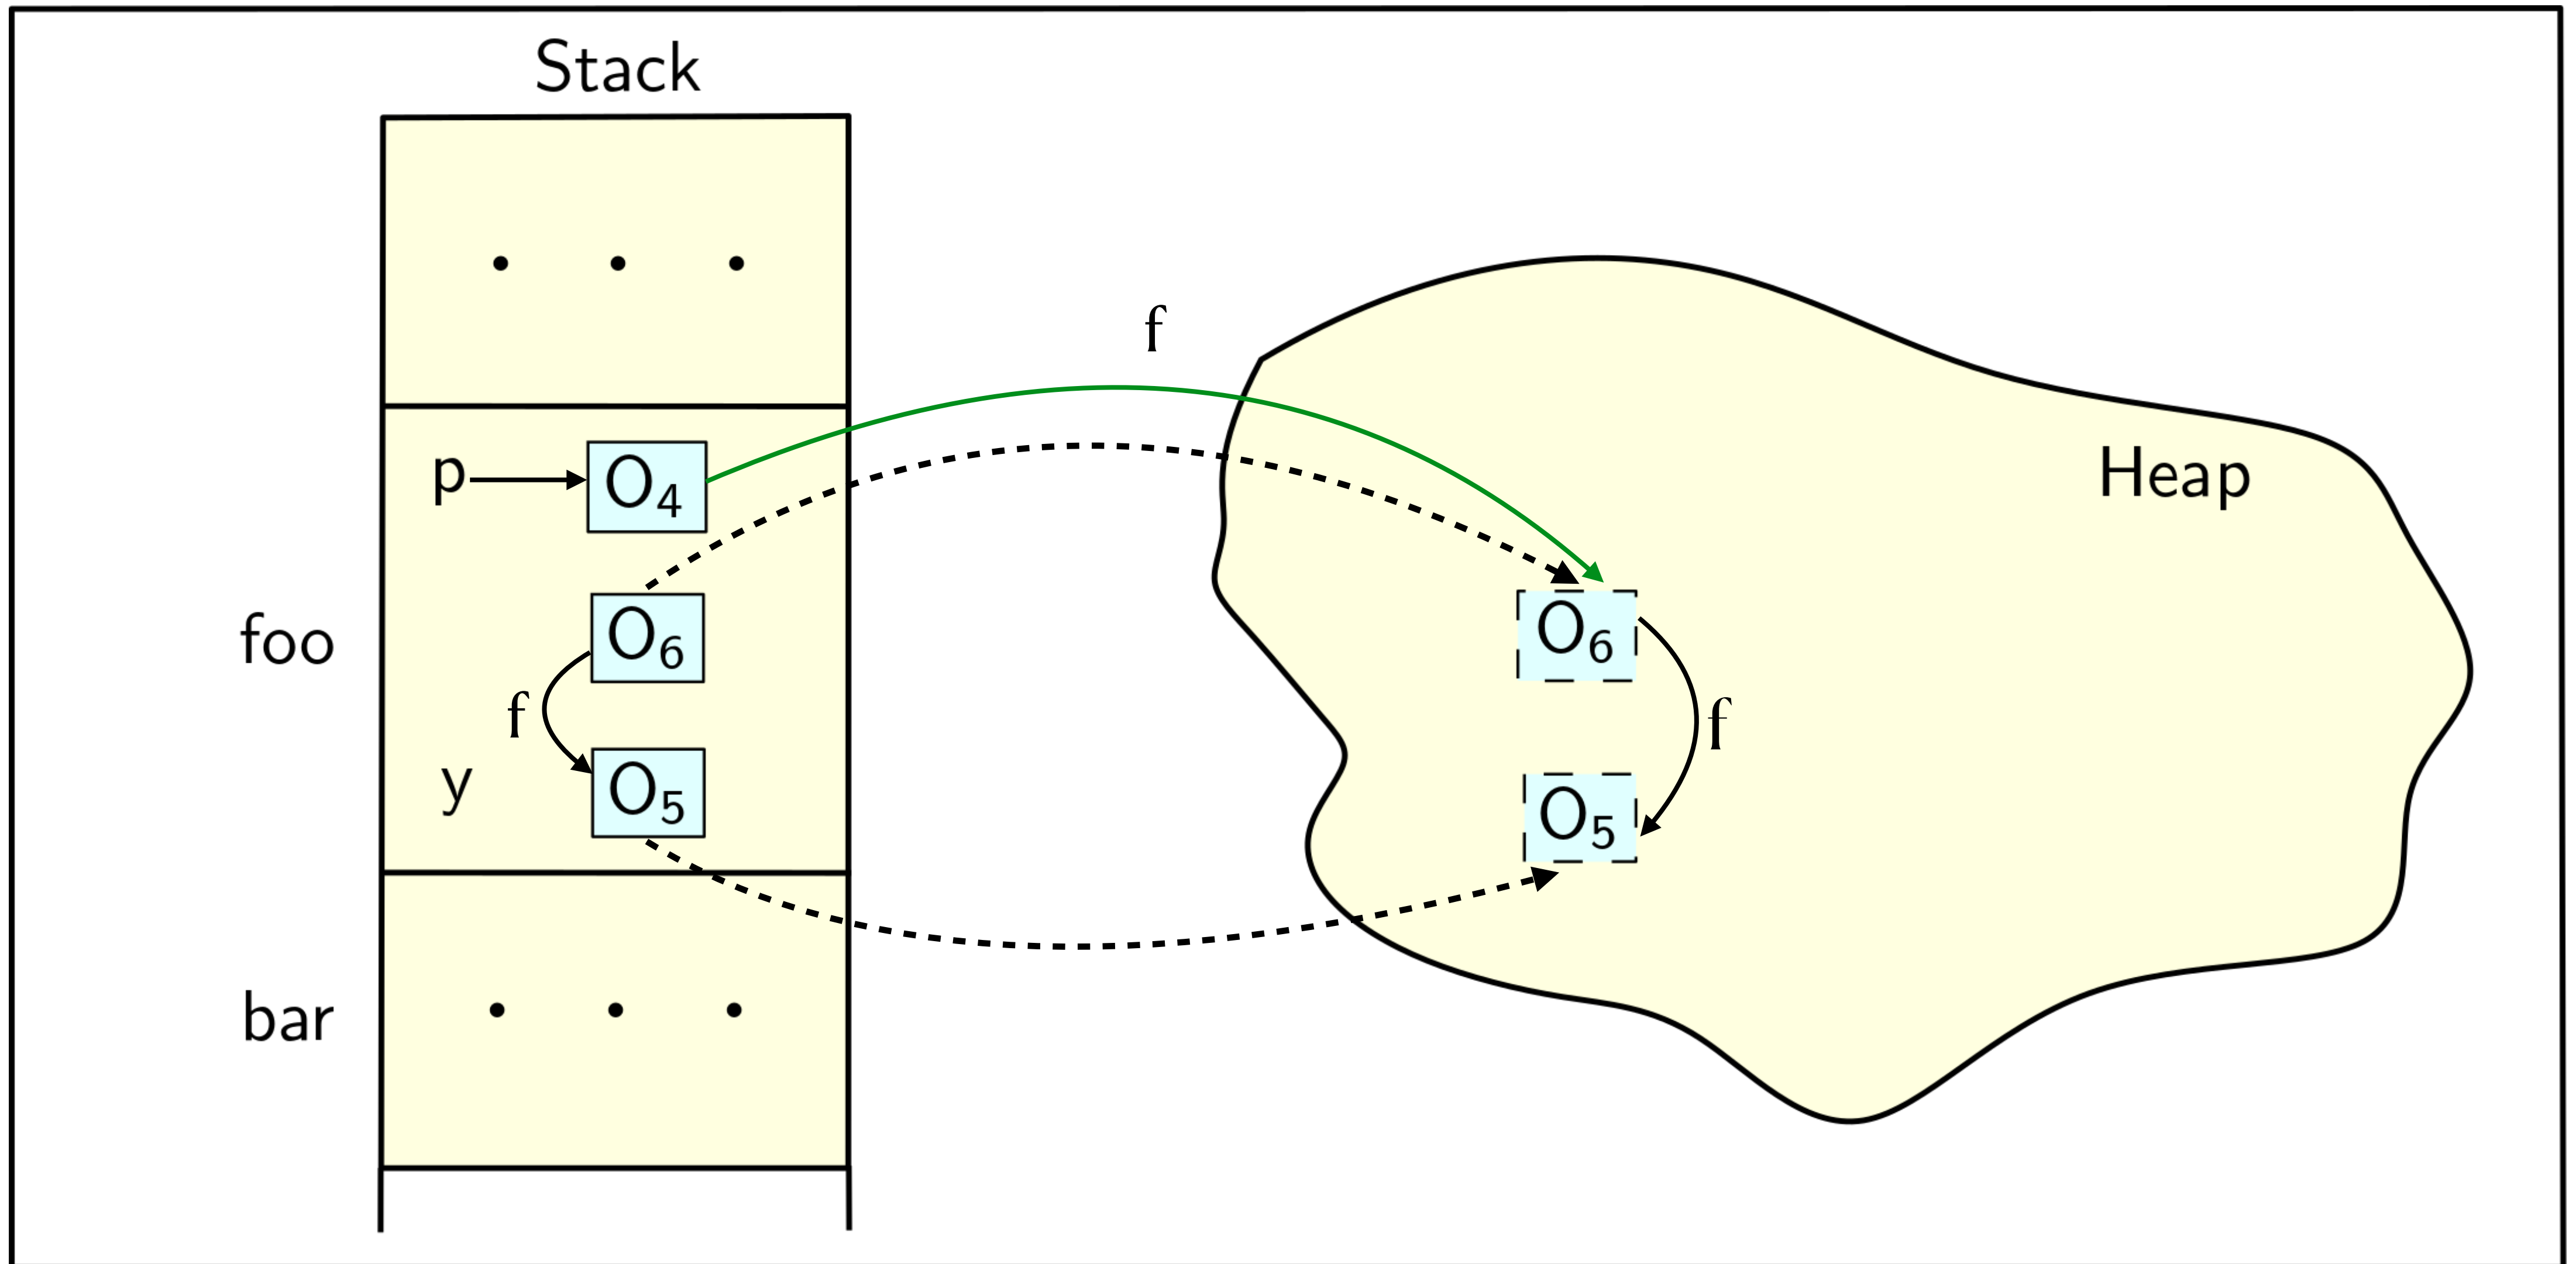
Heapification



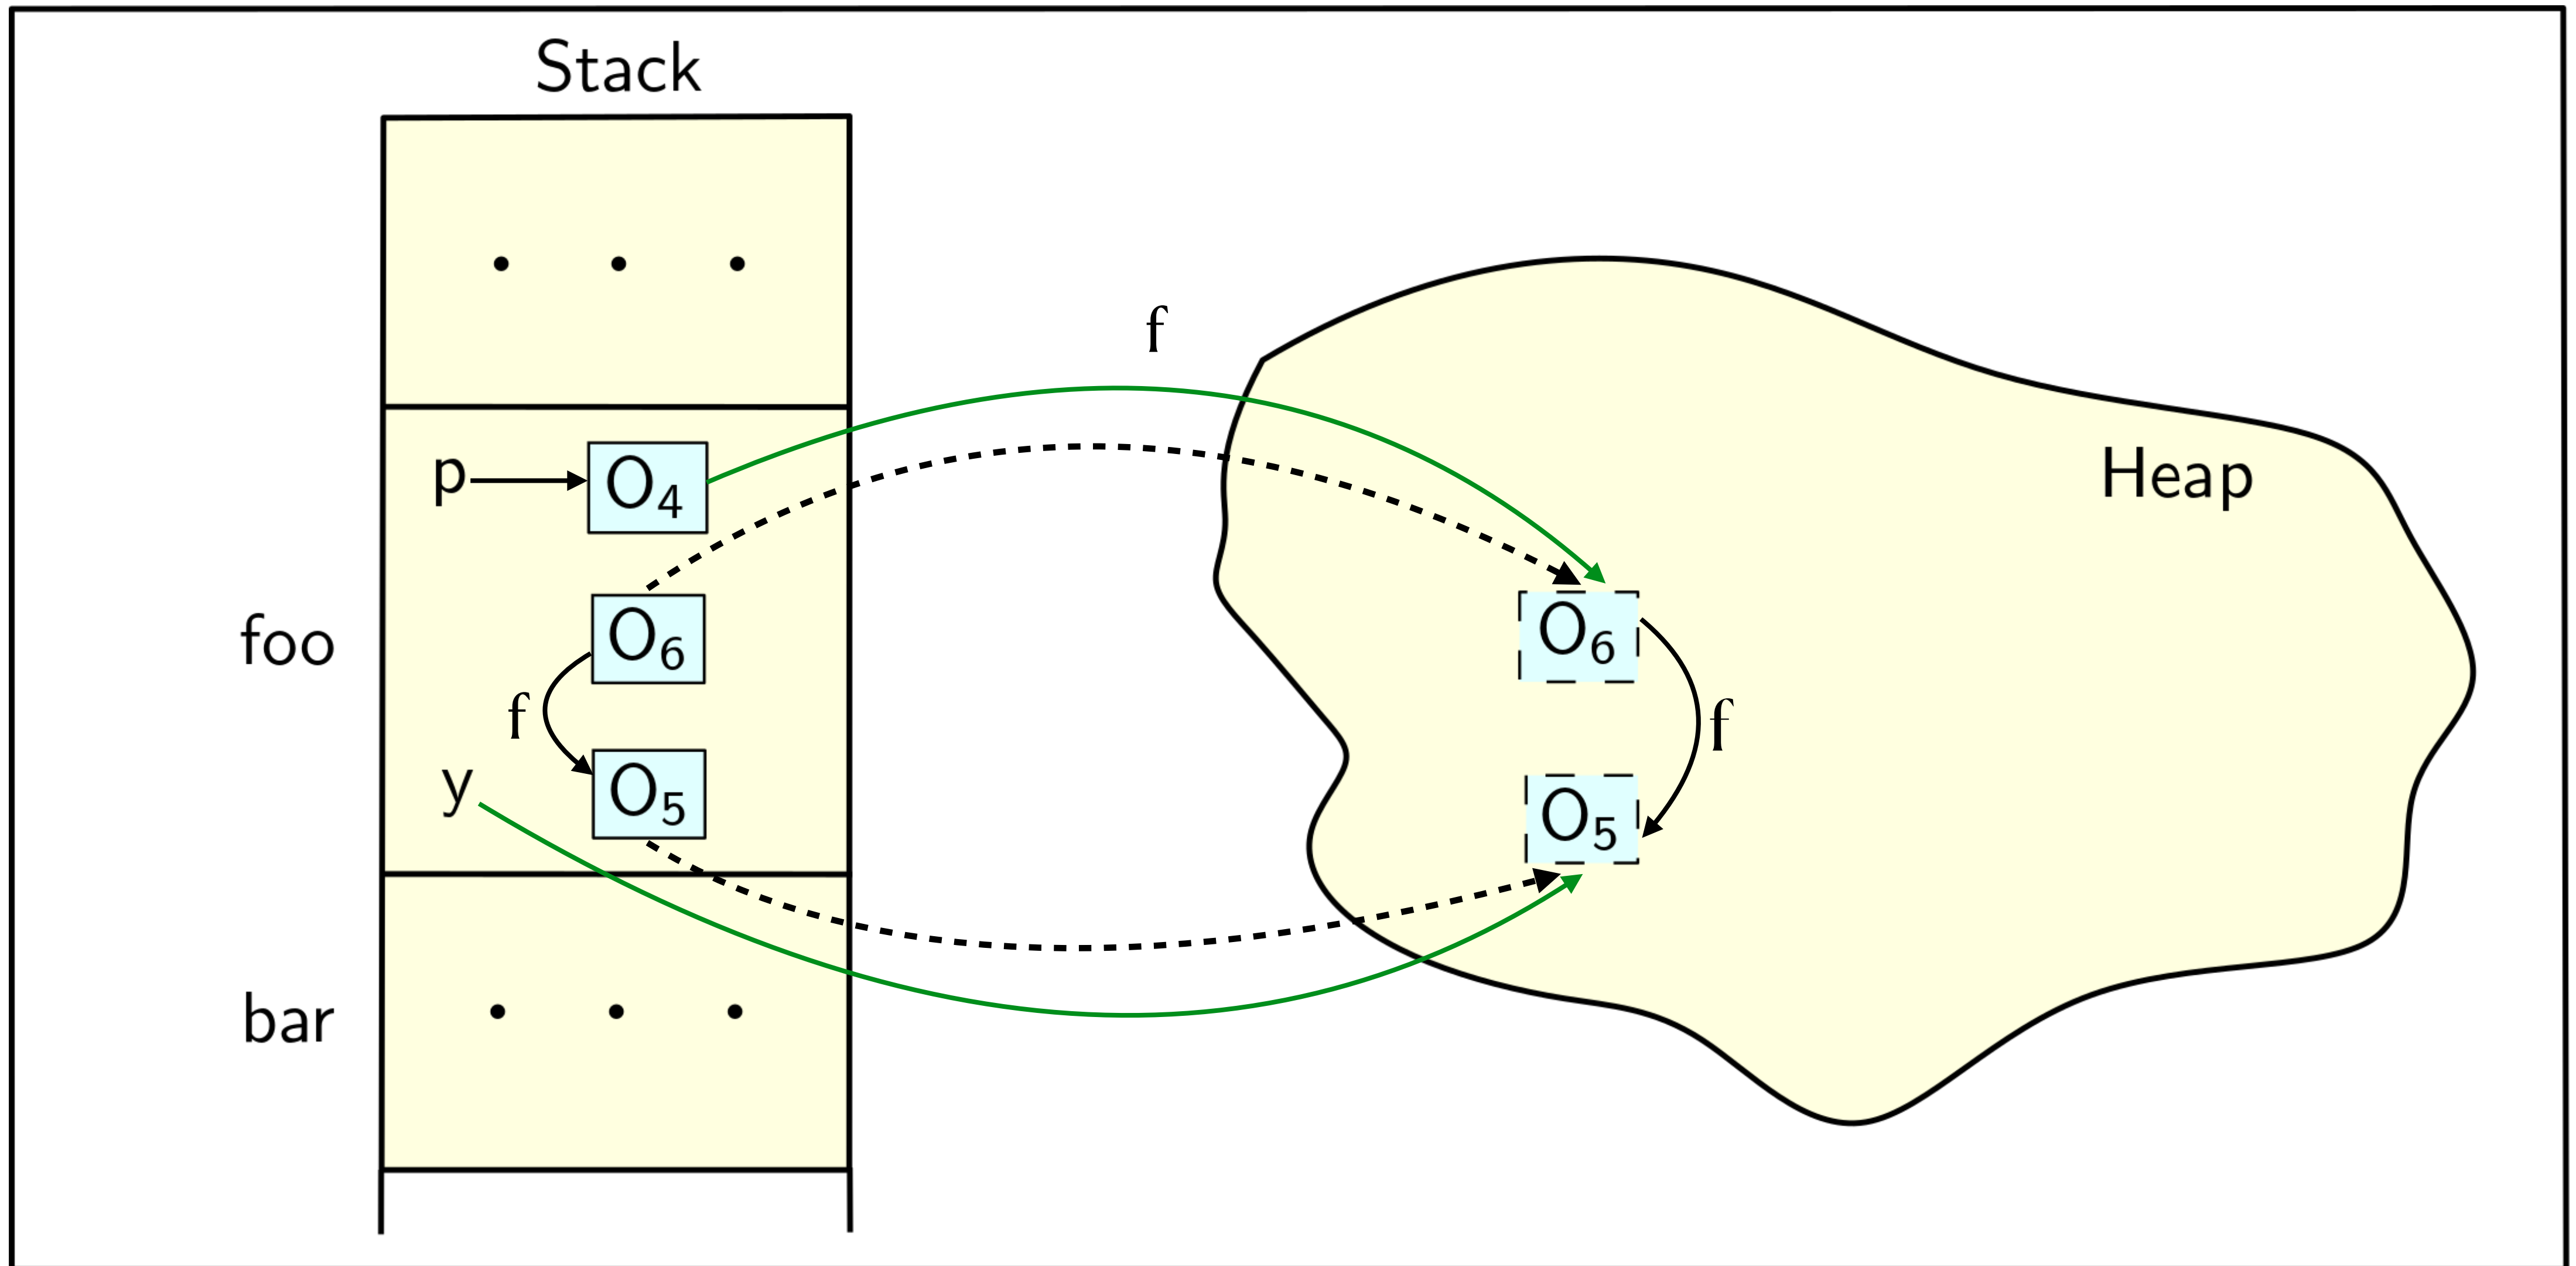
Heapification



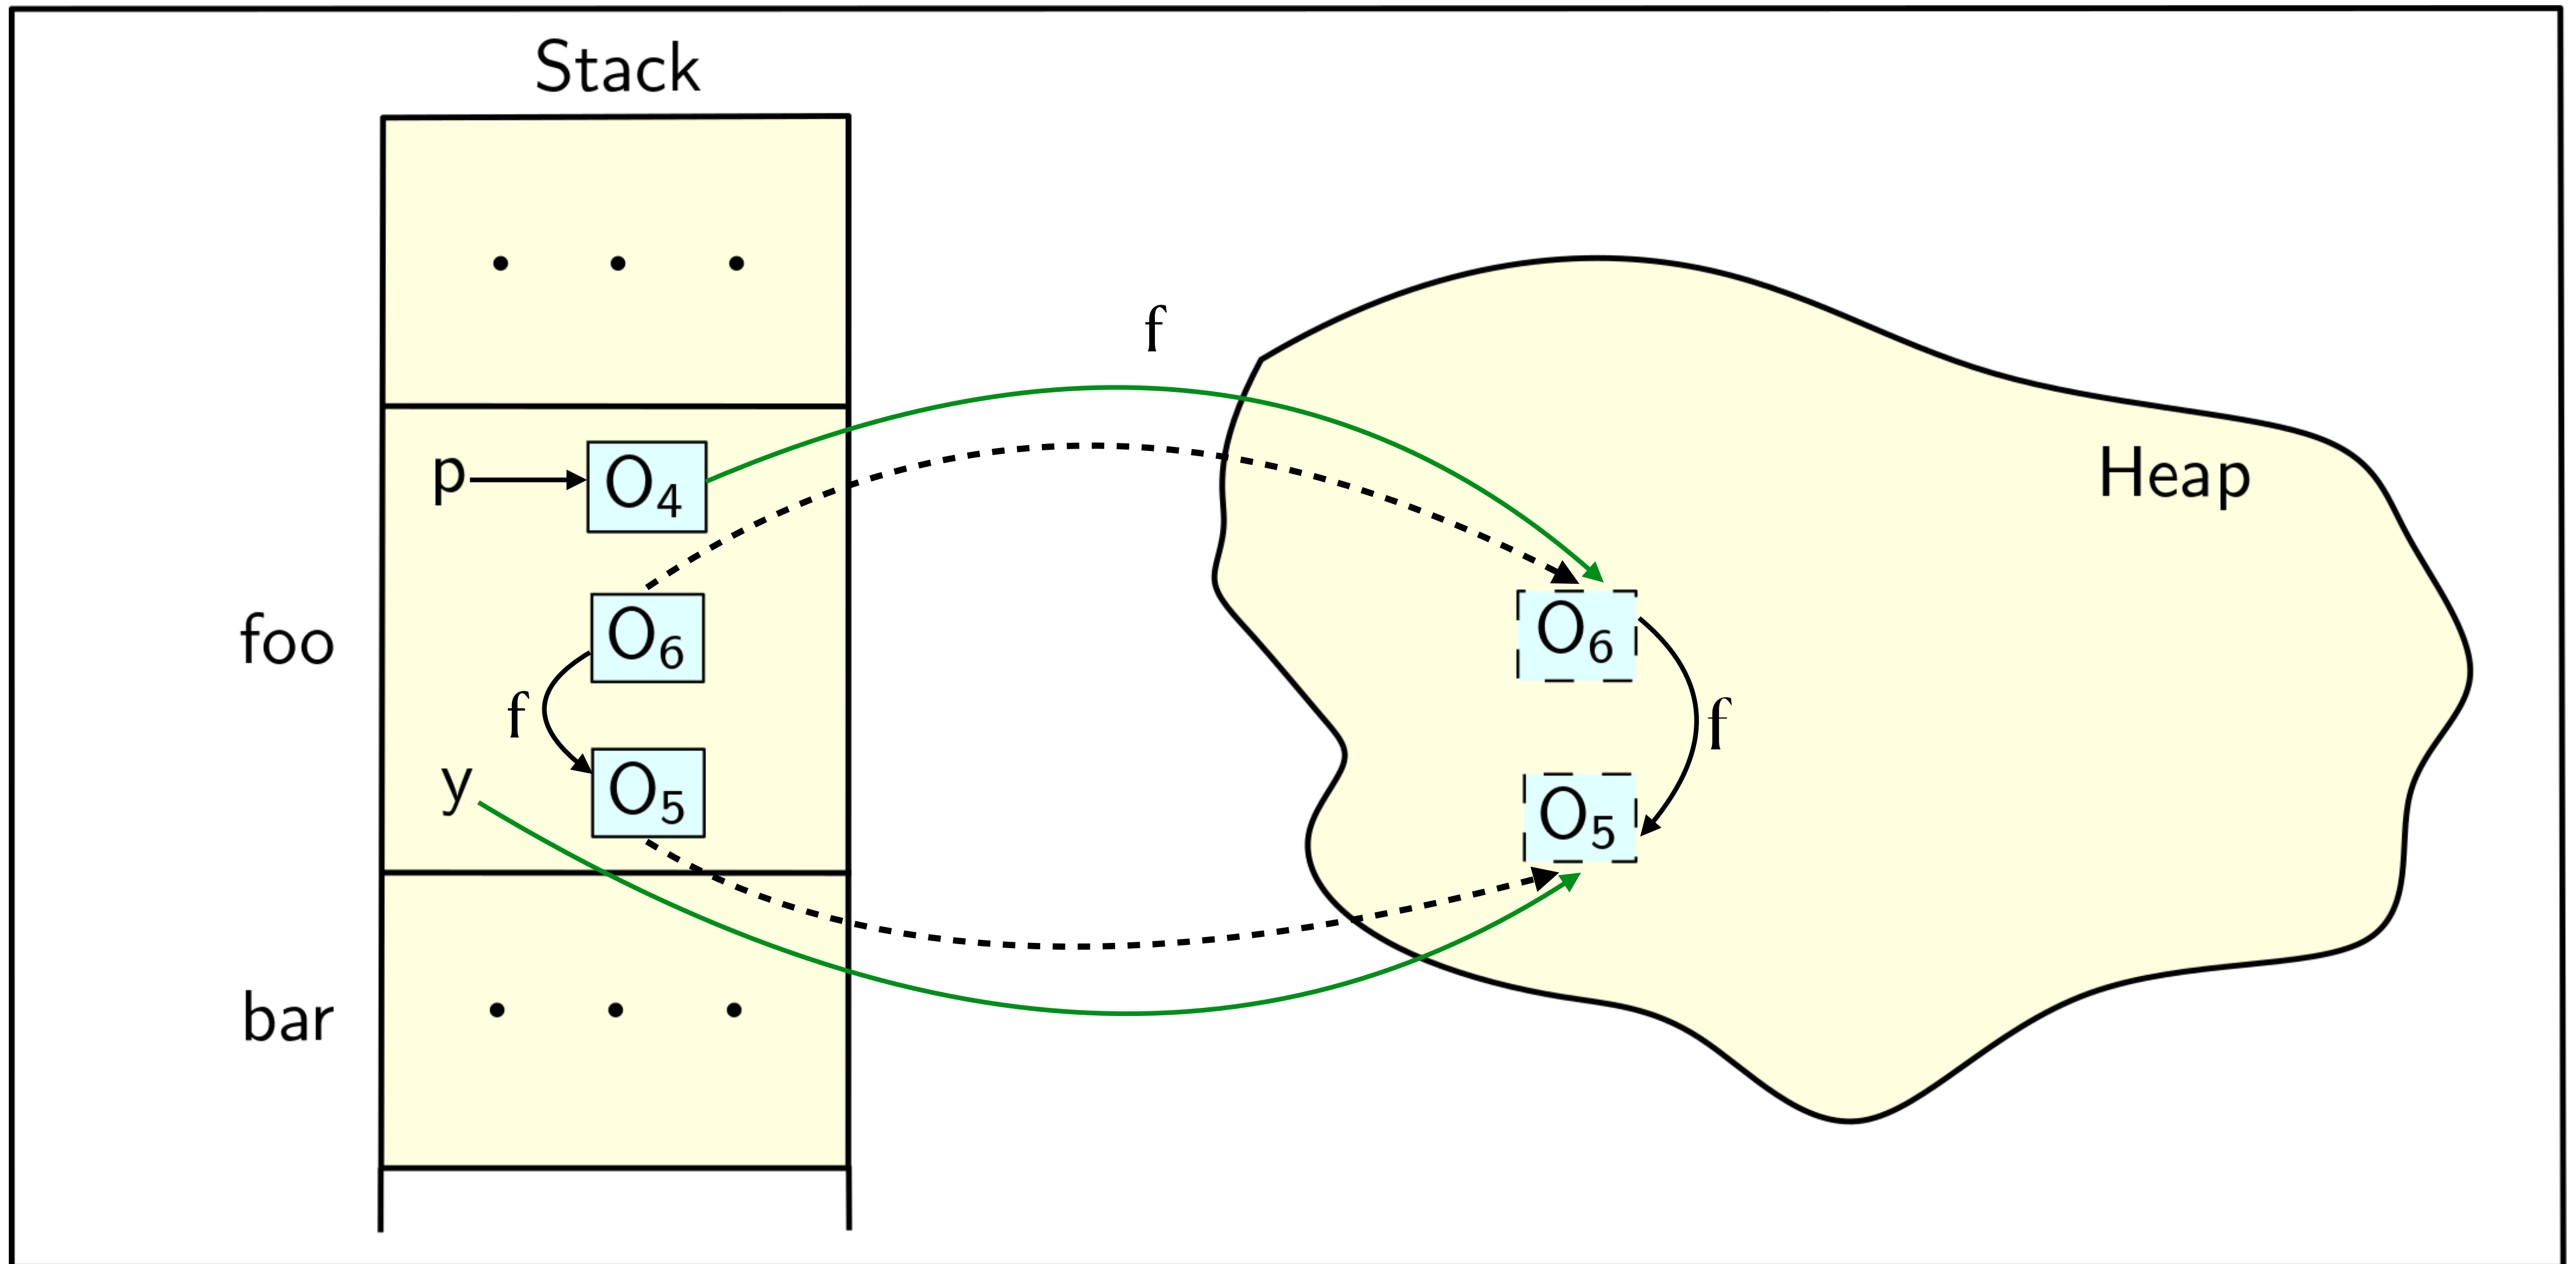
Heapification



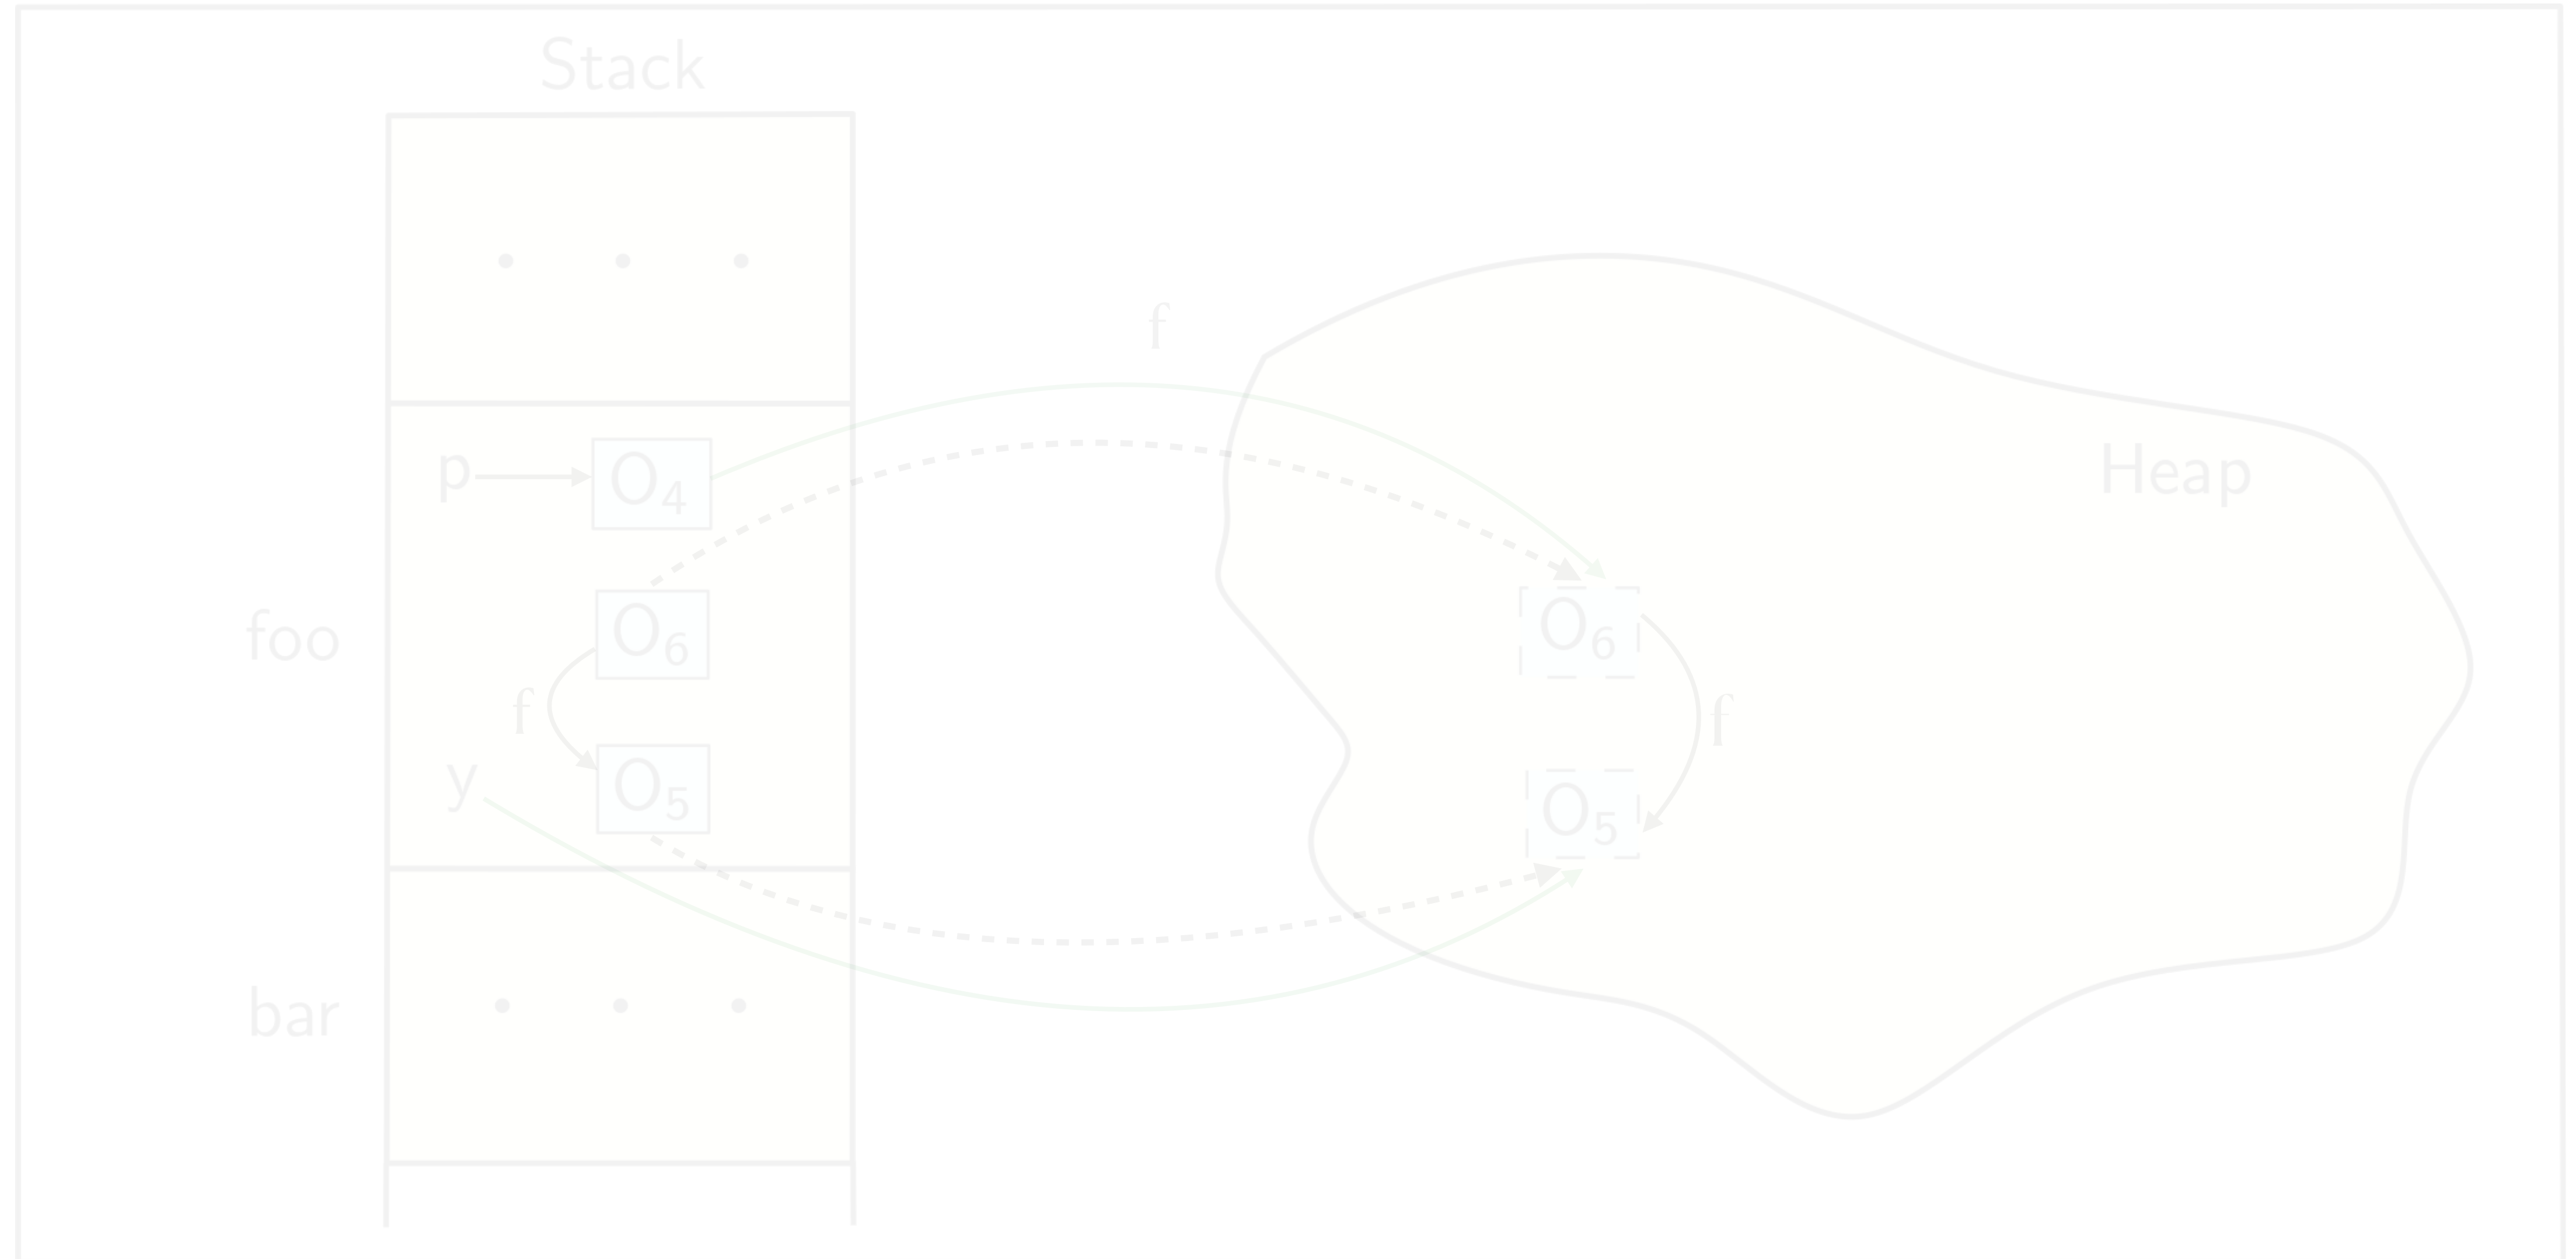
Heapification



Heapification



Heapification



Heapification

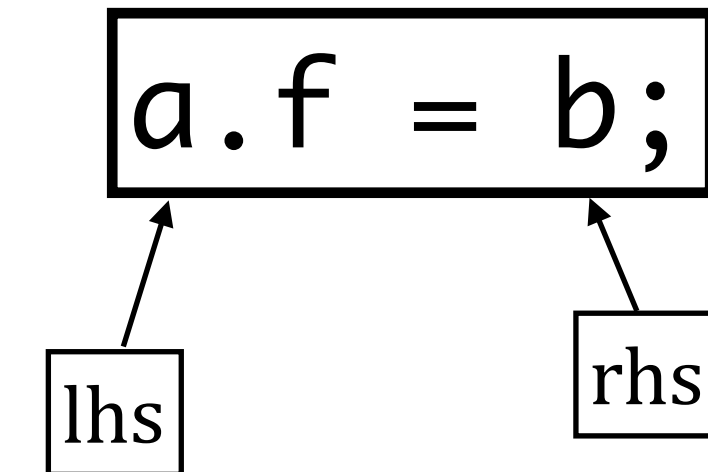
How to identify the need for heapification?

Checking the Need for Heapification

```
1 Procedure HeapificationCheckAtStore(lhs, rhs)
2   if rhs object is outside stack bounds then
3     | No heapification required.
4   else
5     | /* The rhs object is present on the stack */
6     | if lhs object is outside stack bounds then
7       | Heapify starting from the rhs object.
8     | else
9       | /* Both lhs and rhs objects are on the stack */
10      | if rhs object has been allocated before the lhs object then
11      | | No heapification required.
12      | | else
13      | | | /* The lhs object has been allocated in either the same frame or a deeper frame as
14      | | | compared to the rhs object */
15      | | | Perform stack-walk and heapify if needed.
```

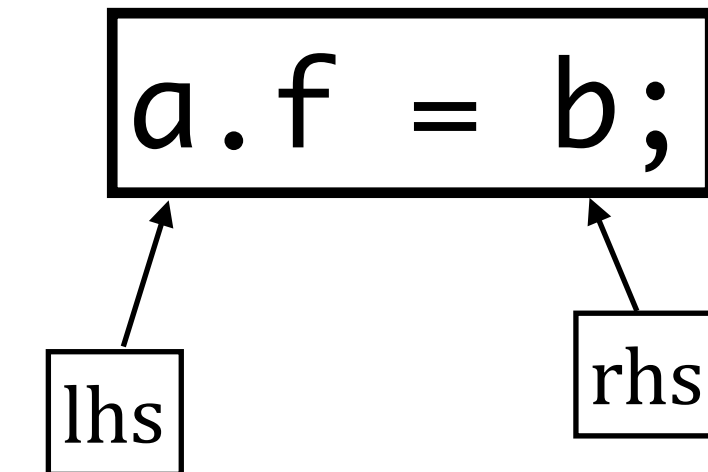
Checking the Need for Heapification

```
1 Procedure HeapificationCheckAtStore(lhs, rhs)
2   if rhs object is outside stack bounds then
3     | No heapification required.
4   else
5     | /* The rhs object is present on the stack */
6     | if lhs object is outside stack bounds then
7       | Heapify starting from the rhs object.
8     | else
9       | /* Both lhs and rhs objects are on the stack */
10      | if rhs object has been allocated before the lhs object then
11        | No heapification required.
12      | else
13        | /* The lhs object has been allocated in either the same frame or a deeper frame as
14          | compared to the rhs object */
15        | Perform stack-walk and heapify if needed.
```



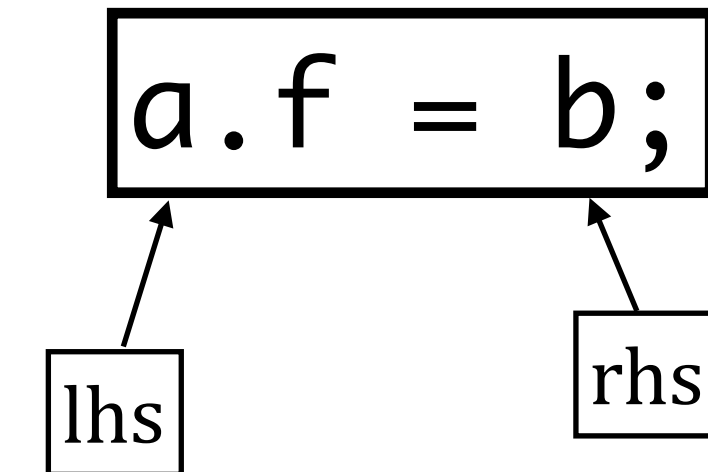
Checking the Need for Heapification

```
1 Procedure HeapificationCheckAtStore(lhs, rhs)
2   if rhs object is outside stack bounds then
3     | No heapification required.
4   else
5     | /* The rhs object is present on the stack */
6     | if lhs object is outside stack bounds then
7       | Heapify starting from the rhs object.
8     | else
9       | /* Both lhs and rhs objects are on the stack */
10      | if rhs object has been allocated before the lhs object then
11        | No heapification required.
12      | else
13        | /* The lhs object has been allocated in either the same frame or a deeper frame as
14          | compared to the rhs object */
15        | Perform stack-walk and heapify if needed.
```



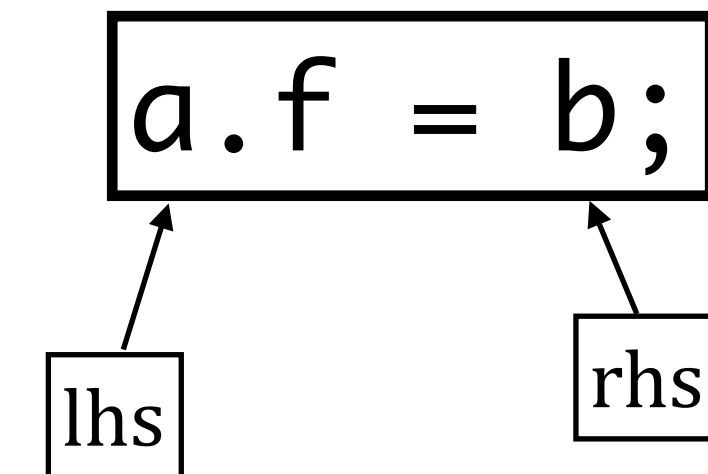
Checking the Need for Heapification

```
1 Procedure HeapificationCheckAtStore(lhs, rhs)
2   if rhs object is outside stack bounds then
3     | No heapification required.
4   else
5     | /* The rhs object is present on the stack */
6     | if lhs object is outside stack bounds then
7       | Heapify starting from the rhs object.
8     | else
9       | /* Both lhs and rhs objects are on the stack */
10      | if rhs object has been allocated before the lhs object then
11        | No heapification required.
12      | else
13        | /* The lhs object has been allocated in either the same frame or a deeper frame as
14          | compared to the rhs object */
15        | Perform stack-walk and heapify if needed.
```



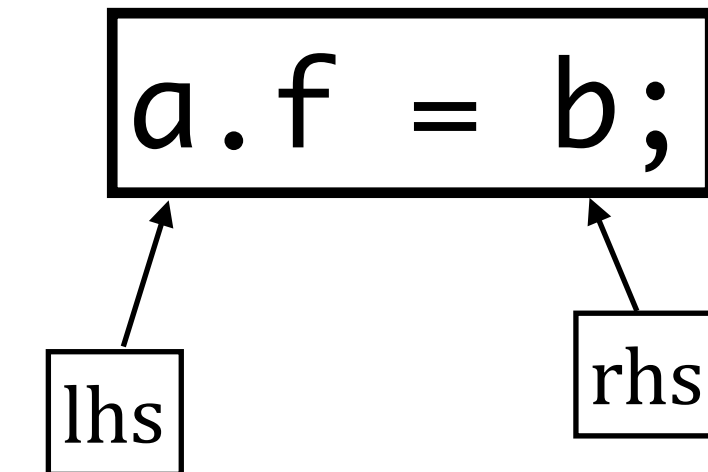
Checking the Need for Heapification

```
1 Procedure HeapificationCheckAtStore(lhs, rhs)
2   if rhs object is outside stack bounds then
3     | No heapification required.
4   else
5     | /* The rhs object is present on the stack */
6     | if (lhs object is outside stack bounds) then
7       | Heapify starting from the rhs object.
8     | else
9       | /* Both lhs and rhs objects are on the stack */
10      | if rhs object has been allocated before the lhs object then
11      | | No heapification required.
12      | else
13      | | /* The lhs object has been allocated in either the same frame or a deeper frame as
14      | | compared to the rhs object */
15      | | Perform stack-walk and heapify if needed.
```



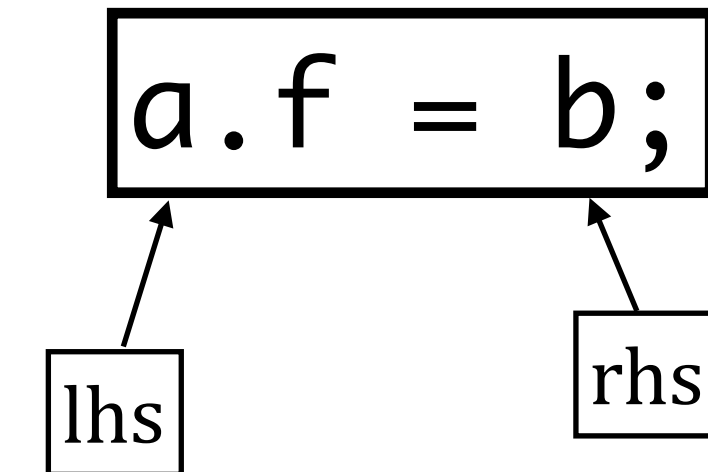
Checking the Need for Heapification

```
1 Procedure HeapificationCheckAtStore(lhs, rhs)
2   if rhs object is outside stack bounds then
3     | No heapification required.
4   else
5     | /* The rhs object is present on the stack */
6     | if lhs object is outside stack bounds then
7       | Heapify starting from the rhs object.
8     | else
9       | /* Both lhs and rhs objects are on the stack */
10      | if rhs object has been allocated before the lhs object then
11        | No heapification required.
12      | else
13        | /* The lhs object has been allocated in either the same frame or a deeper frame as
14          | compared to the rhs object */
15        | Perform stack-walk and heapify if needed.
```



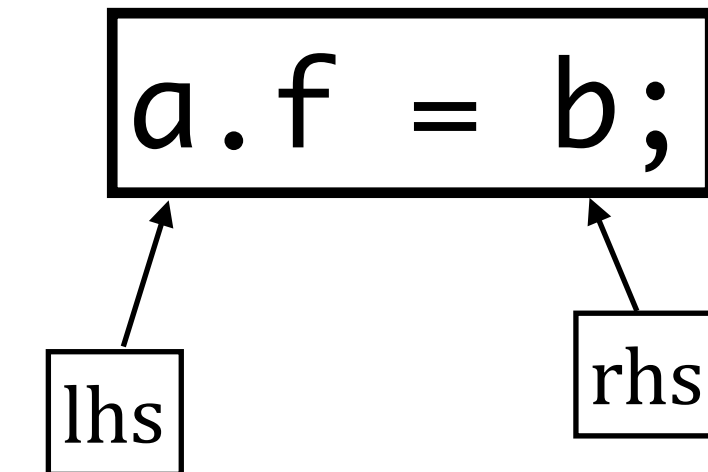
Checking the Need for Heapification

```
1 Procedure HeapificationCheckAtStore(lhs, rhs)
2   if rhs object is outside stack bounds then
3     | No heapification required.
4   else
5     | /* The rhs object is present on the stack */
6     | if lhs object is outside stack bounds then
7       | Heapify starting from the rhs object.
8     | else
9       | /* Both lhs and rhs objects are on the stack */
10      | if rhs object has been allocated before the lhs object then
11        | No heapification required.
12      | else
13        | /* The lhs object has been allocated in either the same frame or a deeper frame as
14          | compared to the rhs object */
15        | Perform stack-walk and heapify if needed.
```



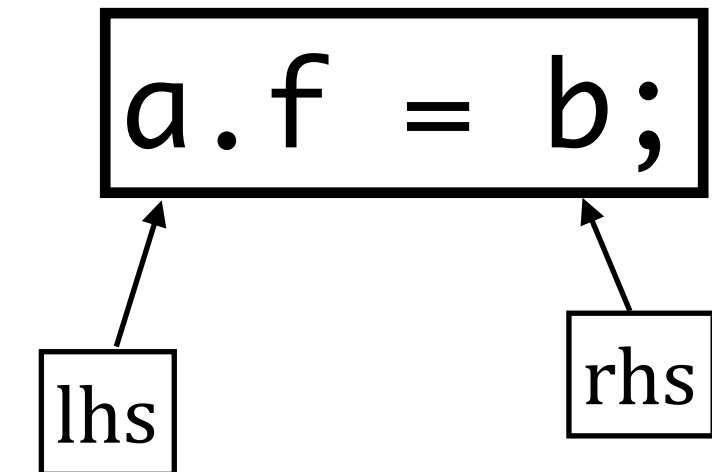
Checking the Need for Heapification

```
1 Procedure HeapificationCheckAtStore(lhs, rhs)
2   if rhs object is outside stack bounds then
3     | No heapification required.
4   else
5     | /* The rhs object is present on the stack */
6     | if lhs object is outside stack bounds then
7       | Heapify starting from the rhs object.
8     | else
9       | /* Both lhs and rhs objects are on the stack */
10      | if rhs object has been allocated before the lhs object then
11        | No heapification required.
12      | else
13        | /* The lhs object has been allocated in either the same frame or a deeper frame as
14          | compared to the rhs object */
15        | Perform stack-walk and heapify if needed.
```



Checking the Need for Heapification

```
1 Procedure HeapificationCheckAtStore(lhs, rhs)
2   if rhs object is outside stack bounds then
3     | No heapification required.
4   else
5     | /* The rhs object is present on the stack */
6     | if lhs object is outside stack bounds then
7       | Heapify starting from the rhs object.
8     | else
9       | /* Both lhs and rhs objects are on the stack */
10      | if rhs object has been allocated before the lhs object then
11        | No heapification required.
12      | else
13        | /* The lhs object has been allocated in either the same frame or a deeper frame as
14          | compared to the rhs object */
15        | Perform stack-walk and heapify if needed.
```

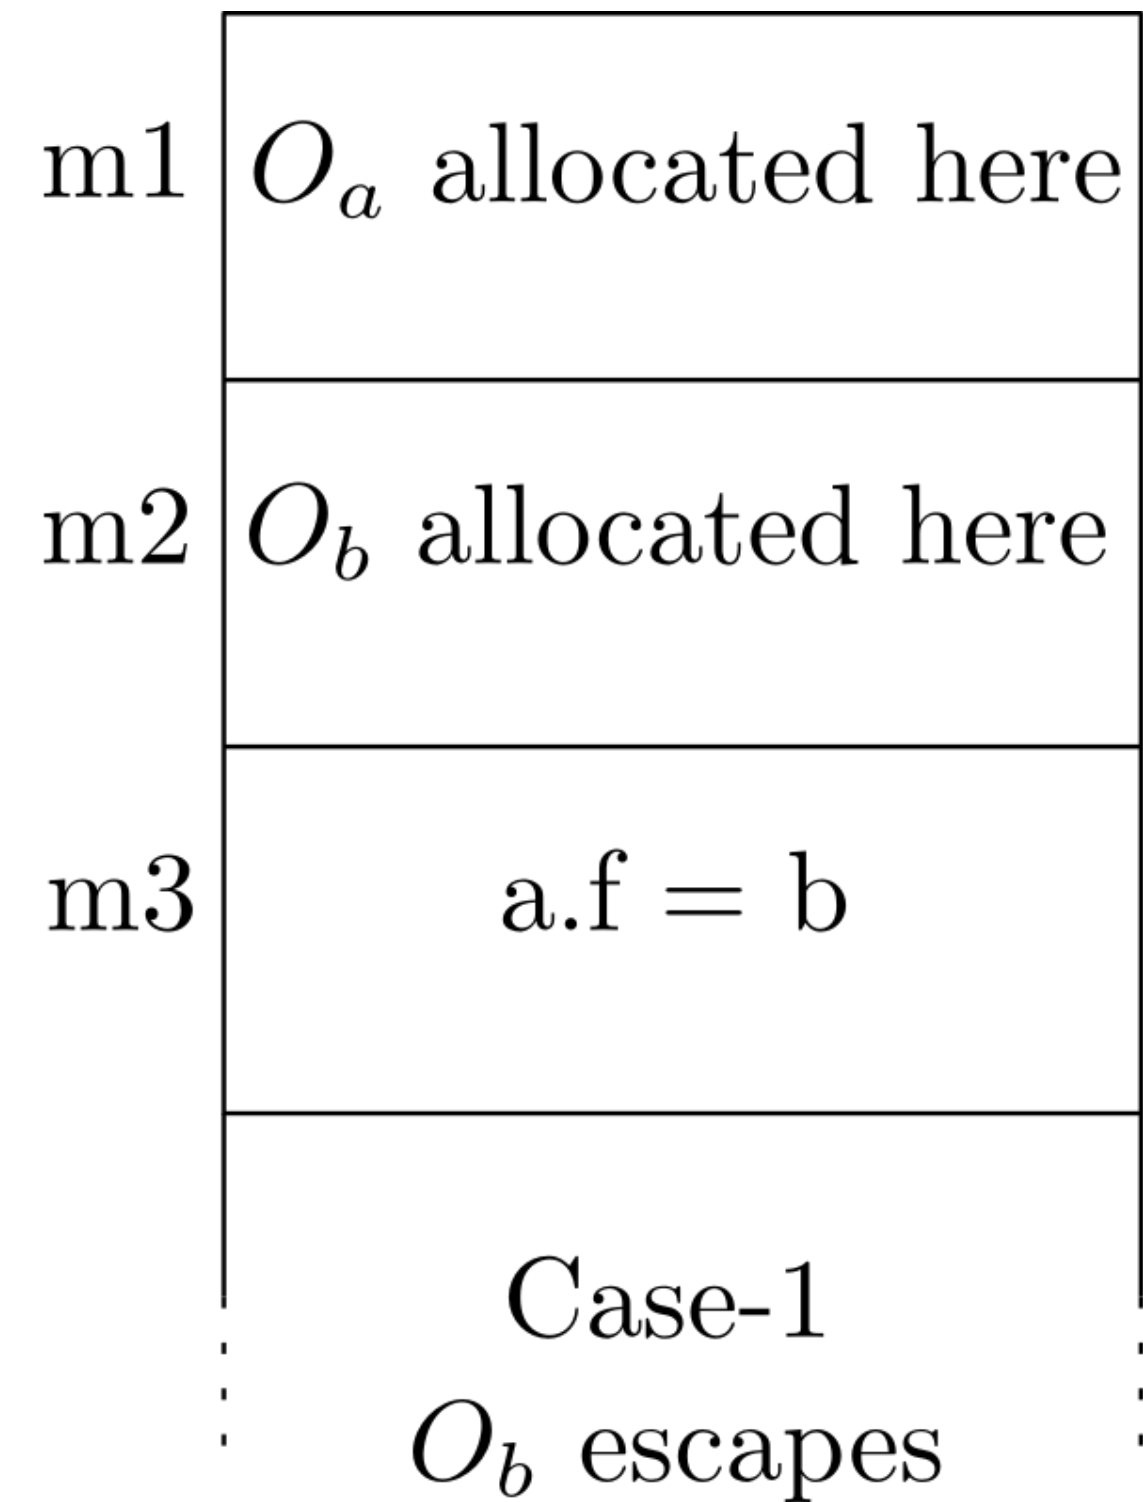


Scenarios at Store Statement

```
1. class T {  
2.     T f;  
3.     void m1() {m2(. . .);} }  
4.     void m2() {m3(. . .);} }  
5.     void m3(T a, T b) {  
6.         a.f = b;  
7.     } /* method m3 */  
8.} /* class T */
```

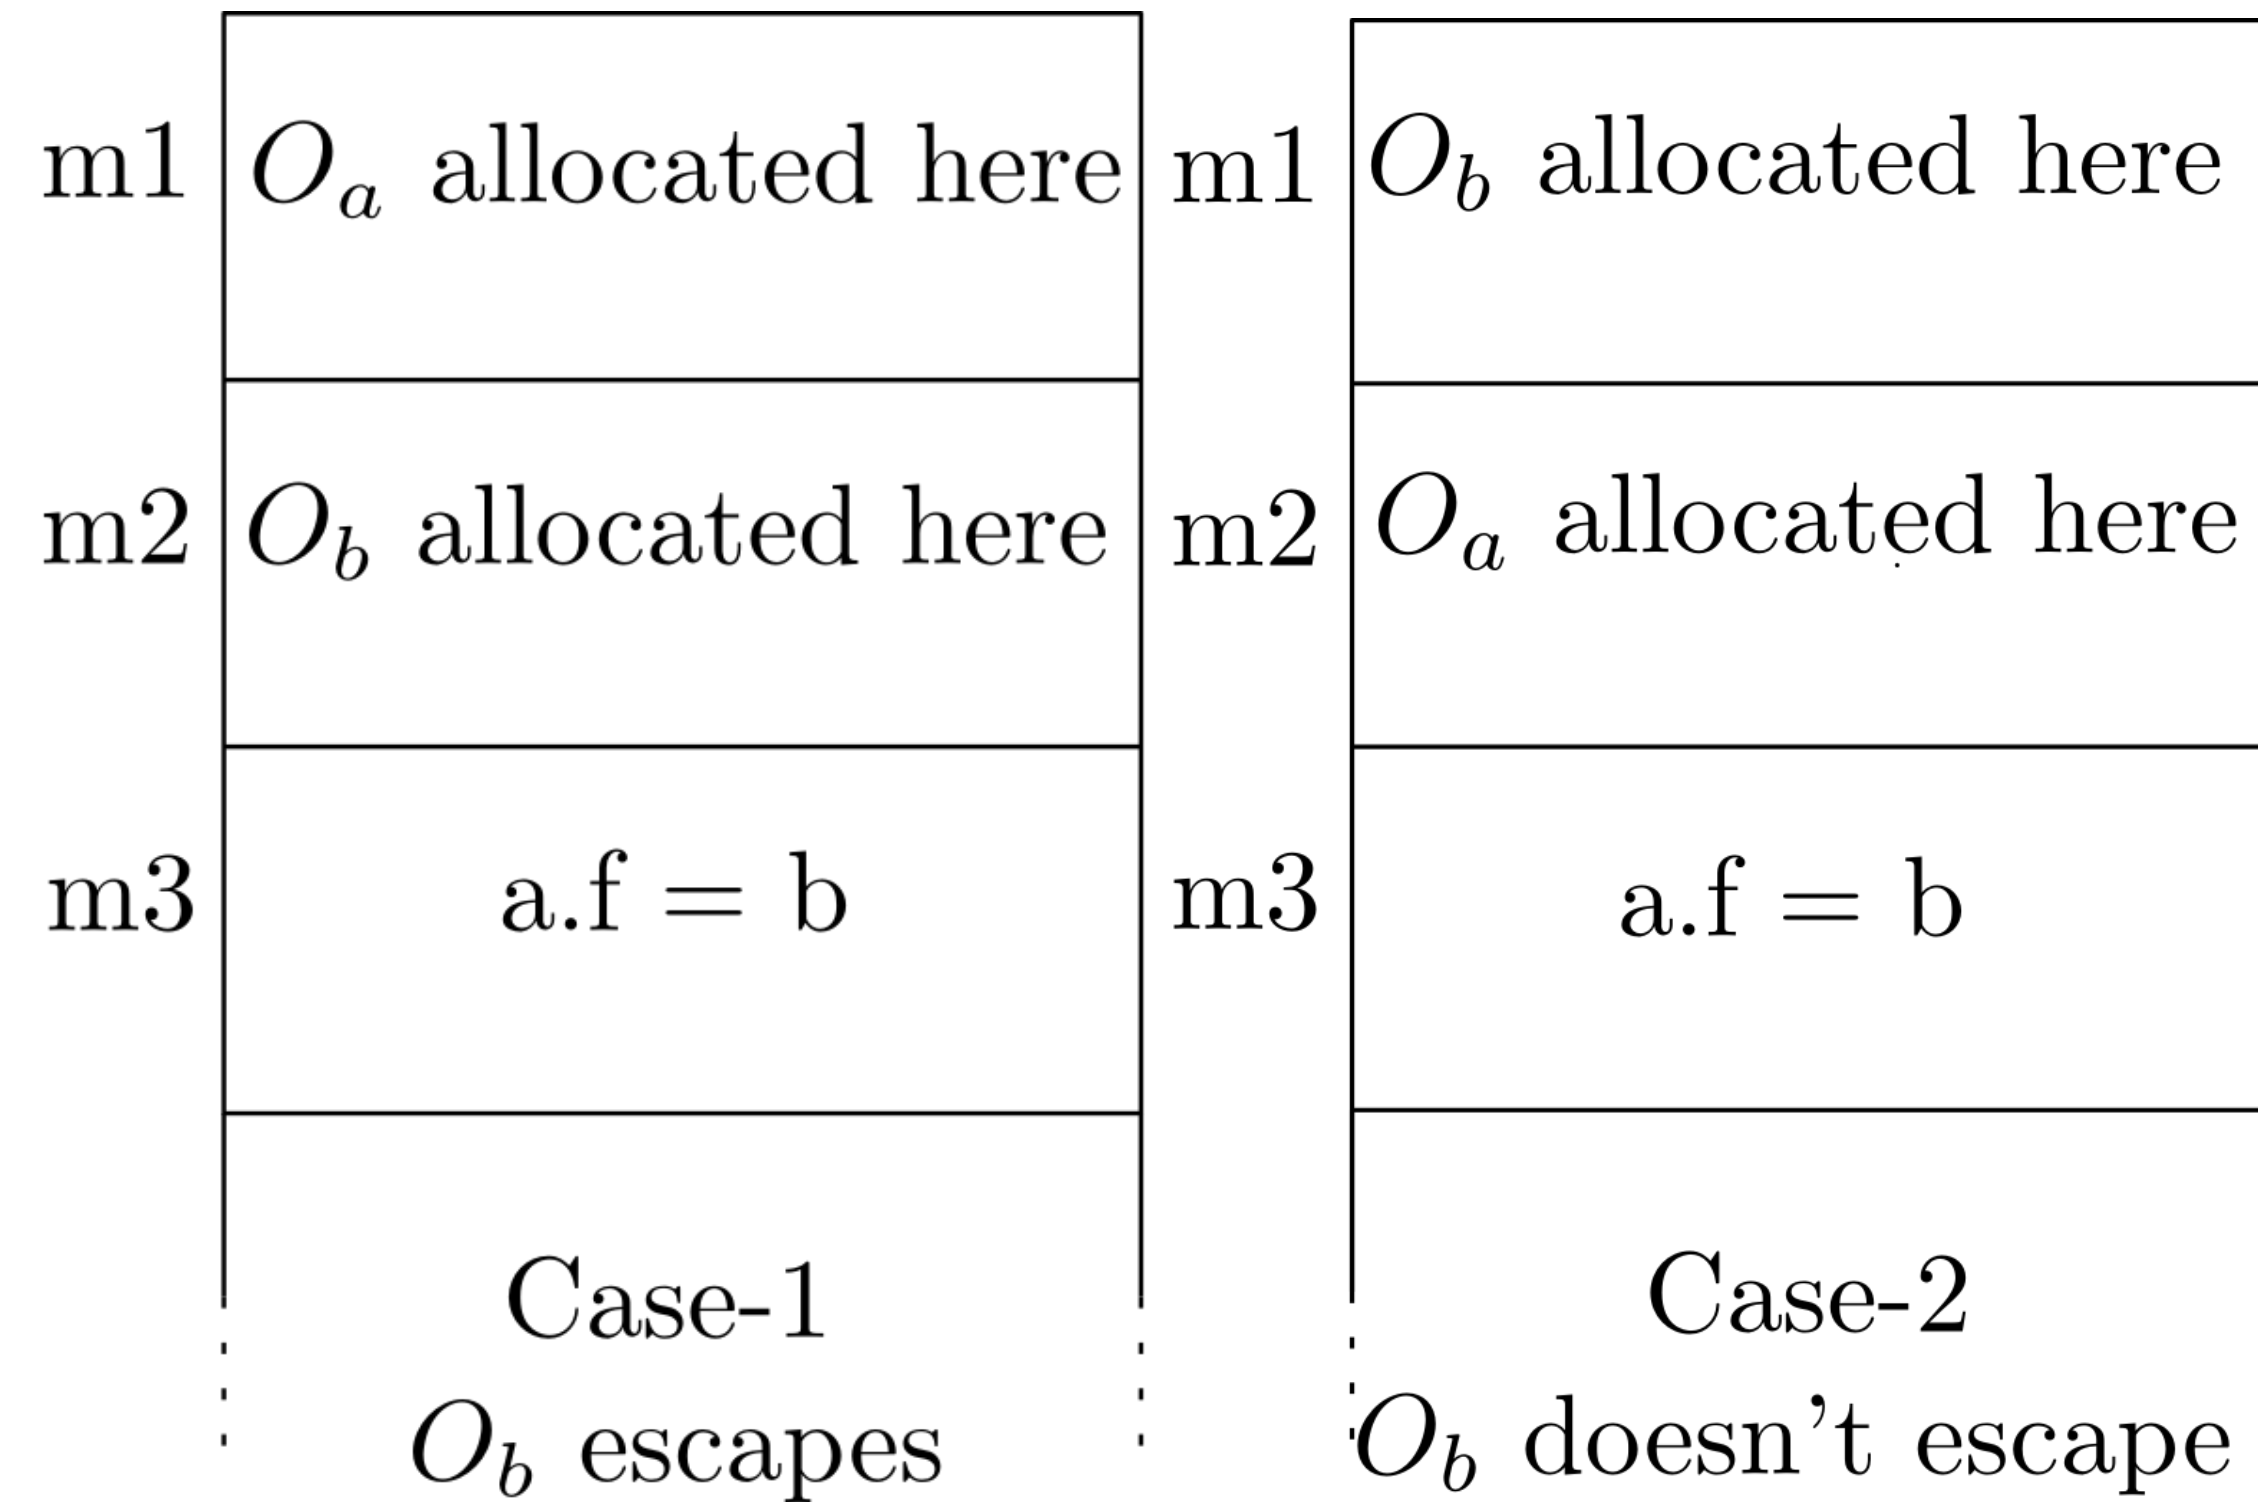
Scenarios at Store Statement

```
1. class T {  
2.     T f;  
3.     void m1() {m2(. . .);} }  
4.     void m2() {m3(. . .);} }  
5.     void m3(T a, T b) {  
6.         a.f = b;  
7.     } /* method m3 */  
8.} /* class T */
```



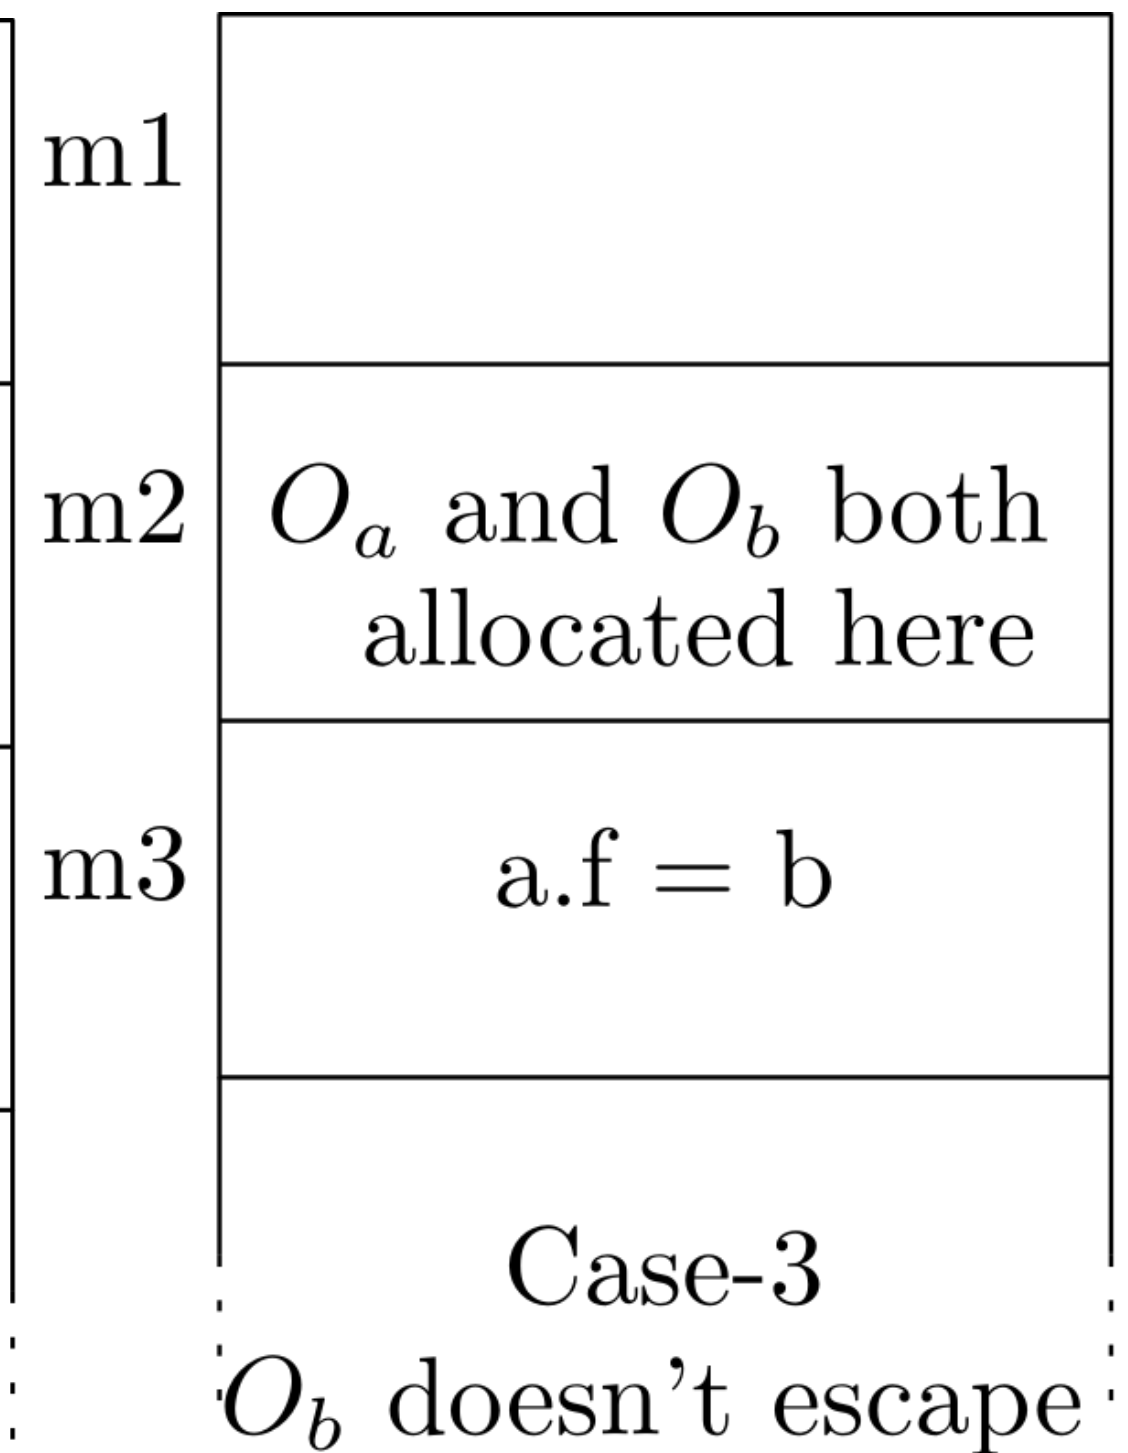
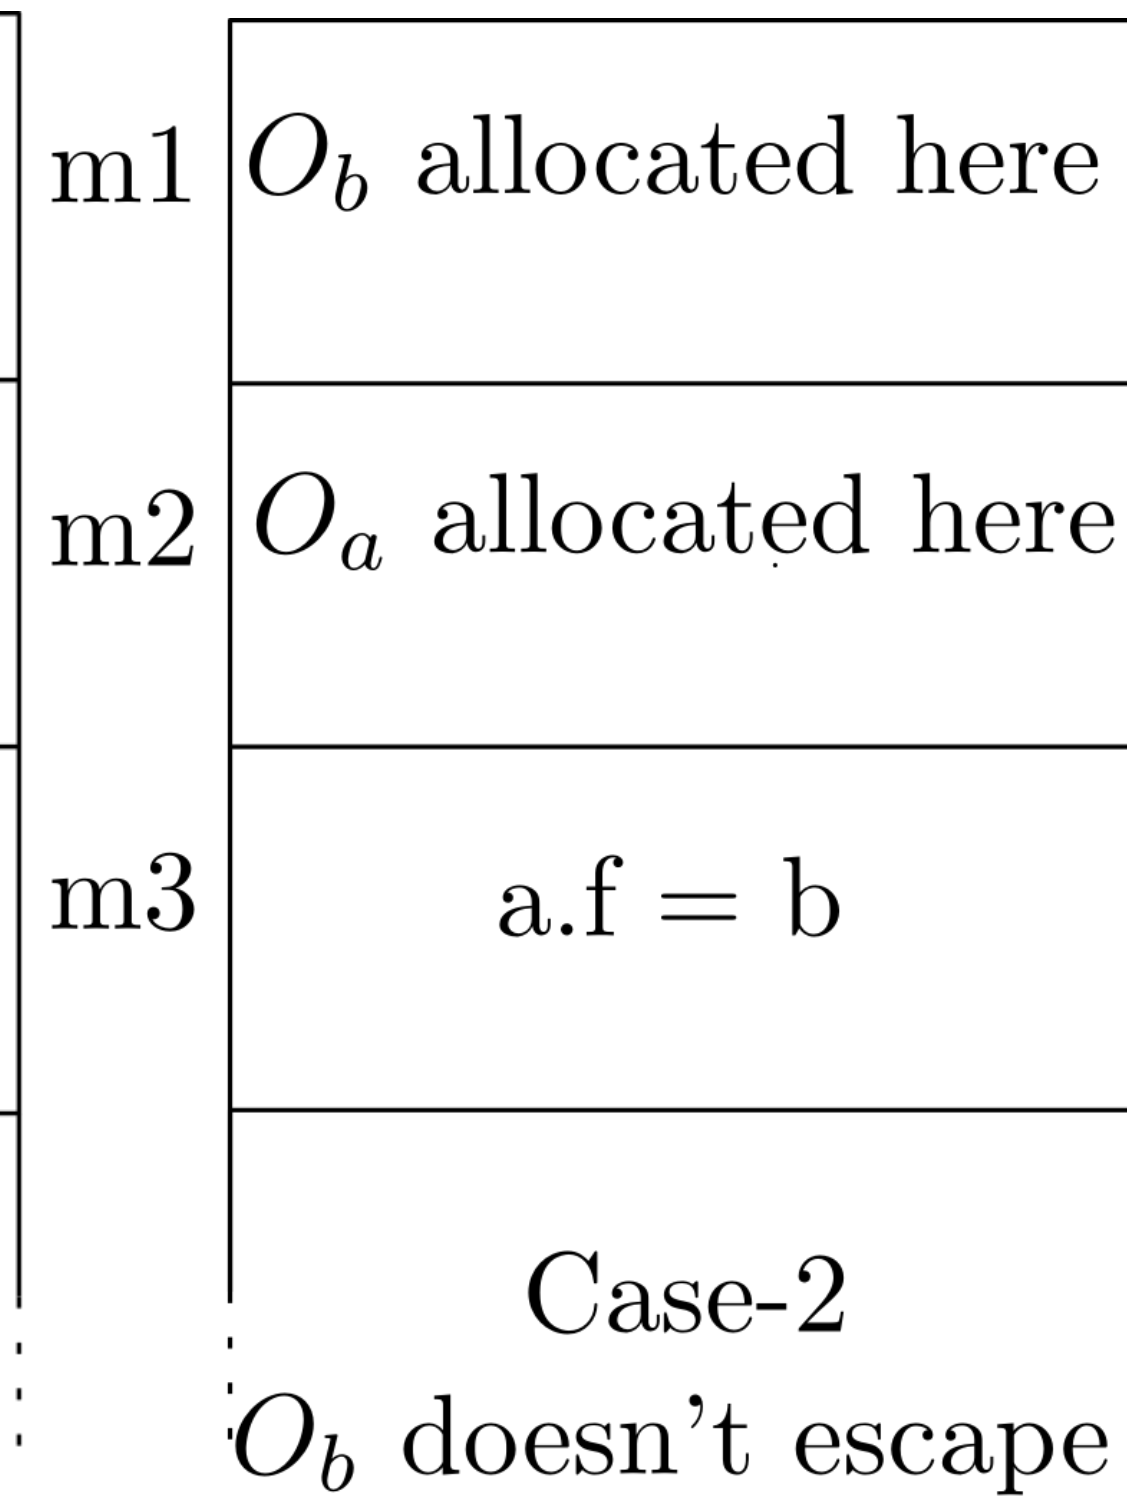
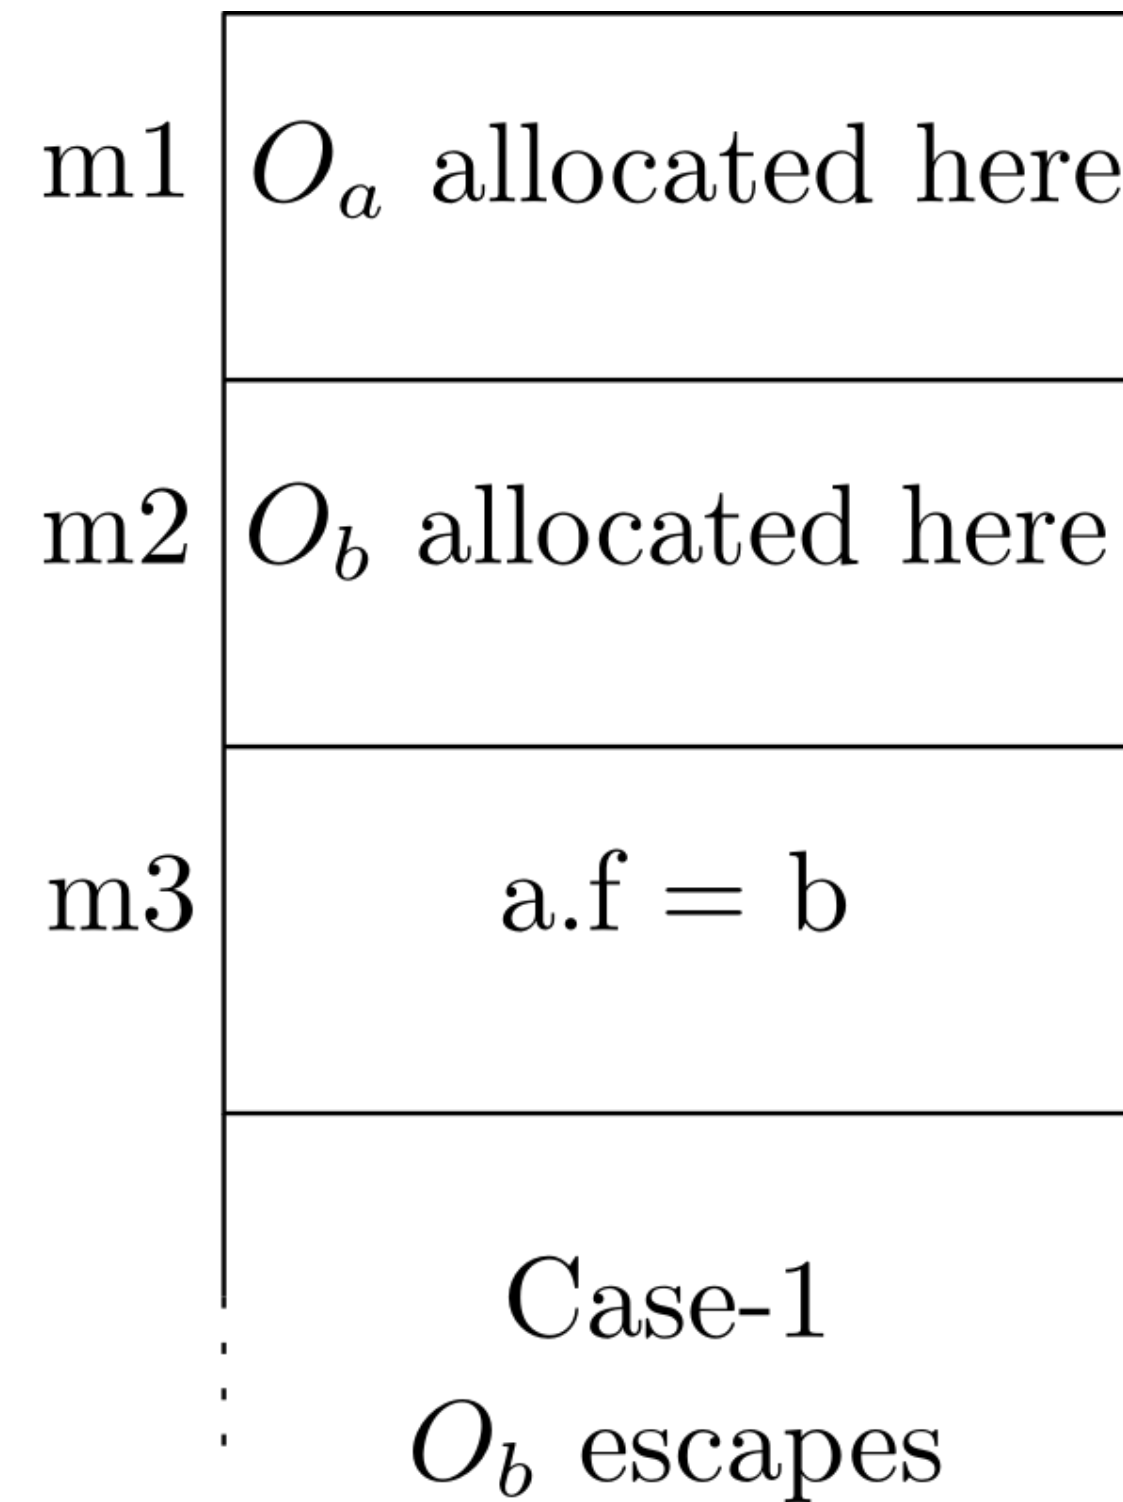
Scenarios at Store Statement

```
1. class T {  
2.     T f;  
3.     void m1() {m2(. . .);} }  
4.     void m2() {m3(. . .);} }  
5.     void m3(T a, T b) {  
6.         a.f = b;  
7.     } /* method m3 */  
8.} /* class T */
```



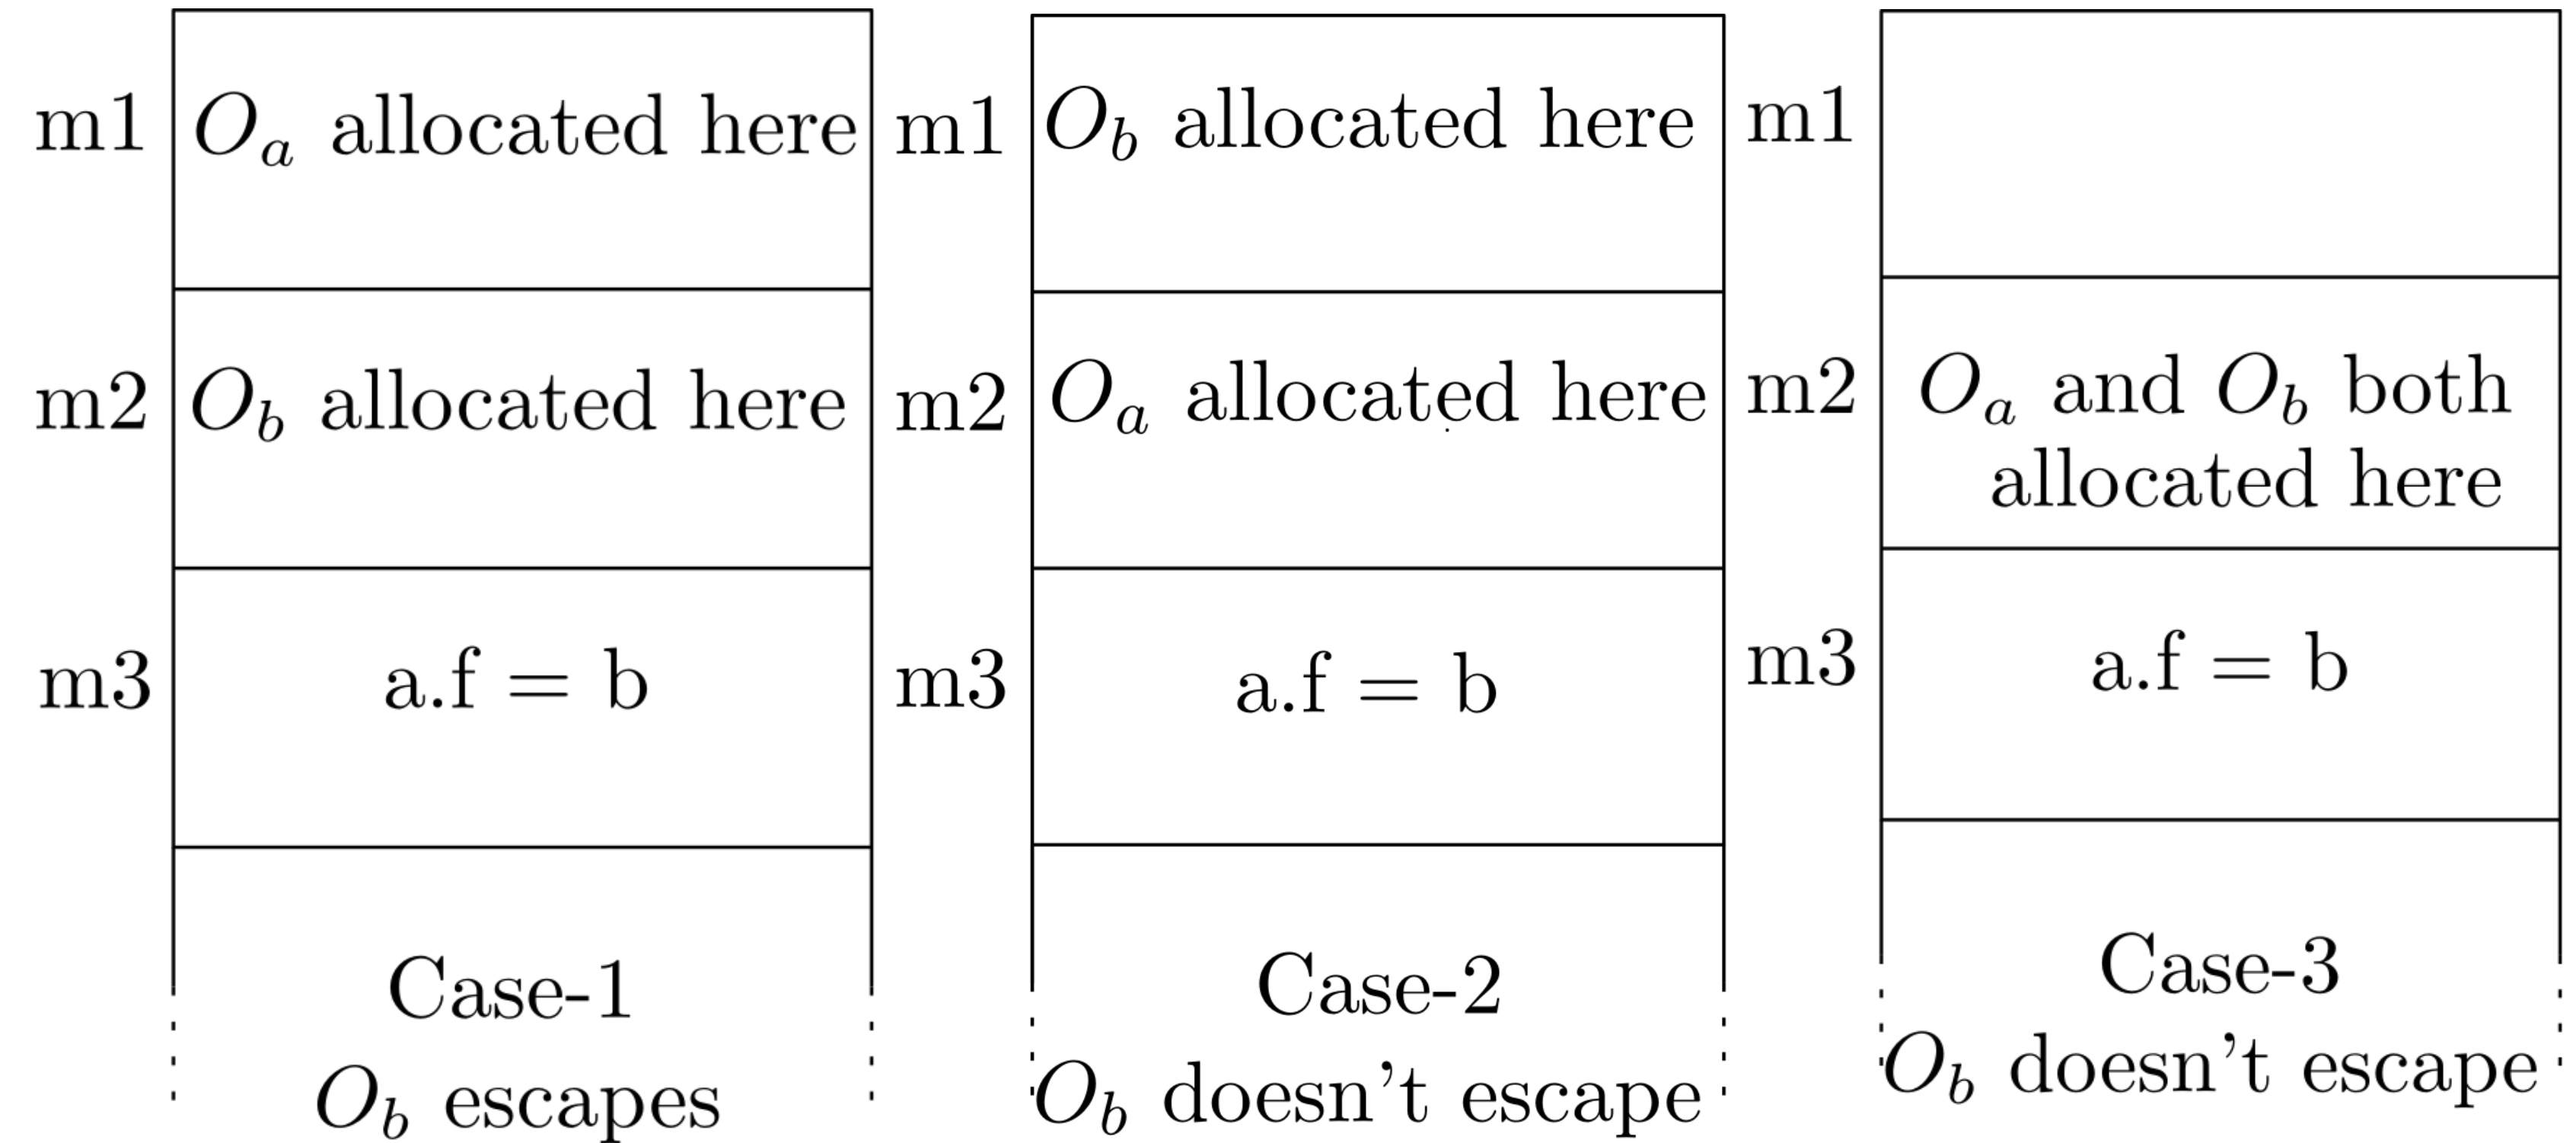
Scenarios at Store Statement

```
1. class T {  
2.   T f;  
3.   void m1() {m2(. . .);} }  
4.   void m2() {m3(. . .);} }  
5.   void m3(T a, T b) {  
6.     a.f = b;  
7.   } /* method m3 */  
8.} /* class T */
```



Scenarios at Store Statement

```
1. class T {  
2.   T f;  
3.   void m1() {m2(. . .);} }  
4.   void m2() {m3(. . .);} }  
5.   void m3(T a, T b) {  
6.     a.f = b;  
7.   } /* method m3 */  
8.} /* class T */
```



Stack Walk — Costly



Ordering Objects on Stack

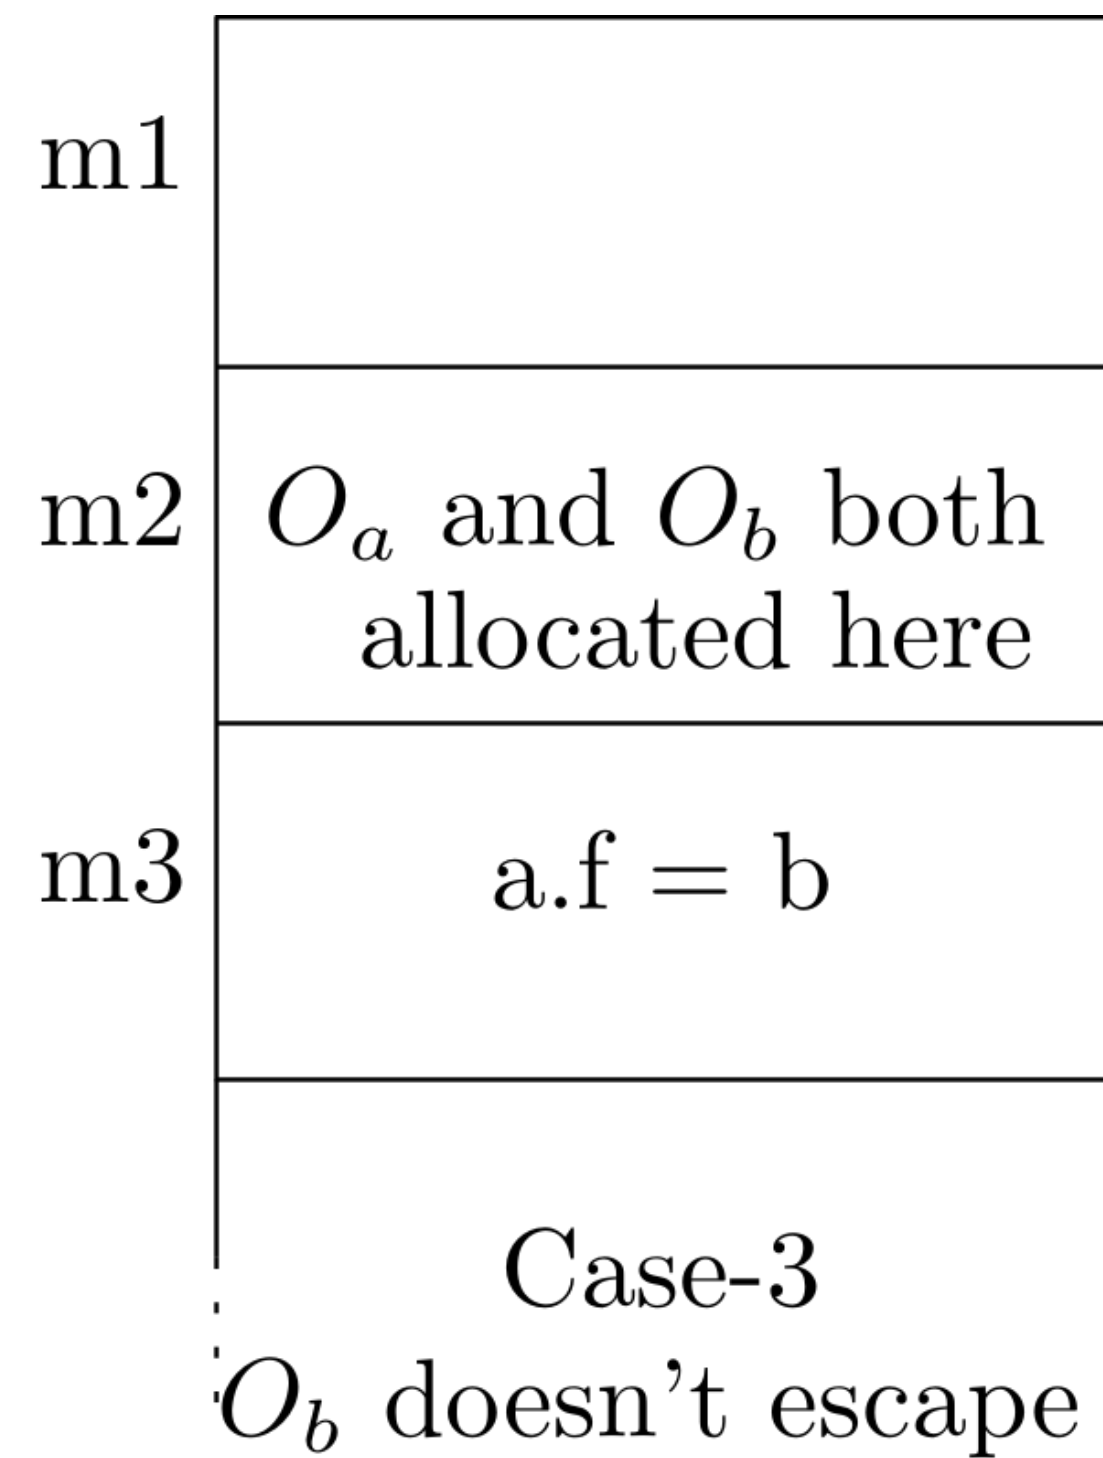
Ordering Objects on Stack

Ordering Objects on Stack

- A simple address-comparison check works majority of times.

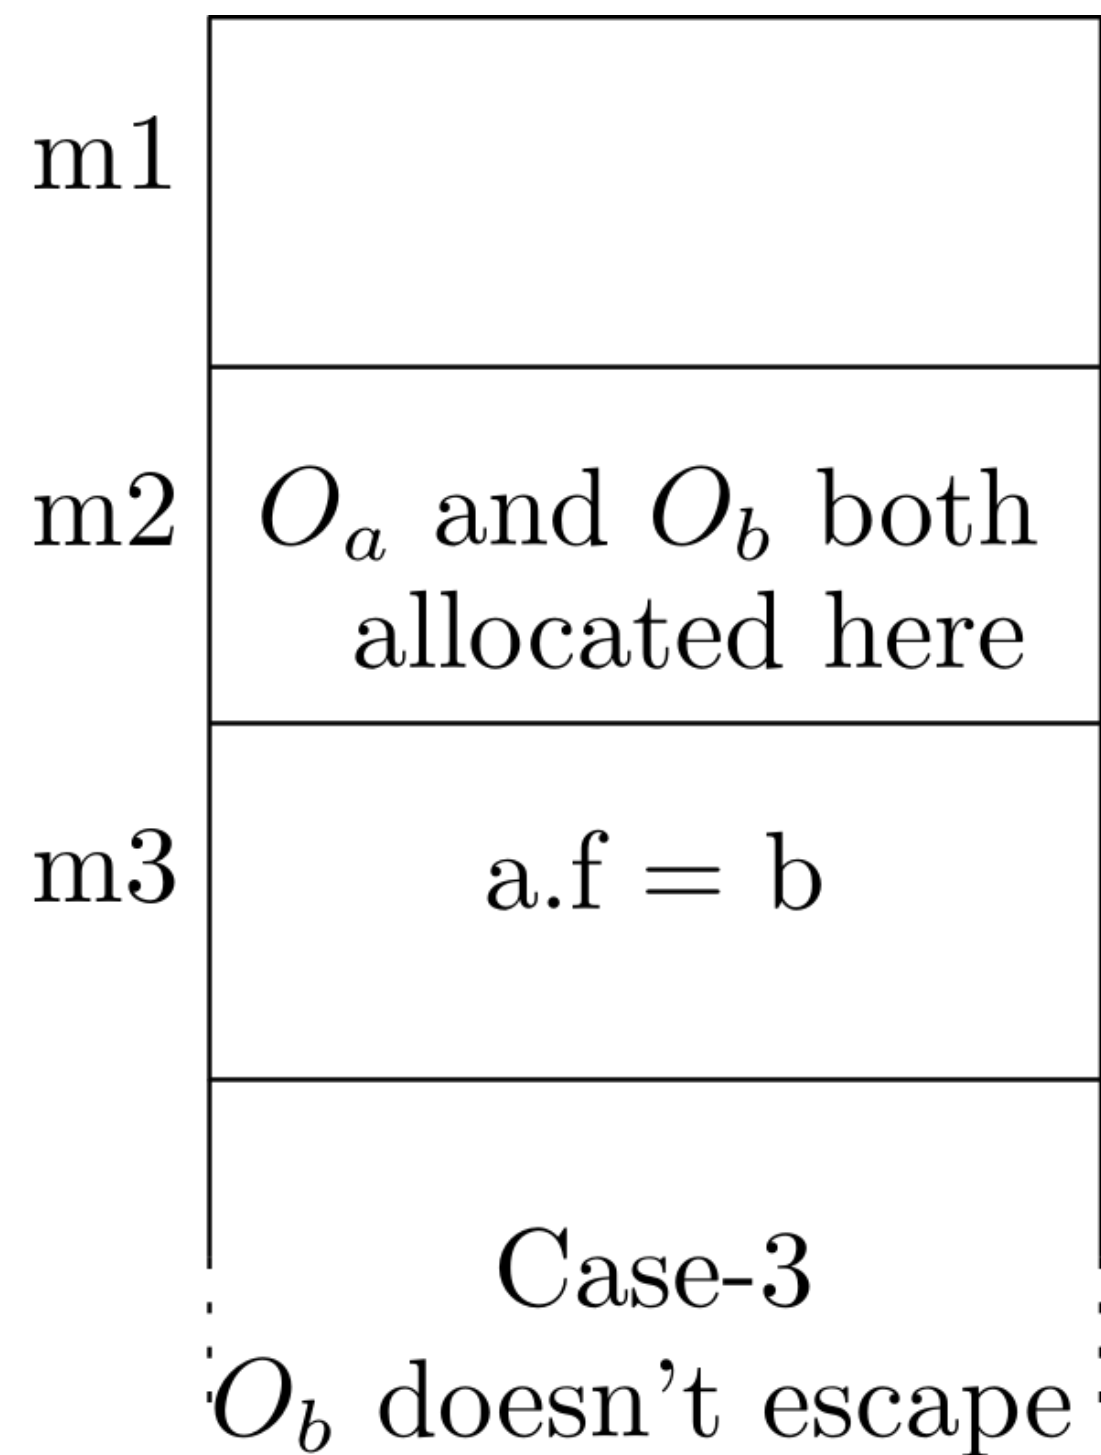
Ordering Objects on Stack

- A simple address-comparison check works majority of times.



Ordering Objects on Stack

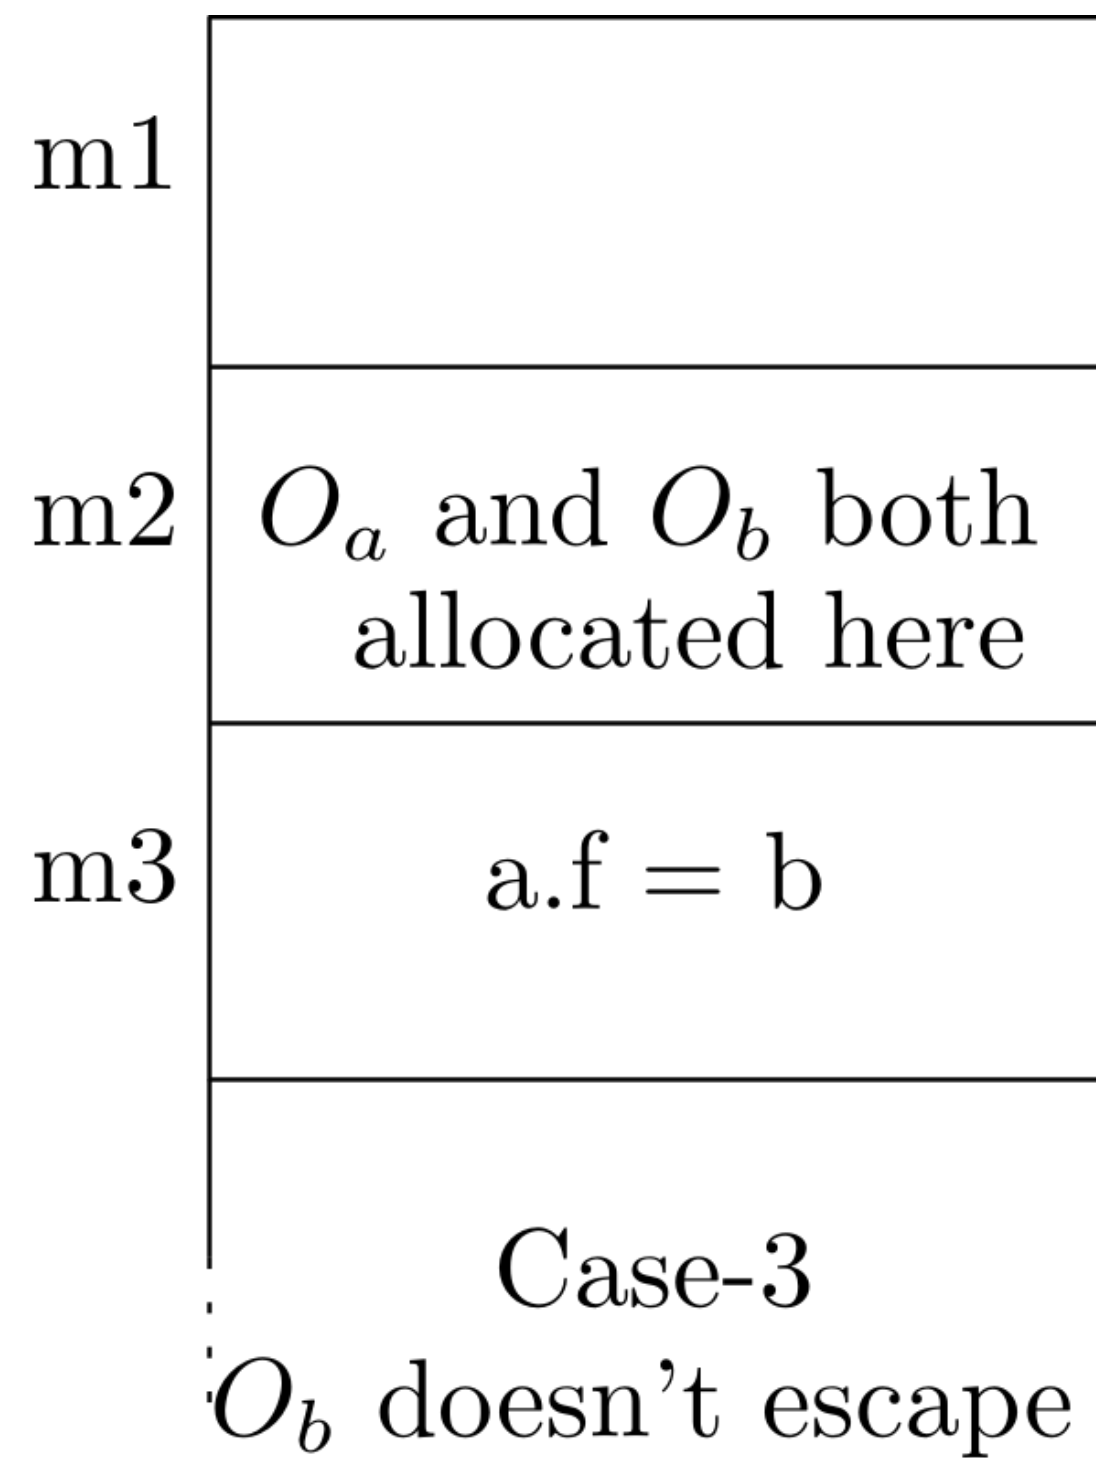
- A simple address-comparison check works majority of times.



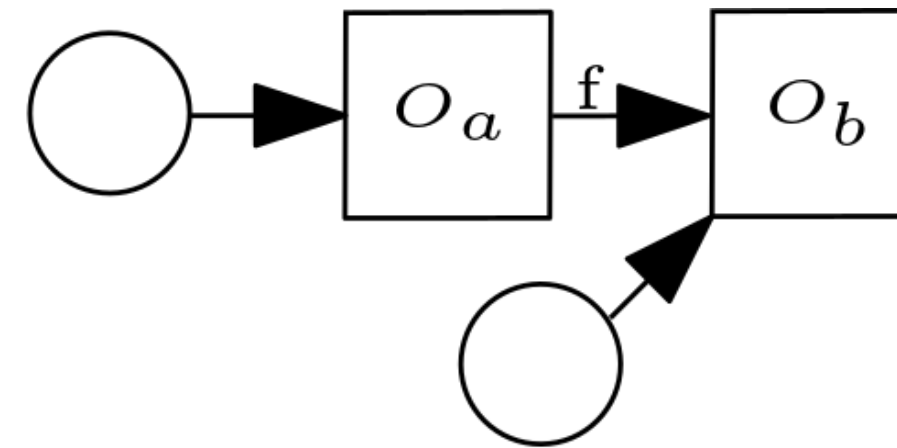
- Statically create a partial order of stack-allocatable objects.

Ordering Objects on Stack

- A simple address-comparison check works majority of times.

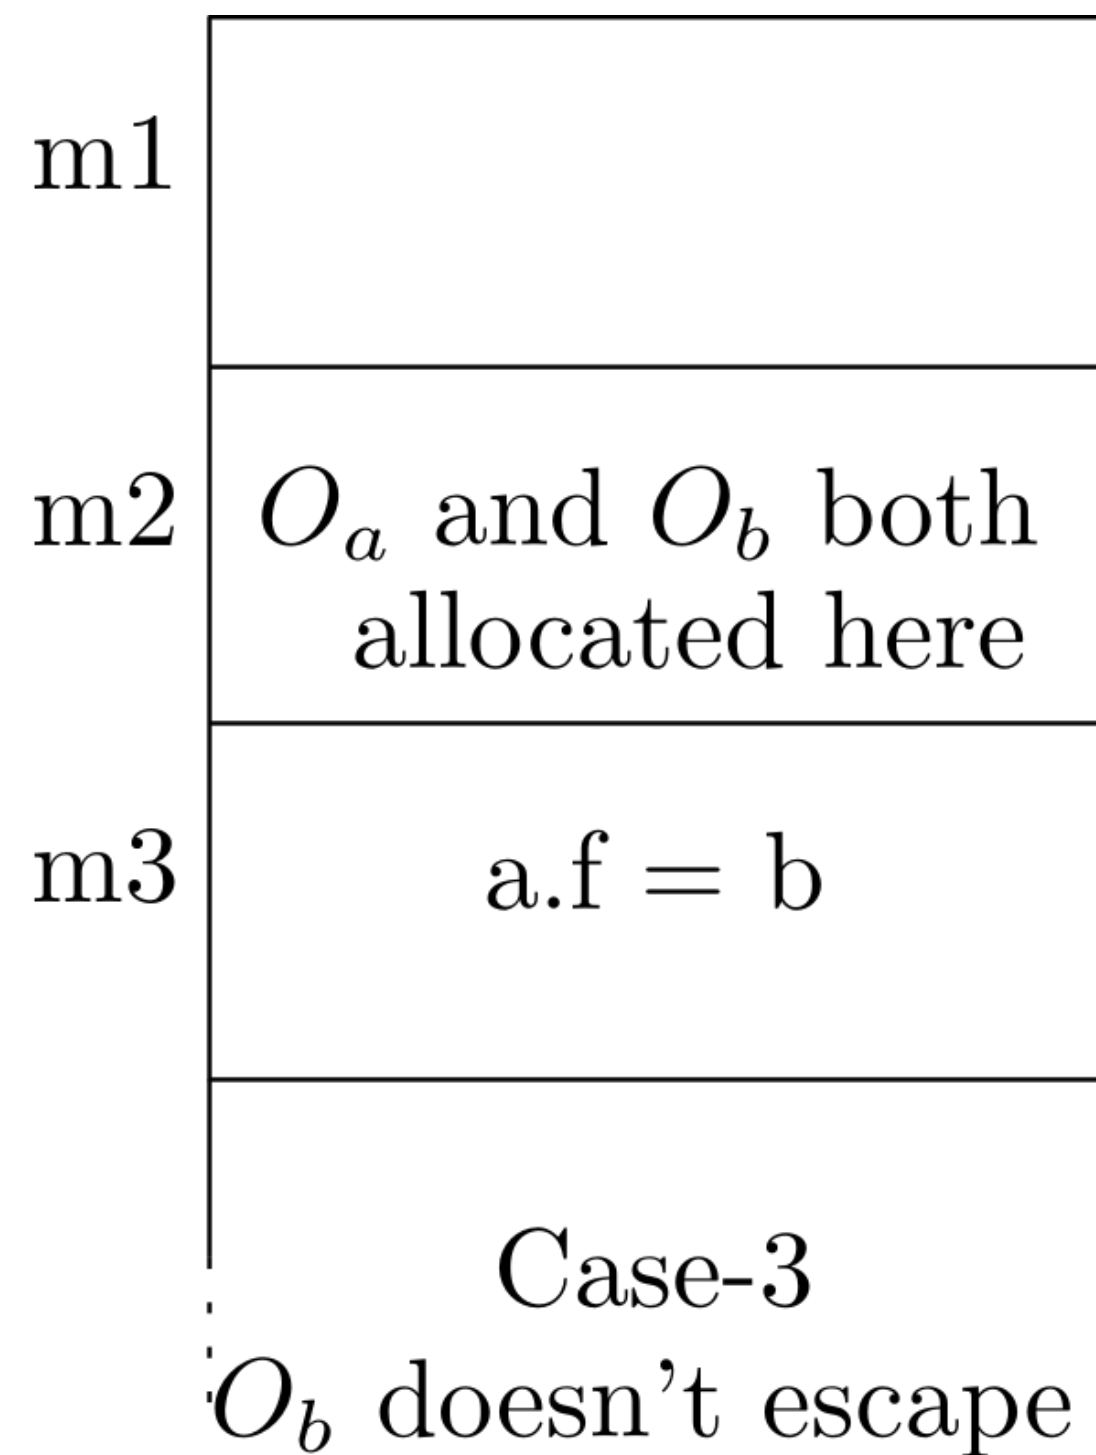


- Statically create a partial order of stack-allocatable objects.

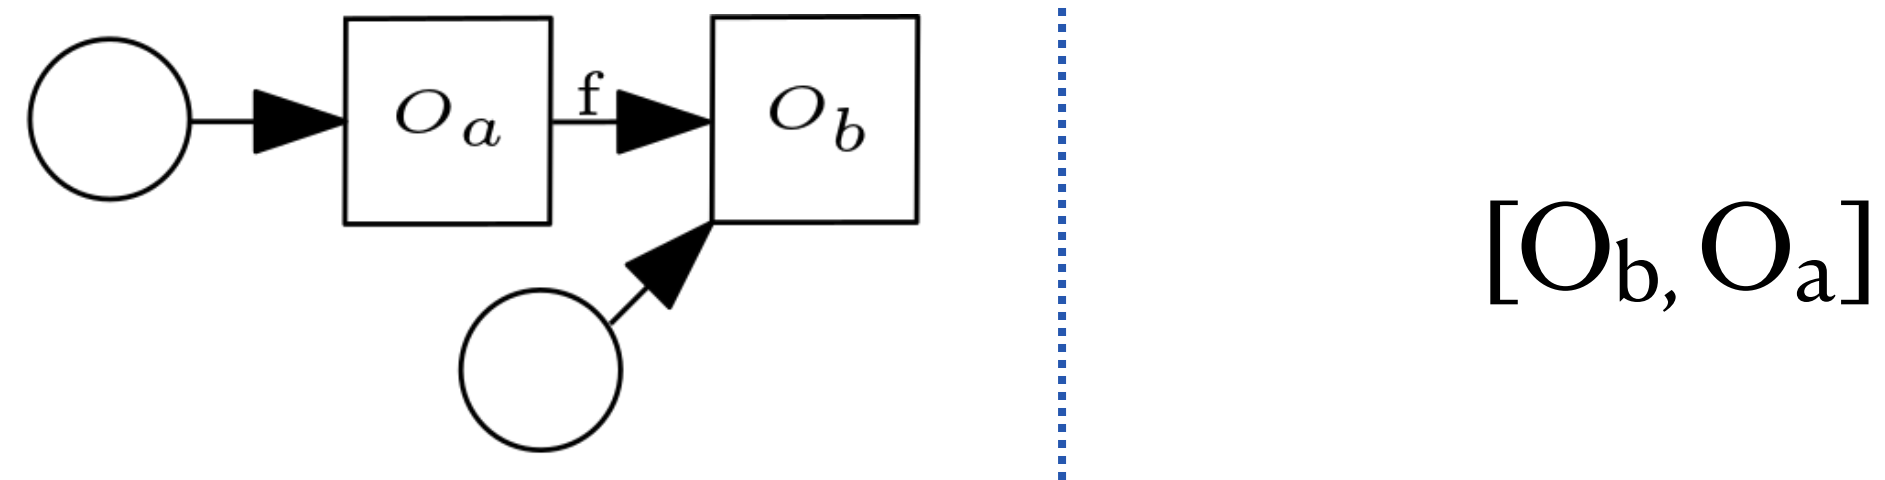


Ordering Objects on Stack

- A simple address-comparison check works majority of times.

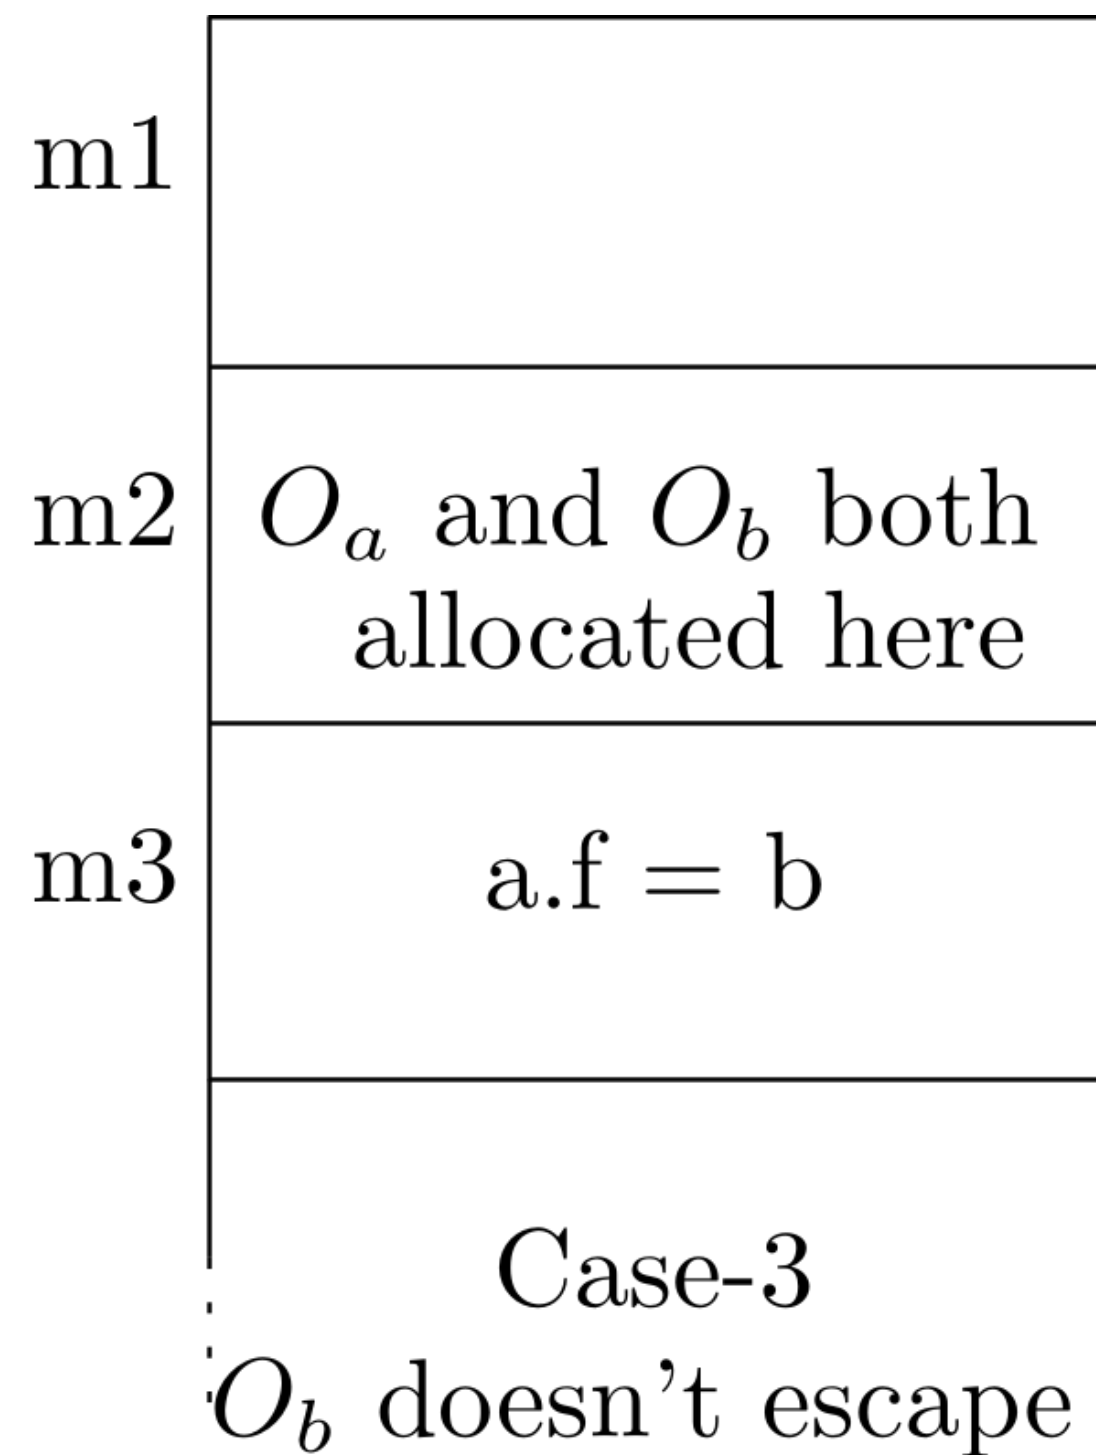


- Statically create a partial order of stack-allocatable objects.

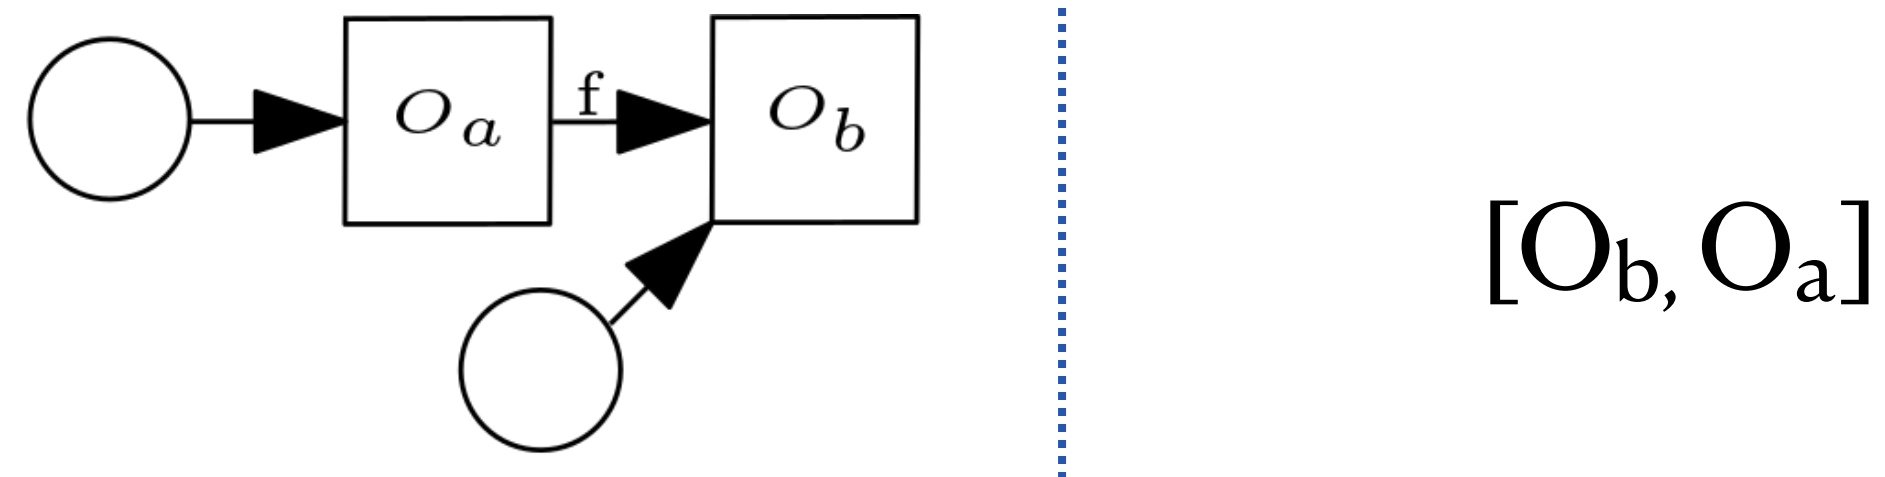


Ordering Objects on Stack

- A simple address-comparison check works majority of times.



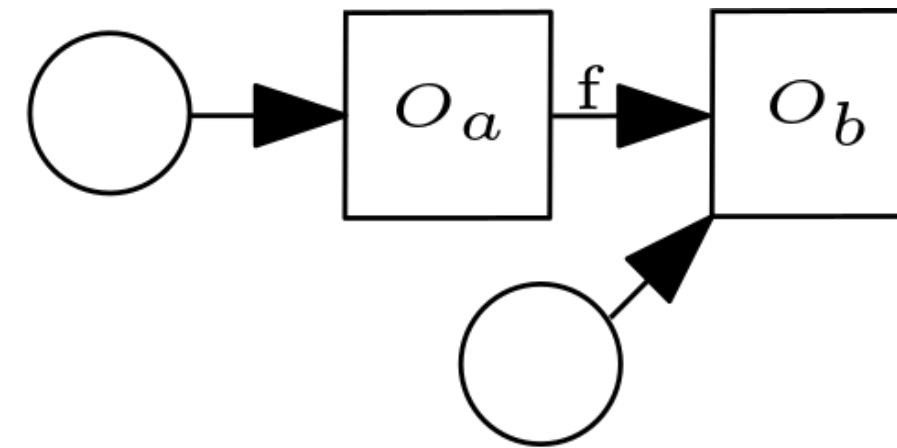
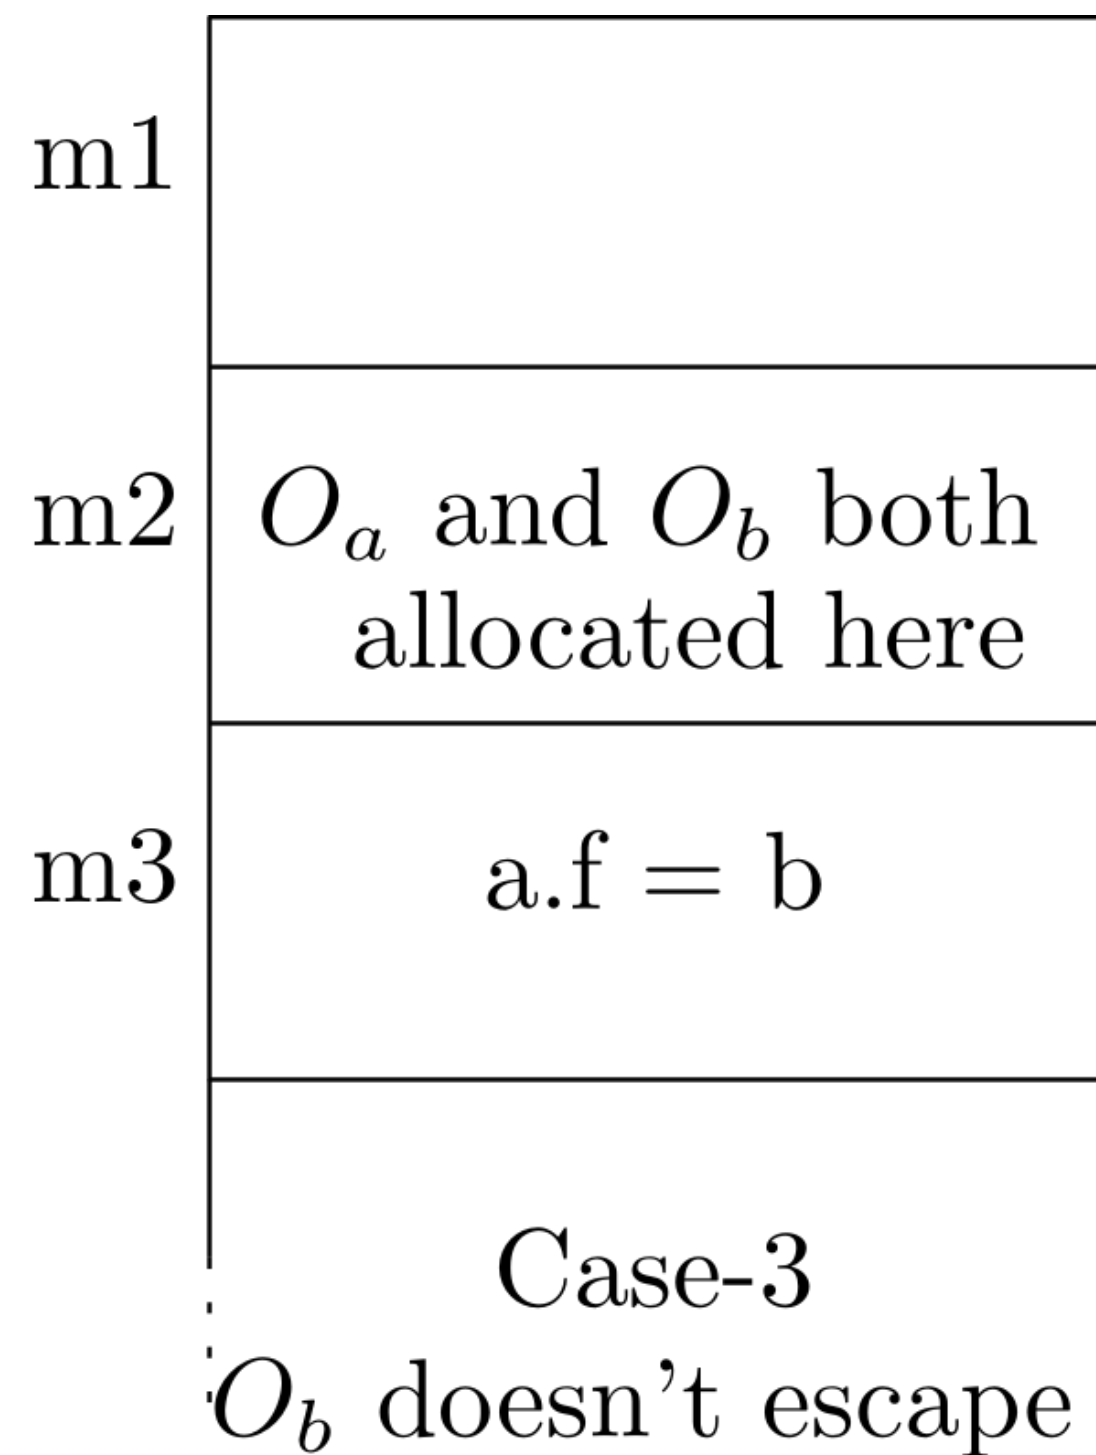
- Statically create a partial order of stack-allocatable objects.



- Use the stack-order in VM to re-order the list of stack allocated objects.

Ordering Objects on Stack

- A simple address-comparison check works majority of times.

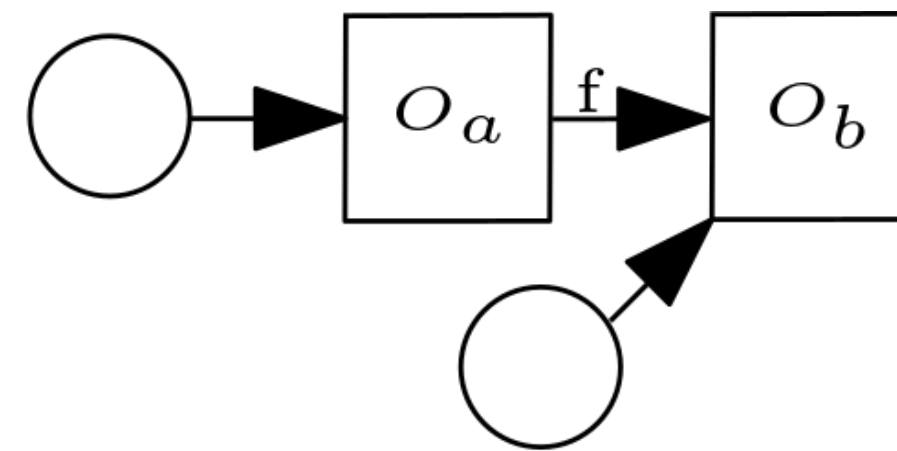
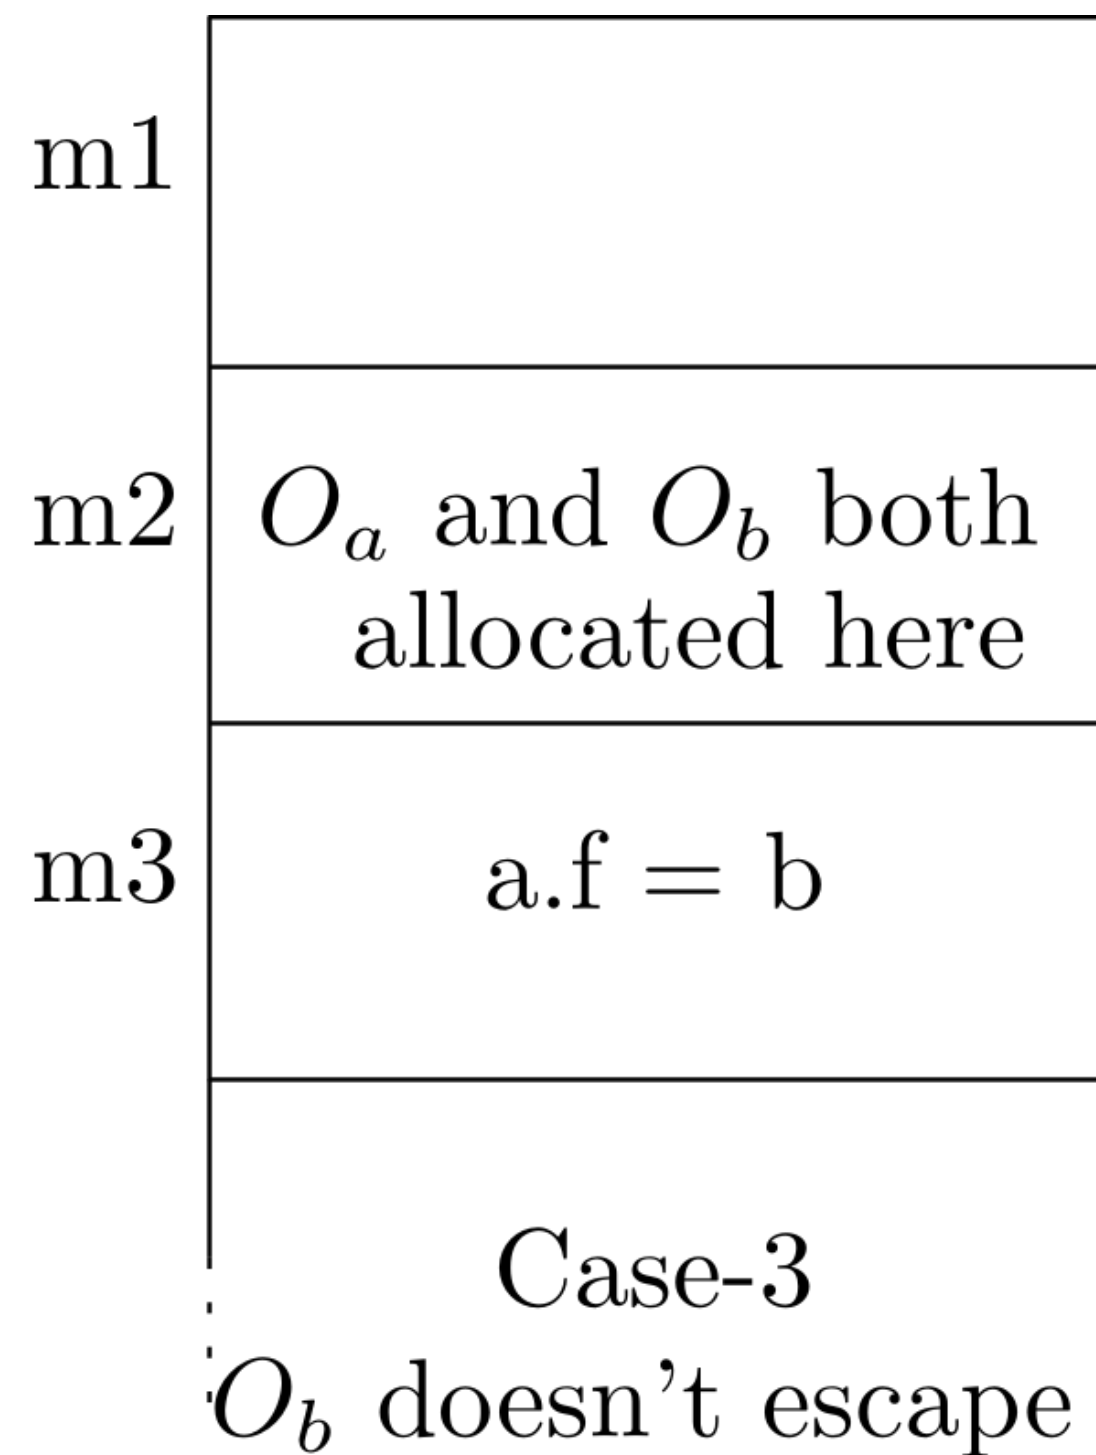


$[O_b, O_a]$

- Statically create a partial order of stack-allocatable objects.
- Use the stack-order in VM to re-order the list of stack allocated objects.
- Reduces cost of heapification checks.

Ordering Objects on Stack

- A simple address-comparison check works majority of times.

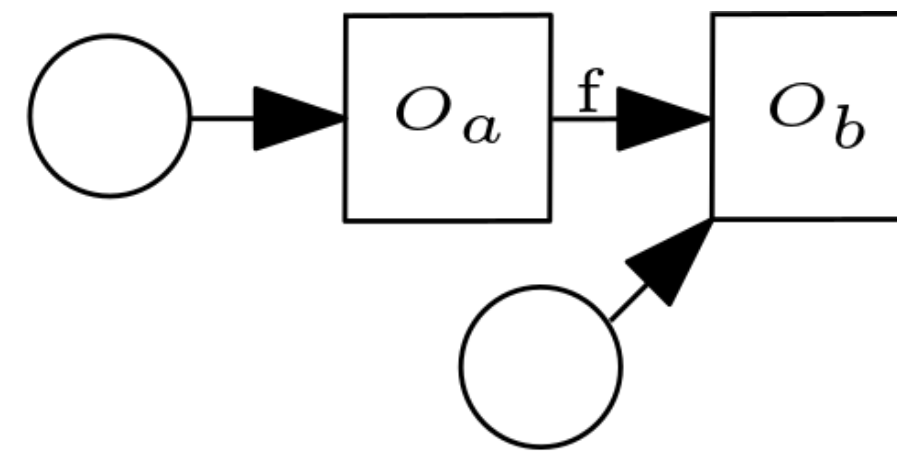
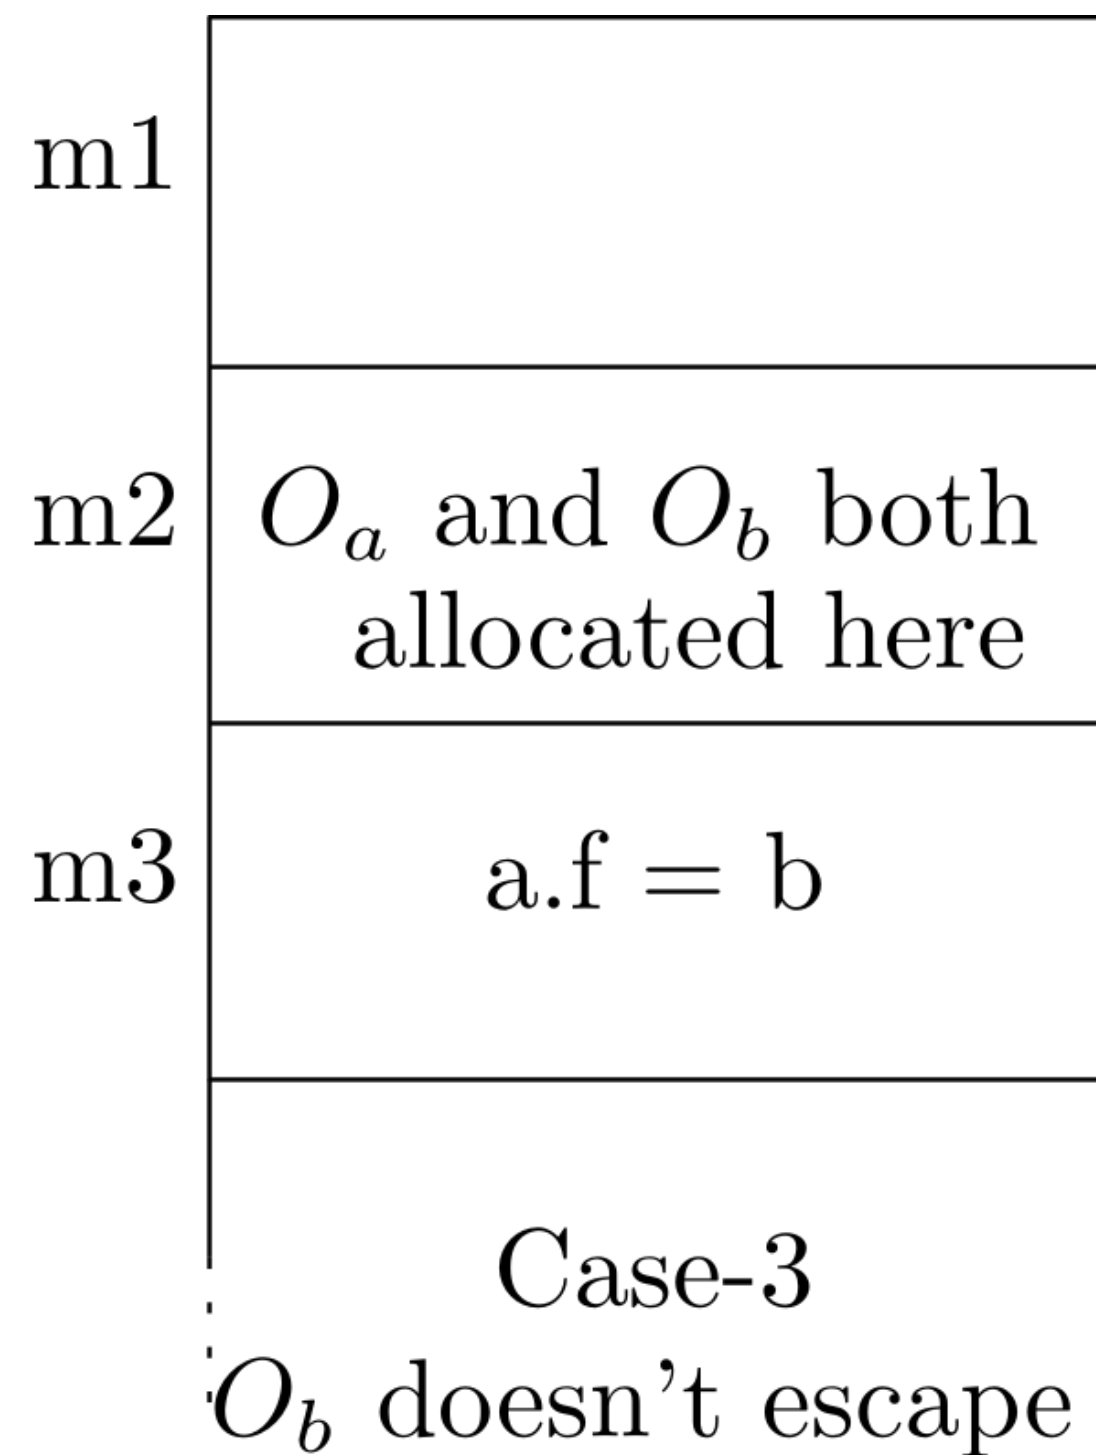


$[O_b, O_a]$

- Statically create a partial order of stack-allocatable objects.
- Use the stack-order in VM to re-order the list of stack allocated objects.
- Reduces cost of heapification checks.
- In case of cycles — result will not be valid only for one store statement.

Ordering Objects on Stack

- A simple address-comparison check works majority of times.



$[O_b, O_a]$

- Statically create a partial order of stack-allocatable objects.
- Use the stack-order in VM to re-order the list of stack allocated objects.
- Reduces cost of heapification checks.
- In case of cycles — result will not be valid only for one store statement. **Stack Walk**

Implementation and Evaluation

Implementation and Evaluation

- **Implementation:**

- Static analysis: Soot
- Runtime components: OpenJ9 VM

Implementation and Evaluation

- **Implementation:**

- Static analysis: Soot
- Runtime components: OpenJ9 VM

- **Benchmarks:**

- DaCapo suites 23.10-chopin and 9.12 MRI.
- SPECjvm 2008.

Implementation and Evaluation

- Implementation:

- Static analysis: Soot
- Runtime components: OpenJ9 VM

- Benchmarks:

- DaCapo suites 23.10-chopin and 9.12 MRI.
- SPECjvm 2008.

- Evaluation schemes:

- **BASE**: Stack allocation with the existing scheme.
- **OPT**: Stack allocation with our optimistic scheme.

Implementation and Evaluation

- **Implementation:**

- Static analysis: Soot
- Runtime components: OpenJ9 VM

- **Benchmarks:**

- DaCapo suites 23.10-chopin and 9.12 MRI.
- SPECjvm 2008.

- **Evaluation schemes:**

- **BASE:** Stack allocation with the existing scheme.
- **OPT:** Stack allocation with our optimistic scheme.

- **Compute:**

- Enhancement in stack allocation.
- Impact on performance and garbage collection.

Evaluation (Stack Allocation)

	Non Optimistic Scheme (BASE)			Optimistic Scheme (OPT)		
Benchmark	Static Count	Dynamic Count	Stack Bytes	Static Count	Dynamic Count	Stack Bytes
graphchi	0 (0.0 %)	0M (0.00%)	0MB	32 (4.15%)	506.3M (6.9%)	9184.6MB
fop	10 (0.15%)	0.04M (0.002%)	1MB	50 (0.77%)	9.8M (0.42%)	161.2MB
h2	61 (2.33%)	29M (0.92%)	523MB	94 (3.87%)	452M (13.92%)	10801MB
luindex	35 (1.35%)	3M (2.39%)	98MB	89 (3.49%)	5M (3.49%)	133MB
lusearch	30 (1.09%)	25M (3.23%)	775MB	78 (3.05%)	59M (7.4%)	1686MB
pmd	89 (1.09%)	52M (7.20%)	1310MB	191 (3.97%)	105M (14.2%)	2465MB
compiler	93 (1.73%)	94M (5.50%)	1720MB	137 (2.75%)	105M (6.17%)	2329MB
rsa	16 (1.13%)	0.1M (1.1%)	46MB	35 (3.18%)	7M (4.62%)	170MB
signverify	15 (0.84%)	0.24M (0.86%)	6.8MB	51 (3.10%)	2.1M (7.24%)	49.4MB

Evaluation (Stack Allocation)

	Non Optimistic Scheme (BASE)			Optimistic Scheme (OPT)		
Benchmark	Static Count	Dynamic Count	Stack Bytes	Static Count	Dynamic Count	Stack Bytes
graphchi	0 (0.0 %)	0M (0.00%)	0MB	32 (4.15%)	506.3M (6.9%)	9184.6MB
fop	10 (0.15%)	0.04M (0.002%)	1MB	50 (0.77%)	9.8M (0.42%)	161.2MB
h2	61 (2.33%)	29M (0.92%)	523MB	94 (3.87%)	452M (13.92%)	10801MB
luindex	35 (1.35%)	3M (2.39%)	98MB	89 (3.49%)	5M (3.49%)	133MB
lusearch	30 (1.09%)	25M (3.23%)	775MB	78 (3.05%)	59M (7.4%)	1686MB
pmd	89 (1.09%)	52M (7.20%)	1310MB	191 (3.97%)	105M (14.2%)	2465MB
compiler	93 (1.73%)	94M (5.50%)	1720MB	137 (2.75%)	105M (6.17%)	2329MB
rsa	16 (1.13%)	0.1M (1.1%)	46MB	35 (3.18%)	7M (4.62%)	170MB
signverify	15 (0.84%)	0.24M (0.86%)	6.8MB	51 (3.10%)	2.1M (7.24%)	49.4MB

Evaluation (Stack Allocation)

	Non Optimistic Scheme (BASE)			Optimistic Scheme (OPT)		
Benchmark	Static Count	Dynamic Count	Stack Bytes	Static Count	Dynamic Count	Stack Bytes
graphchi	0 (0.0 %)	0M (0.00%)	0MB	32 (4.15%)	506.3M (6.9%)	9184.6MB
fop	10 (0.15%)	0.04M (0.002%)	1MB	50 (0.77%)	9.8M (0.42%)	161.2MB
h2	61 (2.33%)	29M (0.92%)	523MB	94 (3.87%)	452M (13.92%)	10801MB
luindex	35 (1.35%)	3M (2.39%)	98MB	89 (3.49%)	5M (3.49%)	133MB
lusearch	30 (1.09%)	25M (3.23%)	775MB	78 (3.05%)	59M (7.4%)	1686MB
pmd	89 (1.09%)	52M (7.20%)	1310MB	191 (3.97%)	105M (14.2%)	2465MB
compiler	93 (1.73%)	94M (5.50%)	1720MB	137 (2.75%)	105M (6.17%)	2329MB
rsa	16 (1.13%)	0.1M (1.1%)	46MB	35 (3.18%)	7M (4.62%)	170MB
signverify	15 (0.84%)	0.24M (0.86%)	6.8MB	51 (3.10%)	2.1M (7.24%)	49.4MB

Evaluation (Stack Allocation)

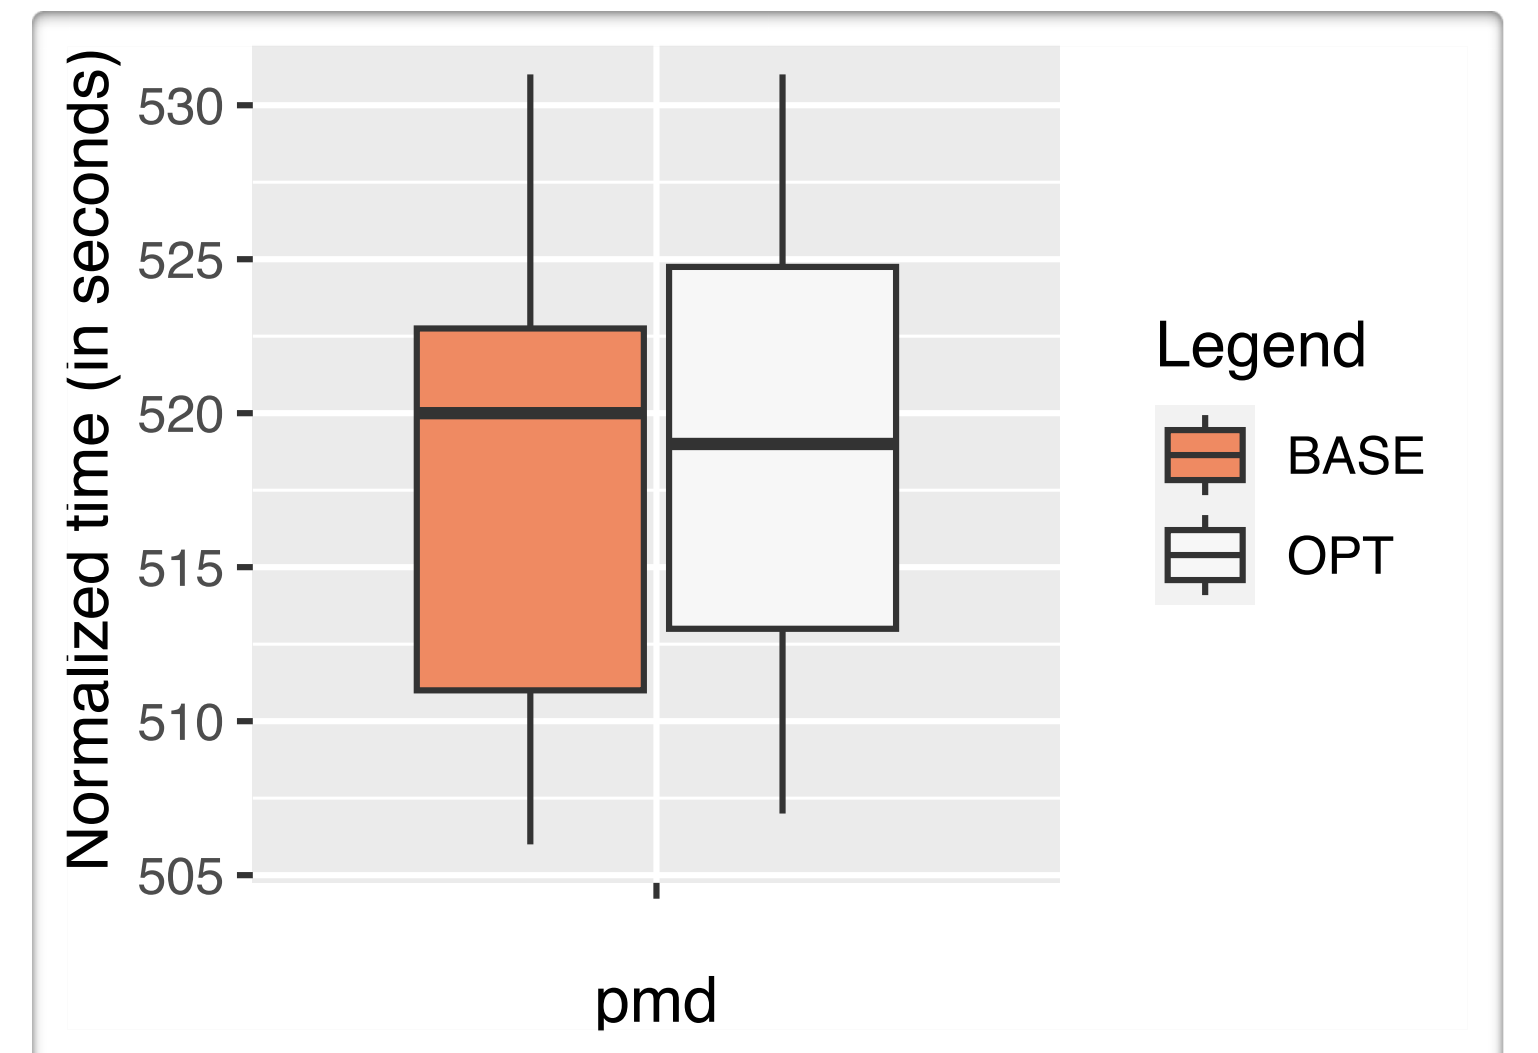
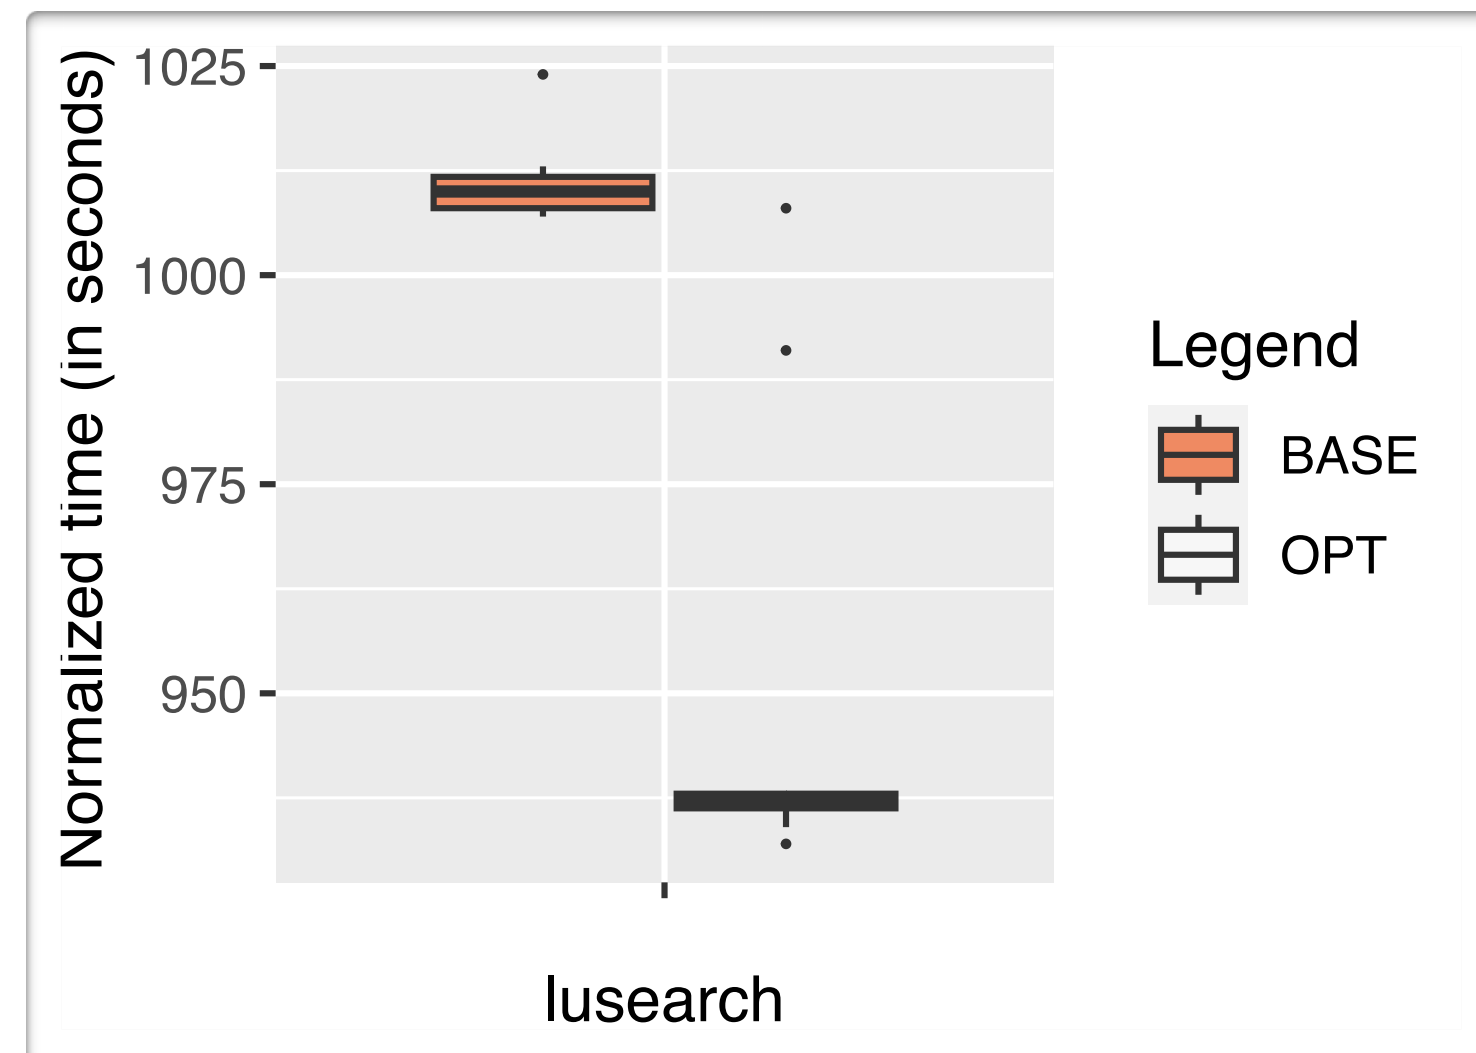
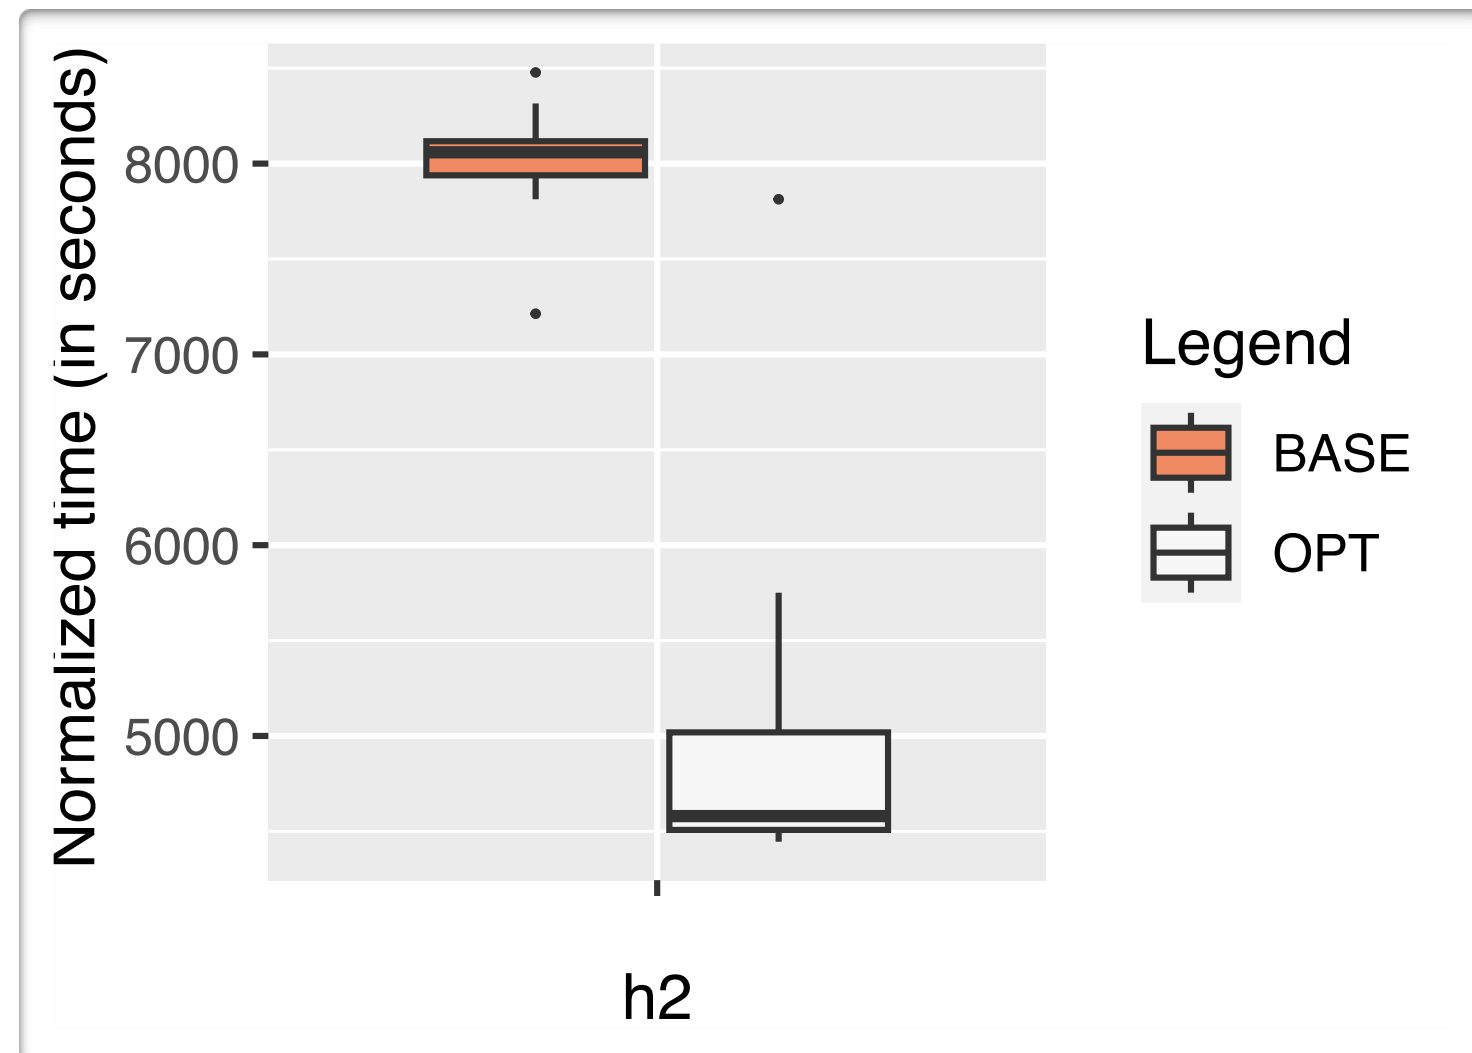
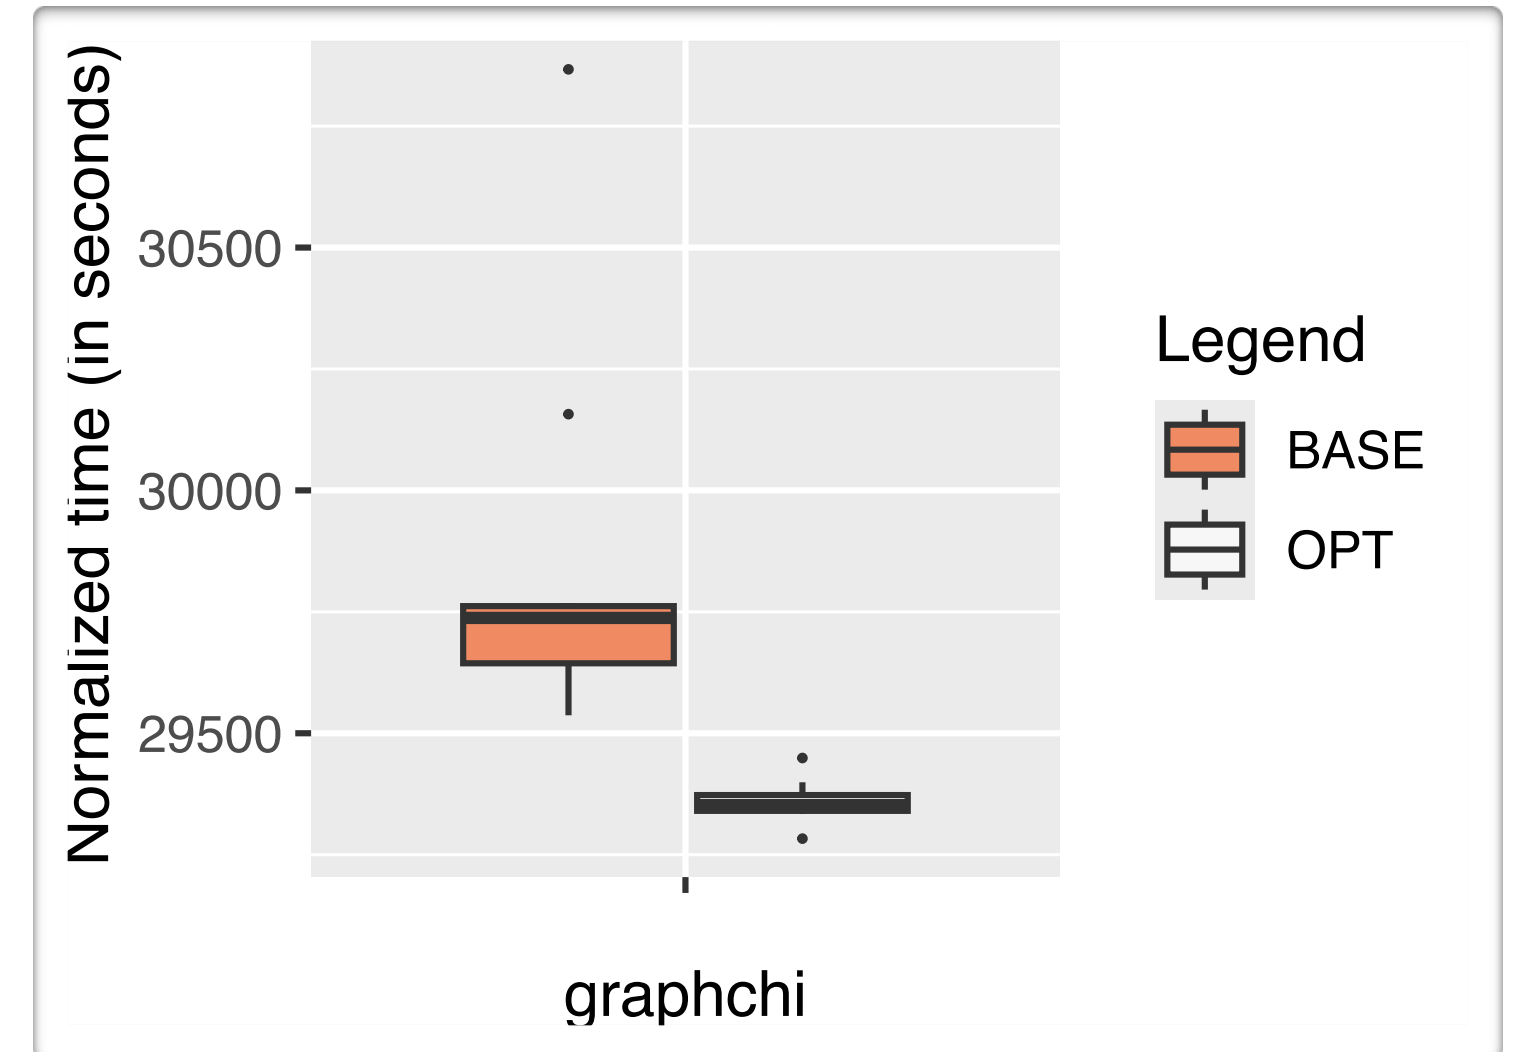
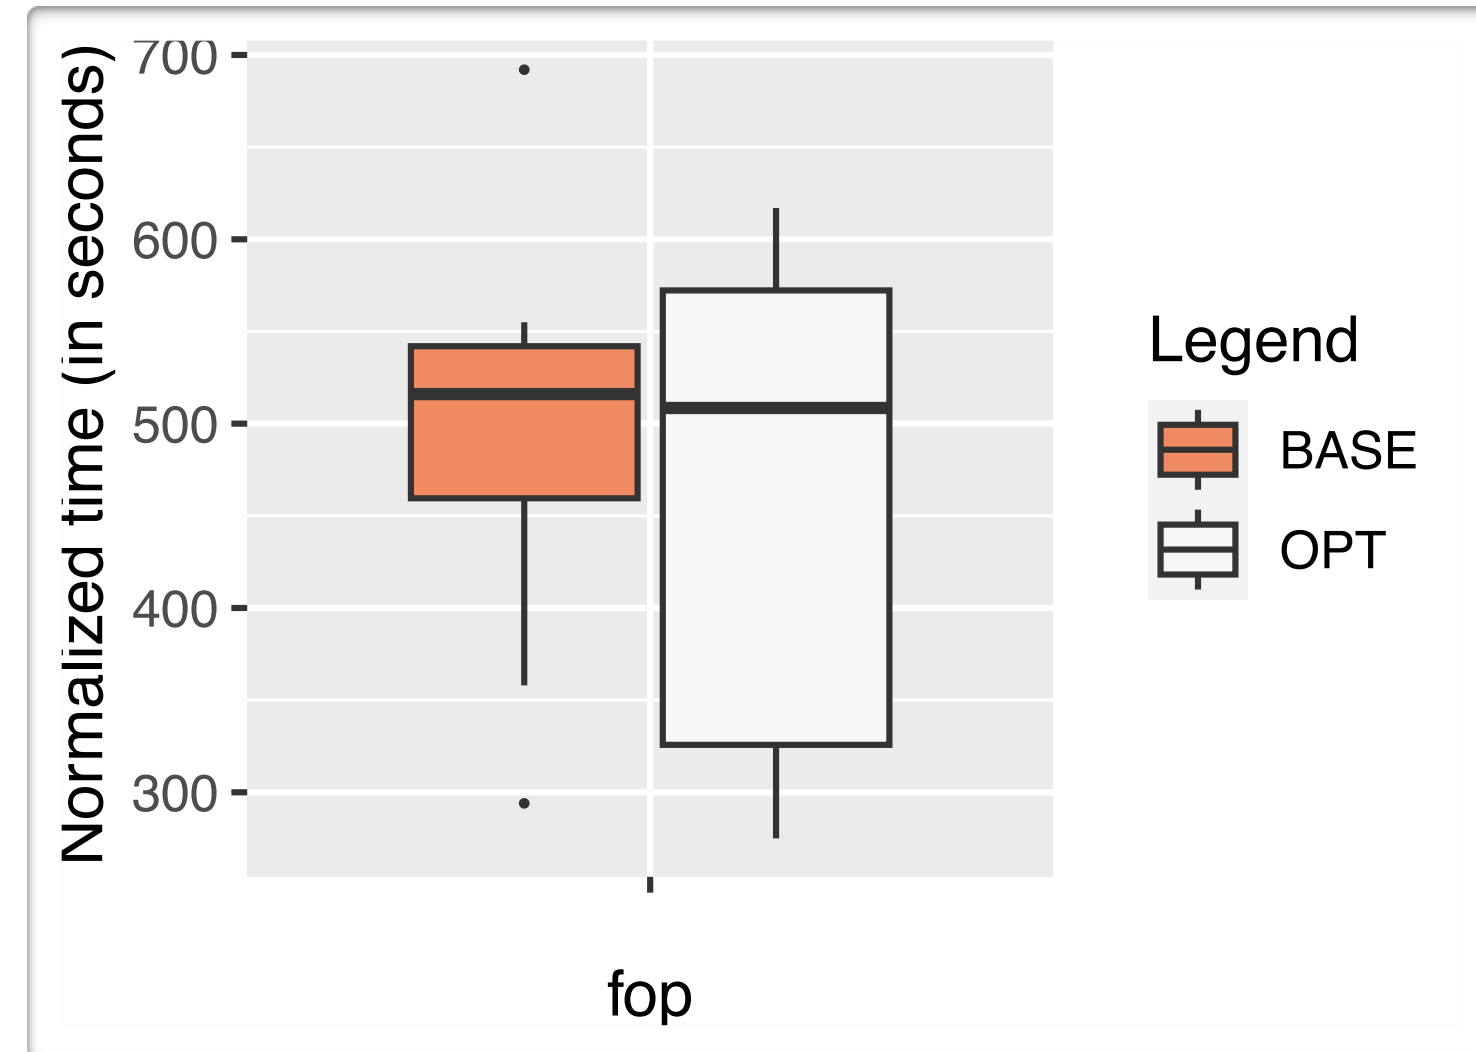
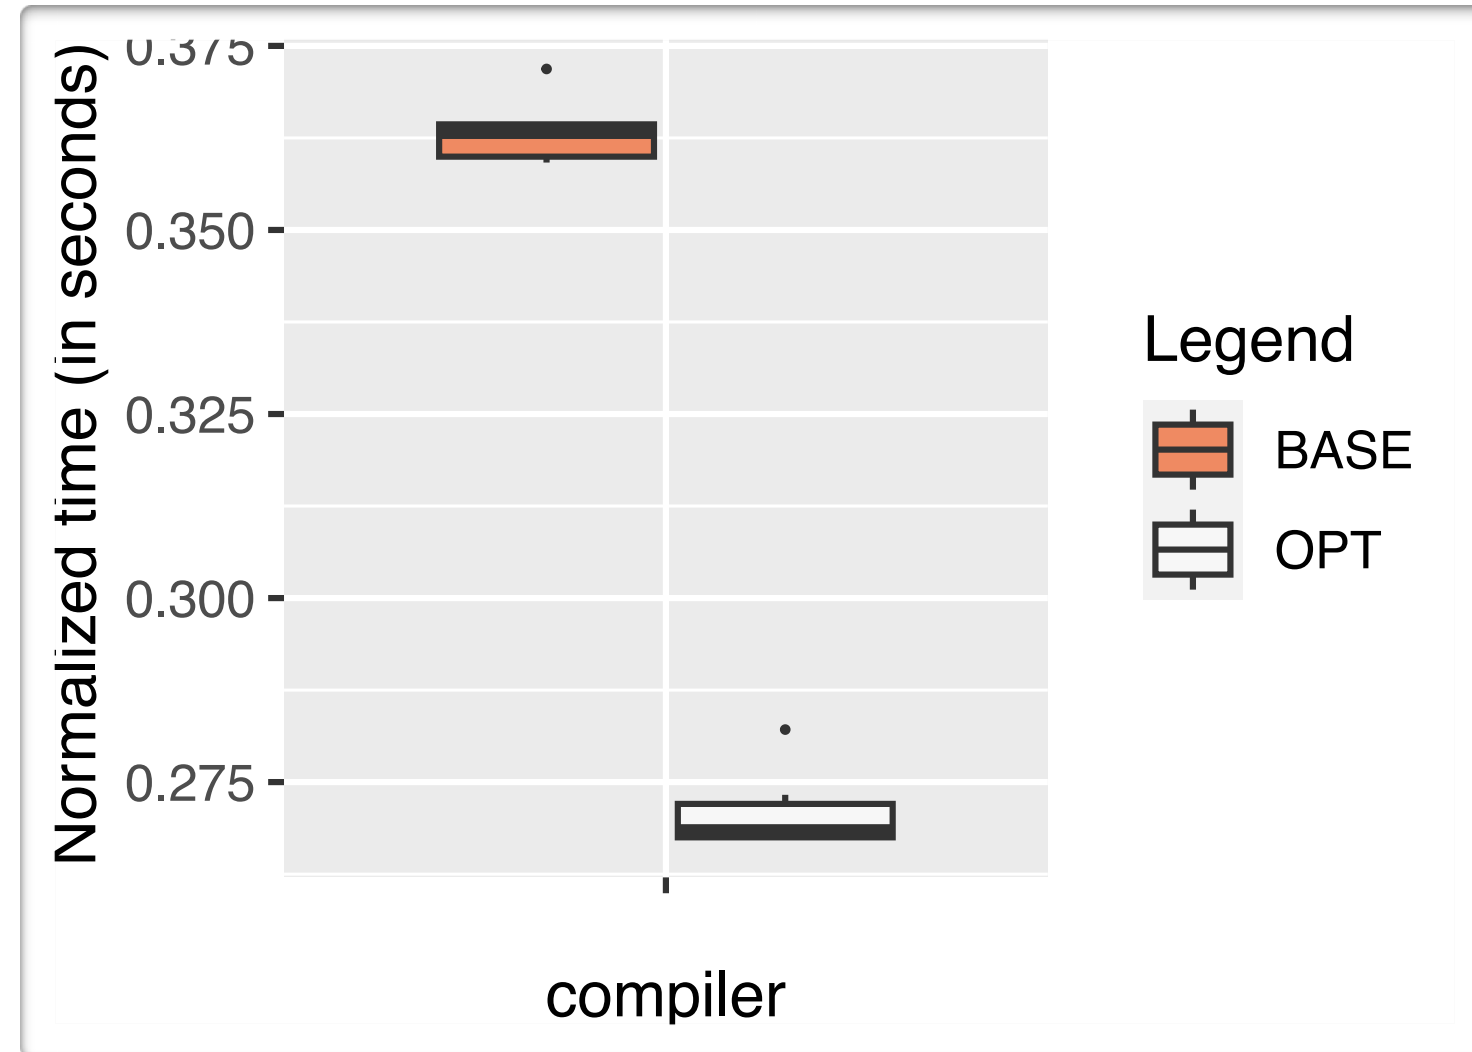
	Non Optimistic Scheme (BASE)			Optimistic Scheme (OPT)		
Benchmark	Static Count	Dynamic Count	Stack Bytes	Static Count	Dynamic Count	Stack Bytes
graphchi	0 (0.0 %)	0M (0.00%)	0MB	32 (4.15%)	506.3M (6.9%)	9184.6MB
fop	10 (0.15%)	0.04M (0.002%)	1MB	50 (0.77%)	9.8M (0.42%)	161.2MB
h2	61 (2.33%)	29M (0.92%)	523MB	94 (3.87%)	452M (13.92%)	10801MB
luindex	35 (1.35%)	3M (2.39%)	98MB	89 (3.49%)	5M (3.49%)	133MB
lusearch	30 (1.09%)	25M (3.23%)	775MB	78 (3.05%)	59M (7.4%)	1686MB
pmd	89 (1.09%)	52M (7.20%)	1310MB	191 (3.97%)	105M (14.2%)	2465MB
compiler	93 (1.73%)	94M (5.50%)	1720MB	137 (2.75%)	105M (6.17%)	2329MB
rsa	16 (1.13%)	0.1M (1.1%)	46MB	35 (3.18%)	7M (4.62%)	170MB
signverify	15 (0.84%)	0.24M (0.86%)	6.8MB	51 (3.10%)	2.1M (7.24%)	49.4MB

Evaluation (Stack Allocation)

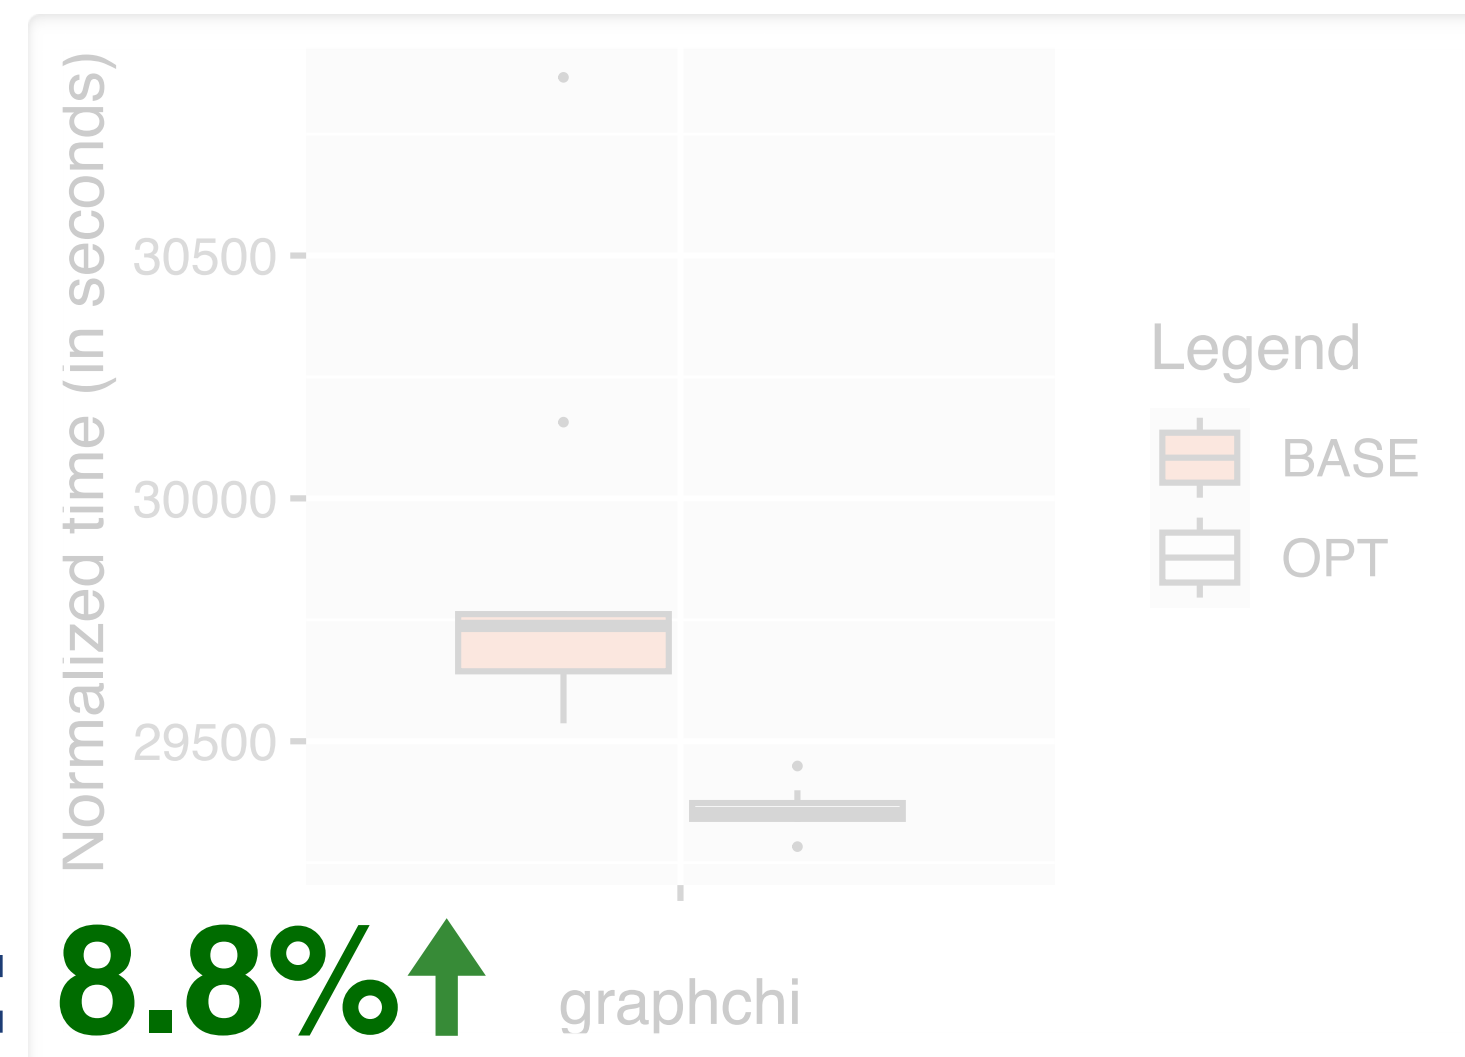
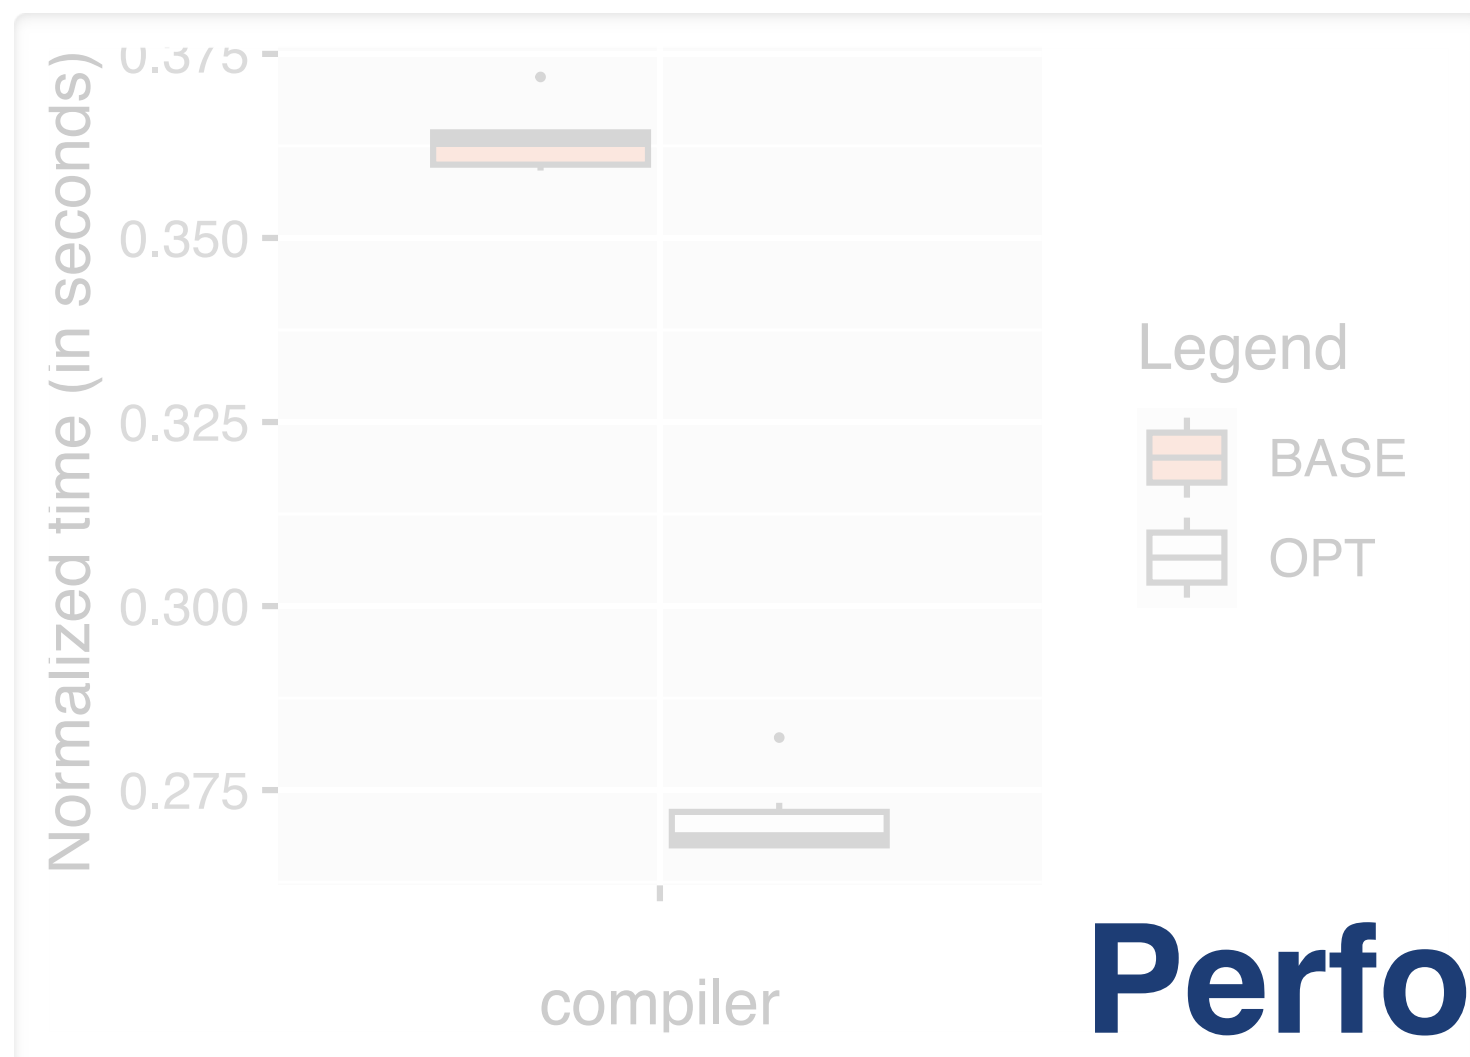
	Non Optimistic Scheme (BASE)			Optimistic Scheme (OPT)		
Benchmark	Static Count	Dynamic Count	Stack Bytes	Static Count	Dynamic Count	Stack Bytes
graphchi	0 (0.0 %)	0M (0.00%)	0MB	32 (4.15%)	506.3M (6.9%)	9184.6MB
fop	10 (0.15%)	0.04M (0.002%)	1MB	50 (0.77%)	9.8M (0.42%)	161.2MB
h2	61 (2.33%)	29M (0.92%)	523MB	94 (3.87%)	452M (13.92%)	10801MB
luindex	35 (1.35%)	3M (2.39%)	98MB	89 (3.49%)	5M (3.49%)	133MB
lusearch	30 (1.09%)	25M (3.23%)	775MB	78 (3.05%)	59M (7.4%)	1686MB
pmd	89 (1.09%)	52M (7.20%)	1310MB	191 (3.97%)	105M (14.2%)	2465MB
compiler	93 (1.73%)	94M (5.50%)	1720MB	137 (2.75%)	105M (6.17%)	2329MB
rsa	16 (1.13%)	0.1M (1.1%)	46MB	35 (3.18%)	7M (4.62%)	170MB
signverify	15 (0.84%)	0.24M (0.86%)	6.8MB	51 (3.10%)	2.1M (7.24%)	49.4MB

Stack Allocation: 71%↑ Stack Bytes: 54%↑
(Less Heap Allocation)

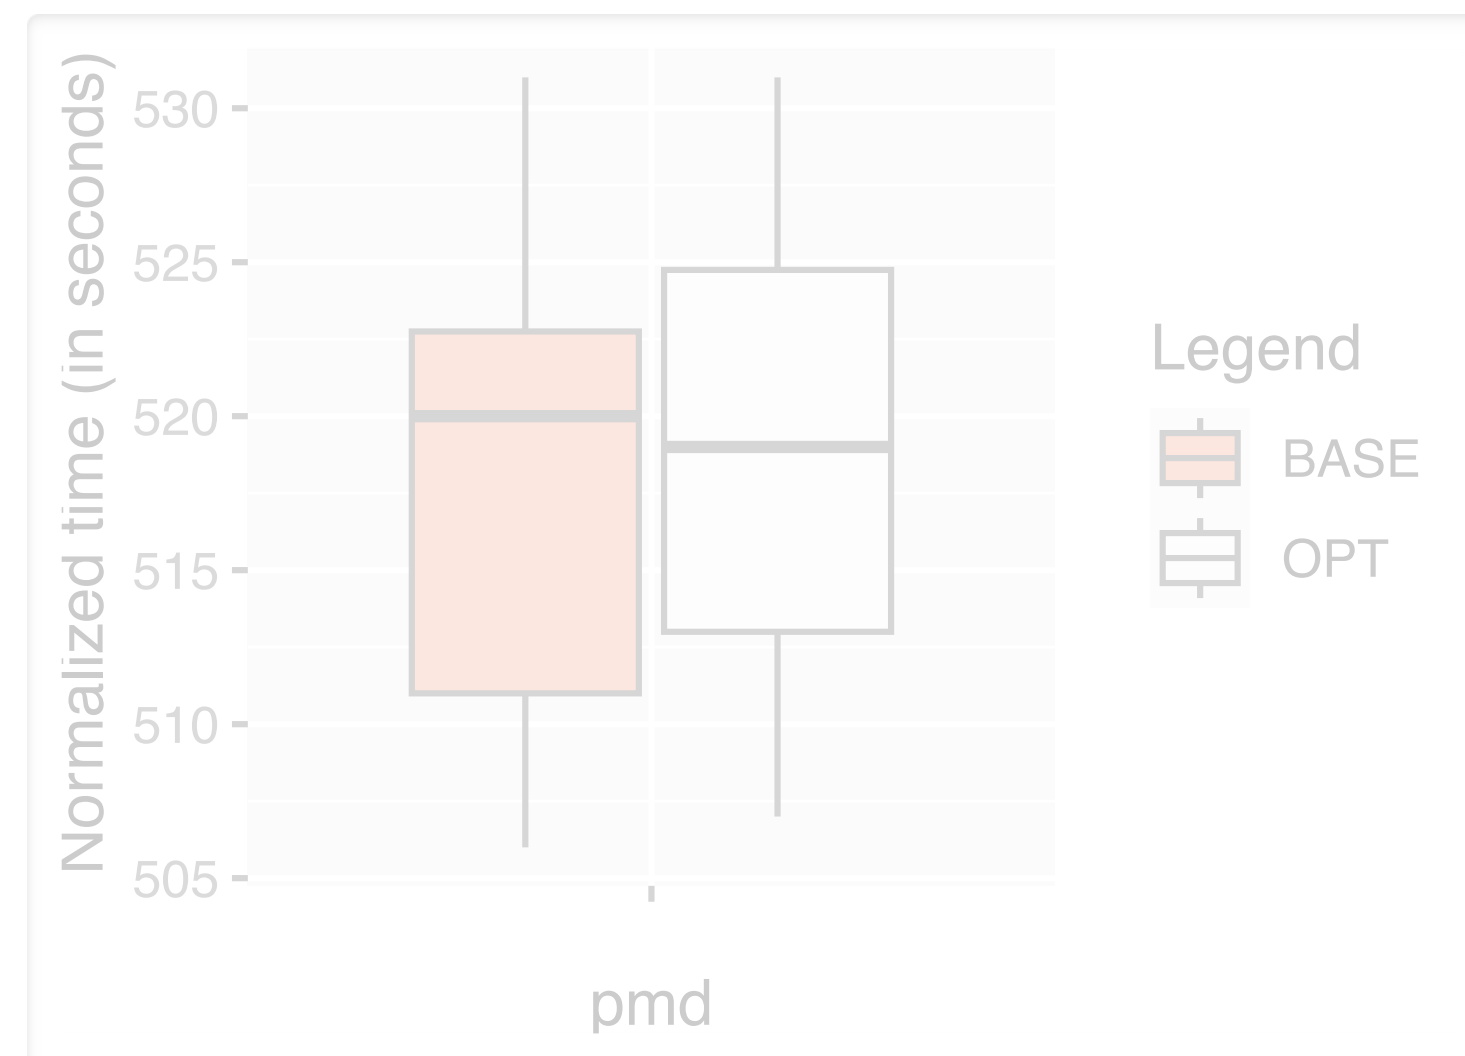
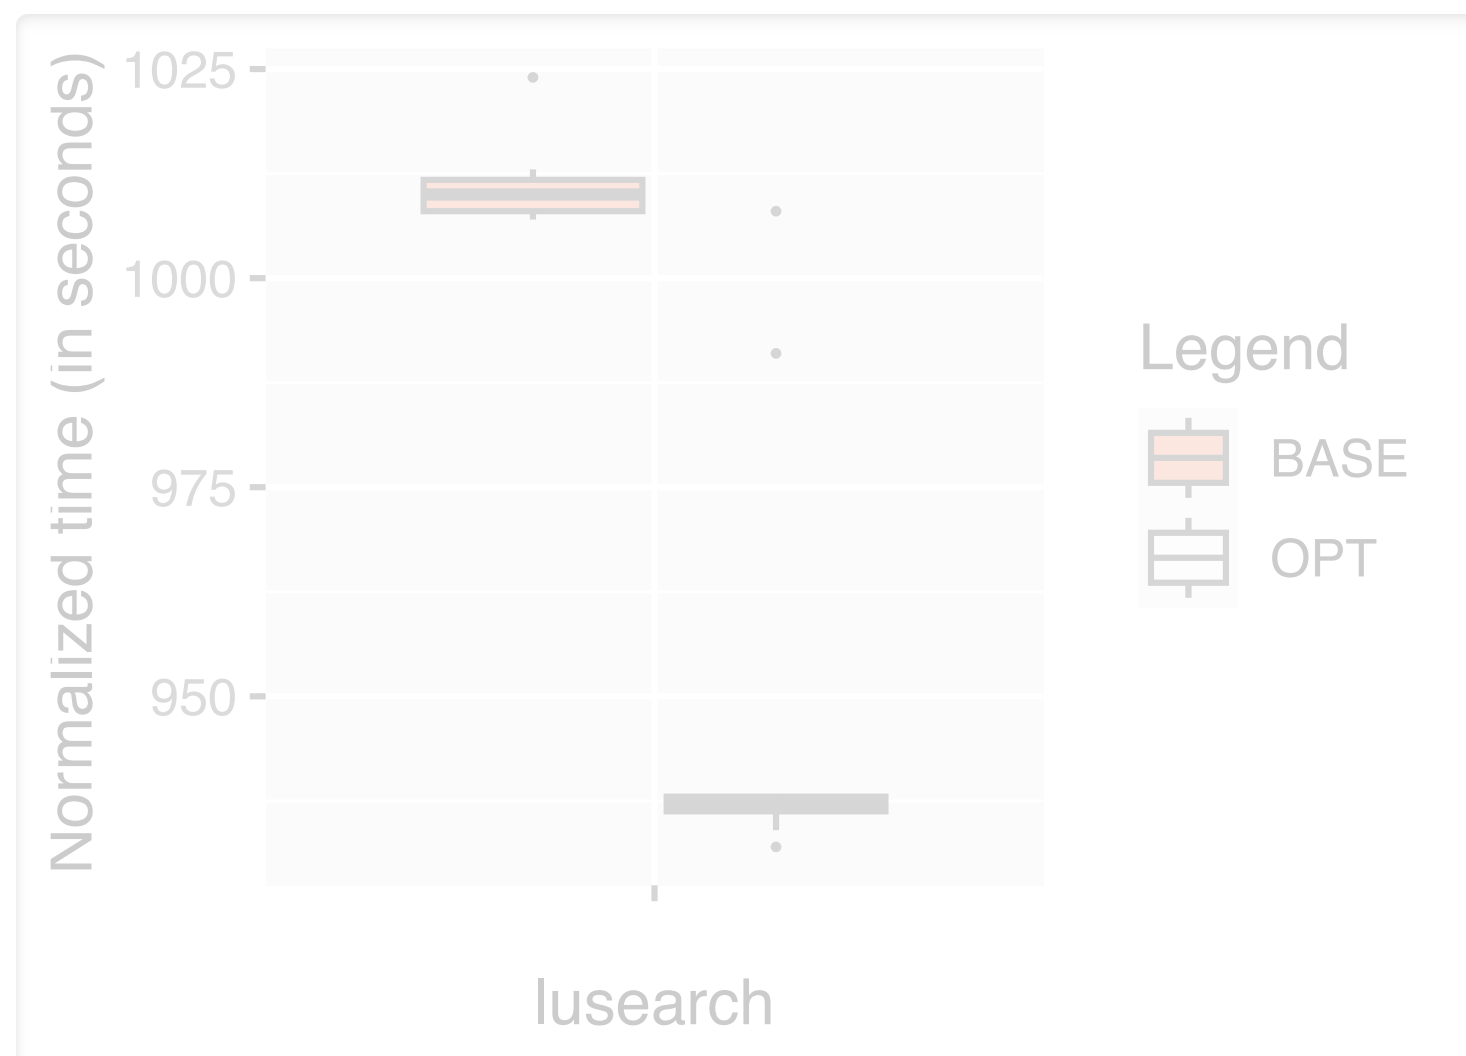
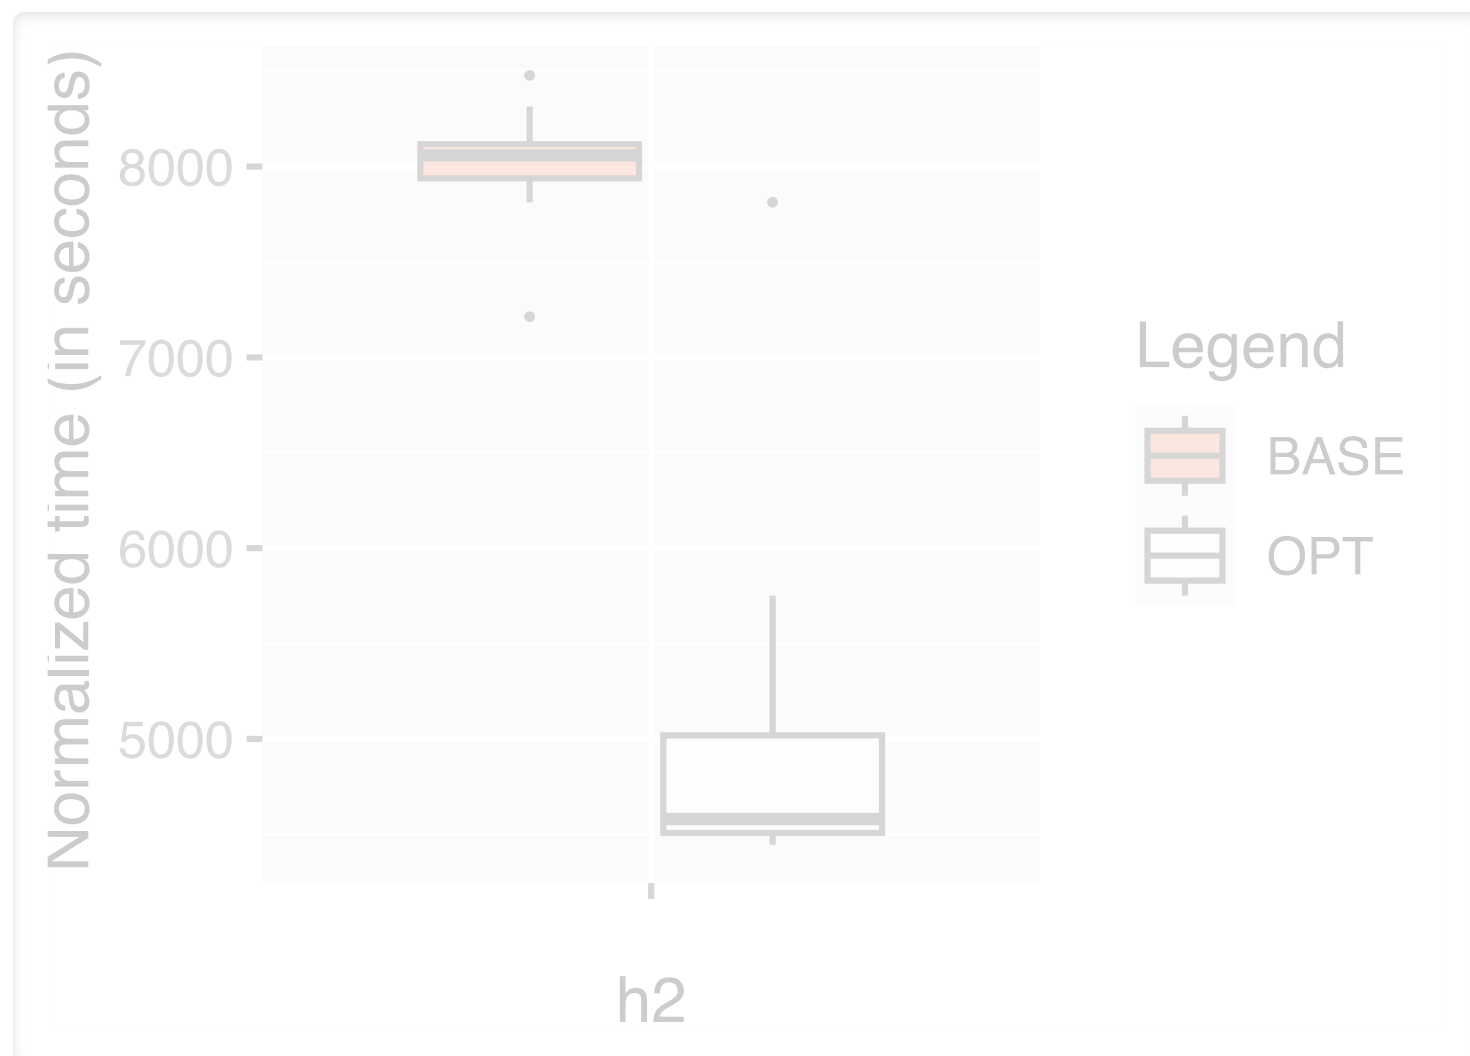
Performance



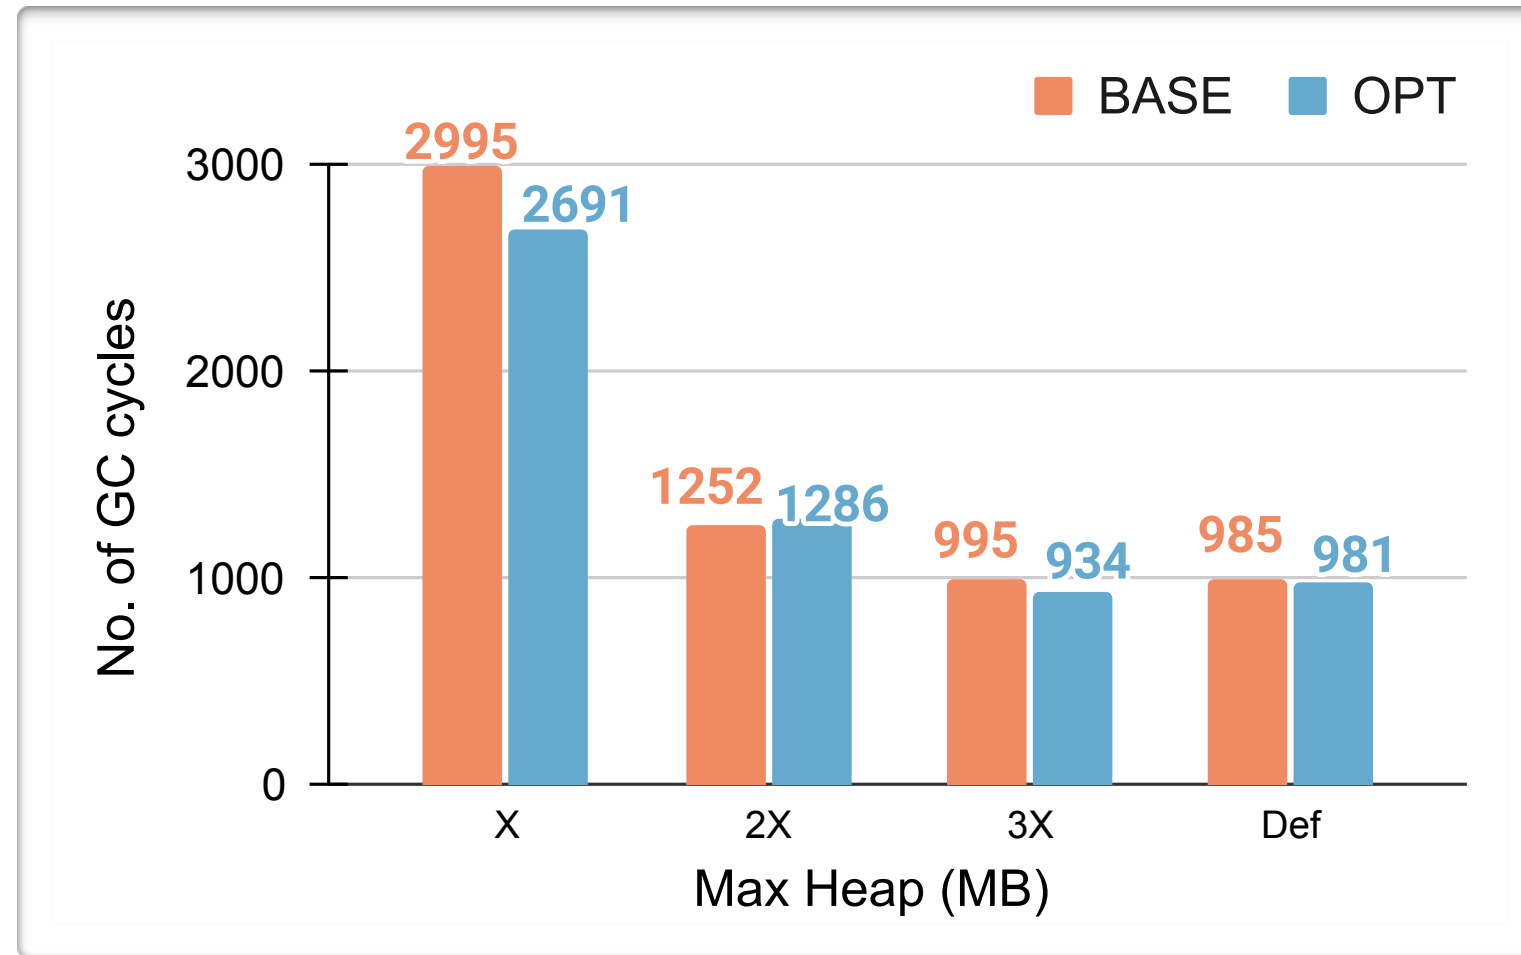
Performance



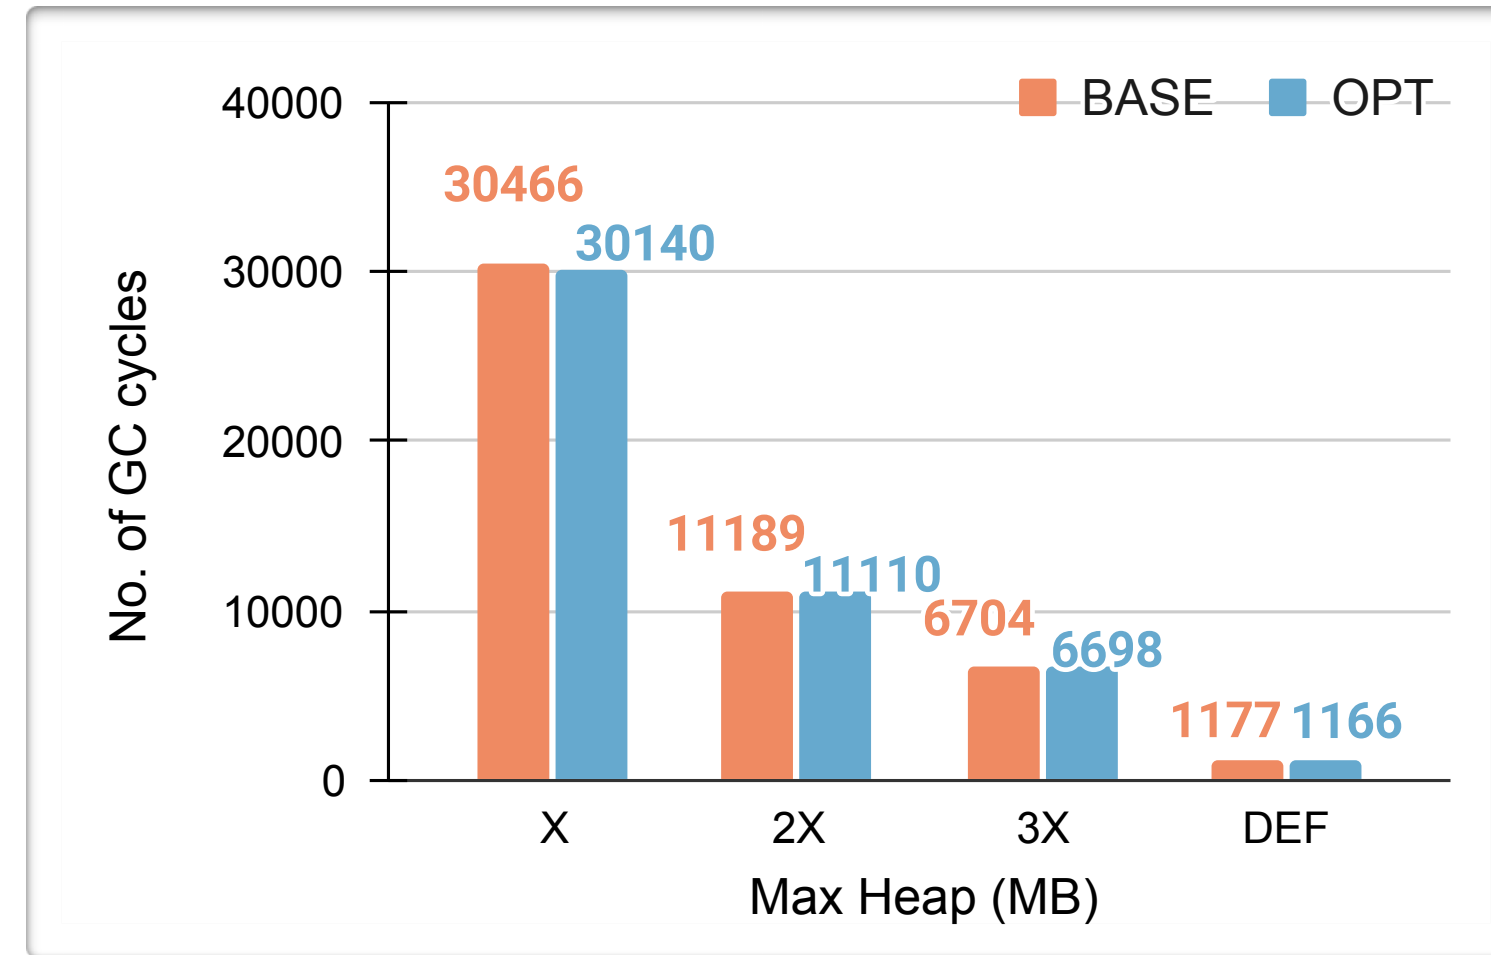
Performance Improvement: **8.8%↑**



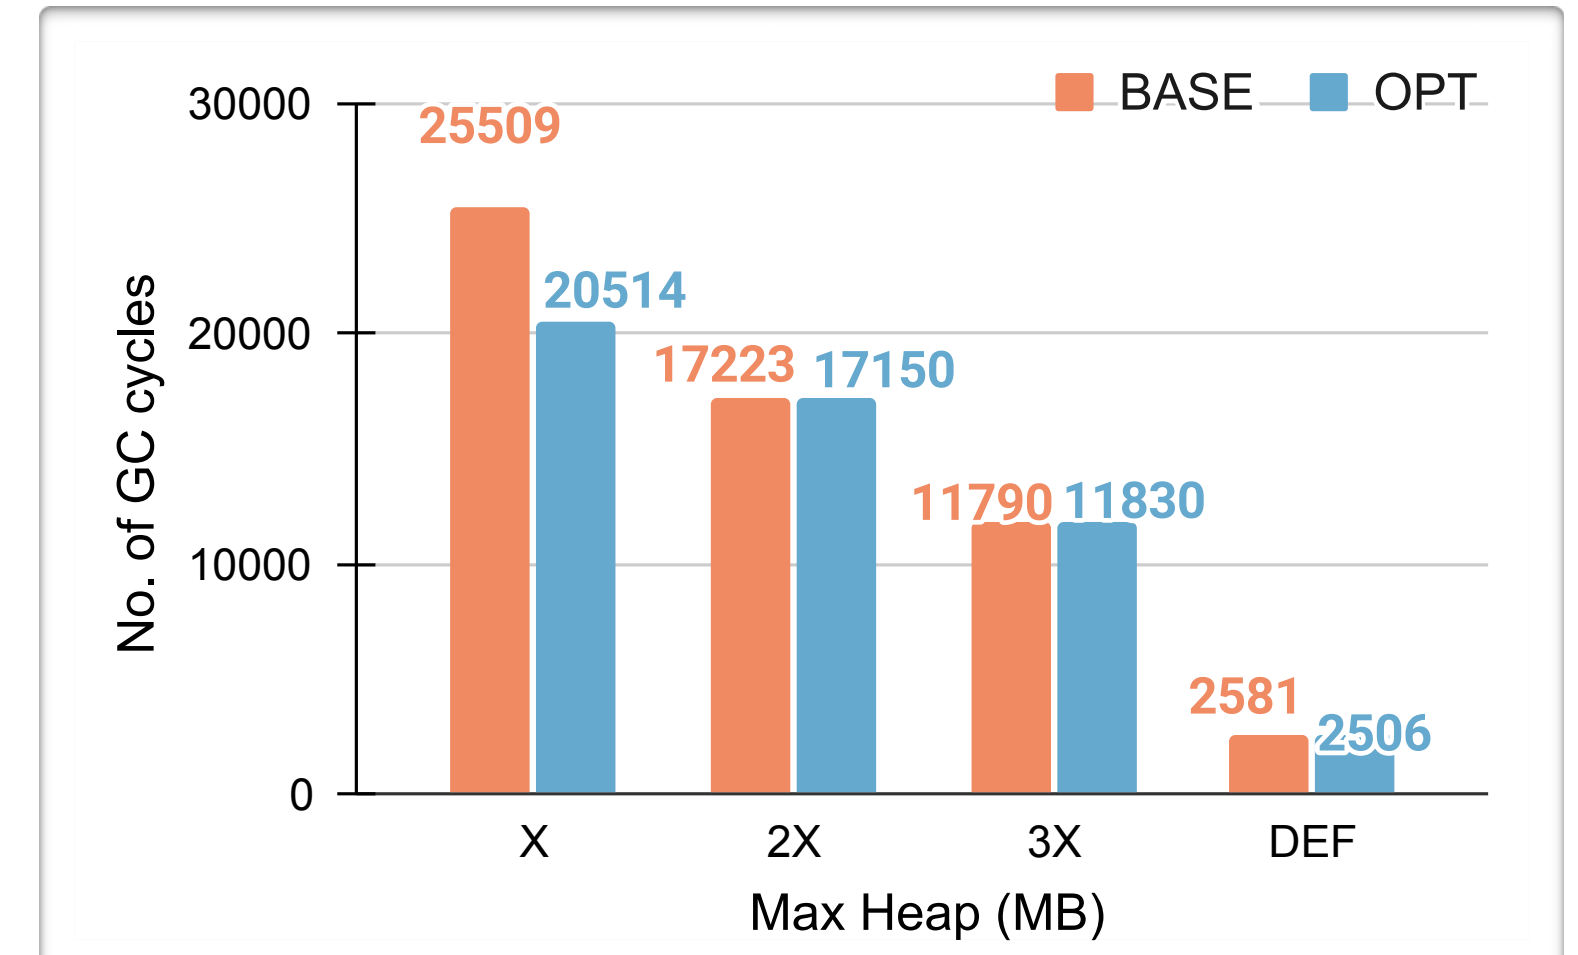
Garbage Collection



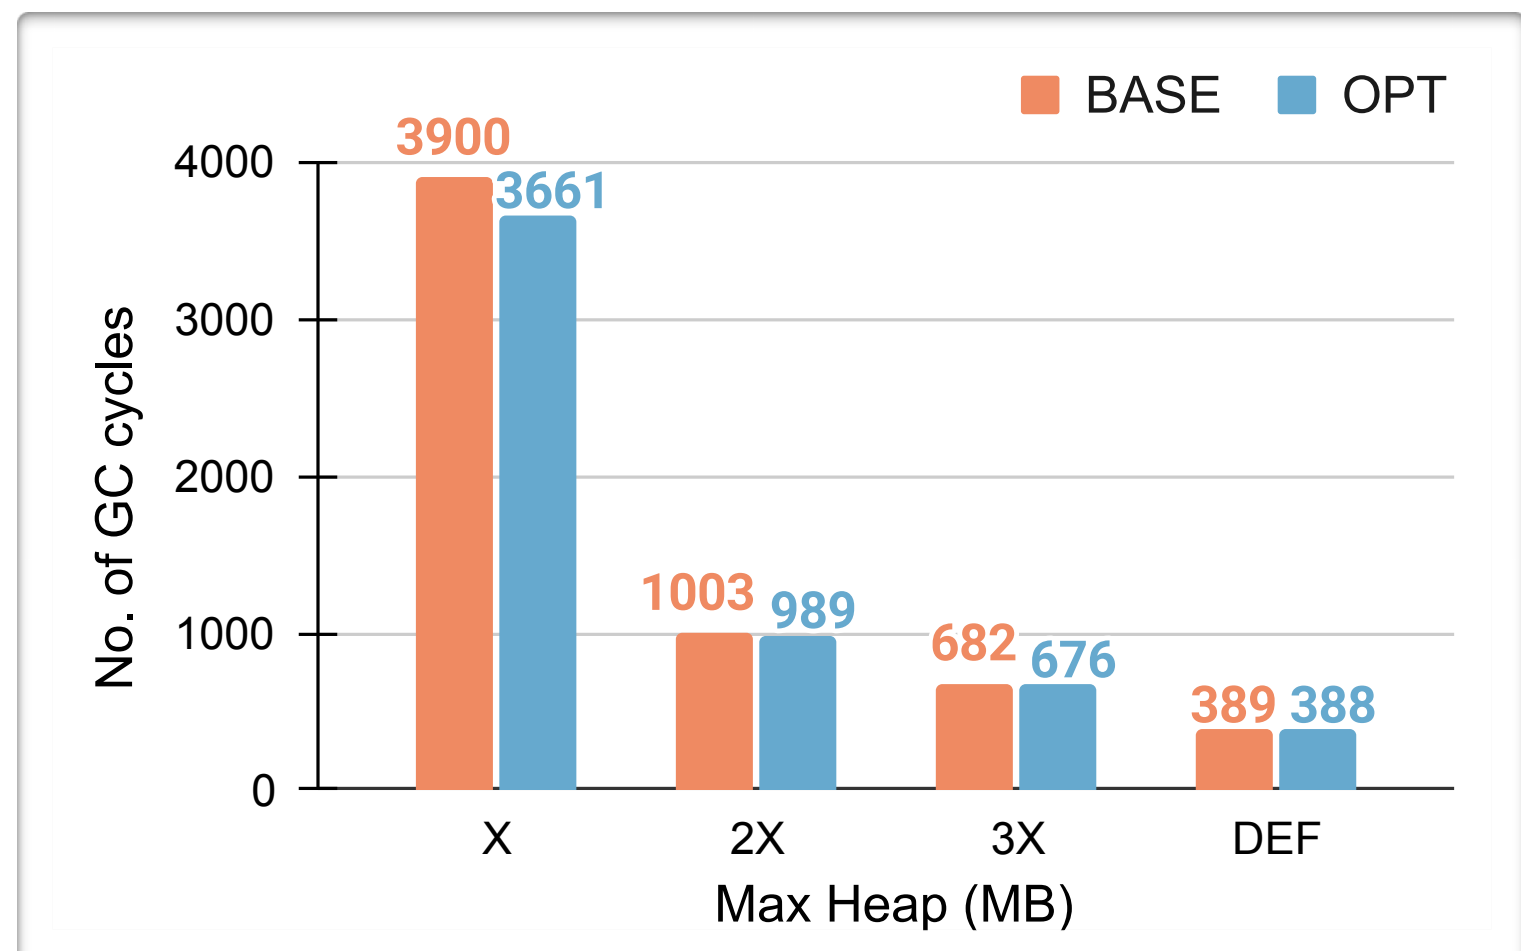
compiler



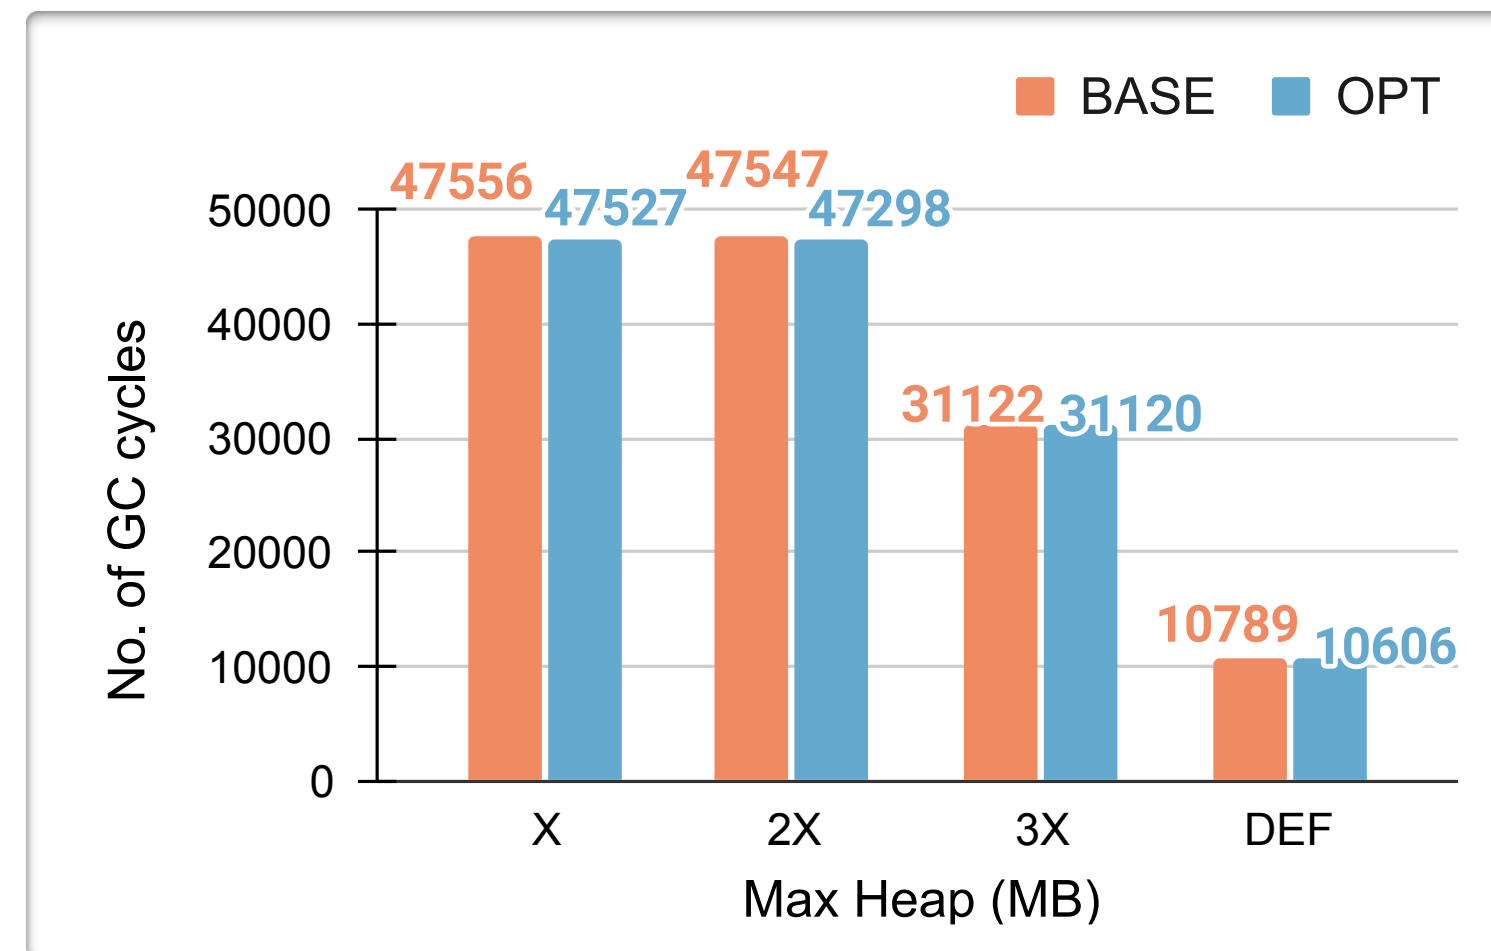
fop



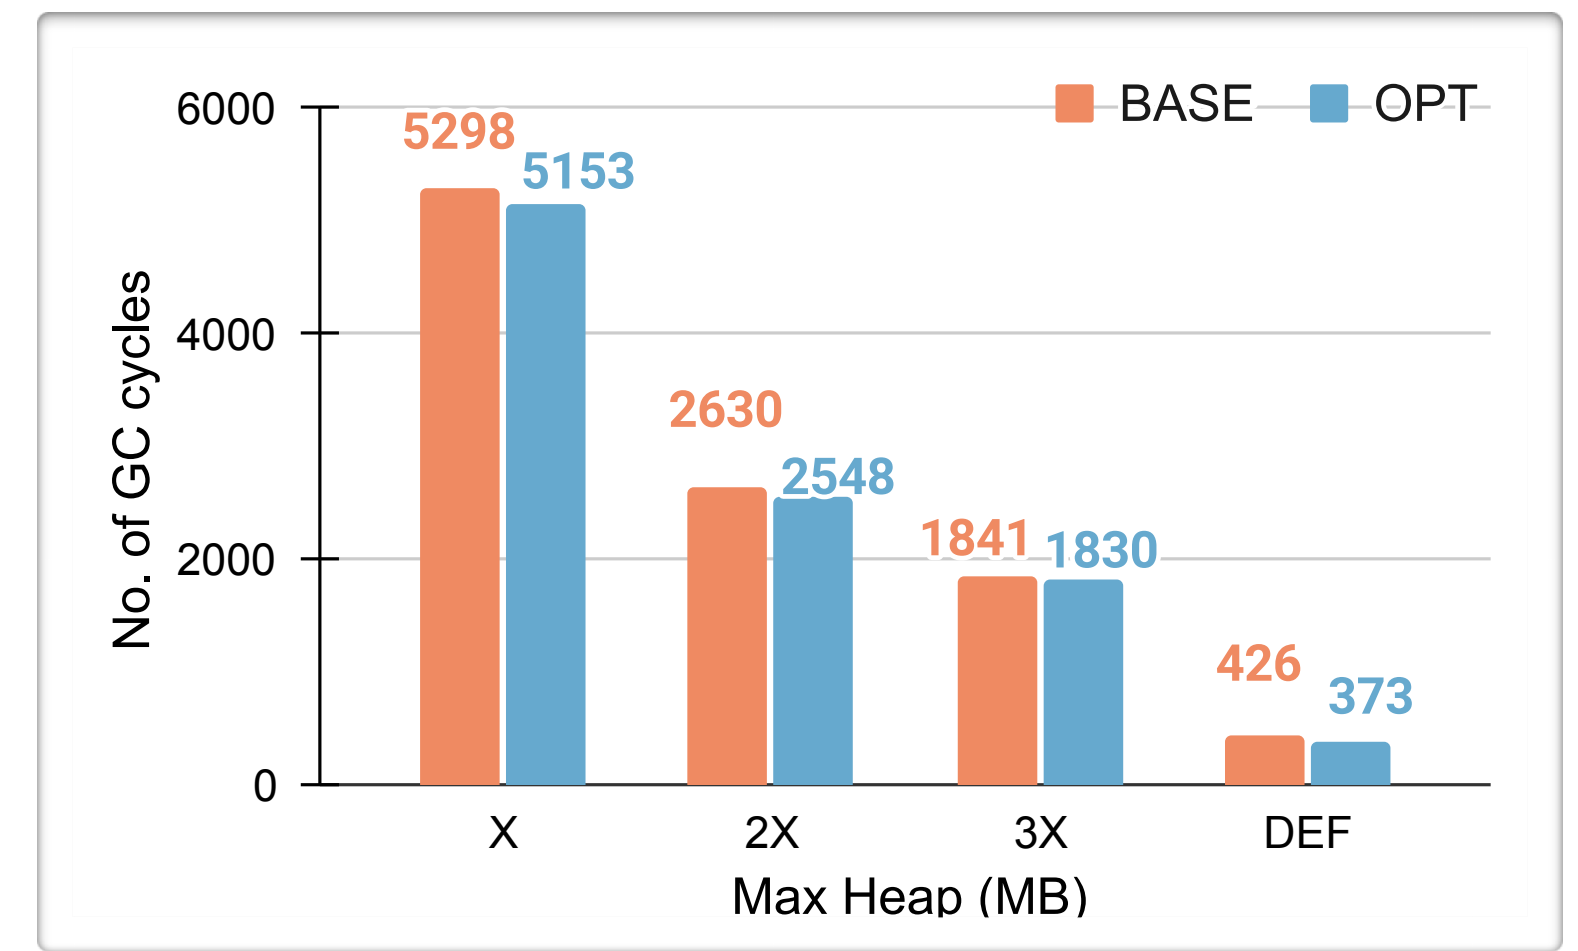
graphchi



h2

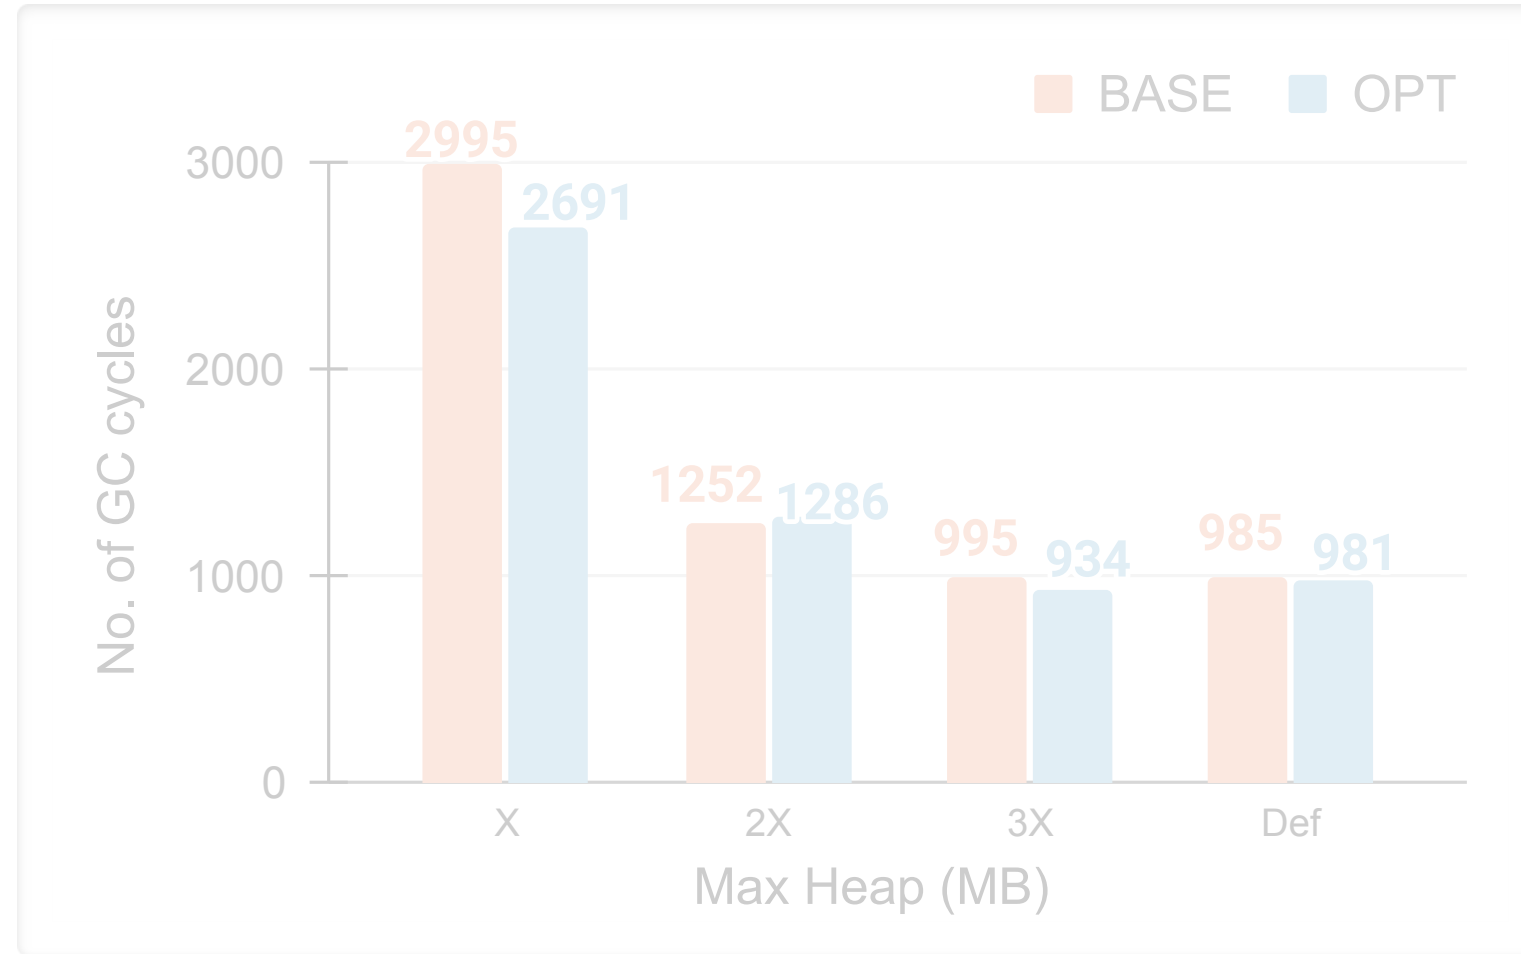


lusearch

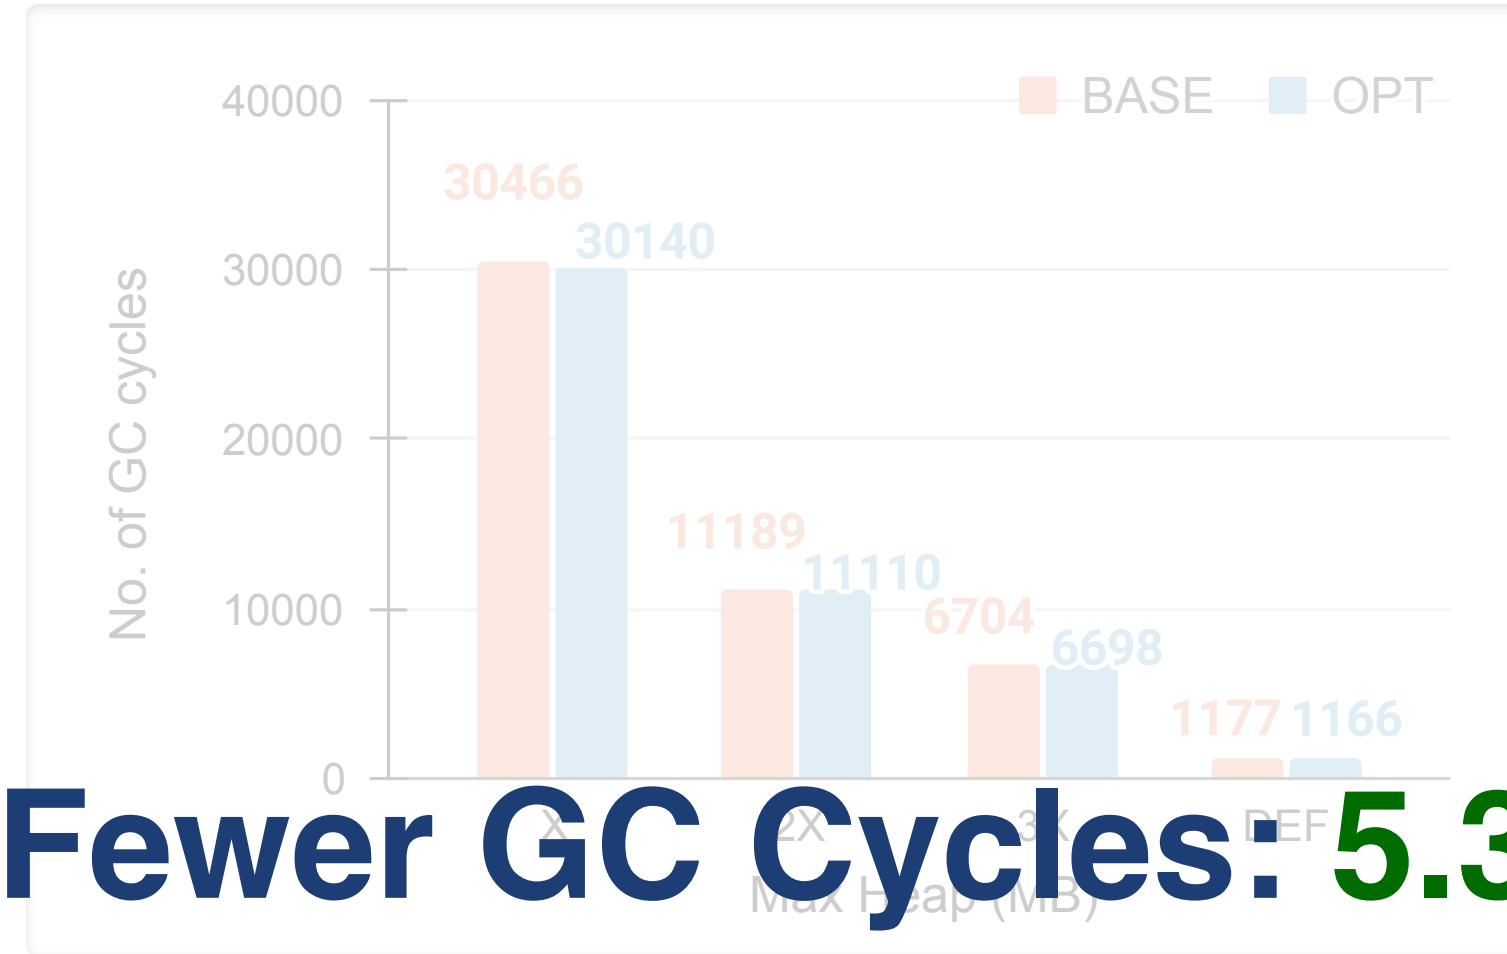


pmd

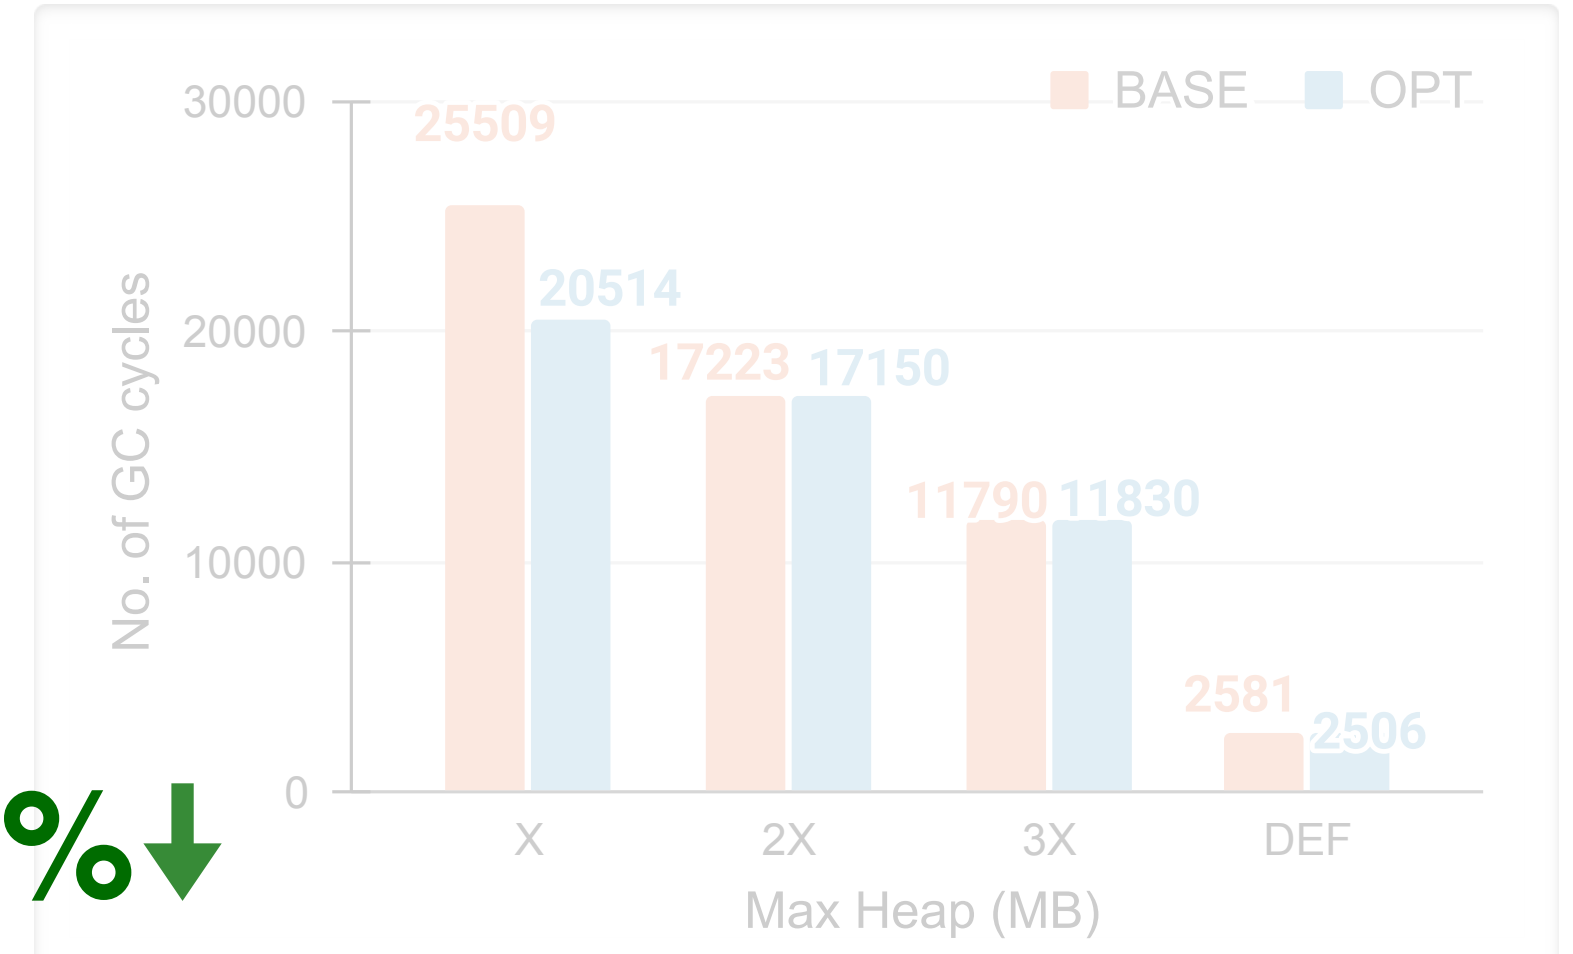
Garbage Collection



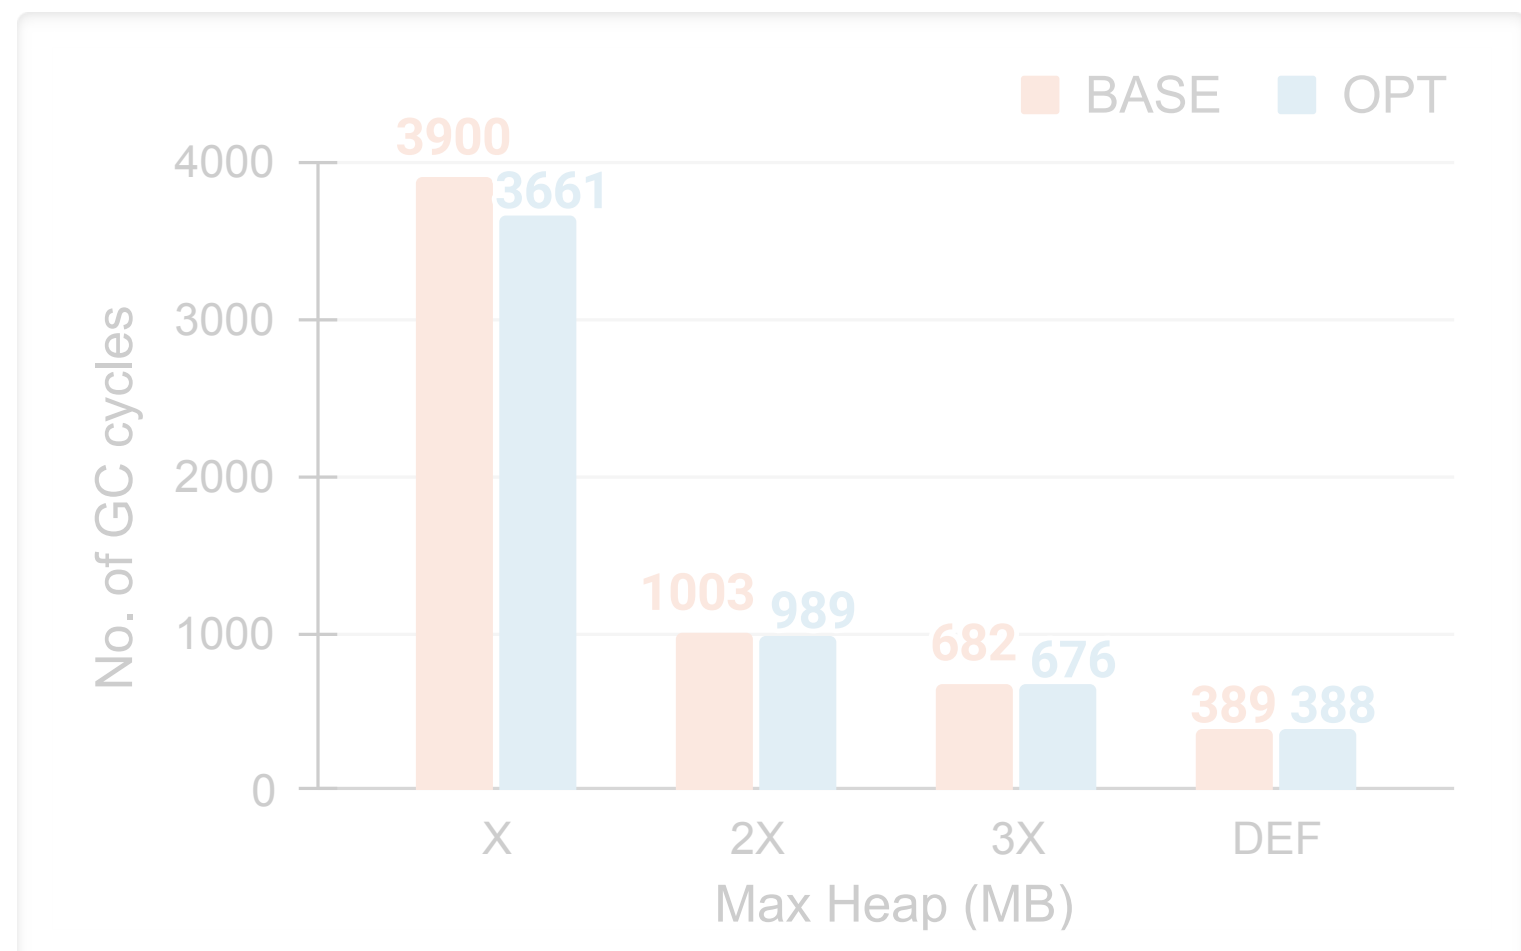
compiler



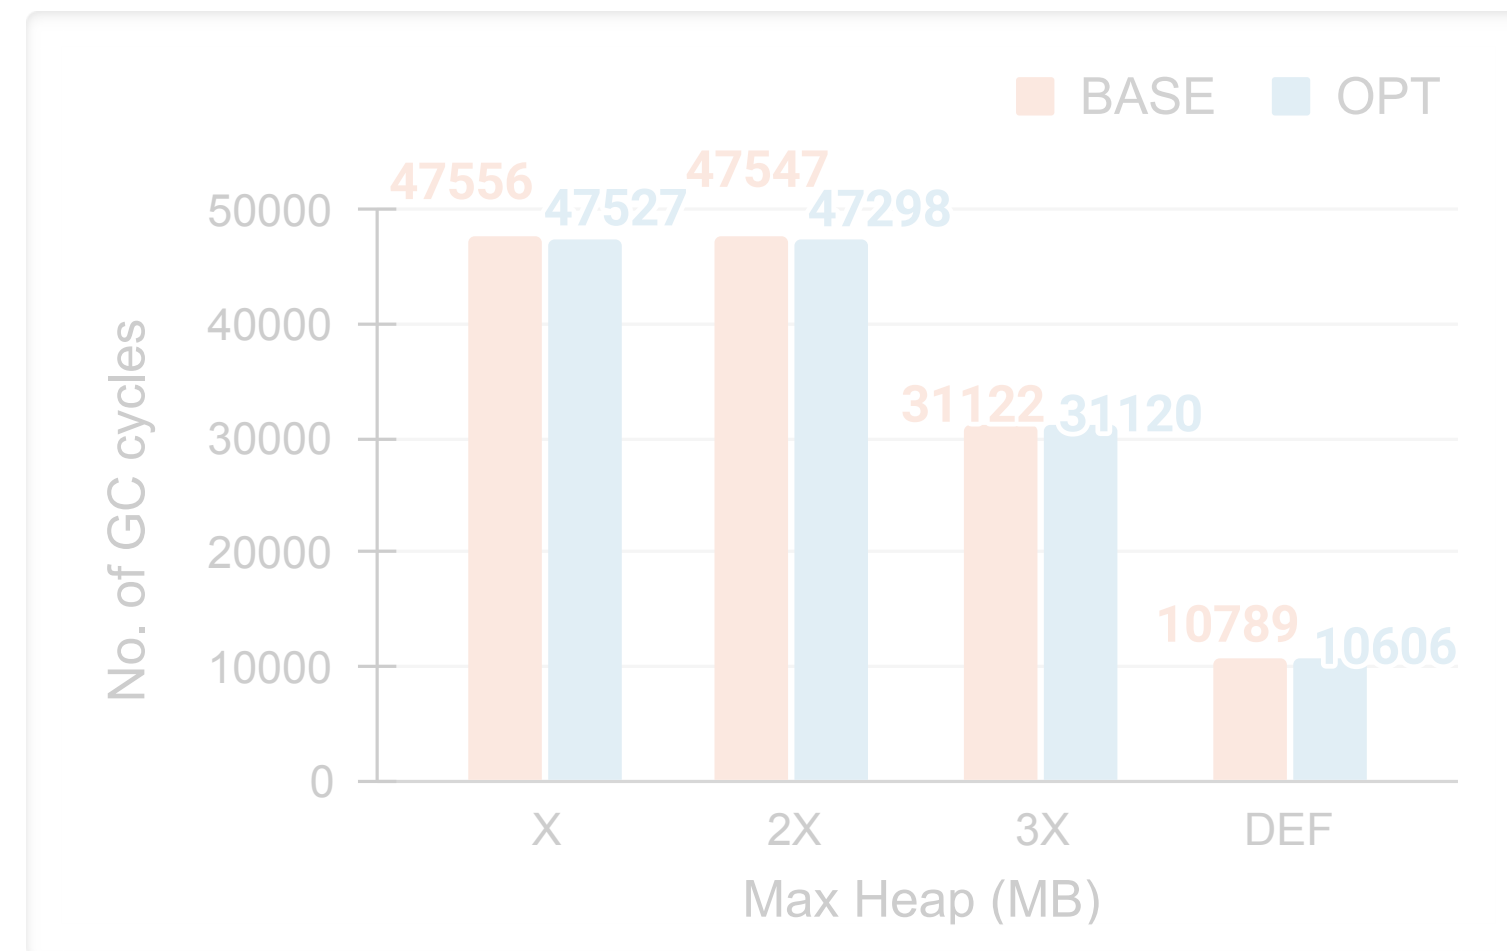
fop



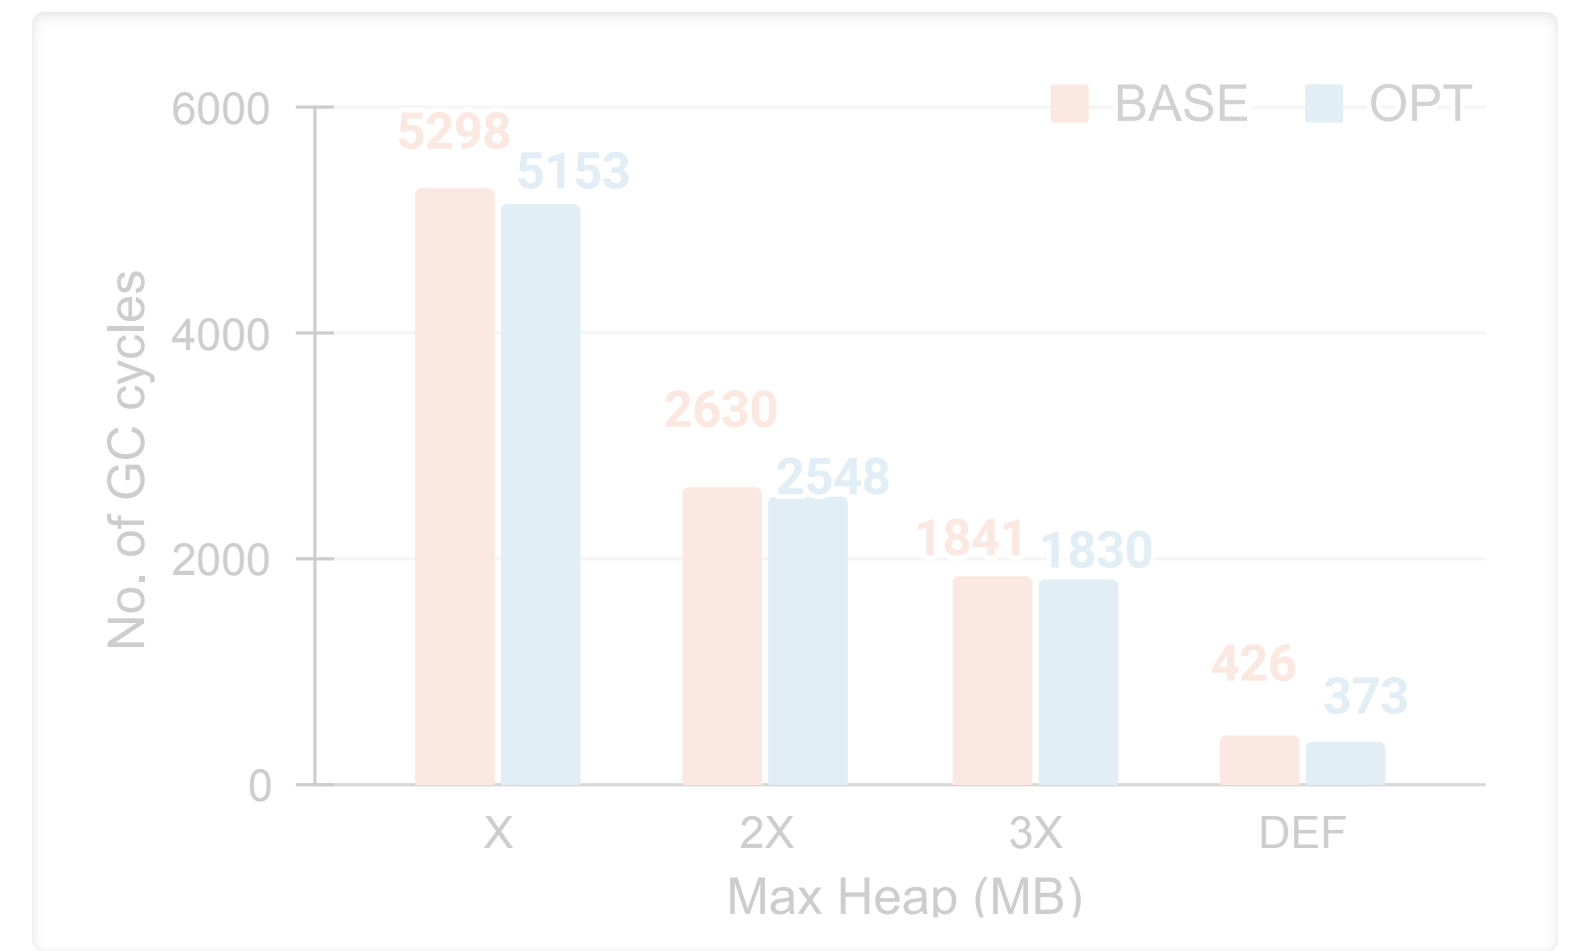
graphchi



h2



lusearch



pmd

Fewer GC Cycles: 5.3%↓

Take Aways

Take Aways

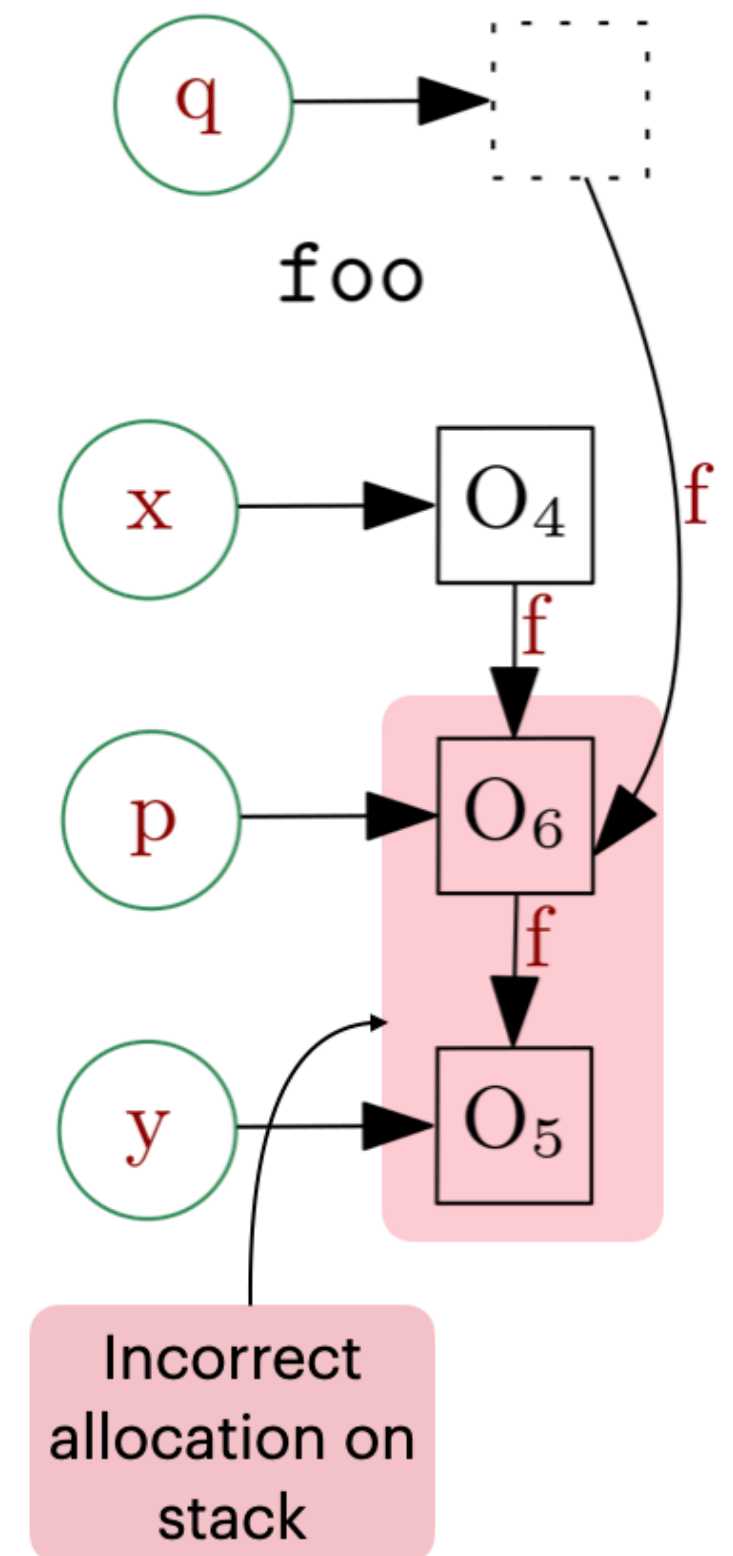
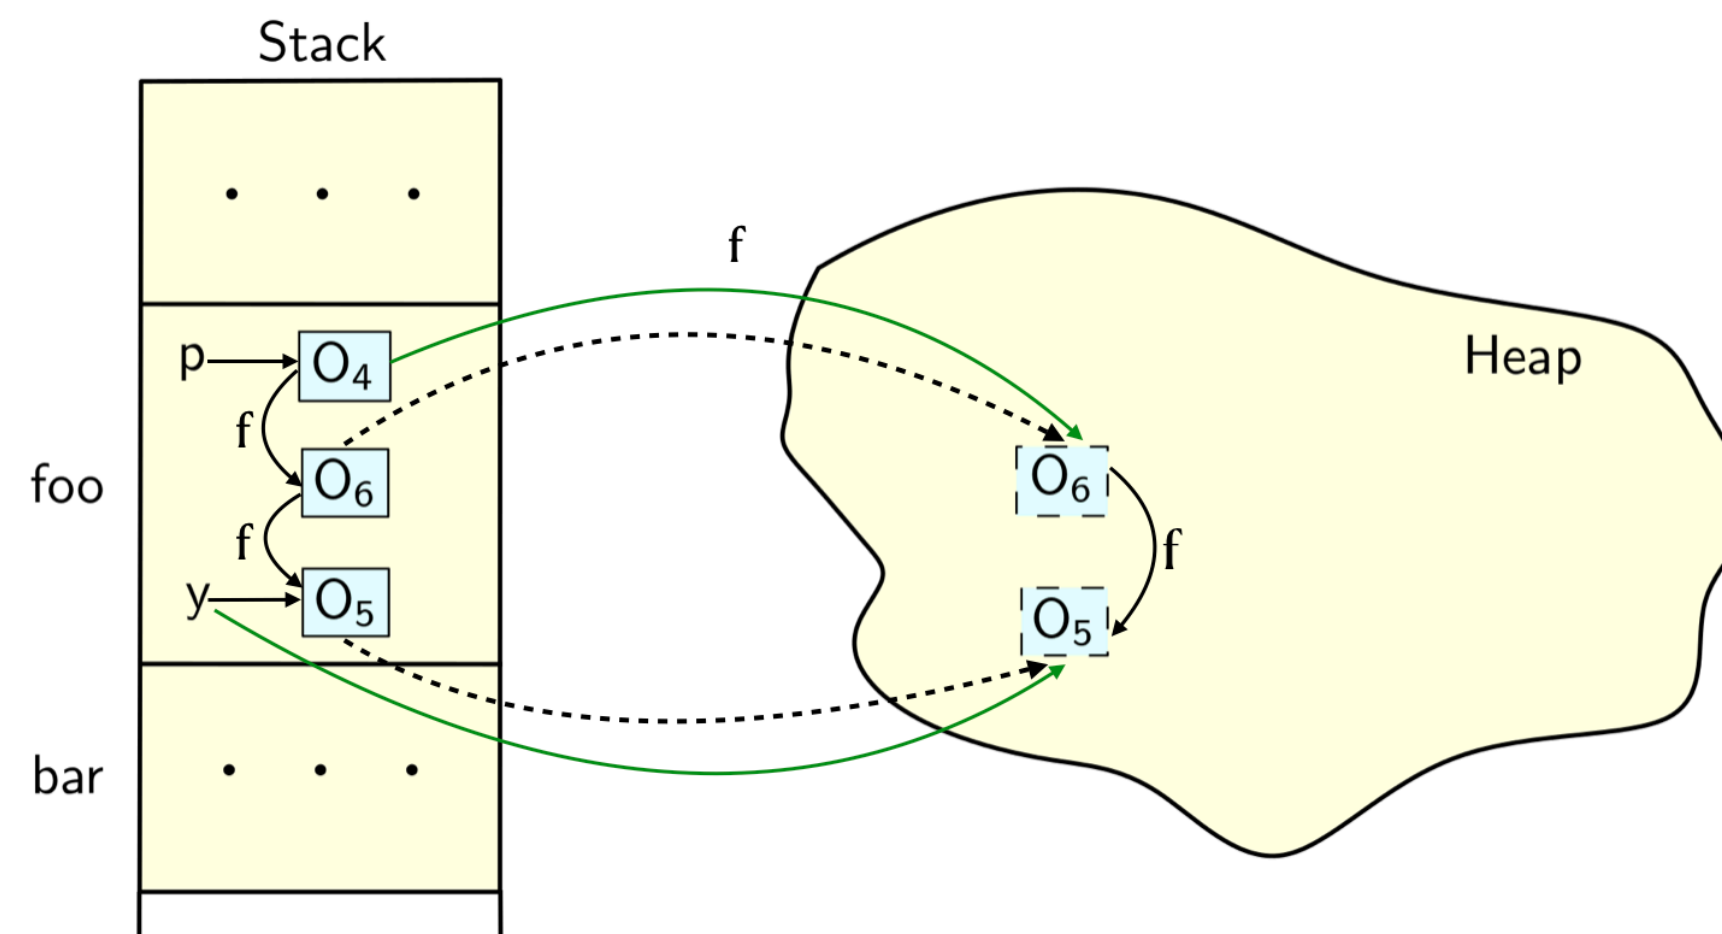
- **An important OO Optimization:**
Allocating method-local objects on the stack frames of their allocating methods.

Take Aways

- **An important OO Optimization:**
Allocating method-local objects on the stack frames of their allocating methods.
- Used static escape analysis to **optimistically allocate** identified objects on stack to improve the precision without thwarting the efficiency.

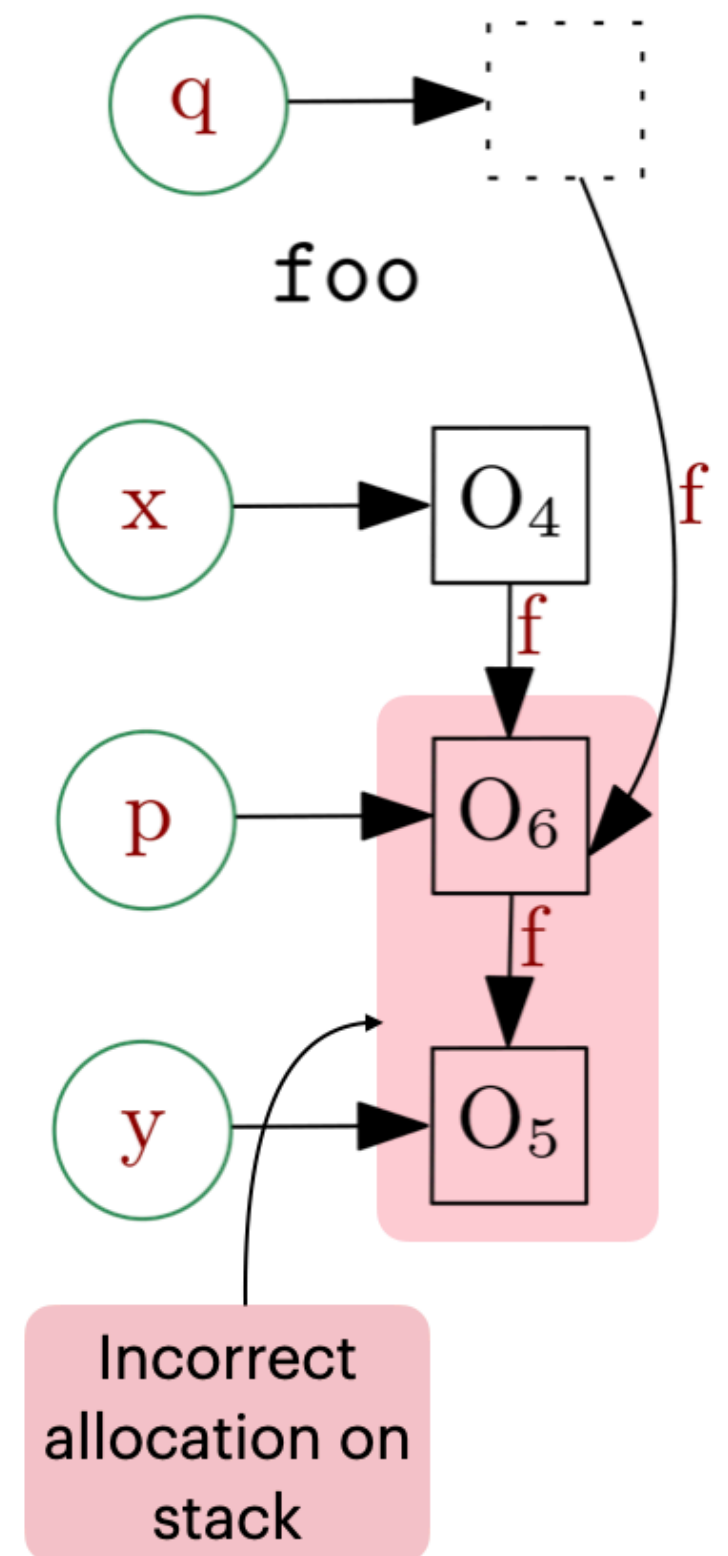
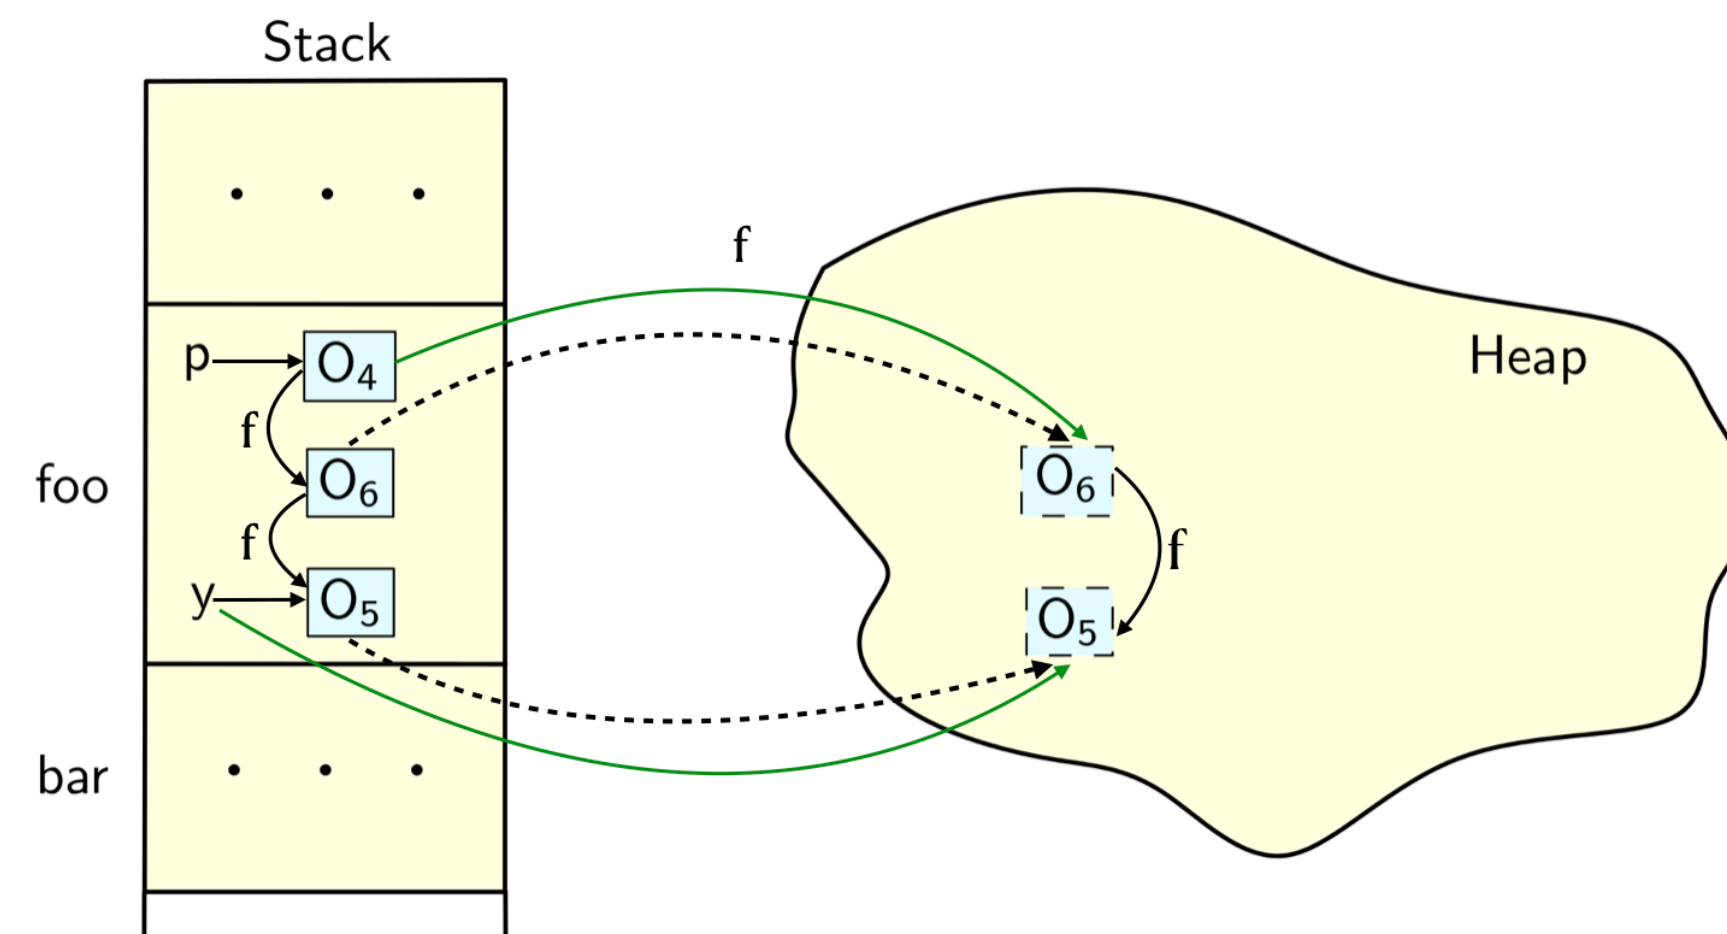
Take Aways

- An important OO Optimization: Allocating method-local objects on the stack frames of their allocating methods.
- Used static escape analysis to **optimistically allocate** identified objects on stack to improve the precision without thwarting the efficiency.
- Ensure **functional correctness** in cases static analysis results do not correspond to the runtime environment.



Take Aways

- **An important OO Optimization:** Allocating method-local objects on the stack frames of their allocating methods.
 - Used static escape analysis to **optimistically allocate** identified objects on stack to improve the precision without thwarting the efficiency.
 - Ensure **functional correctness** in cases static analysis results do not correspond to the runtime environment.
- **Overall, one of the first approaches to soundly and efficiently use static (offline) analysis results in a JIT compiler!**



Take Aways



Optimistic Stack Allocation and Dynamic Heapification for Managed Runtimes

ADITYA ANAND, Indian Institute of Technology Bombay, India

SOLAI ADITHYA, Indian Institute of Technology Mandi, India

SWAPNIL RUSTAGI, Indian Institute of Technology Mandi, India

PRIYAM SETH, Indian Institute of Technology Mandi, India

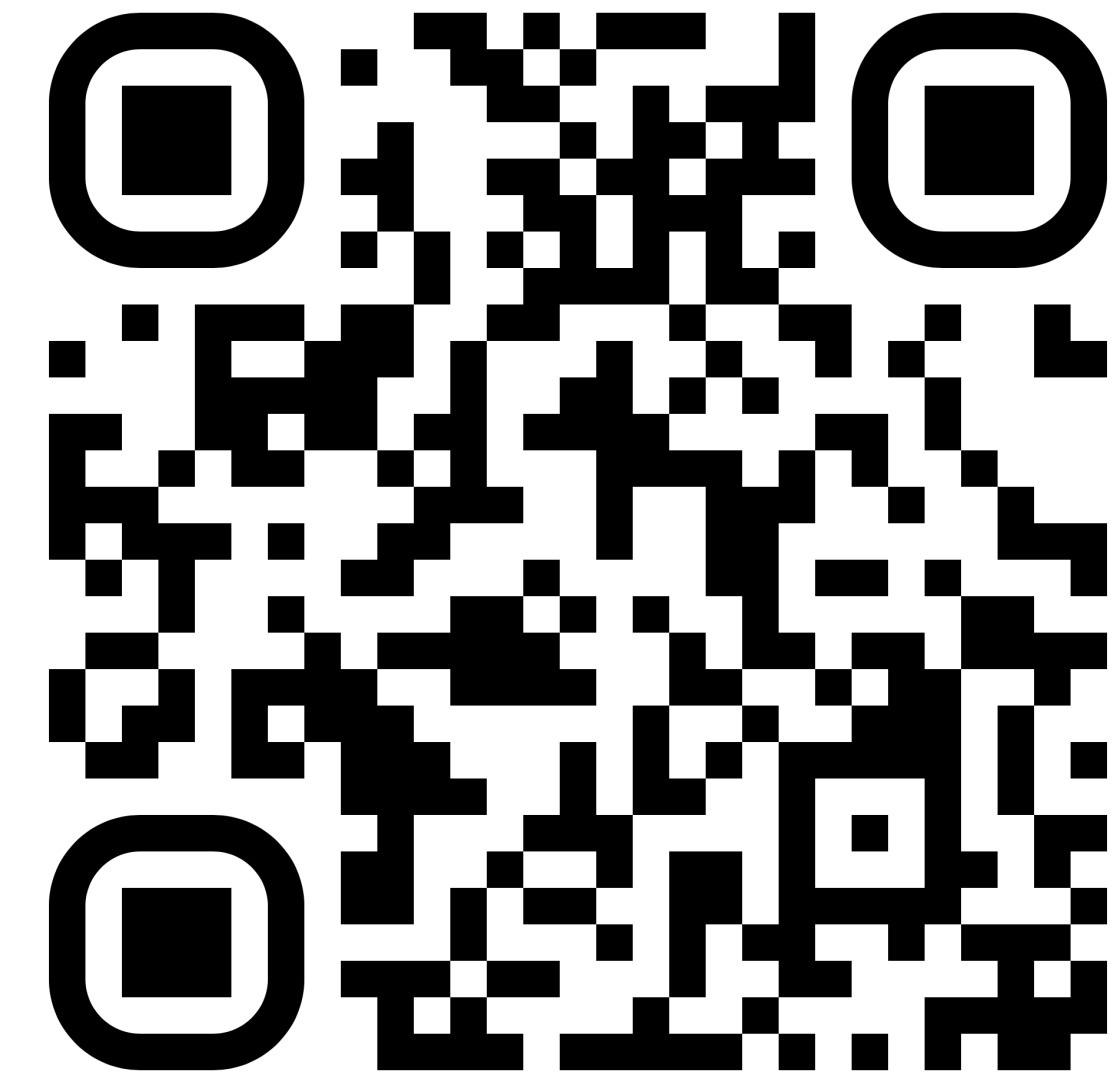
VIJAY SUNDARESAN, IBM Canada Lab, Canada

DARYL MAIER, IBM Canada Lab, Canada

V. KRISHNA NANDIVADA, Indian Institute of Technology Madras, India

MANAS THAKUR, Indian Institute of Technology Bombay, India

The runtimes of managed object-oriented languages such as Java allocate objects on the heap, and rely on automatic garbage collection (GC) techniques for freeing up unused objects. Most such runtimes also consist of just-in-time (JIT) compilers that optimize memory access and GC times by employing *escape analysis*: an object that does not escape (outlive) its allocating method can be allocated on (and freed up with) the stack frame of the corresponding method. However, in order to minimize the time spent in JIT compilation, the scope of such useful analyses is quite limited, thereby restricting their precision significantly. On the contrary, even though it is feasible to perform precise program analyses statically, it is not possible to use their results in a managed runtime without a closed-world assumption. In this paper, we propose a static+dynamic scheme that allows one to harness the results of a precise static escape analysis for allocating objects on stack, while taking care of both soundness and efficiency concerns in the runtime.



Paper Link !!

Take Aways



Optimistic Stack Allocation and Dynamic Heapification for Managed Runtimes

ADITYA ANAND, Indian Institute of Technology Bombay, India

SOLAI ADITHYA, Indian Institute of Technology Mandi, India

SWAPNIL RUSTAGI, Indian Institute of Technology Mandi, India

PRIYAM SETH, Indian Institute of Technology Mandi, India

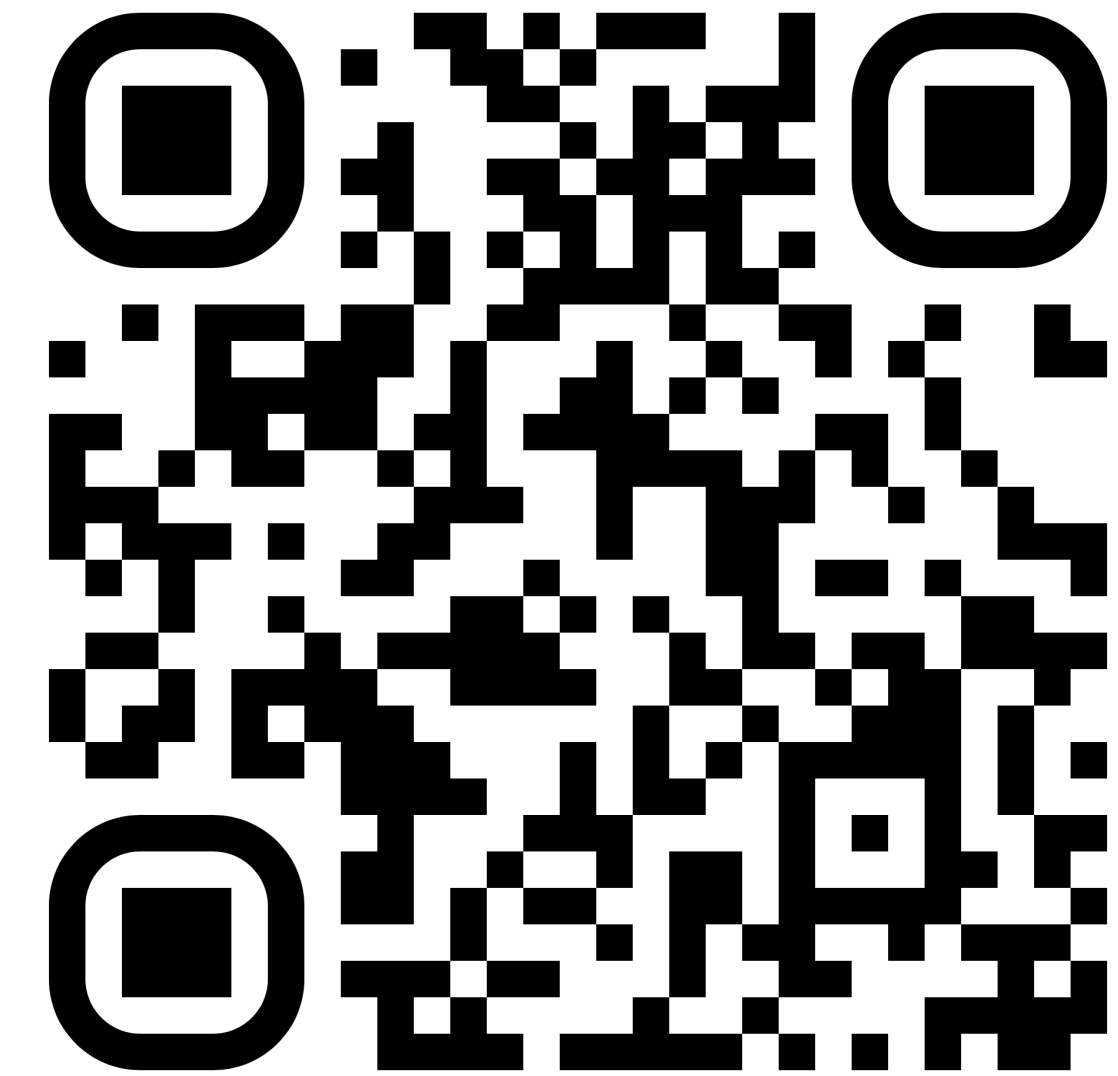
VIJAY SUNDARESAN, IBM Canada Lab, Canada

DARYL MAIER, IBM Canada Lab, Canada

V. KRISHNA NANDIVADA, Indian Institute of Technology Madras, India

MANAS THAKUR, Indian Institute of Technology Bombay, India

The runtimes of managed object-oriented languages such as Java allocate objects on the heap, and rely on automatic garbage collection (GC) techniques for freeing up unused objects. Most such runtimes also consist of just-in-time (JIT) compilers that optimize memory access and GC times by employing *escape analysis*: an object that does not escape (outlive) its allocating method can be allocated on (and freed up with) the stack frame of the corresponding method. However, in order to minimize the time spent in JIT compilation, the scope of such useful analyses is quite limited, thereby restricting their precision significantly. On the contrary, even though it is feasible to perform precise program analyses statically, it is not possible to use their results in a managed runtime without a closed-world assumption. In this paper, we propose a static+dynamic scheme that allows one to harness the results of a precise static escape analysis for allocating objects on stack, while taking care of both soundness and efficiency concerns in the runtime.



Paper Link !!

Thank You!!

Take Aways



Optimistic Stack Allocation and Dynamic Heapification for Managed Runtimes

ADITYA ANAND, Indian Institute of Technology Bombay, India

SOLAI ADITHYA, Indian Institute of Technology Mandi, India

SWAPNIL RUSTAGI, Indian Institute of Technology Mandi, India

PRIYAM SETH, Indian Institute of Technology Mandi, India

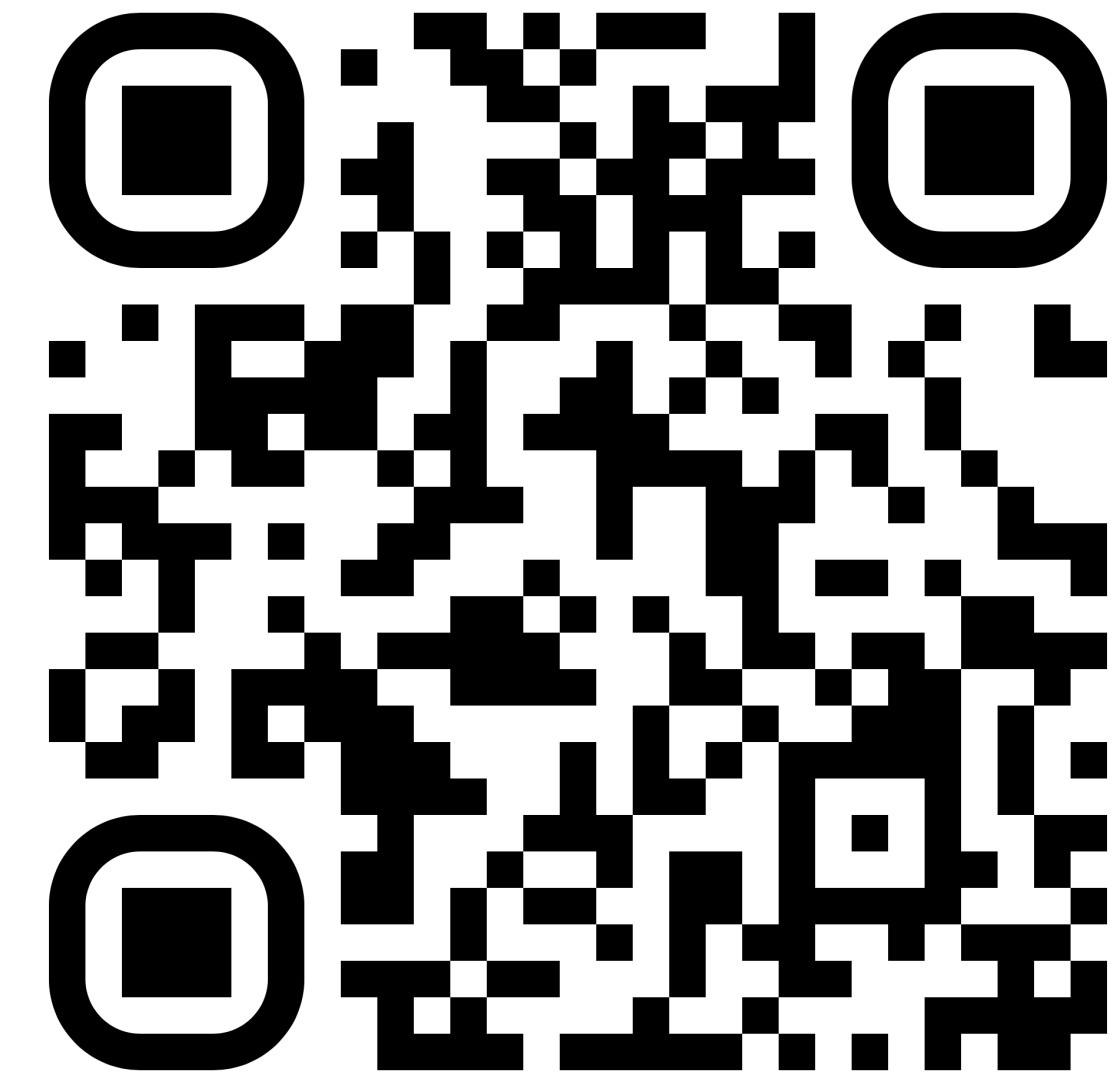
VIJAY SUNDARESAN, IBM Canada Lab, Canada

DARYL MAIER, IBM Canada Lab, Canada

V. KRISHNA NANDIVADA, Indian Institute of Technology Madras, India

MANAS THAKUR, Indian Institute of Technology Bombay, India

The runtimes of managed object-oriented languages such as Java allocate objects on the heap, and rely on automatic garbage collection (GC) techniques for freeing up unused objects. Most such runtimes also consist of just-in-time (JIT) compilers that optimize memory access and GC times by employing *escape analysis*: an object that does not escape (outlive) its allocating method can be allocated on (and freed up with) the stack frame of the corresponding method. However, in order to minimize the time spent in JIT compilation, the scope of such useful analyses is quite limited, thereby restricting their precision significantly. On the contrary, even though it is feasible to perform precise program analyses statically, it is not possible to use their results in a managed runtime without a closed-world assumption. In this paper, we propose a static+dynamic scheme that allows one to harness the results of a precise static escape analysis for allocating objects on stack, while taking care of both soundness and efficiency concerns in the runtime.



Paper Link !!

Thank You!! Questions?

Take Aways

- **An important OO Optimization:** Allocating method-local objects on the stack frames of their allocating methods.
- Used static escape analysis to **optimistically allocate** identified objects on stack to improve the precision without thwarting the efficiency.
- Ensure **functional correctness** in cases static analysis results do not correspond to the runtime environment.
- **Overall, one of the first approaches to soundly and efficiently use static (offline) analysis results in a JIT compiler!**

